

# Formalizing Connections Between Motion Planning and Machine Learning



Siddhartha Srinivasa  
Boeing Endowed Professor  
University of Washington

**W** PAUL G. ALLEN SCHOOL  
OF COMPUTER SCIENCE & ENGINEERING

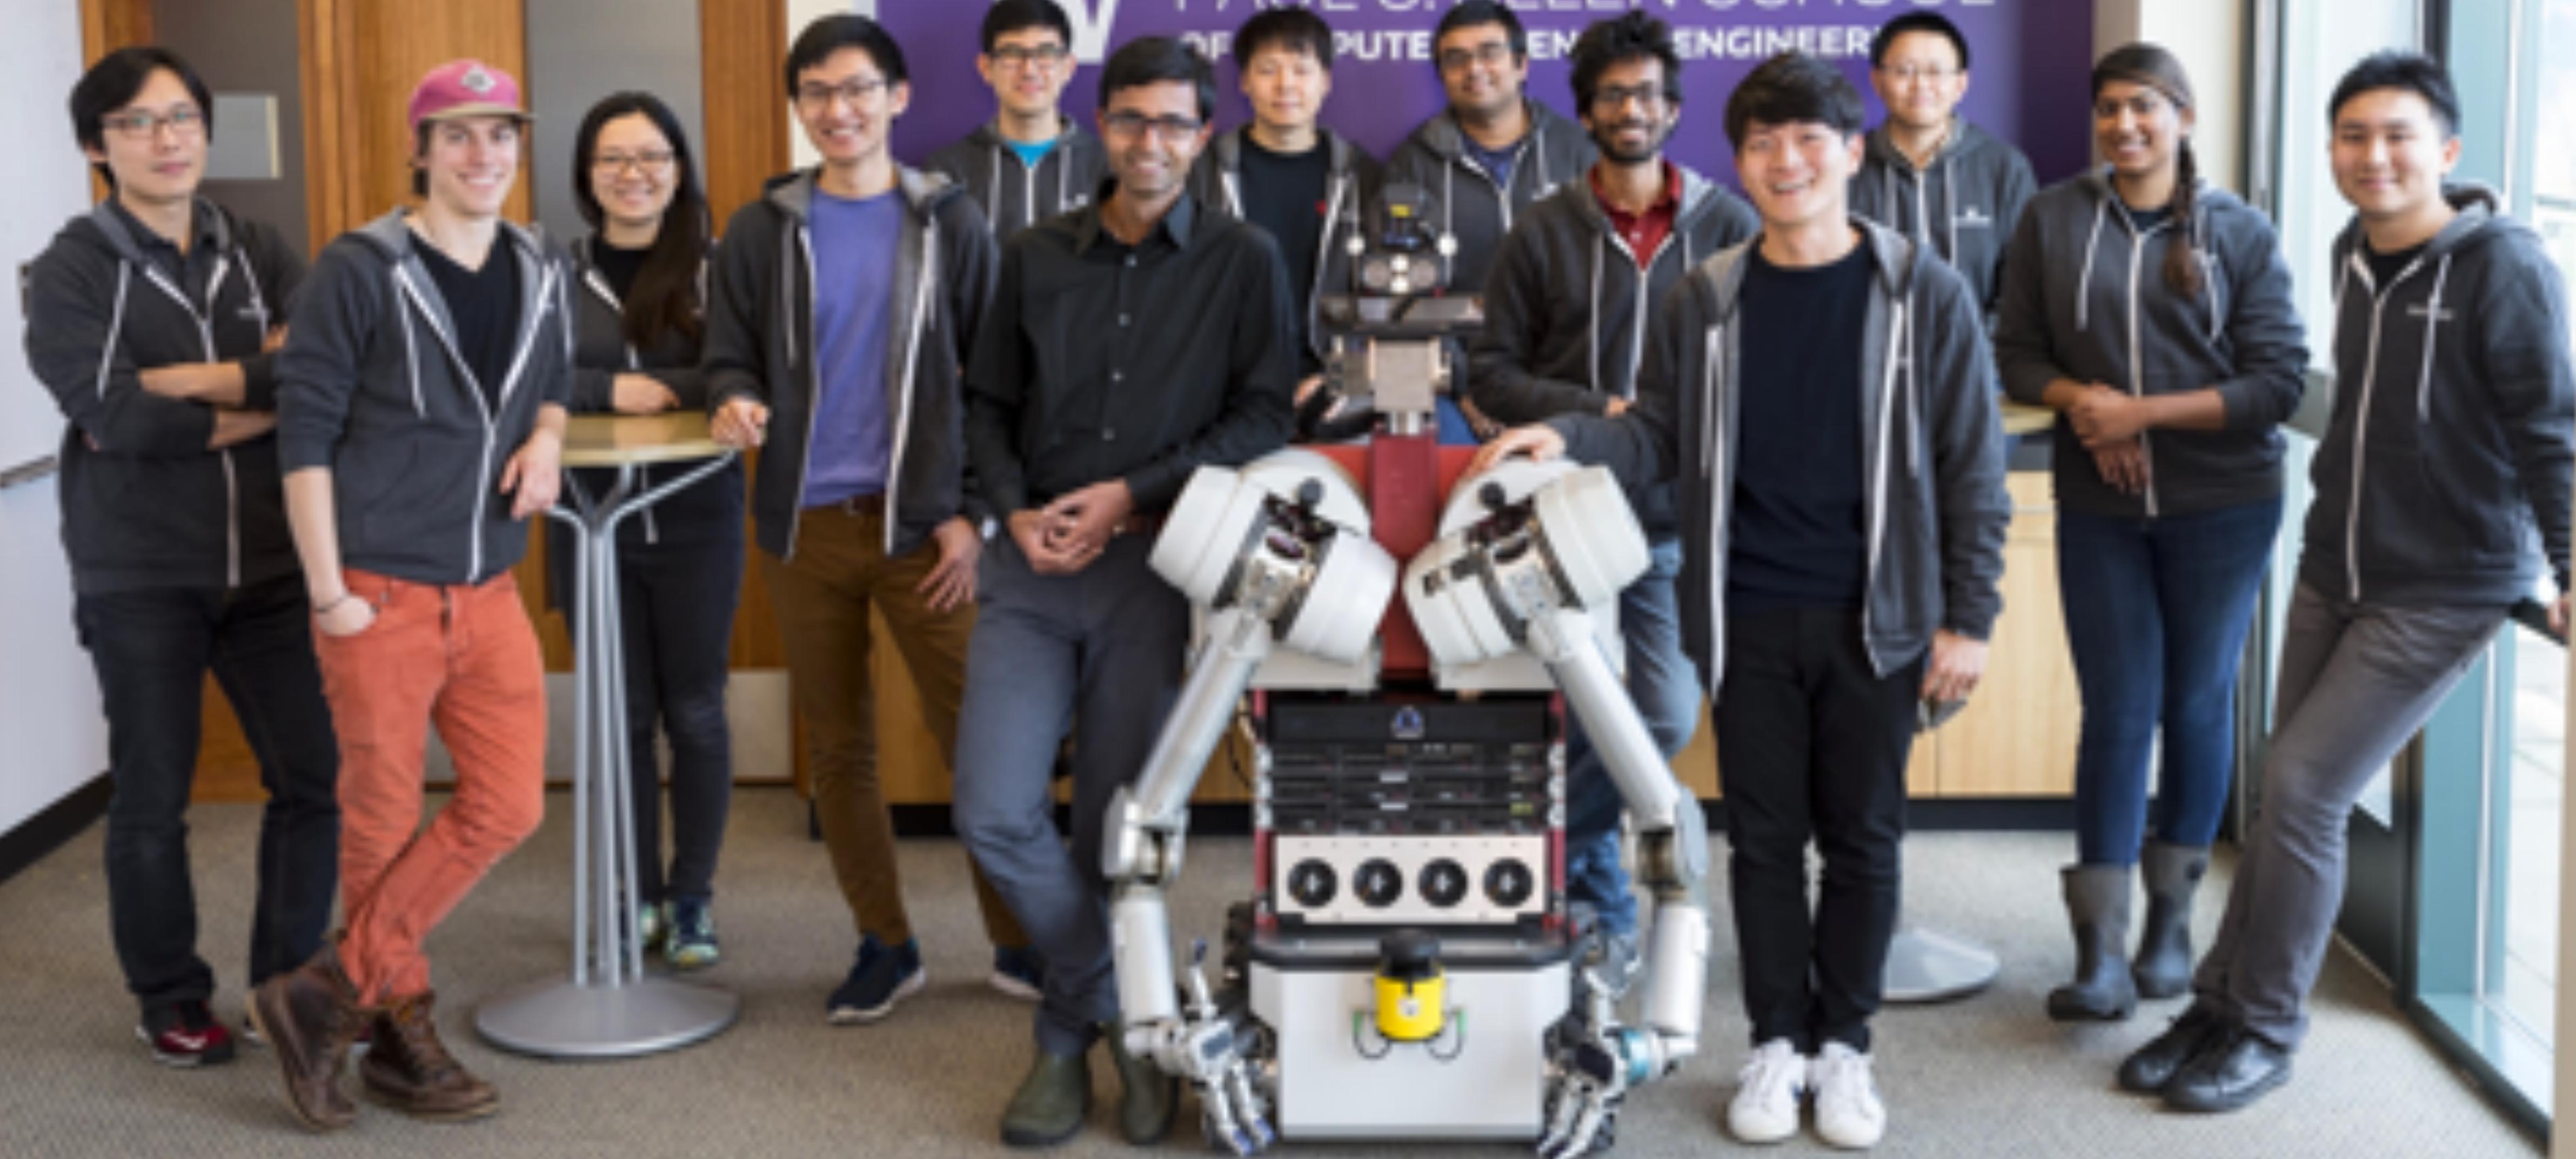


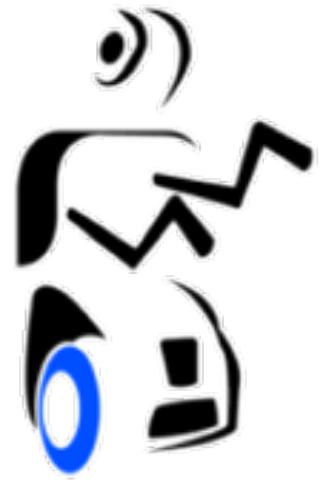
# Problems I Want You to Solve So I can Retire



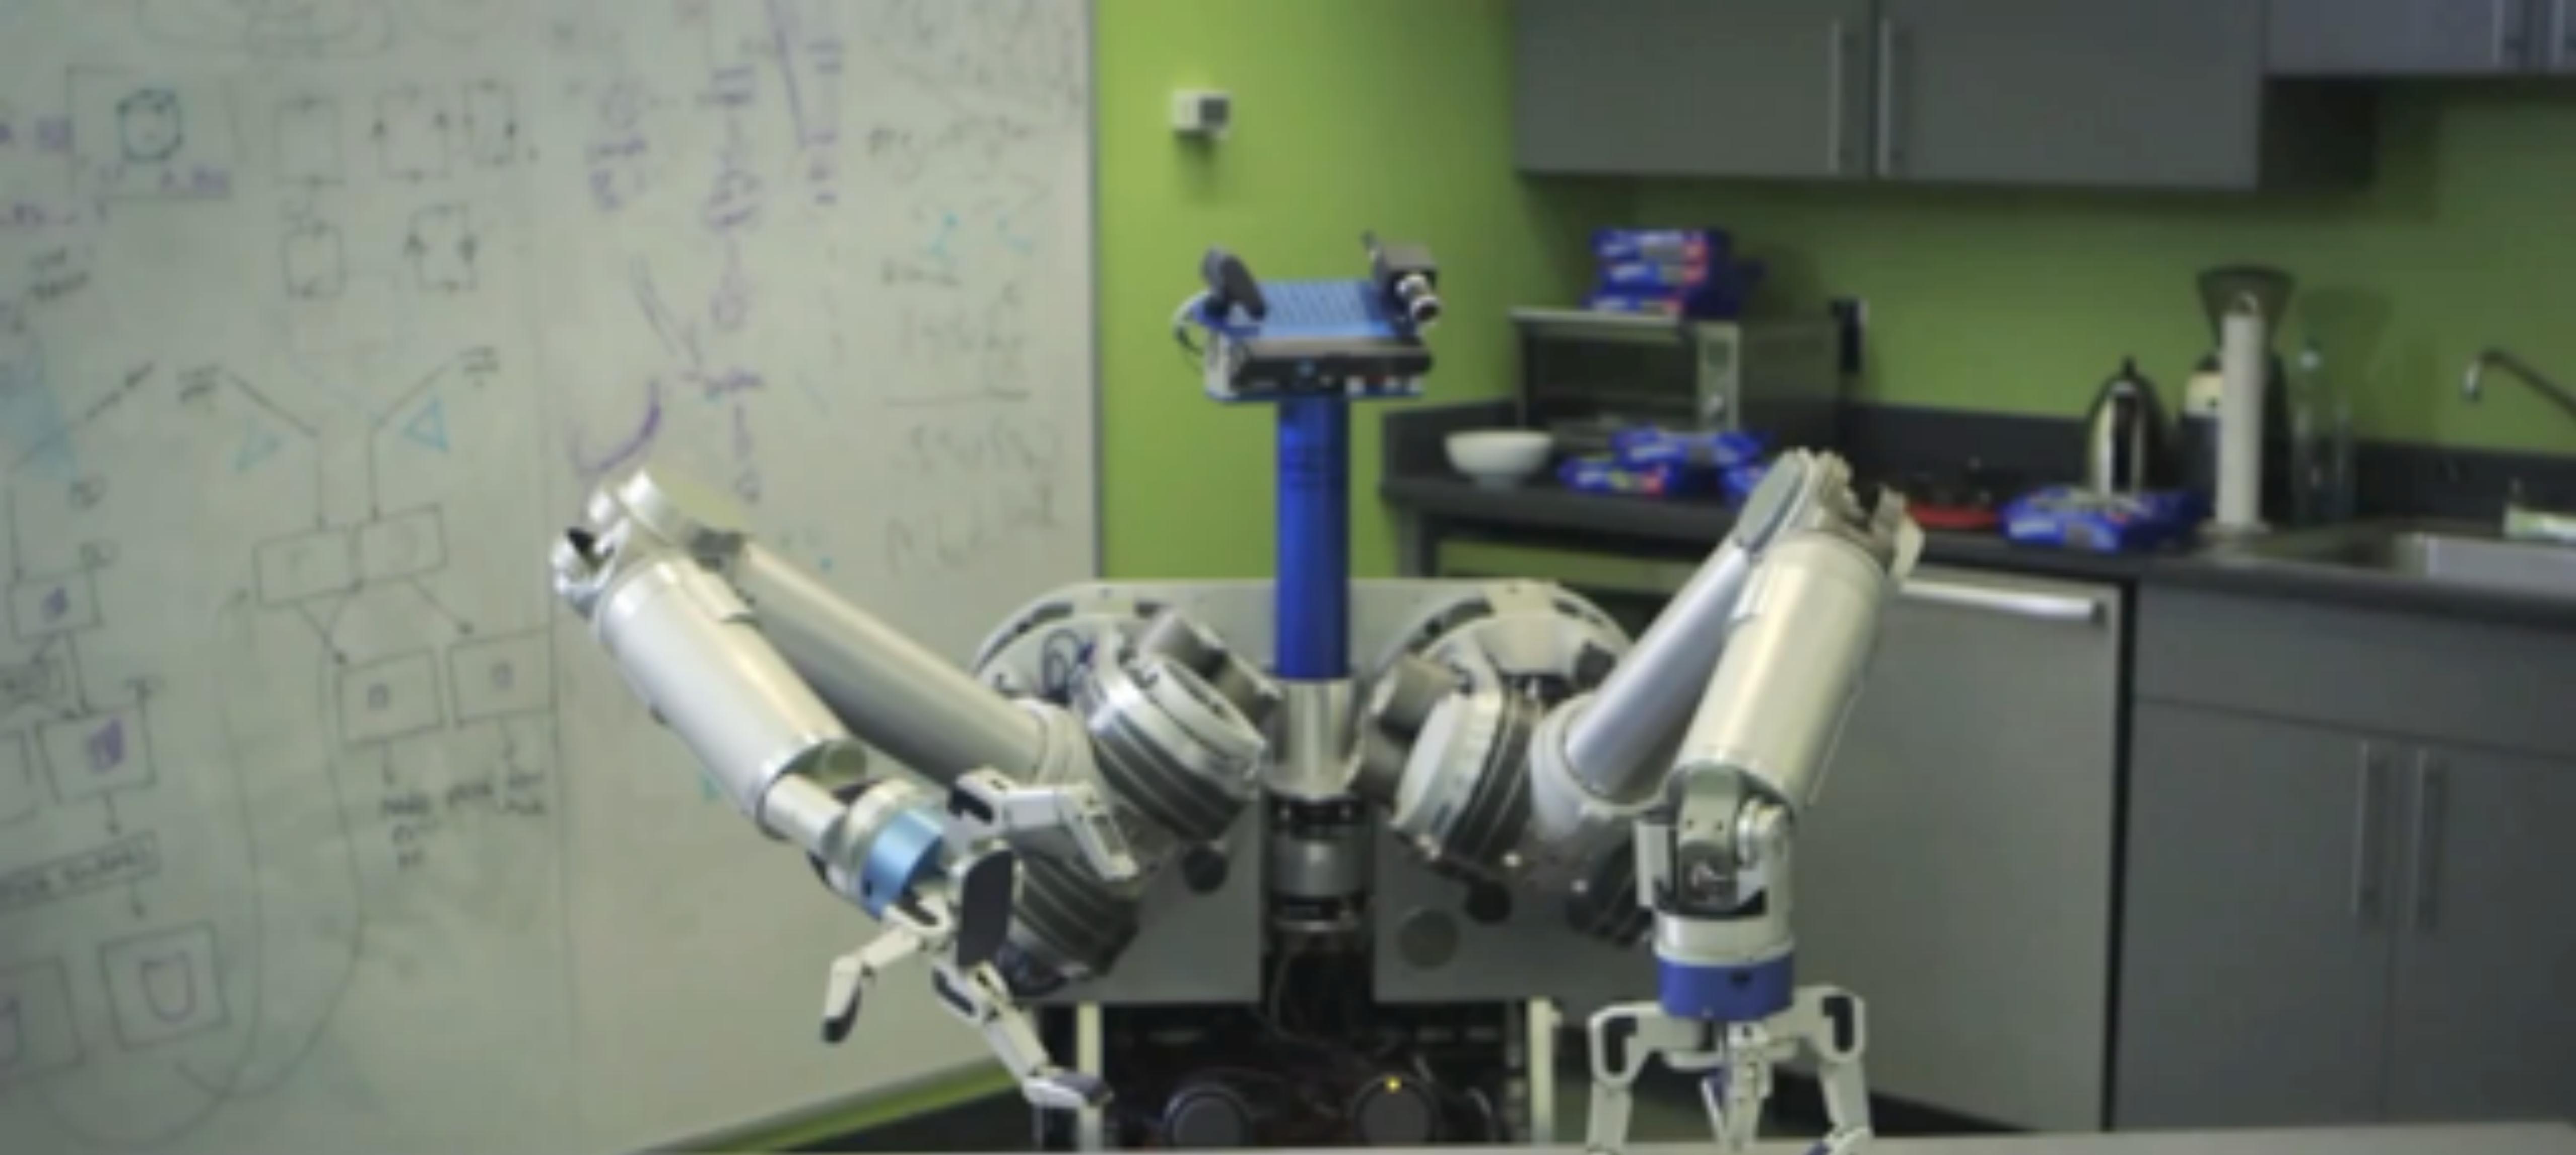
Siddhartha Srinivasa  
Retired Boeing Endowed Professor  
University of Washington

**PAUL G. ALLEN SCHOOL**  
OF COMPUTER SCIENCE & ENGINEERING





# Motion Planning









# Motion Planning is a technology

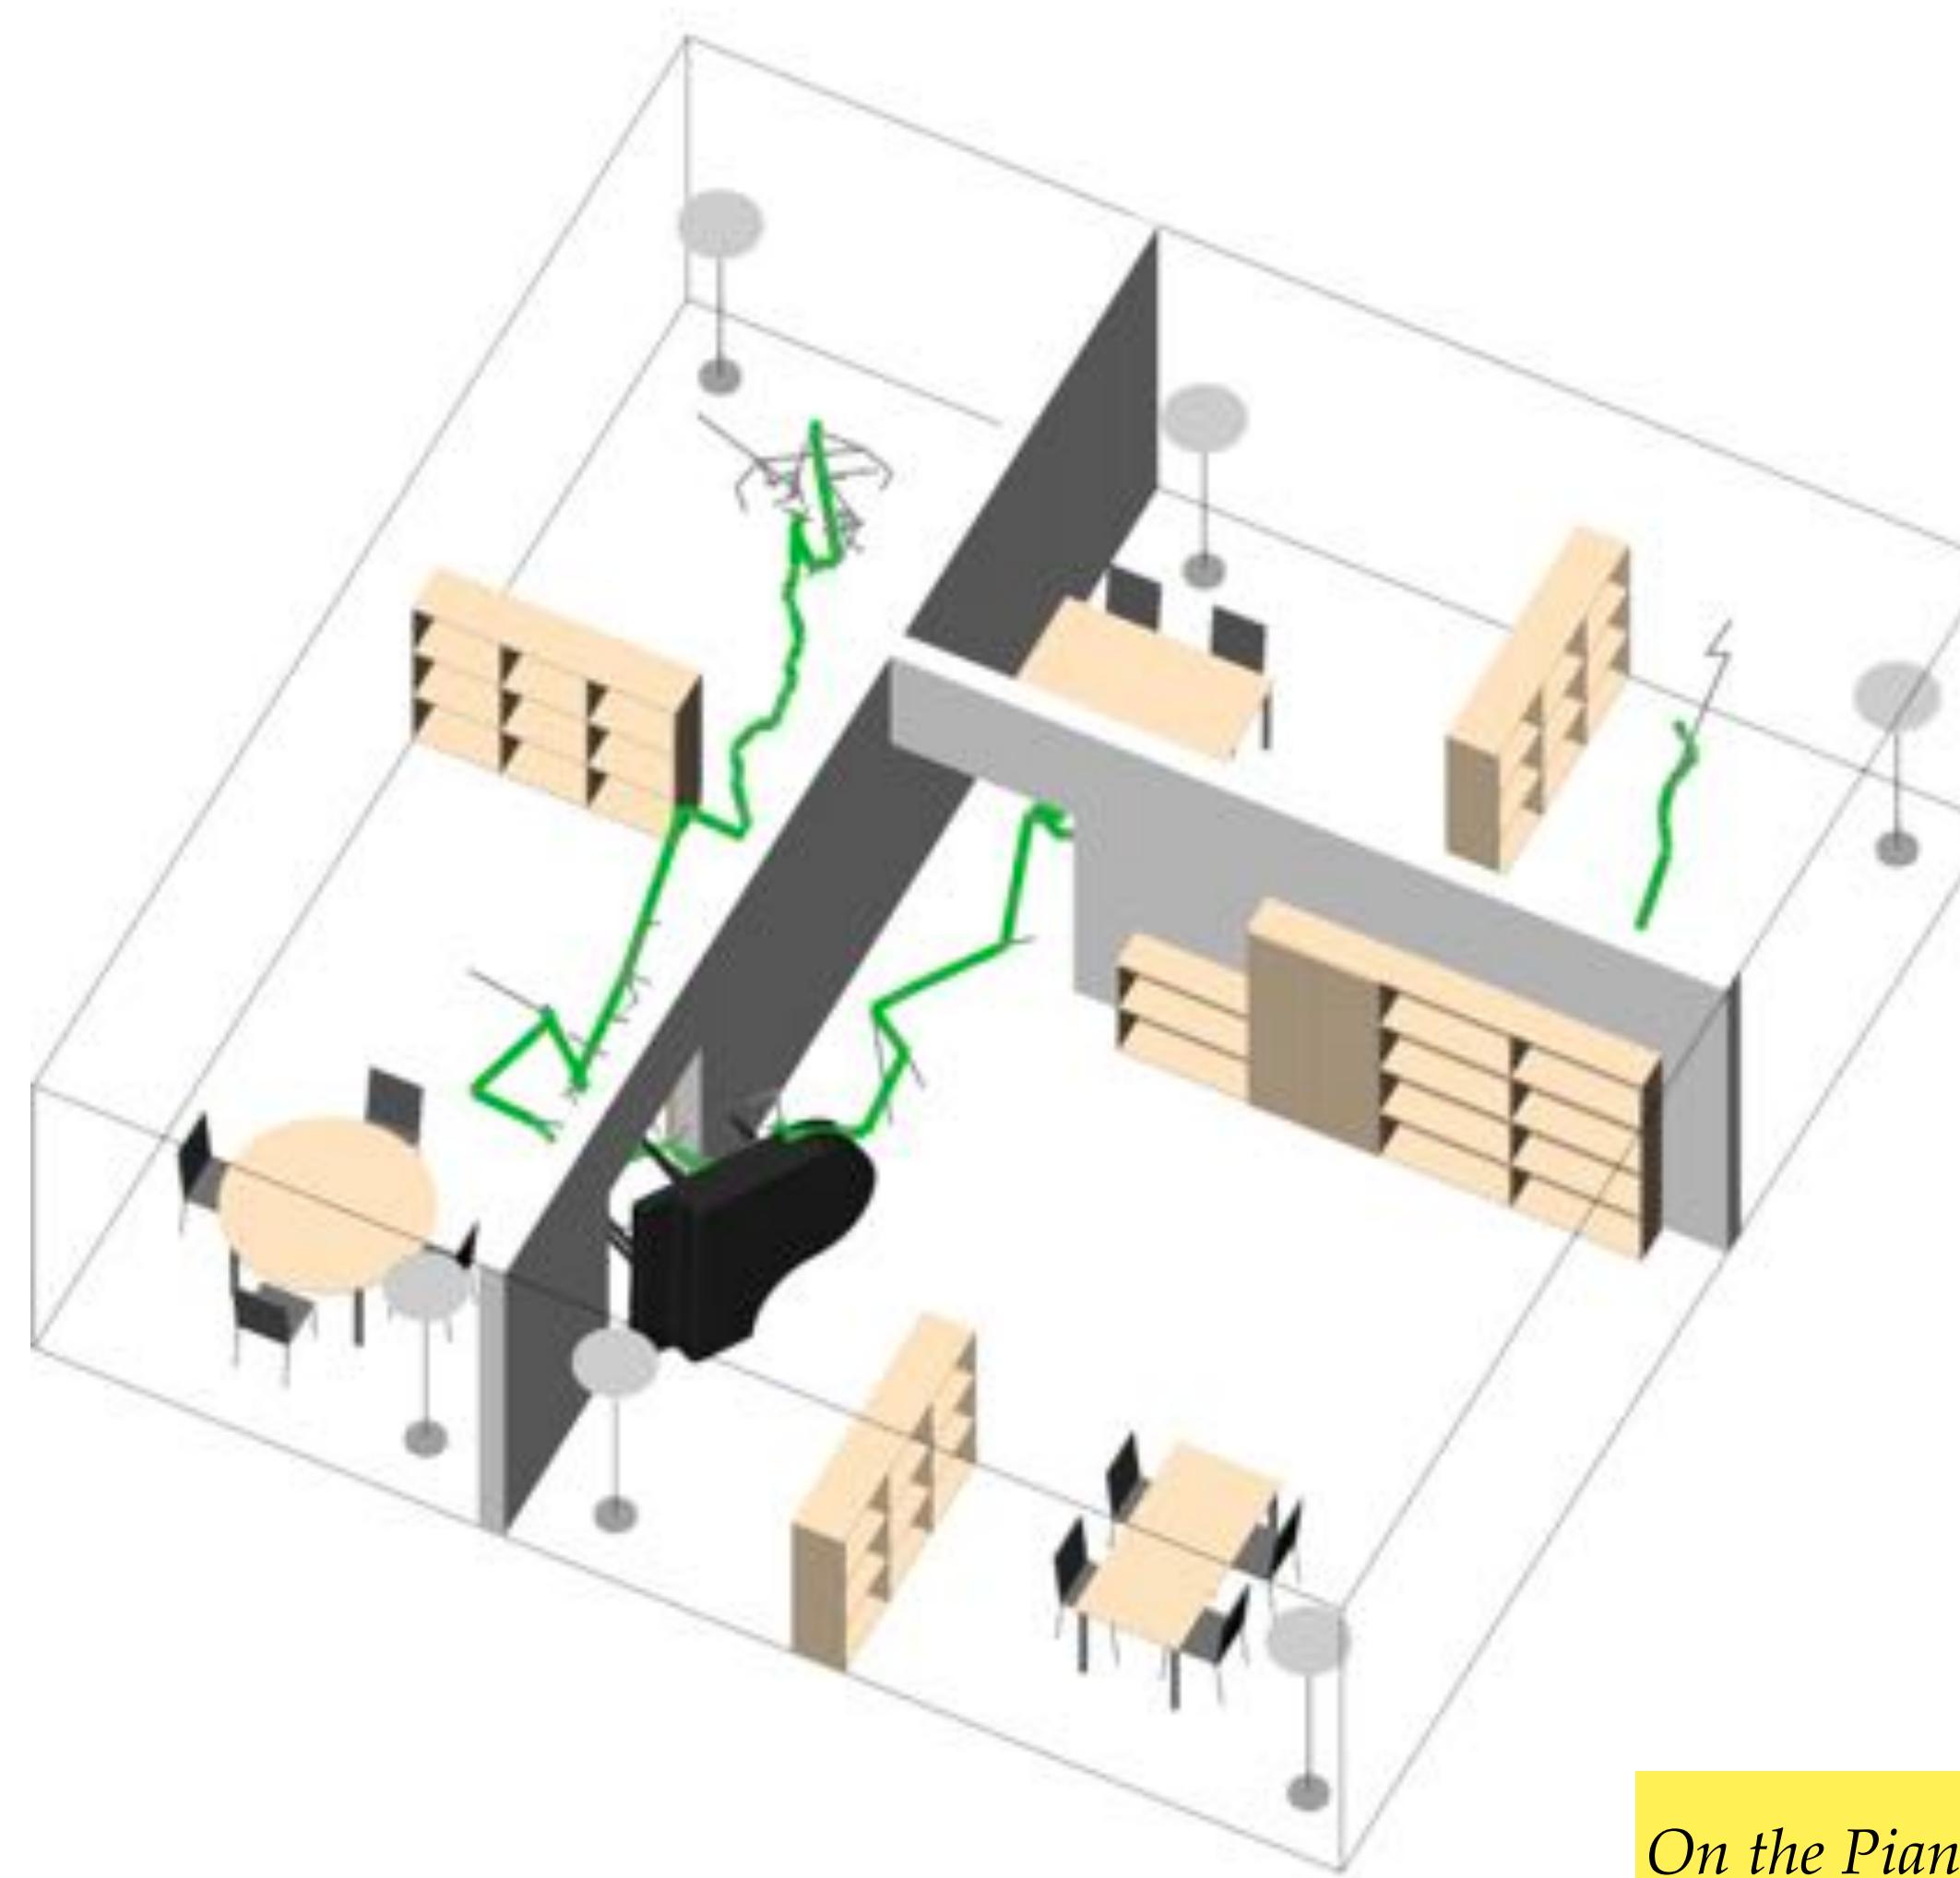


The screenshot shows the homepage of the Open Motion Planning Library (OMPL) at [ompl.kavrakilis.org](http://ompl.kavrakilis.org). The page title is "The Open Motion Planning Library". Below the title, there is a sequence of five 3D wireframe models of a parallel manipulator, each with blue dots representing joints or vertices, illustrating the progression of a motion planning algorithm. A small arrow icon points from left to right above the first model. Below this visual, a section titled "Constrained Motion Planning" contains the text: "OMPL supports planning for kinematically constrained robots. This parallel manipulator (included as a [demo program](#)) has over 150 degrees of freedom, but feasible motions can still be computed in seconds." At the bottom of the page, there is a brief description of OMPL's architecture: "OMPL, the Open Motion Planning Library, consists of many state-of-the-art sampling-based motion planning algorithms. OMPL itself does not contain any code related to, e.g., collision checking or visualization. This is a deliberate design choice, so that OMPL is not tied to a particular collision checker or visualization front end. The library is designed so it can be easily integrated into [systems](#) that provide the additional needed components." To the right, there are two blue buttons: "Download version 1.5.0" (released Jun 12, 2020) and "Cite for citation" (with the note "If you use OMPL in your work").

10-100X

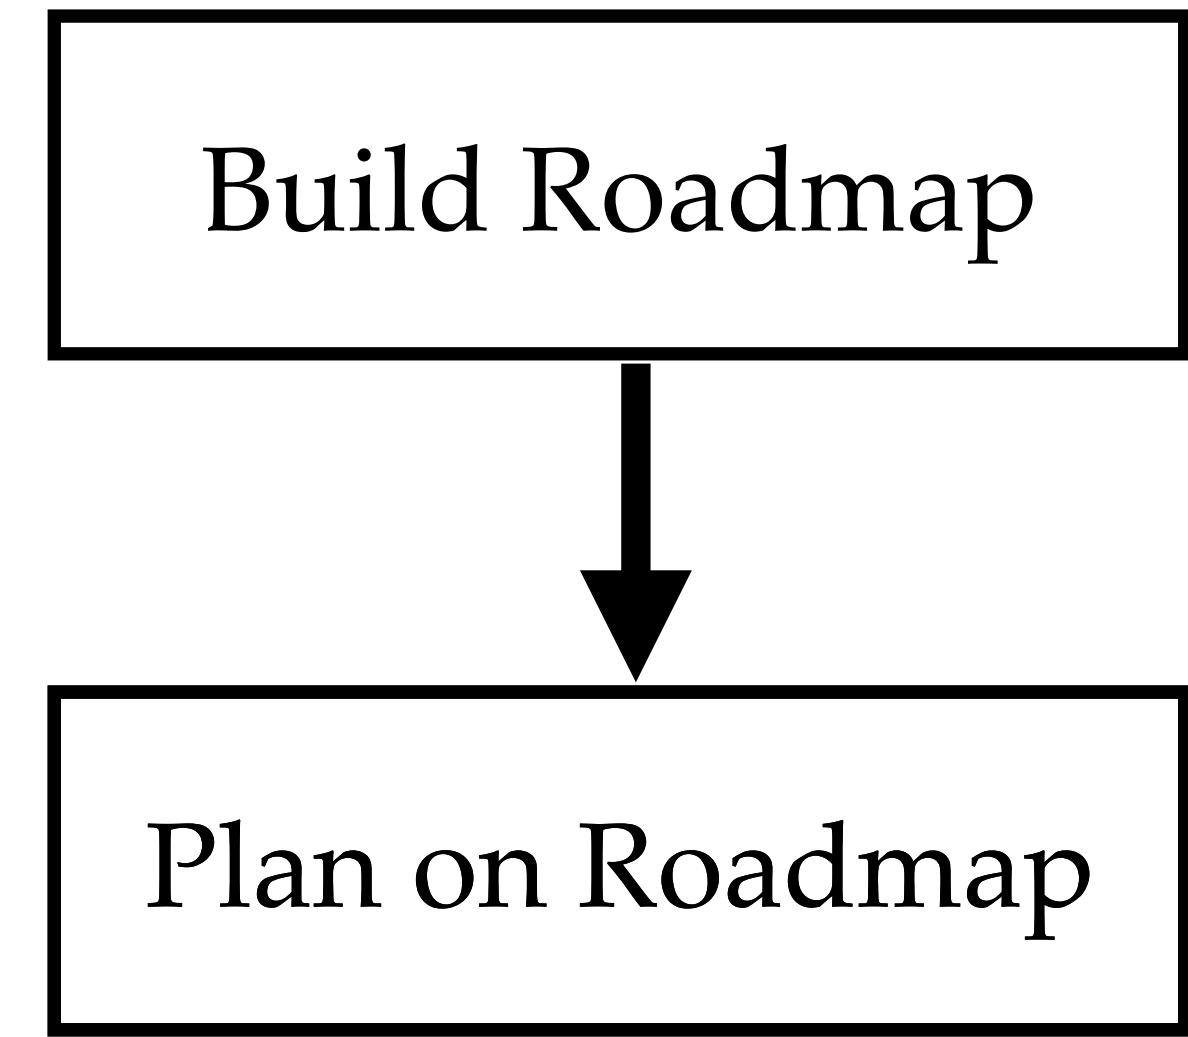
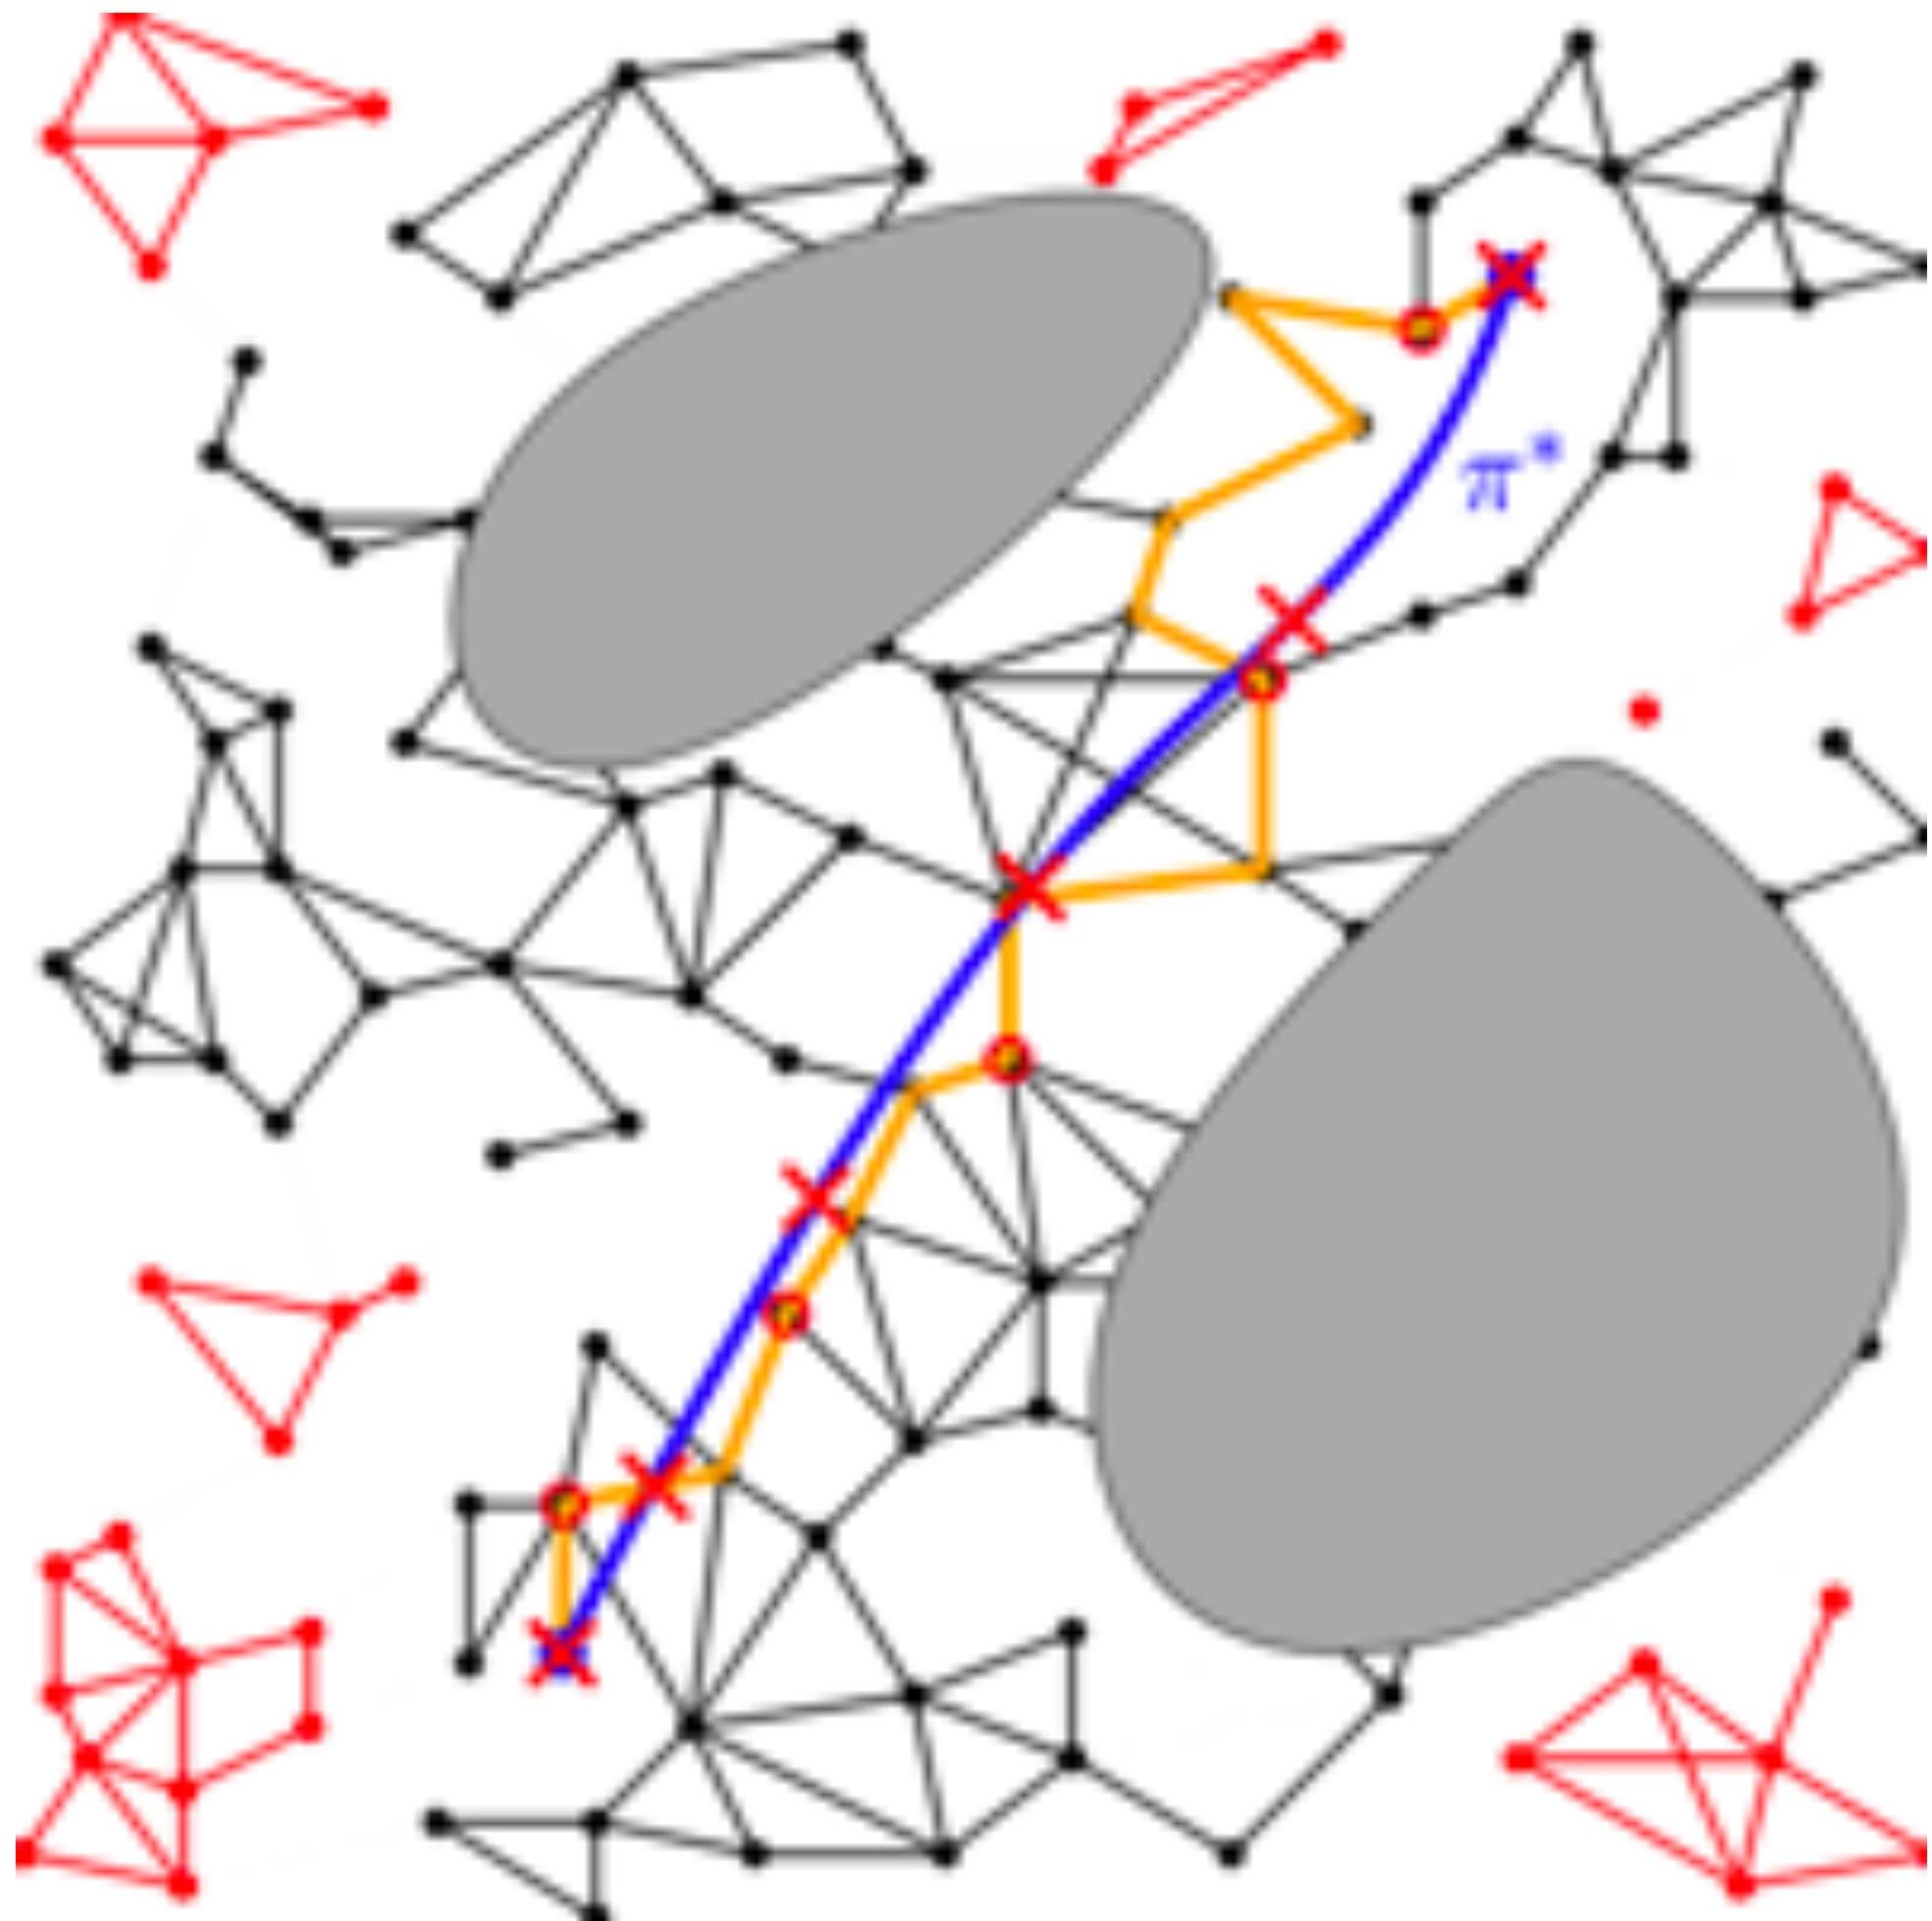
Improvement

# The Piano Movers' Problem



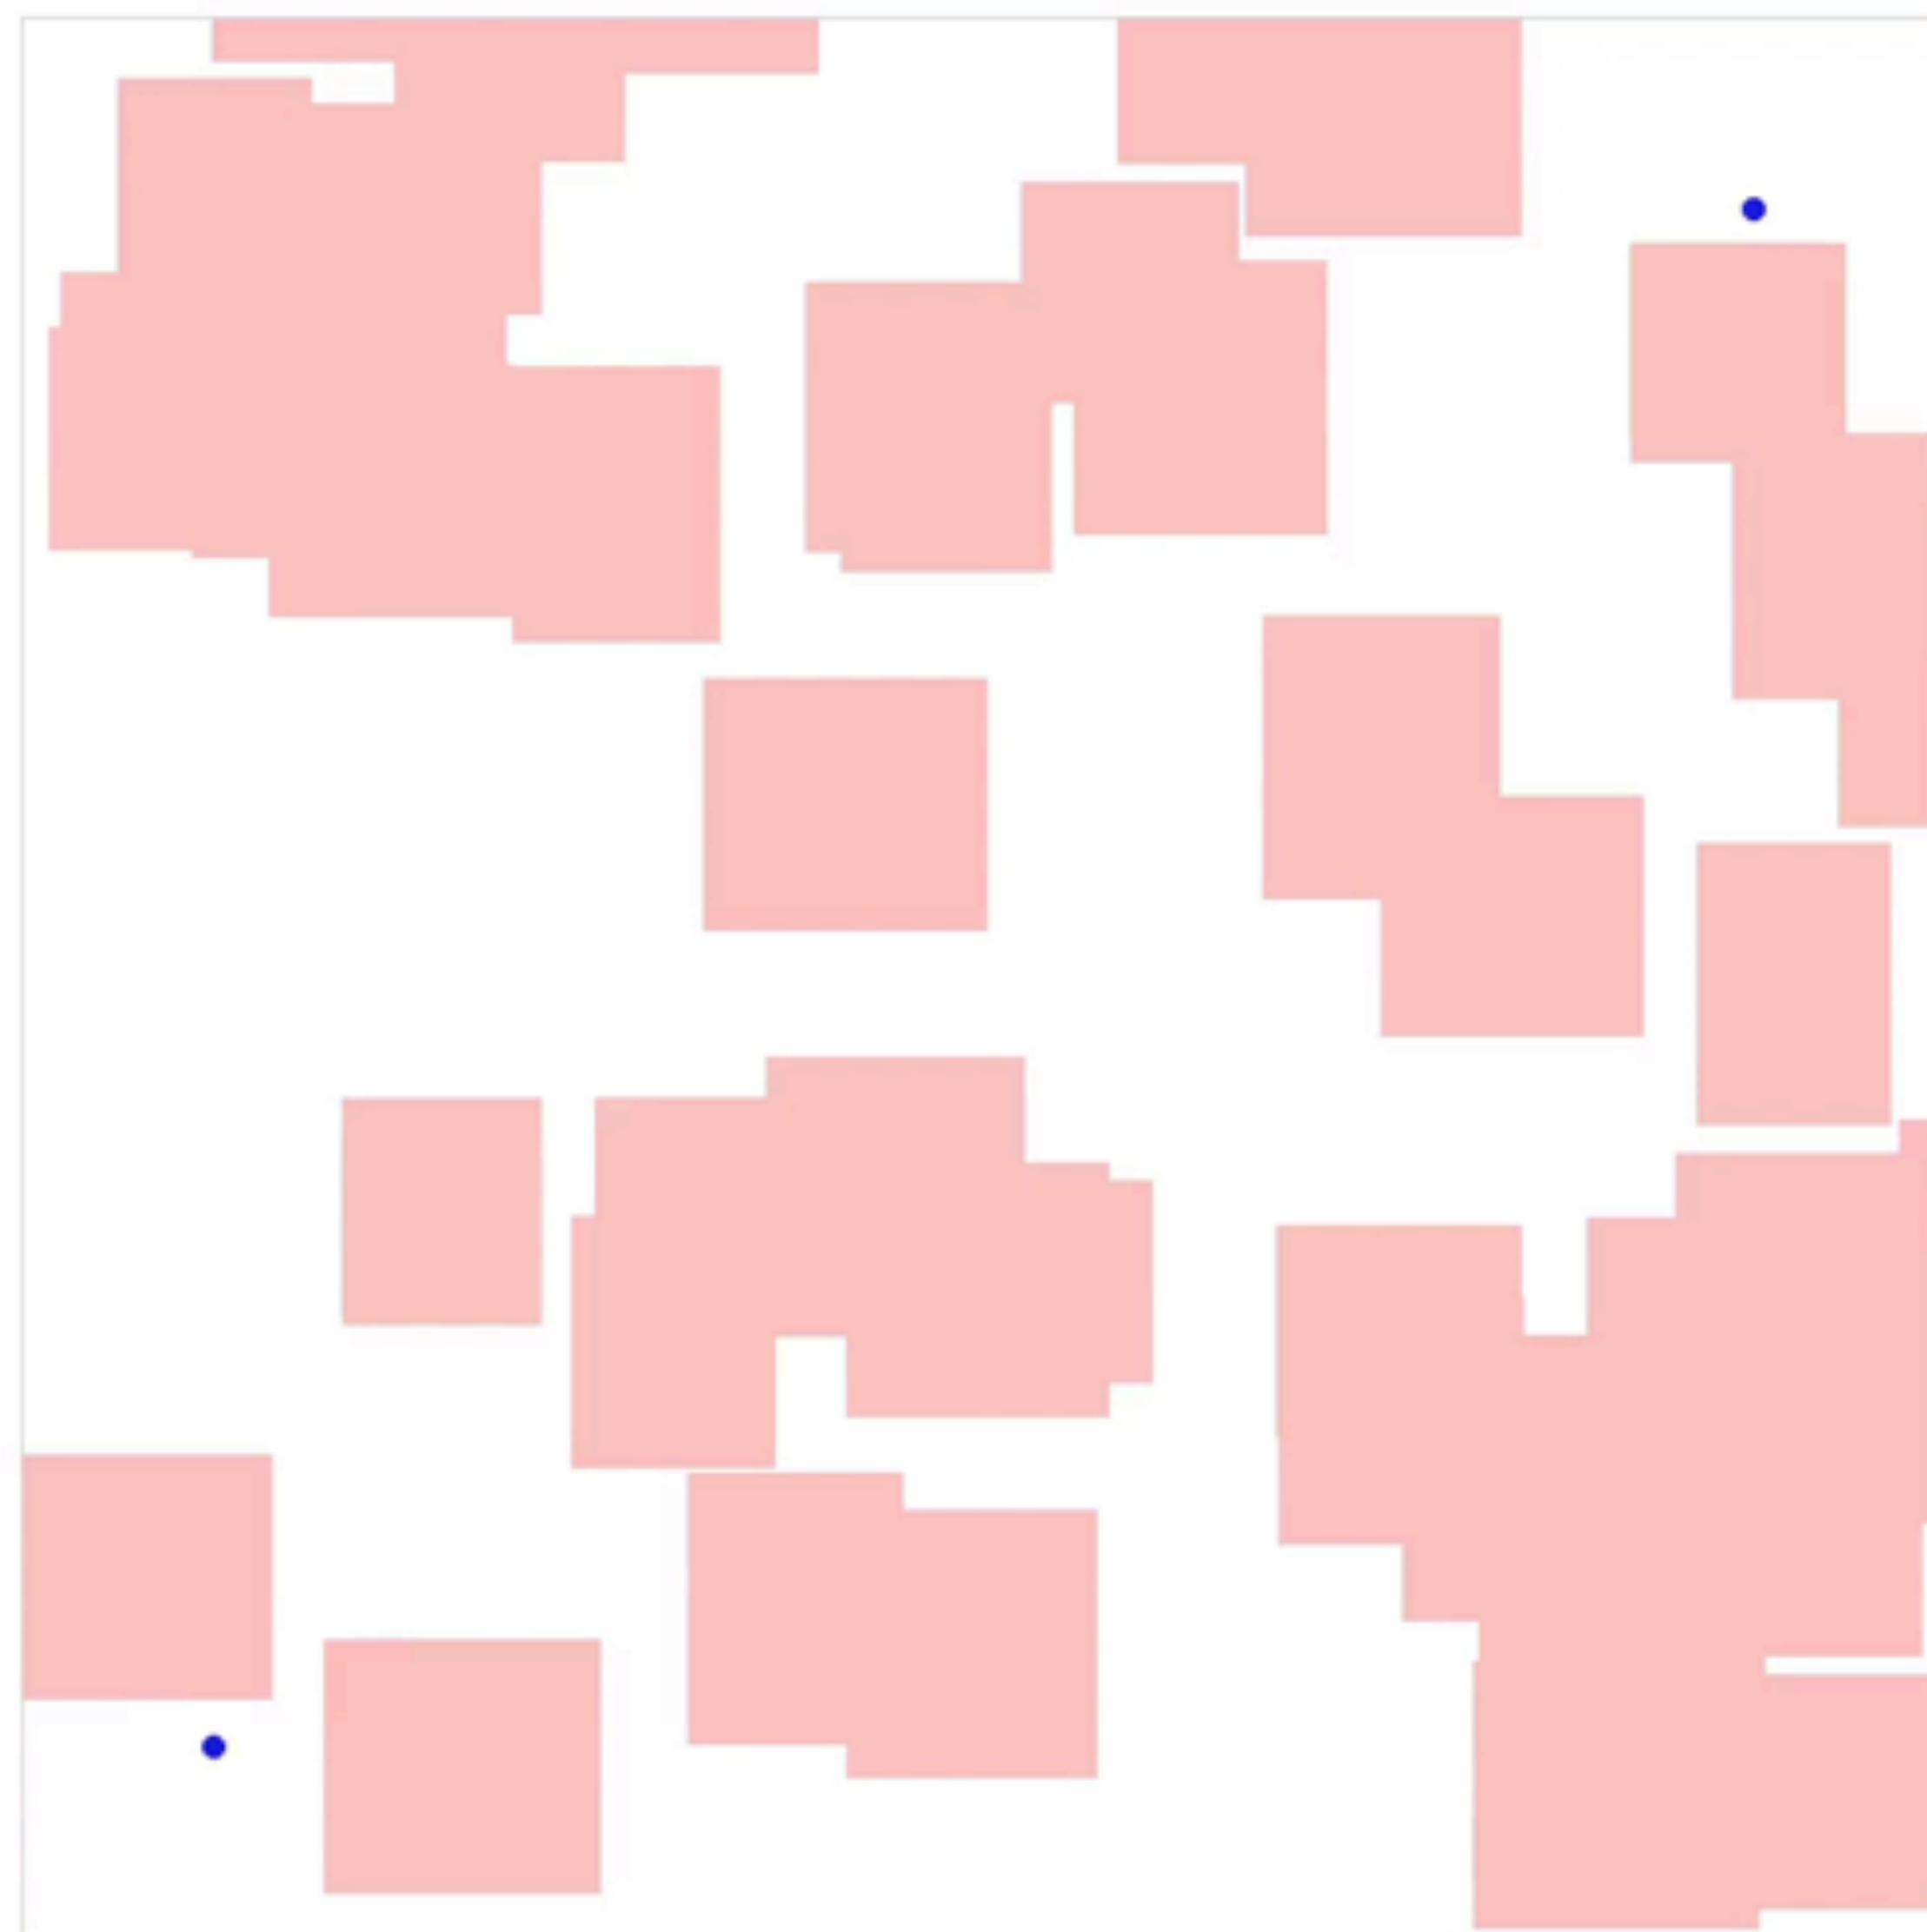
*On the Piano Movers problem. I-III,* Schwartz and Sharir,  
Comm. on Pure and Applied Math., 1983

# Roadmaps



*Probabilistic roadmaps for path planning in high-dimensional configuration spaces, Kavraki et al., IEEE TRO, 1996.*

# A\* Search



# A\* Search



OPTIMAL!!

Is it optimal over something we care about?

# A\* Search: A Personal Journey

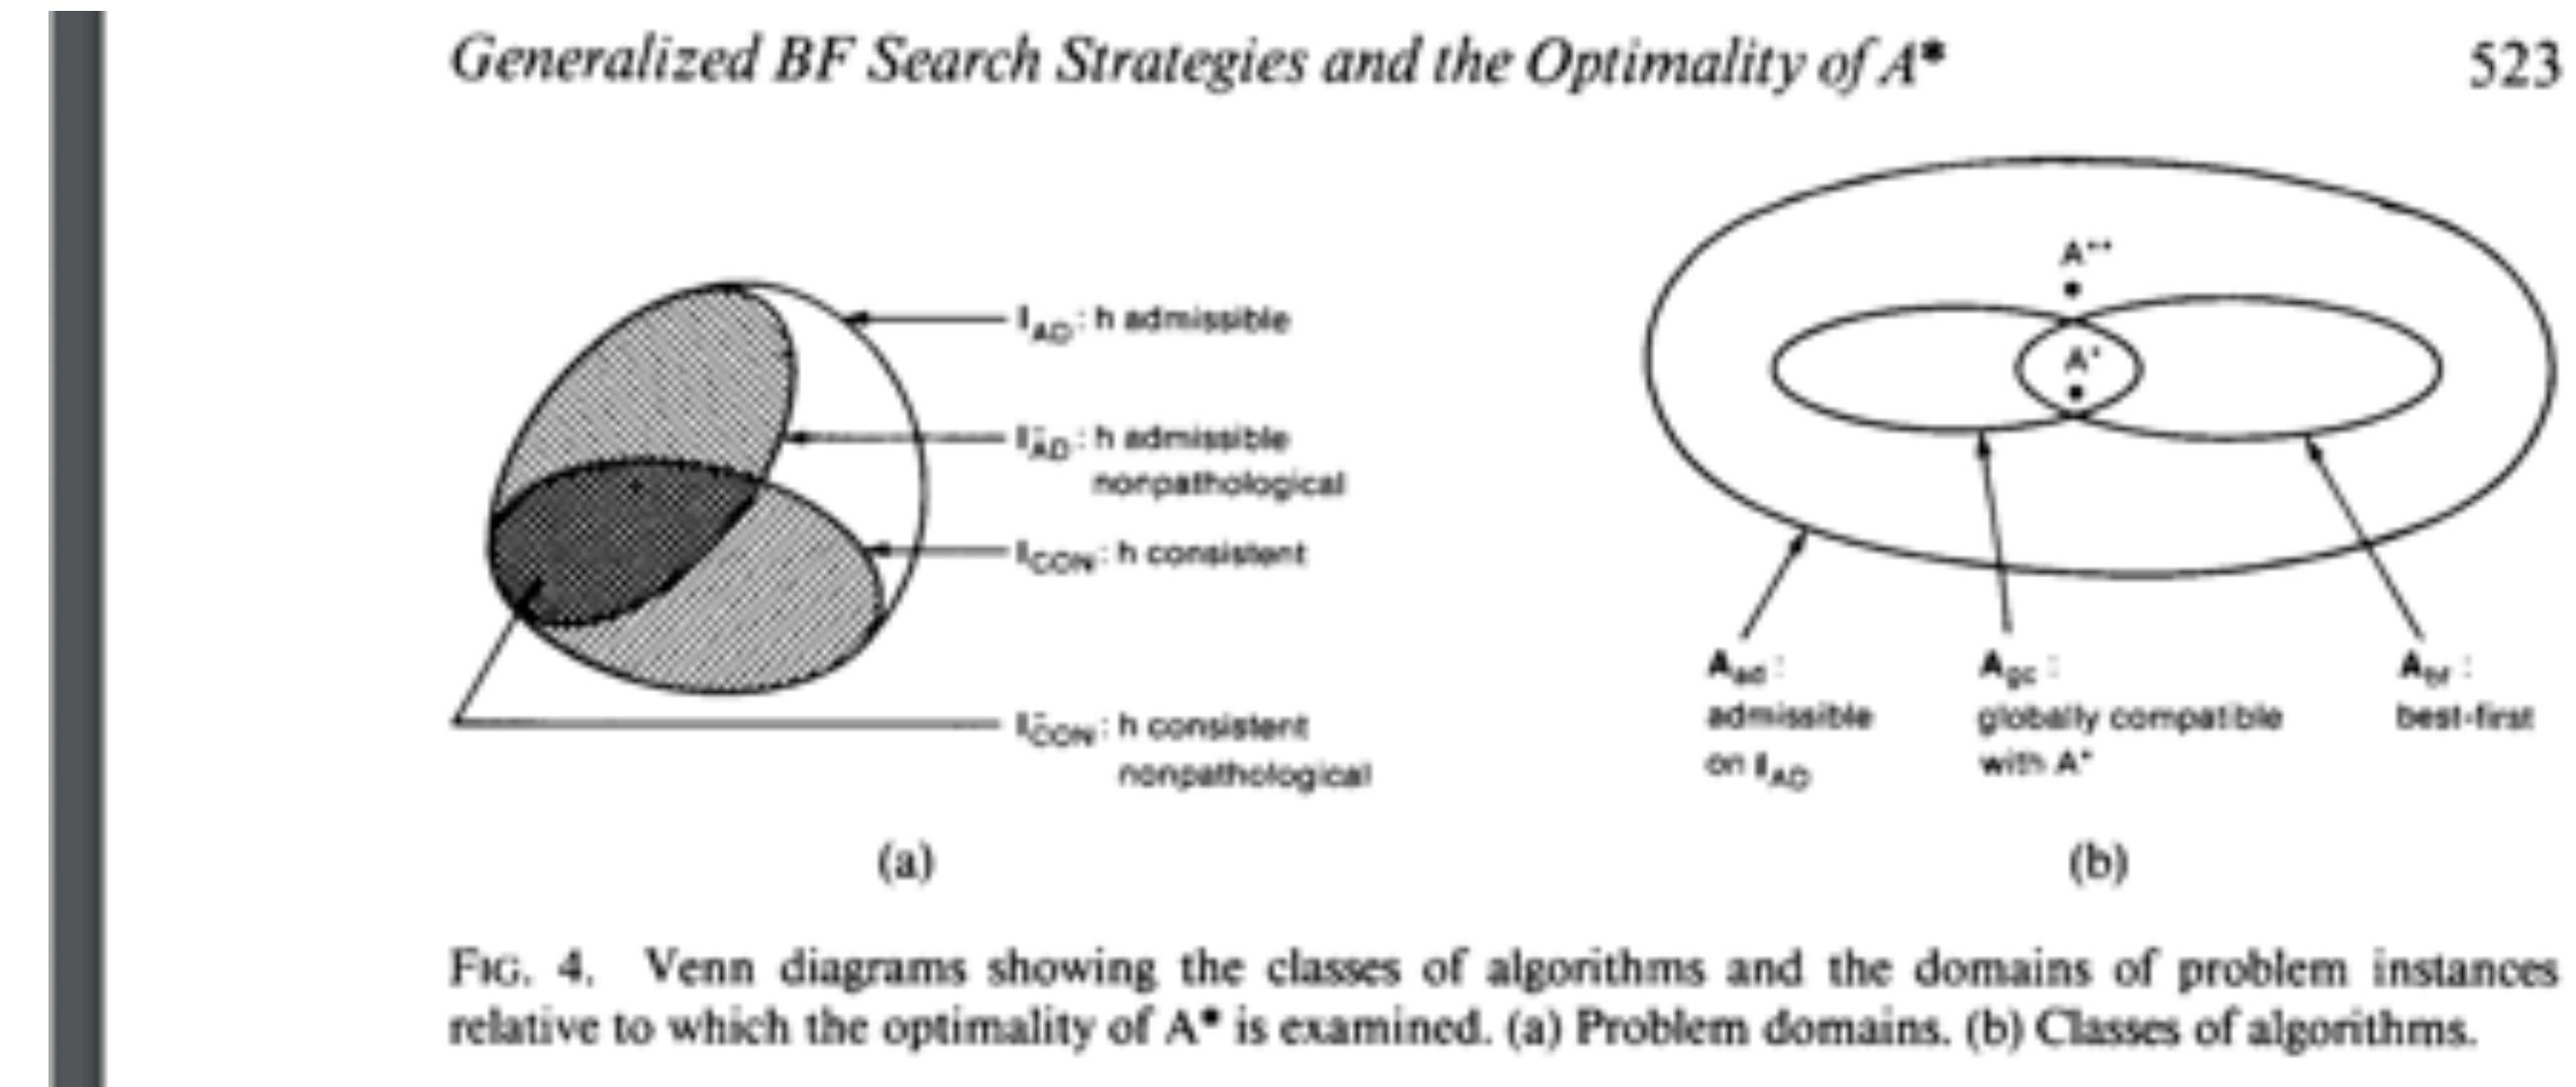
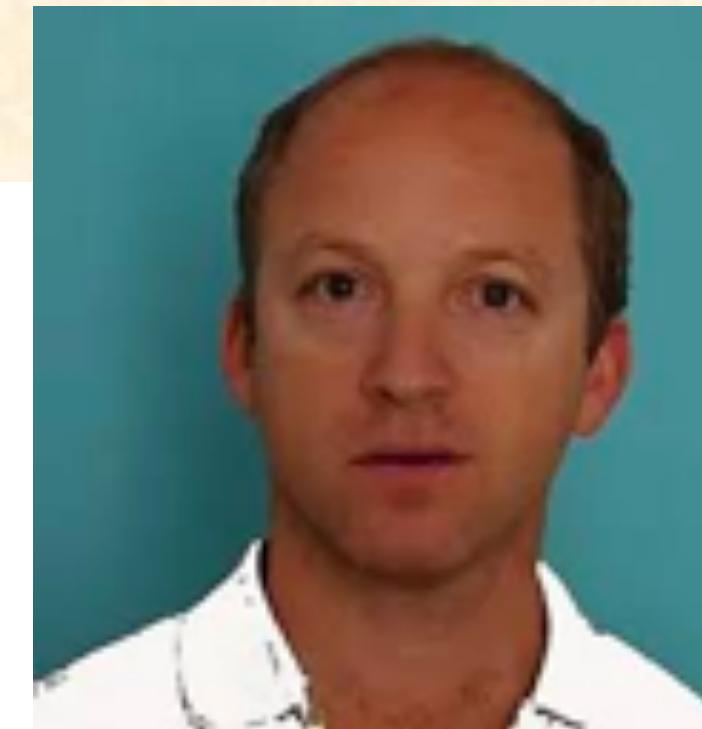
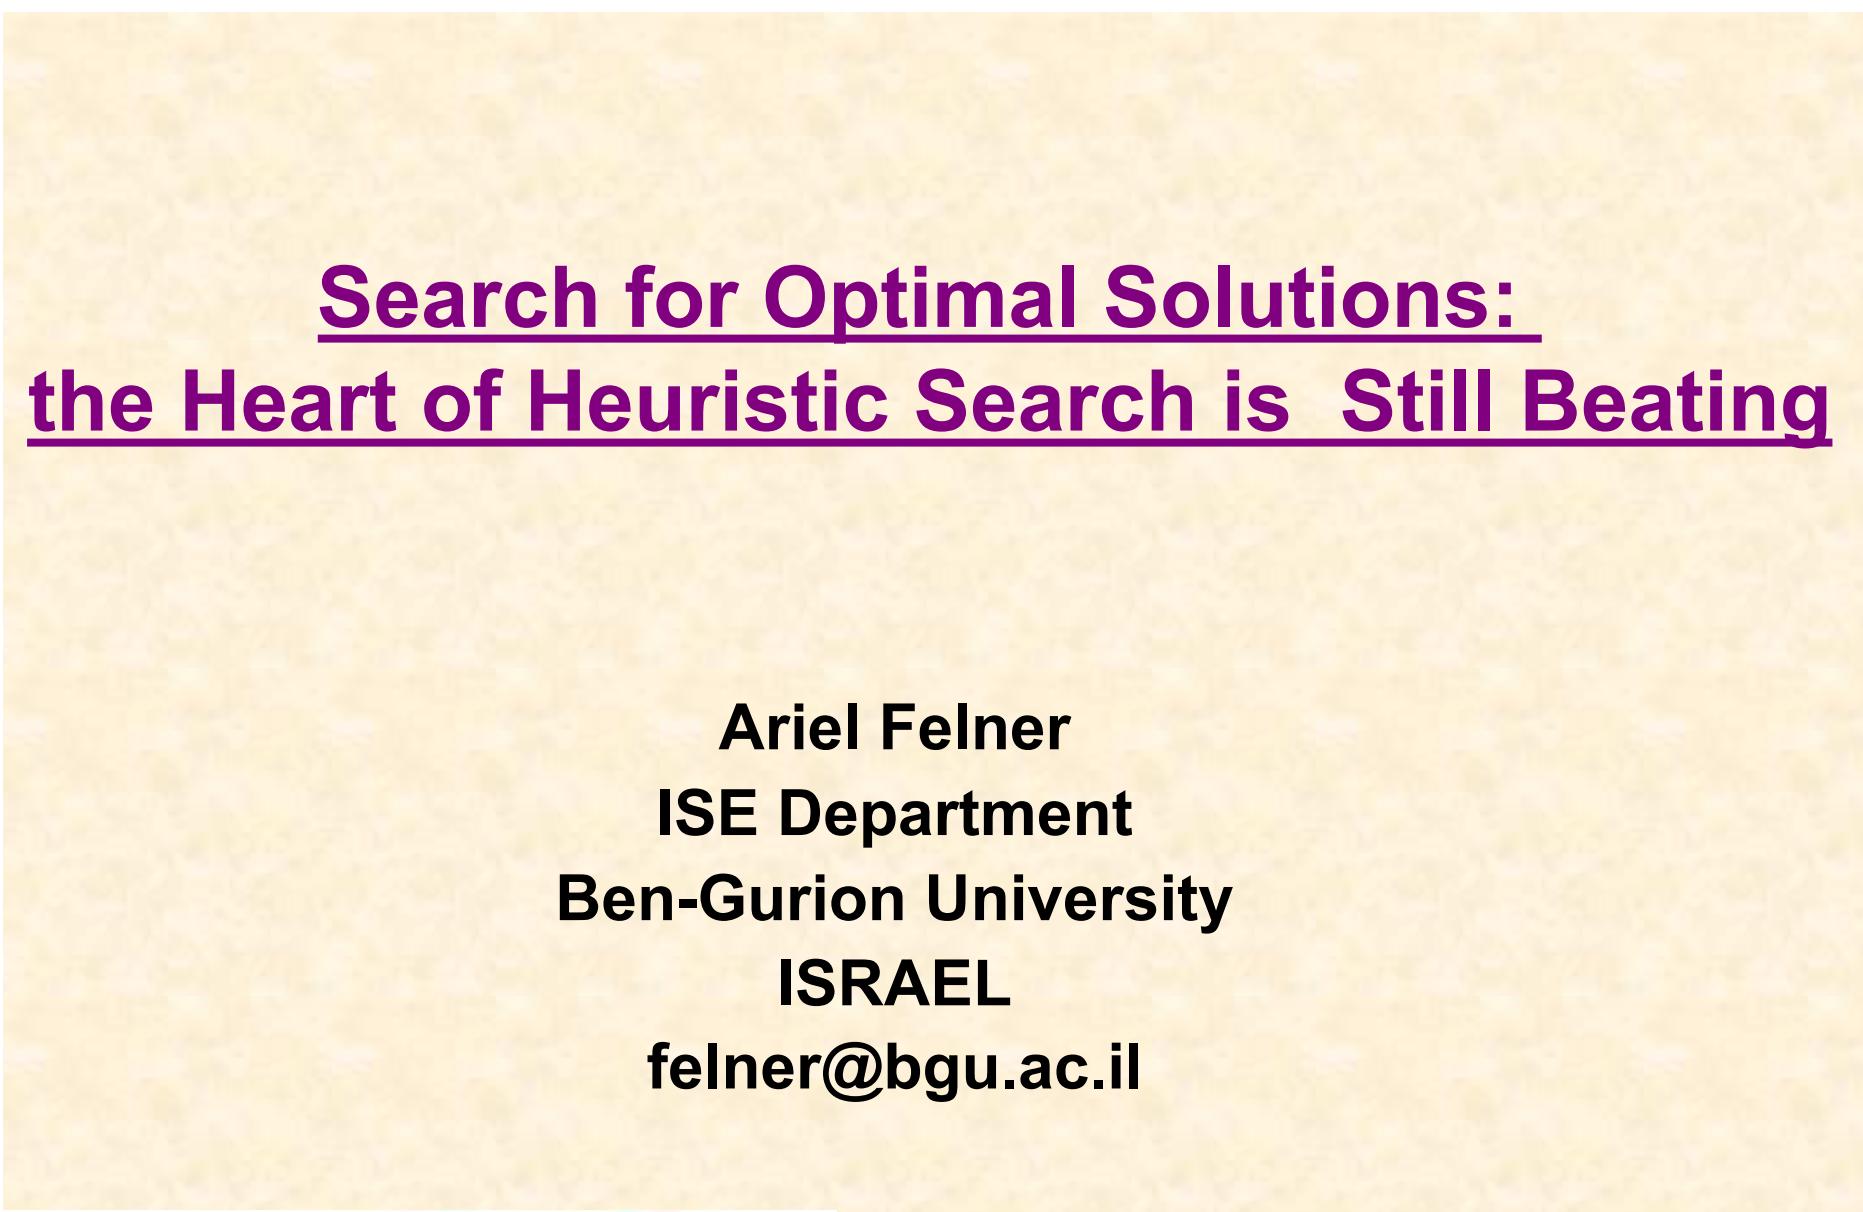
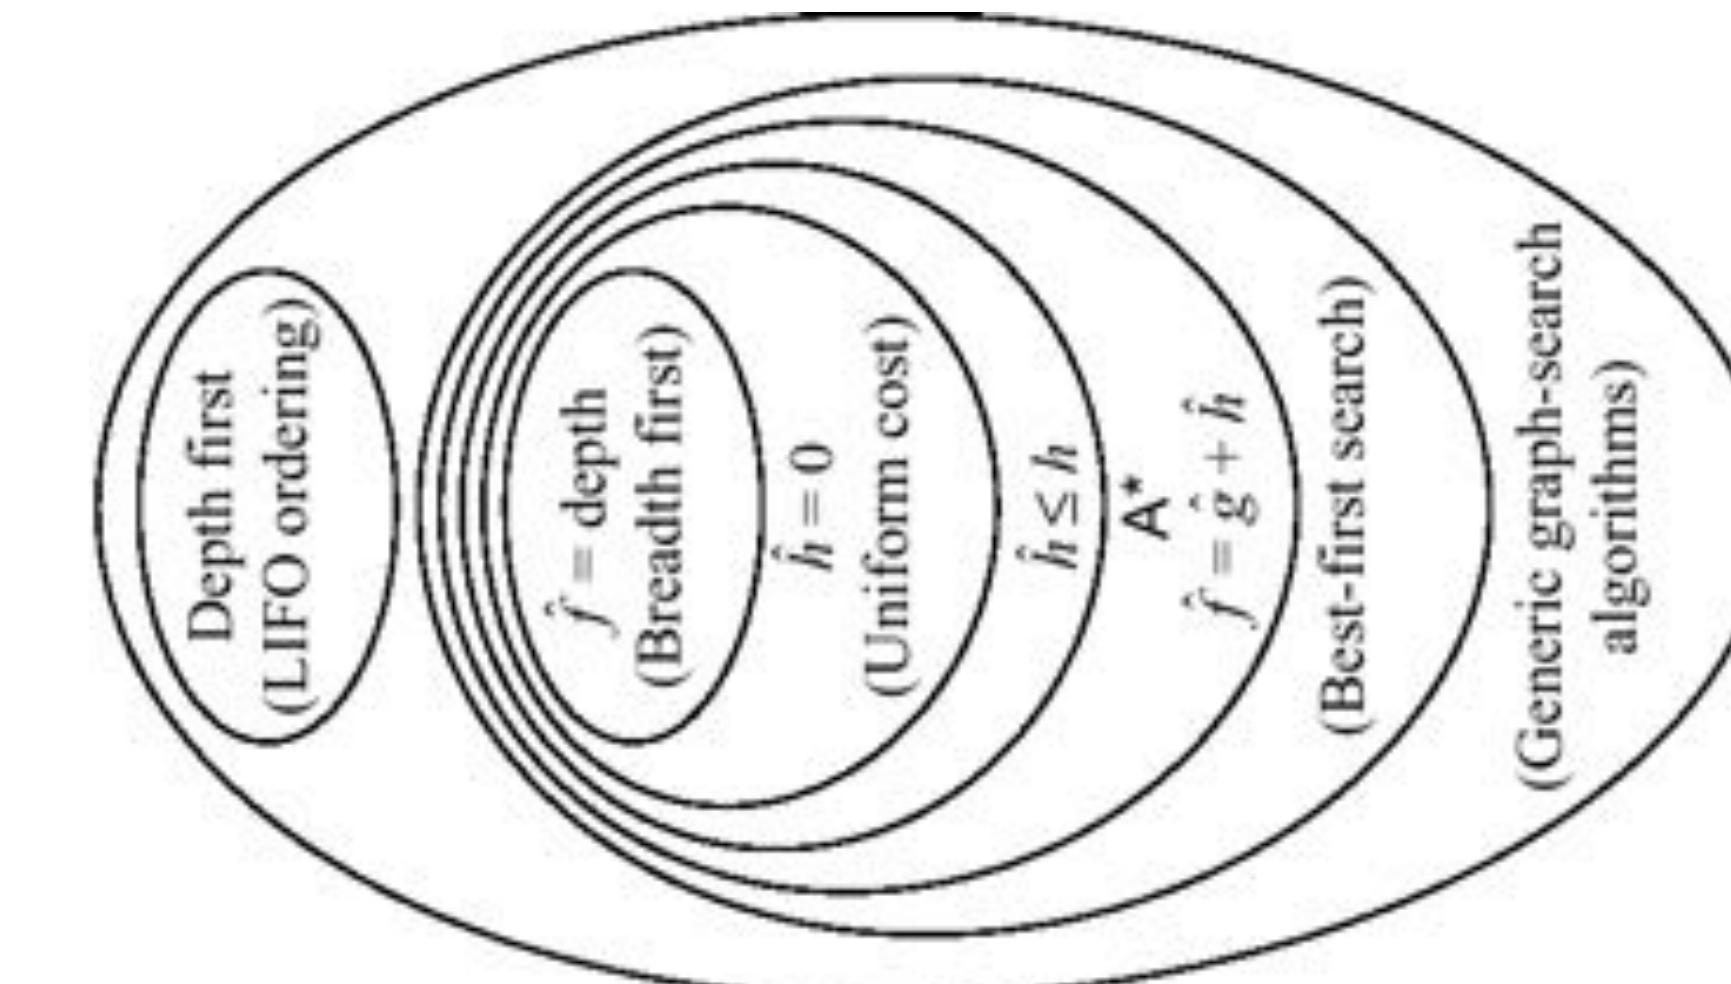
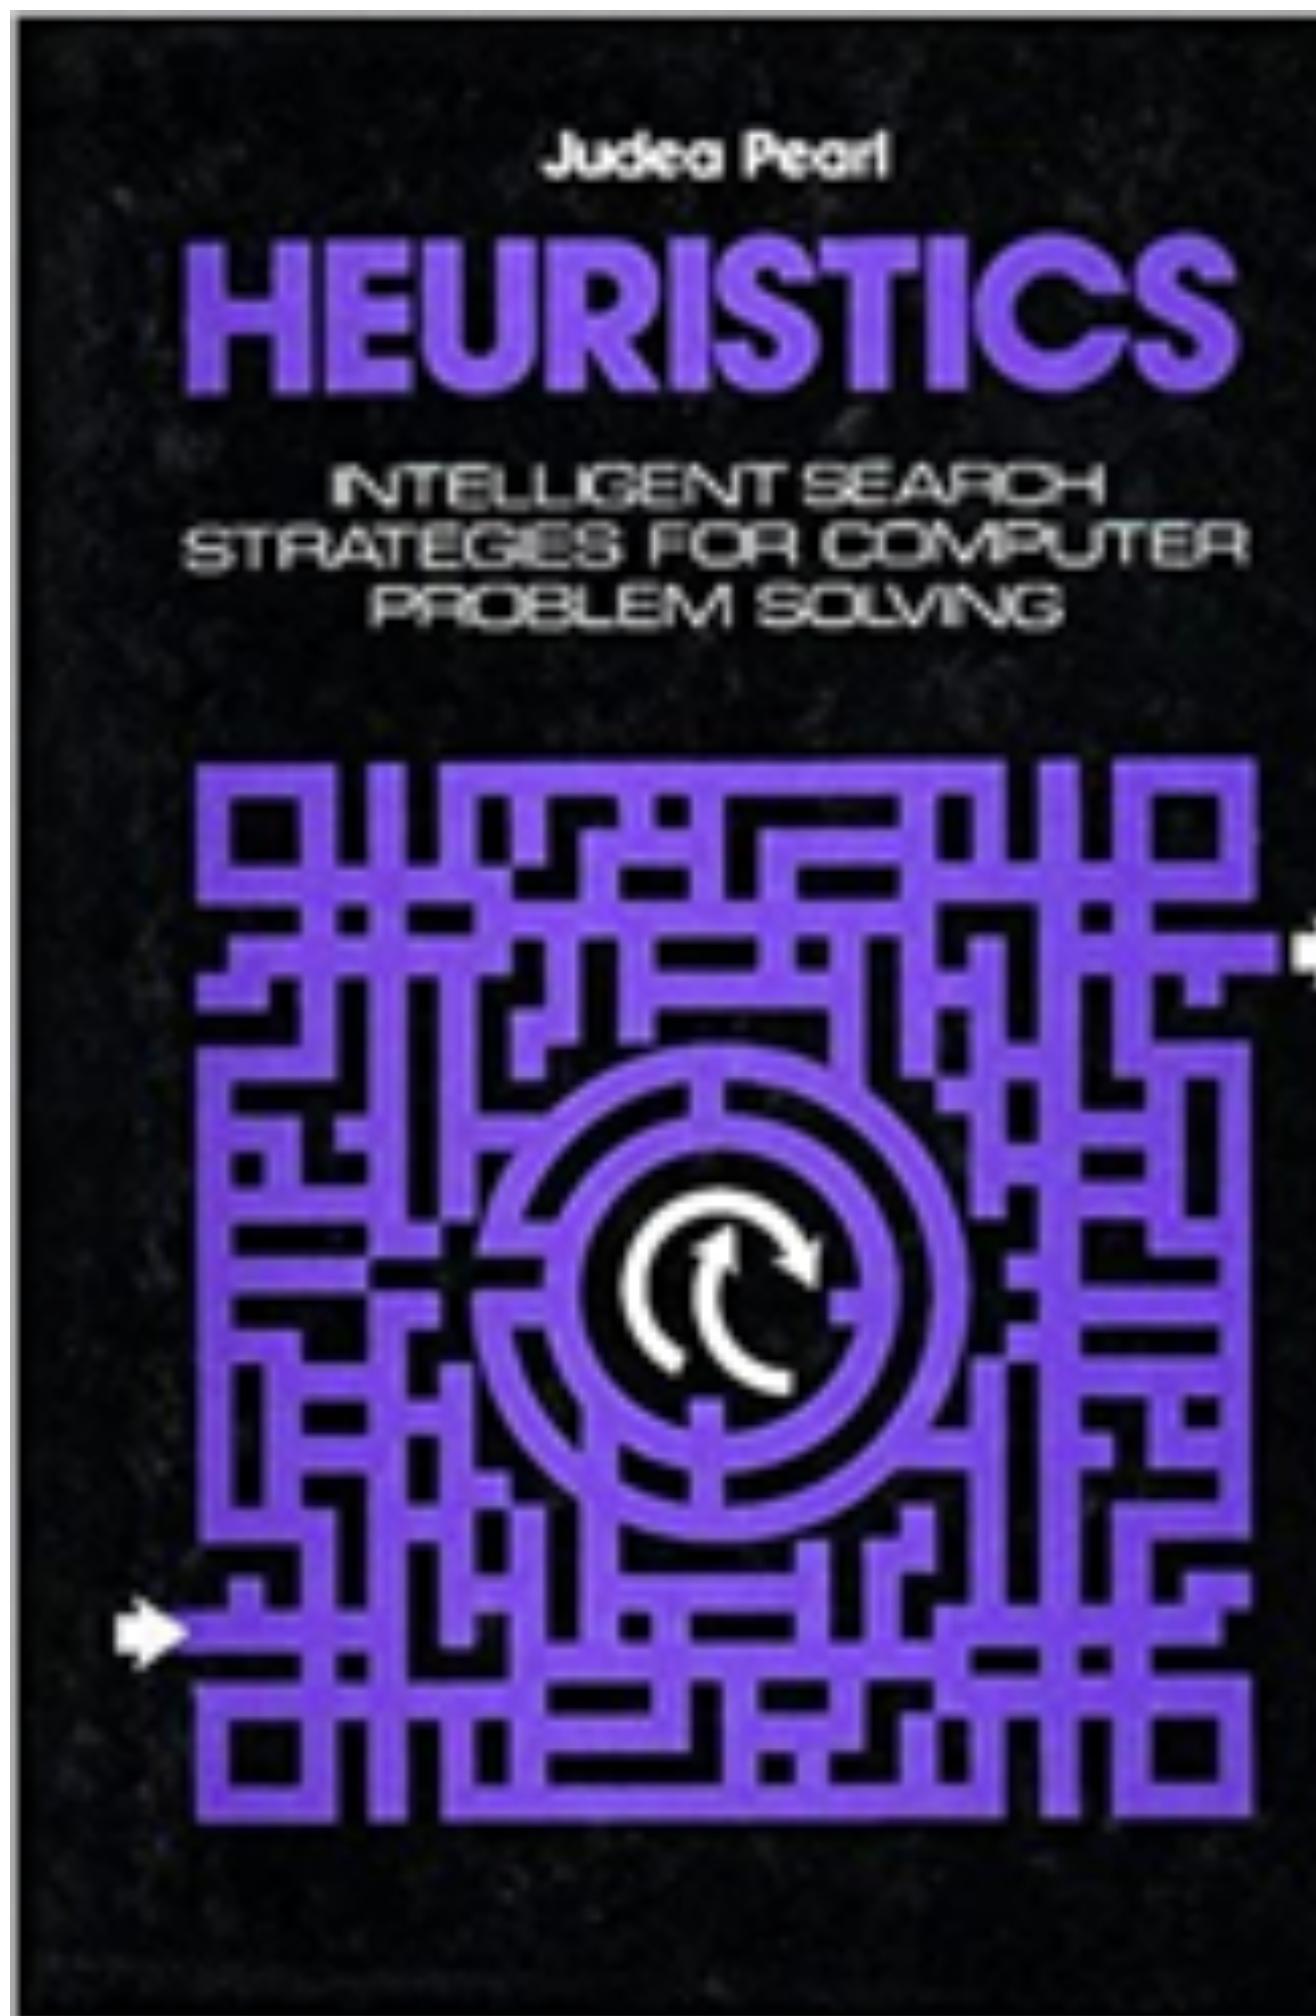


FIG. 4. Venn diagrams showing the classes of algorithms and the domains of problem instances relative to which the optimality of A\* is examined. (a) Problem domains. (b) Classes of algorithms.





# A\* Search: A Personal Journey



## CHAPTER 2

### GENERALIZATION OF THE SHORTEST PATH ALGORITHM

#### 2.1 Intuitive Description

Picture two amoeba, one dyed red and the other dyed blue. The red one is placed on the starting node  $s$ , and the blue one is placed on the terminating node  $t$ . Only the behavior of the red amoeba will be described in detail as the blue amoeba behaves analogously. The red amoeba moves at a velocity  $V_r$ . If the red amoeba reaches a node, it splits into the number of outgoing edges (edges where the initial node is the node where the amoeba is), with one progeny traveling each edge. The red amoeba and its progeny all travel at the same speed  $V_r$ . The blue amoeba and progeny have speed  $V_b$  and are performing in the same fashion with respect to ingoing edges. The first two amoeba of different lineage to meet have traveled the shortest path from  $s$  to  $t$ . Let  $dr(t)$  = the distance covered by red amoeba in time  $t$  and let  $db(t)$  = the distance covered by a blue amoeba in time  $t$ . At any time  $t$  these functions represent the distance covered by all amoeba of the corresponding color. If  $t^*$  is the time at which two amoeba of different color meet, then the distance traveled would be  $dr(t^*) + db(t^*)$ . Since  $dr$  and  $db$  are both monotonic functions with respect to time, any pair of amoeba meeting at  $t^* + \epsilon$ ,  $\epsilon > 0$  would have traveled more than the pair that met first. Therefore the above procedure is correct.

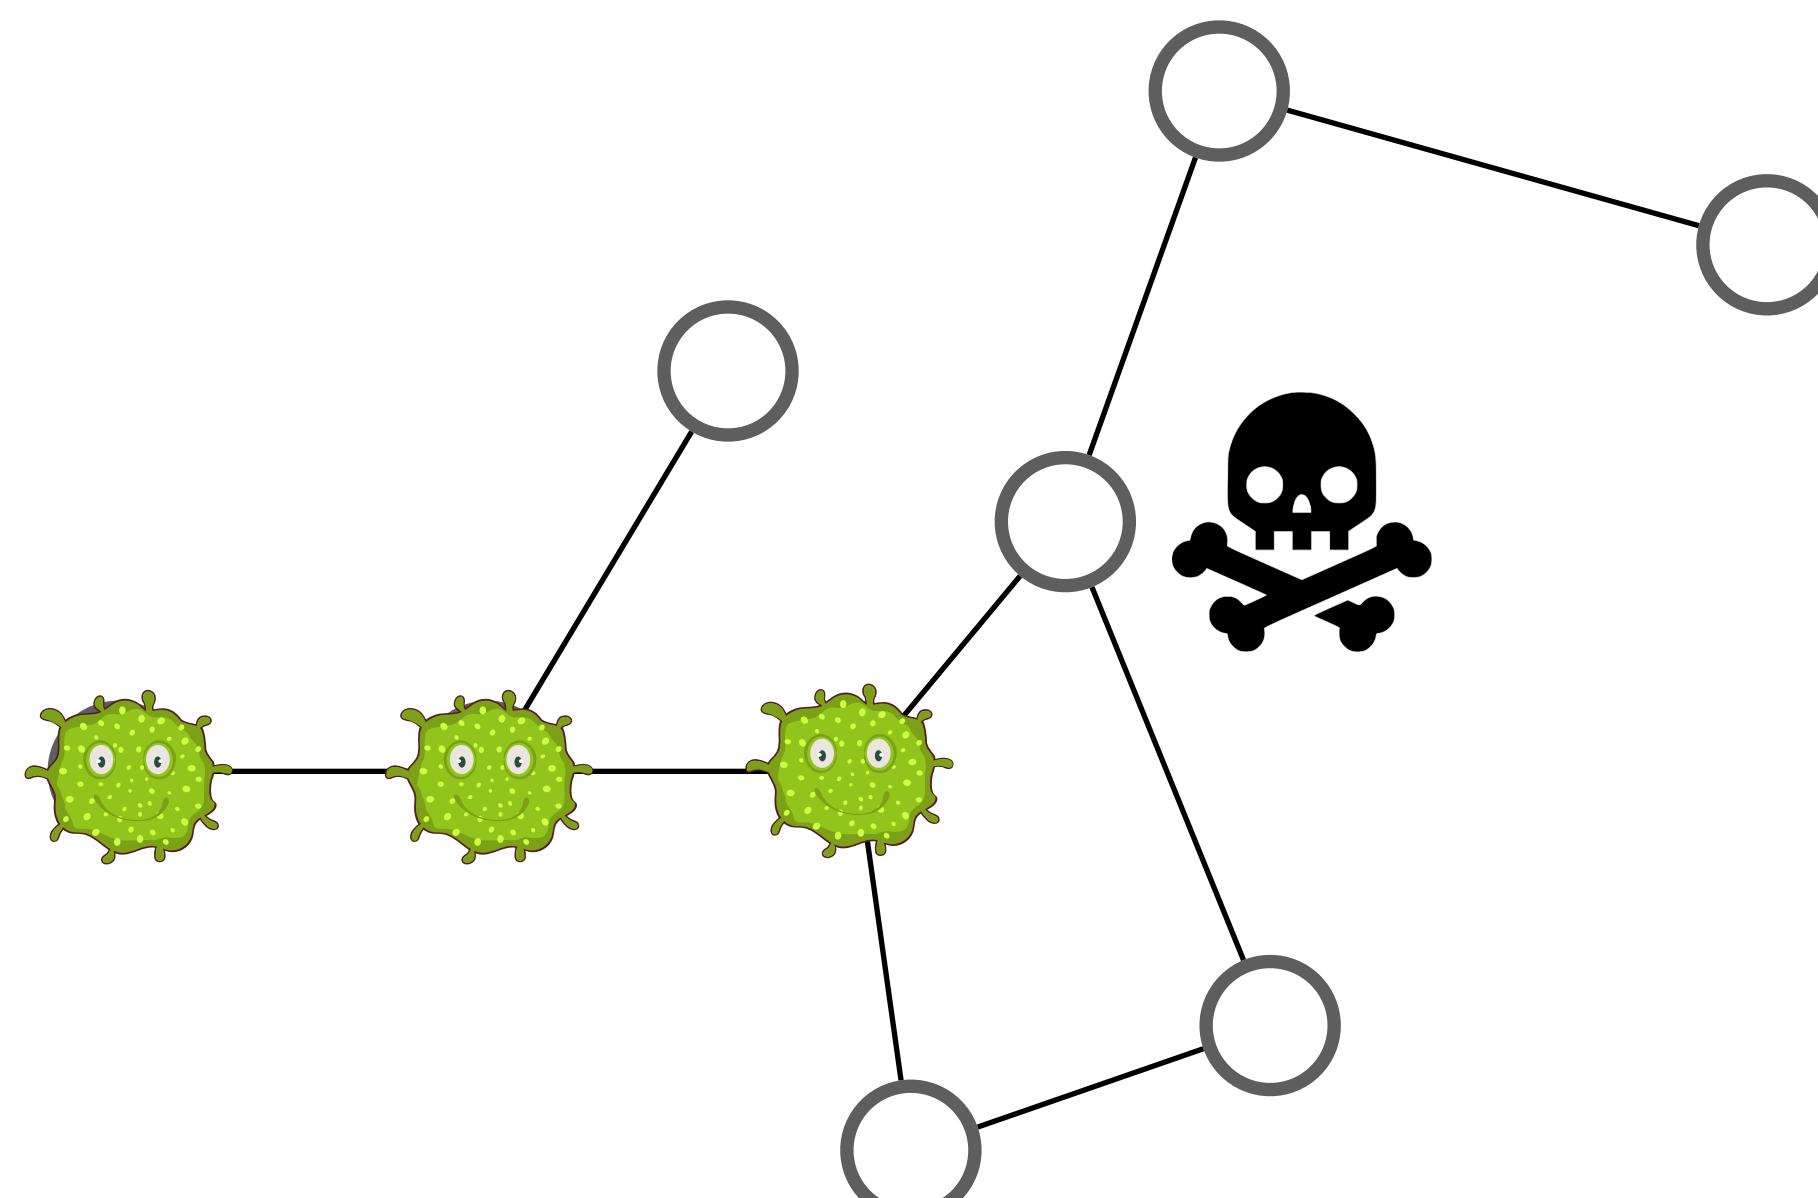
The major complication in implementing the above algorithm by a discrete process is to have a correct stopping criterion.<sup>†</sup>

This intuitive description covers all the standard shortest path methods and one type not previously proposed.

- a.  $V_r \neq 0, V_b = 0$  forward uni-directional algorithm<sup>17</sup>
- b.  $V_r = 0, V_b \neq 0$  backward uni-directional algorithm

<sup>†</sup>Reference 8, p. 174 gives a bi-directional algorithm with an incorrect termination criterion (also see Ref. 7).

# A\* Search: Amoebas!



Optimal Substructure

$$f(a) < f(b) \implies f(a \circ x) < f(b \circ x) \forall x$$

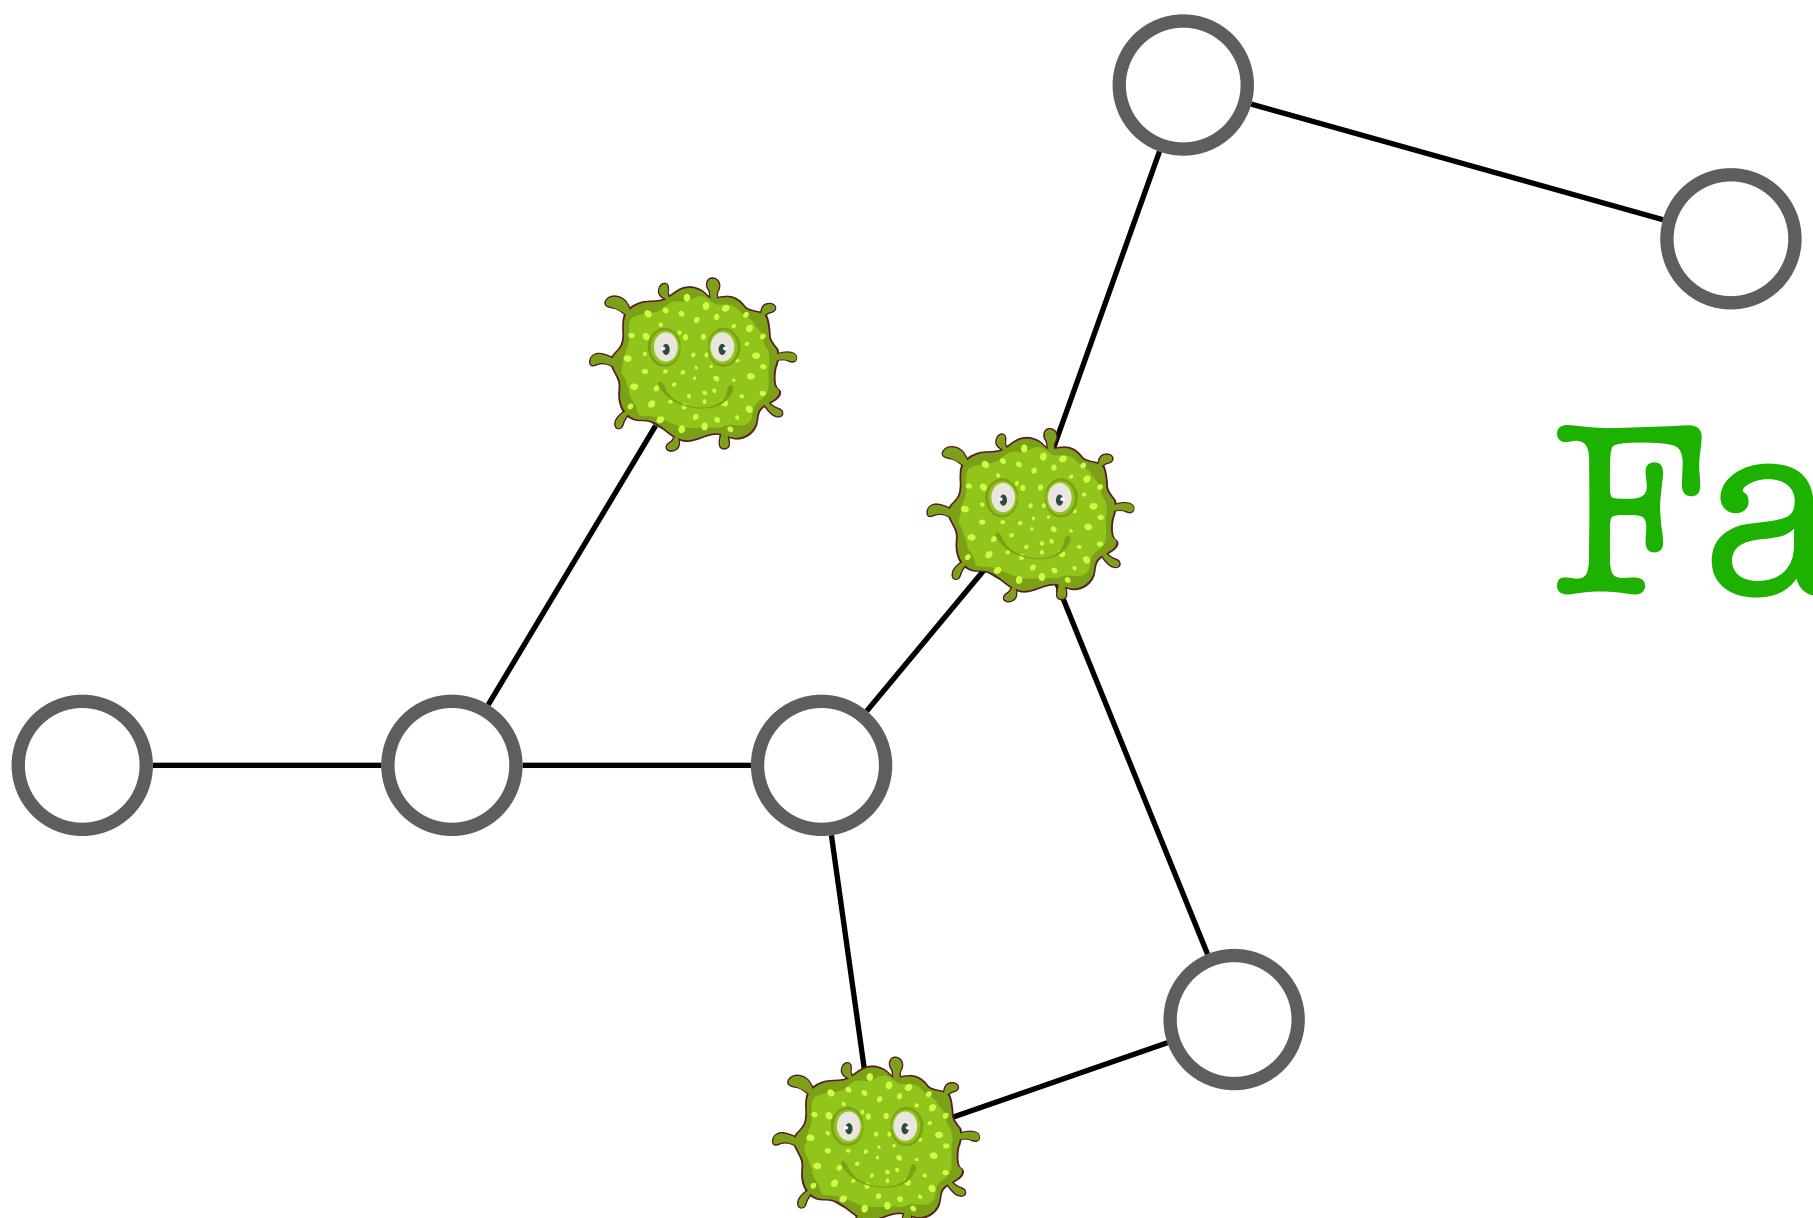
*You will never catch up.*

Bellman Condition

$$f^*(a) = \min_{x \in \text{SUCC}} \{c(a, x) + f^*(b)\}$$

*Be best, locally.*

# A\* Search: Favoritism





# Optimism in the Face of Uncertainty (OFU)

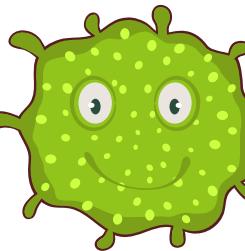
$$\min_{x \in \text{open}} g(x) + h(x)$$



*Always be optimistic under uncertainty.  
You'll either be correct,  
or learn something important if you're wrong.*

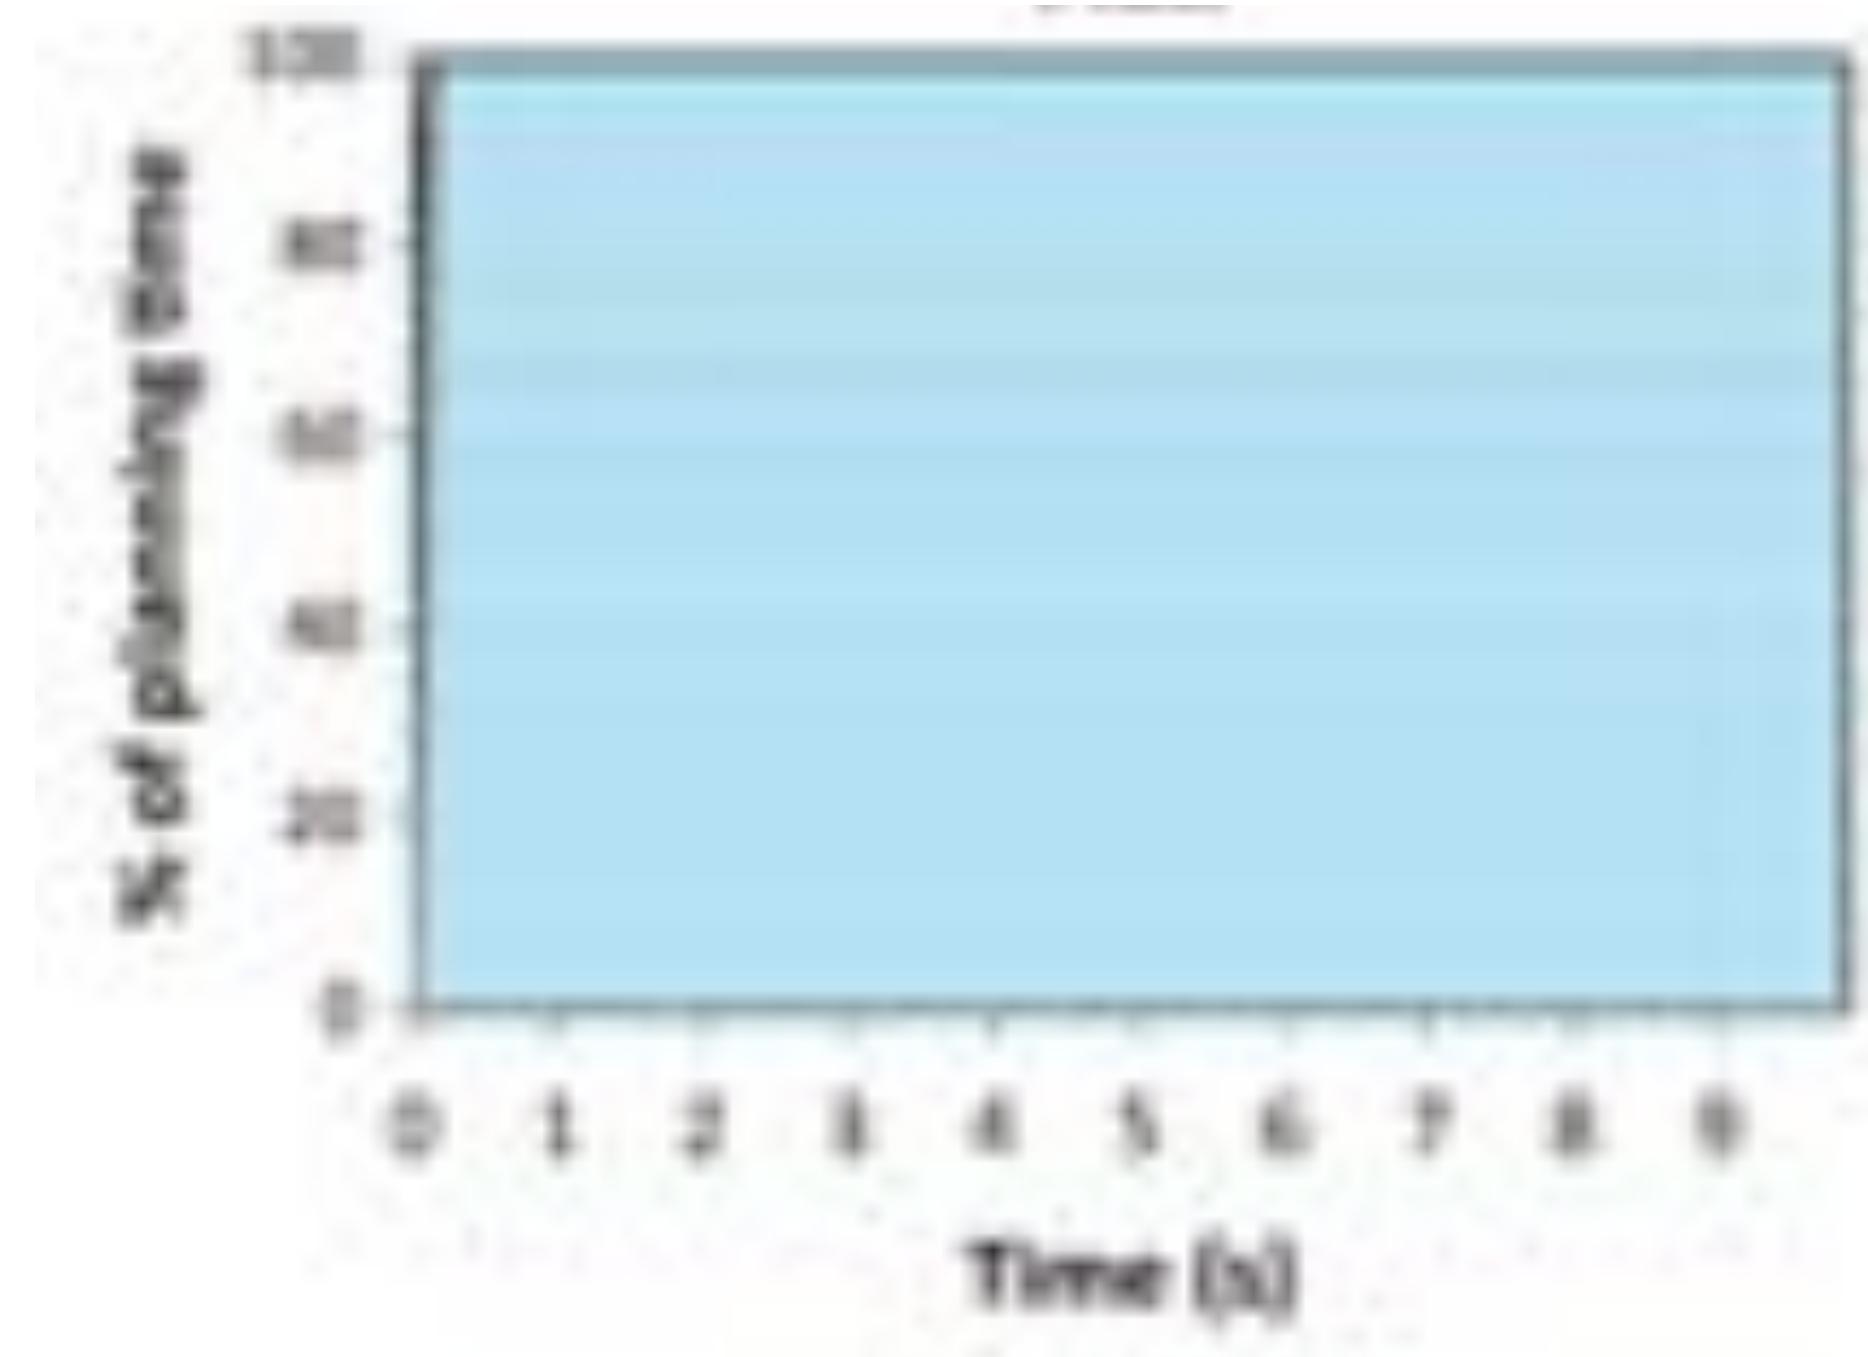
# A\* Search is Optimal ...

Expands the Fewest Number of Vertices



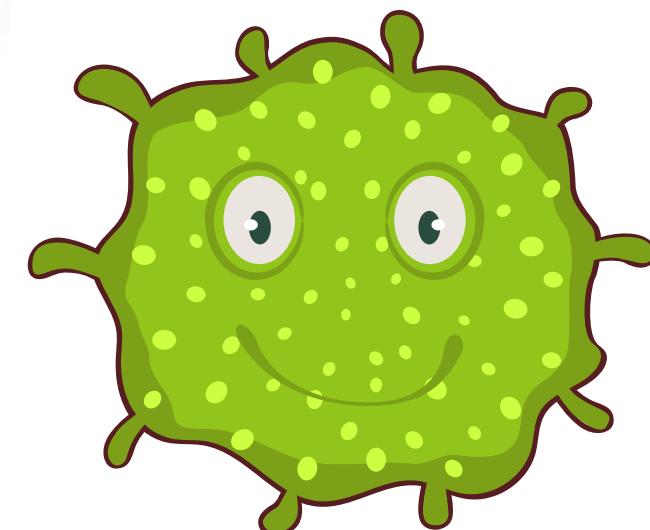
But is this what we  
*really* want in Motion Planning?

# Edge Evaluation Dominates Planning Time



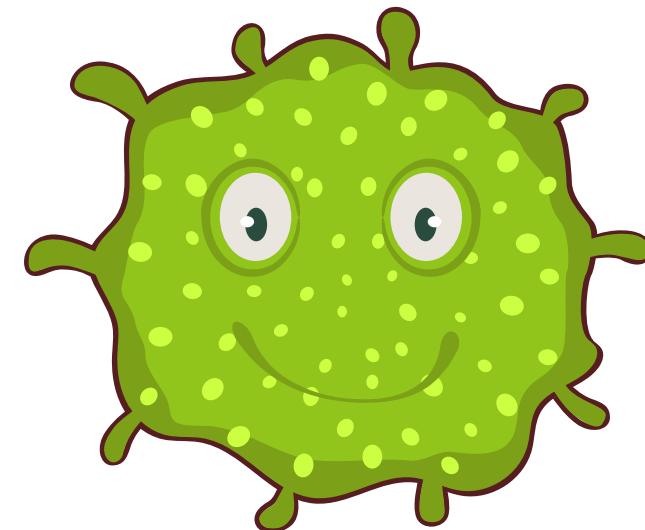
**Edge Evaluations**

**Other**



Amoebas are Cheap  
Slime is Expensive

Is there a Search Algorithm  
that **Minimizes**  
the Number of Edge Evaluations?



I don't care about amoebas.  
What algorithm minimizes slime?

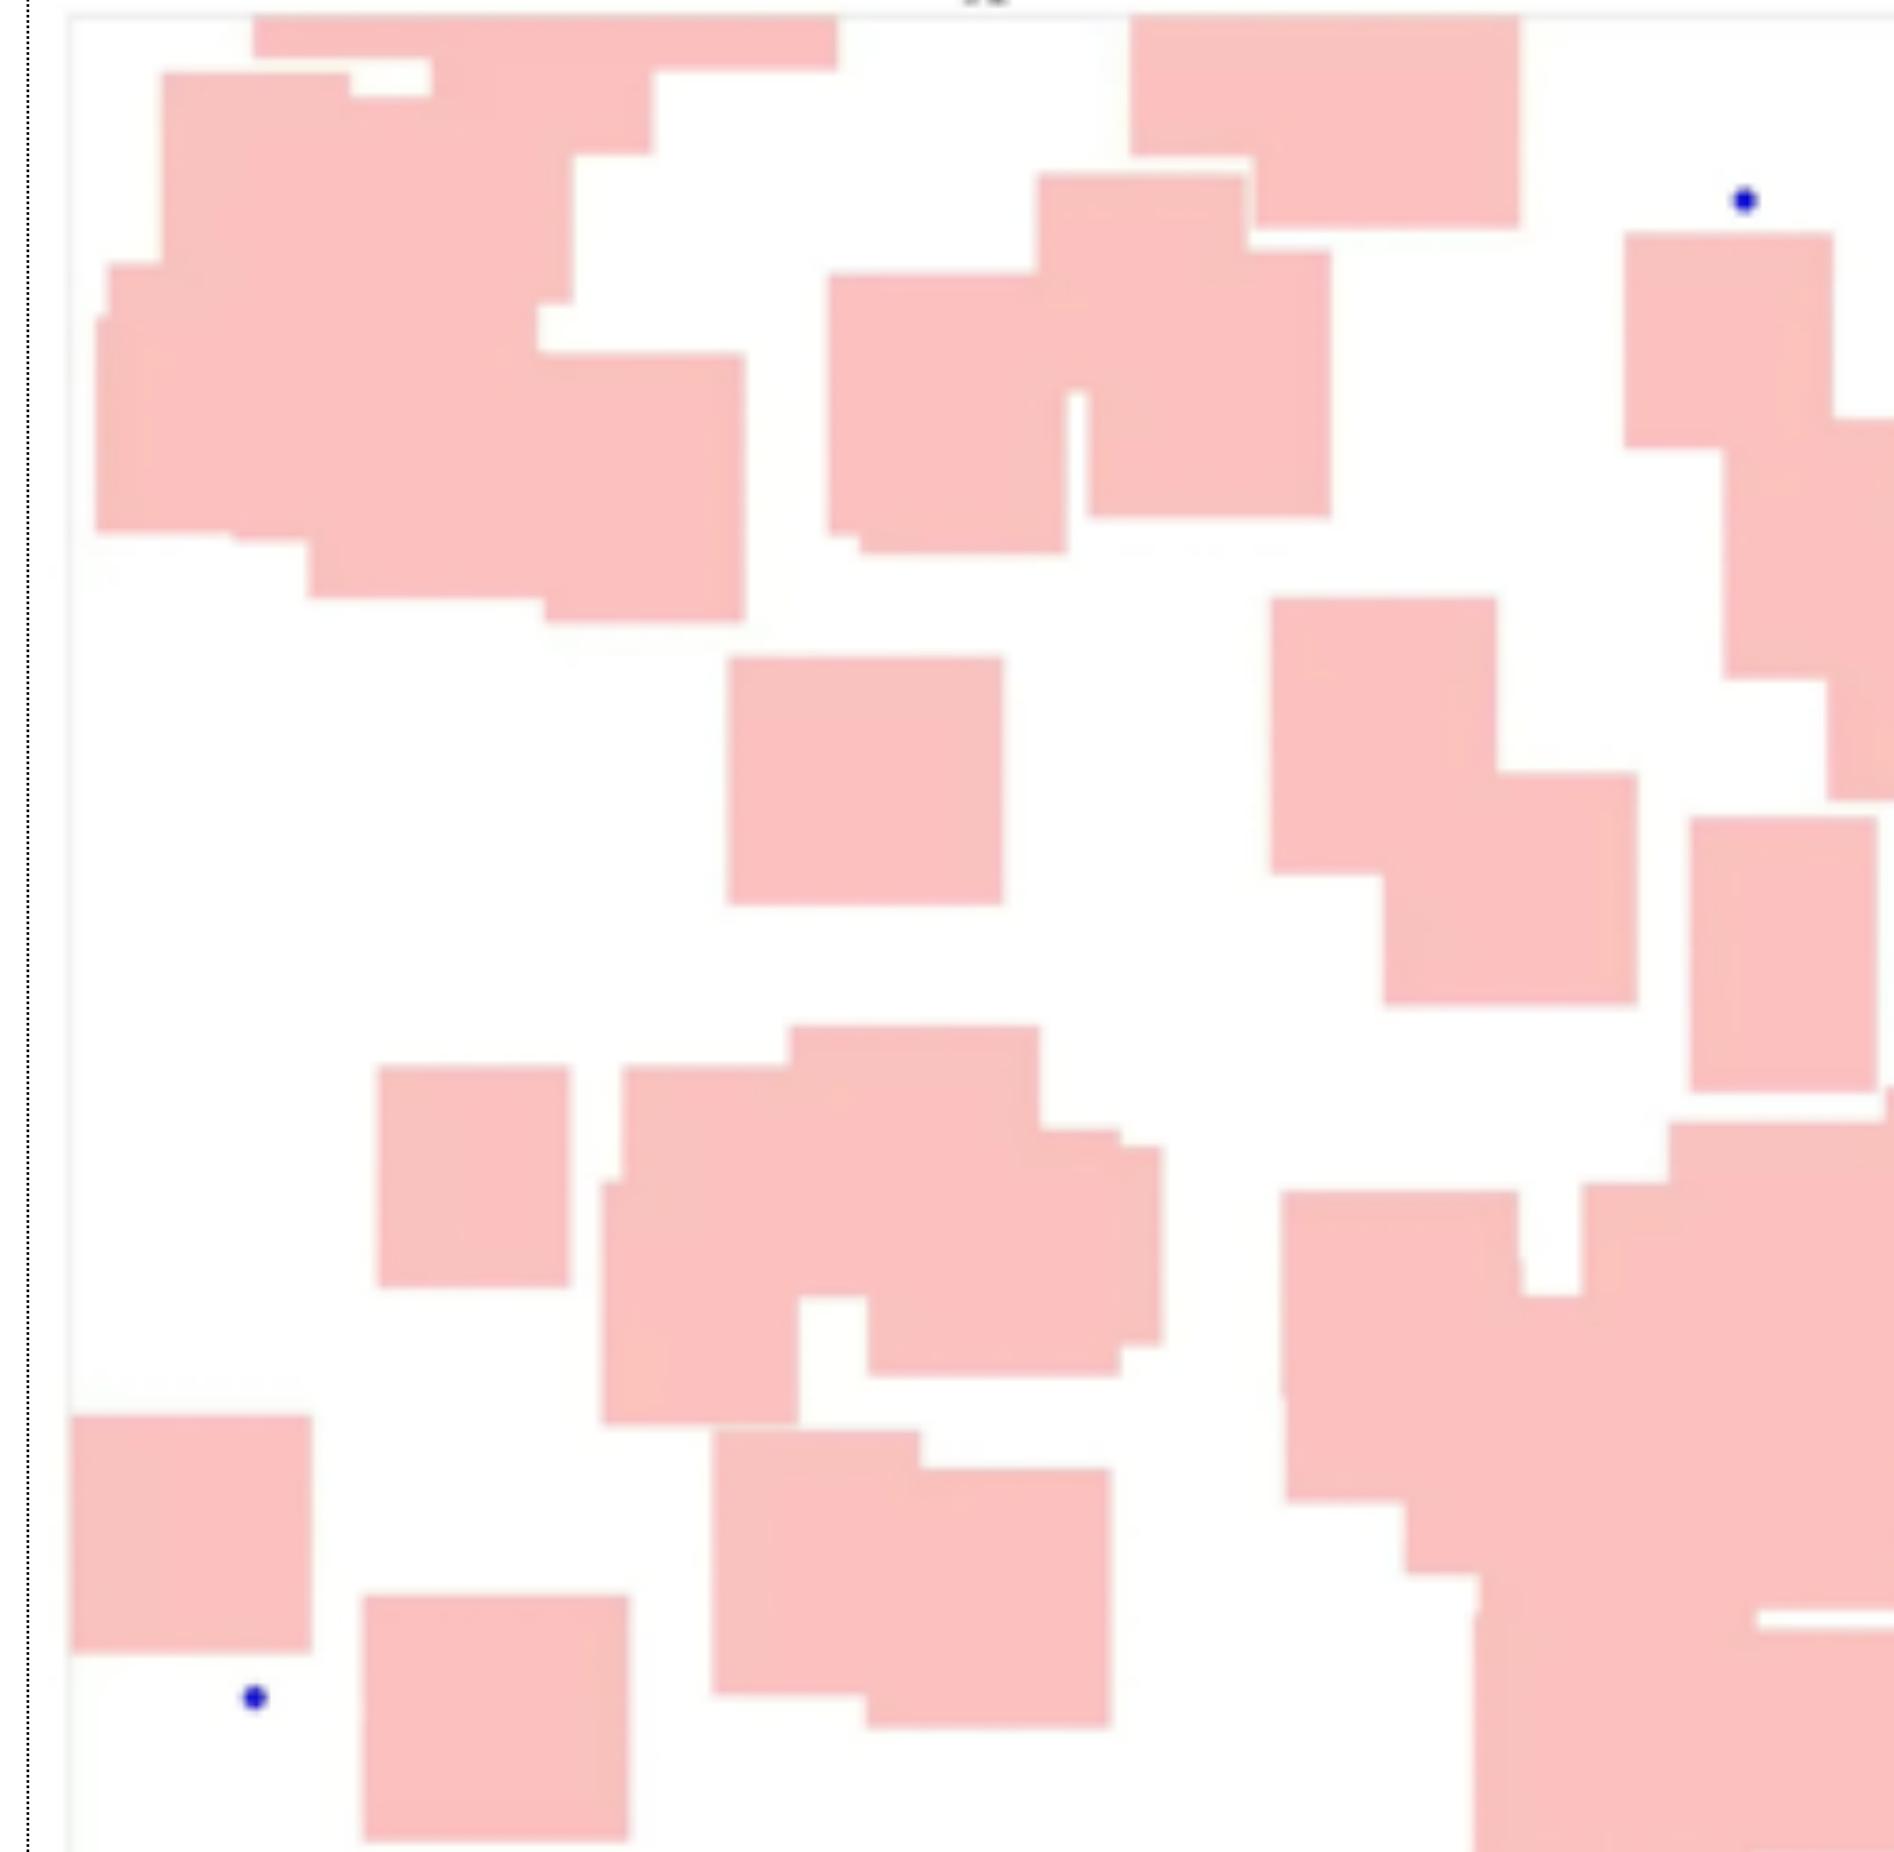
# LazySP

ICAPS 2018 [Best Conference Paper Award Winner]

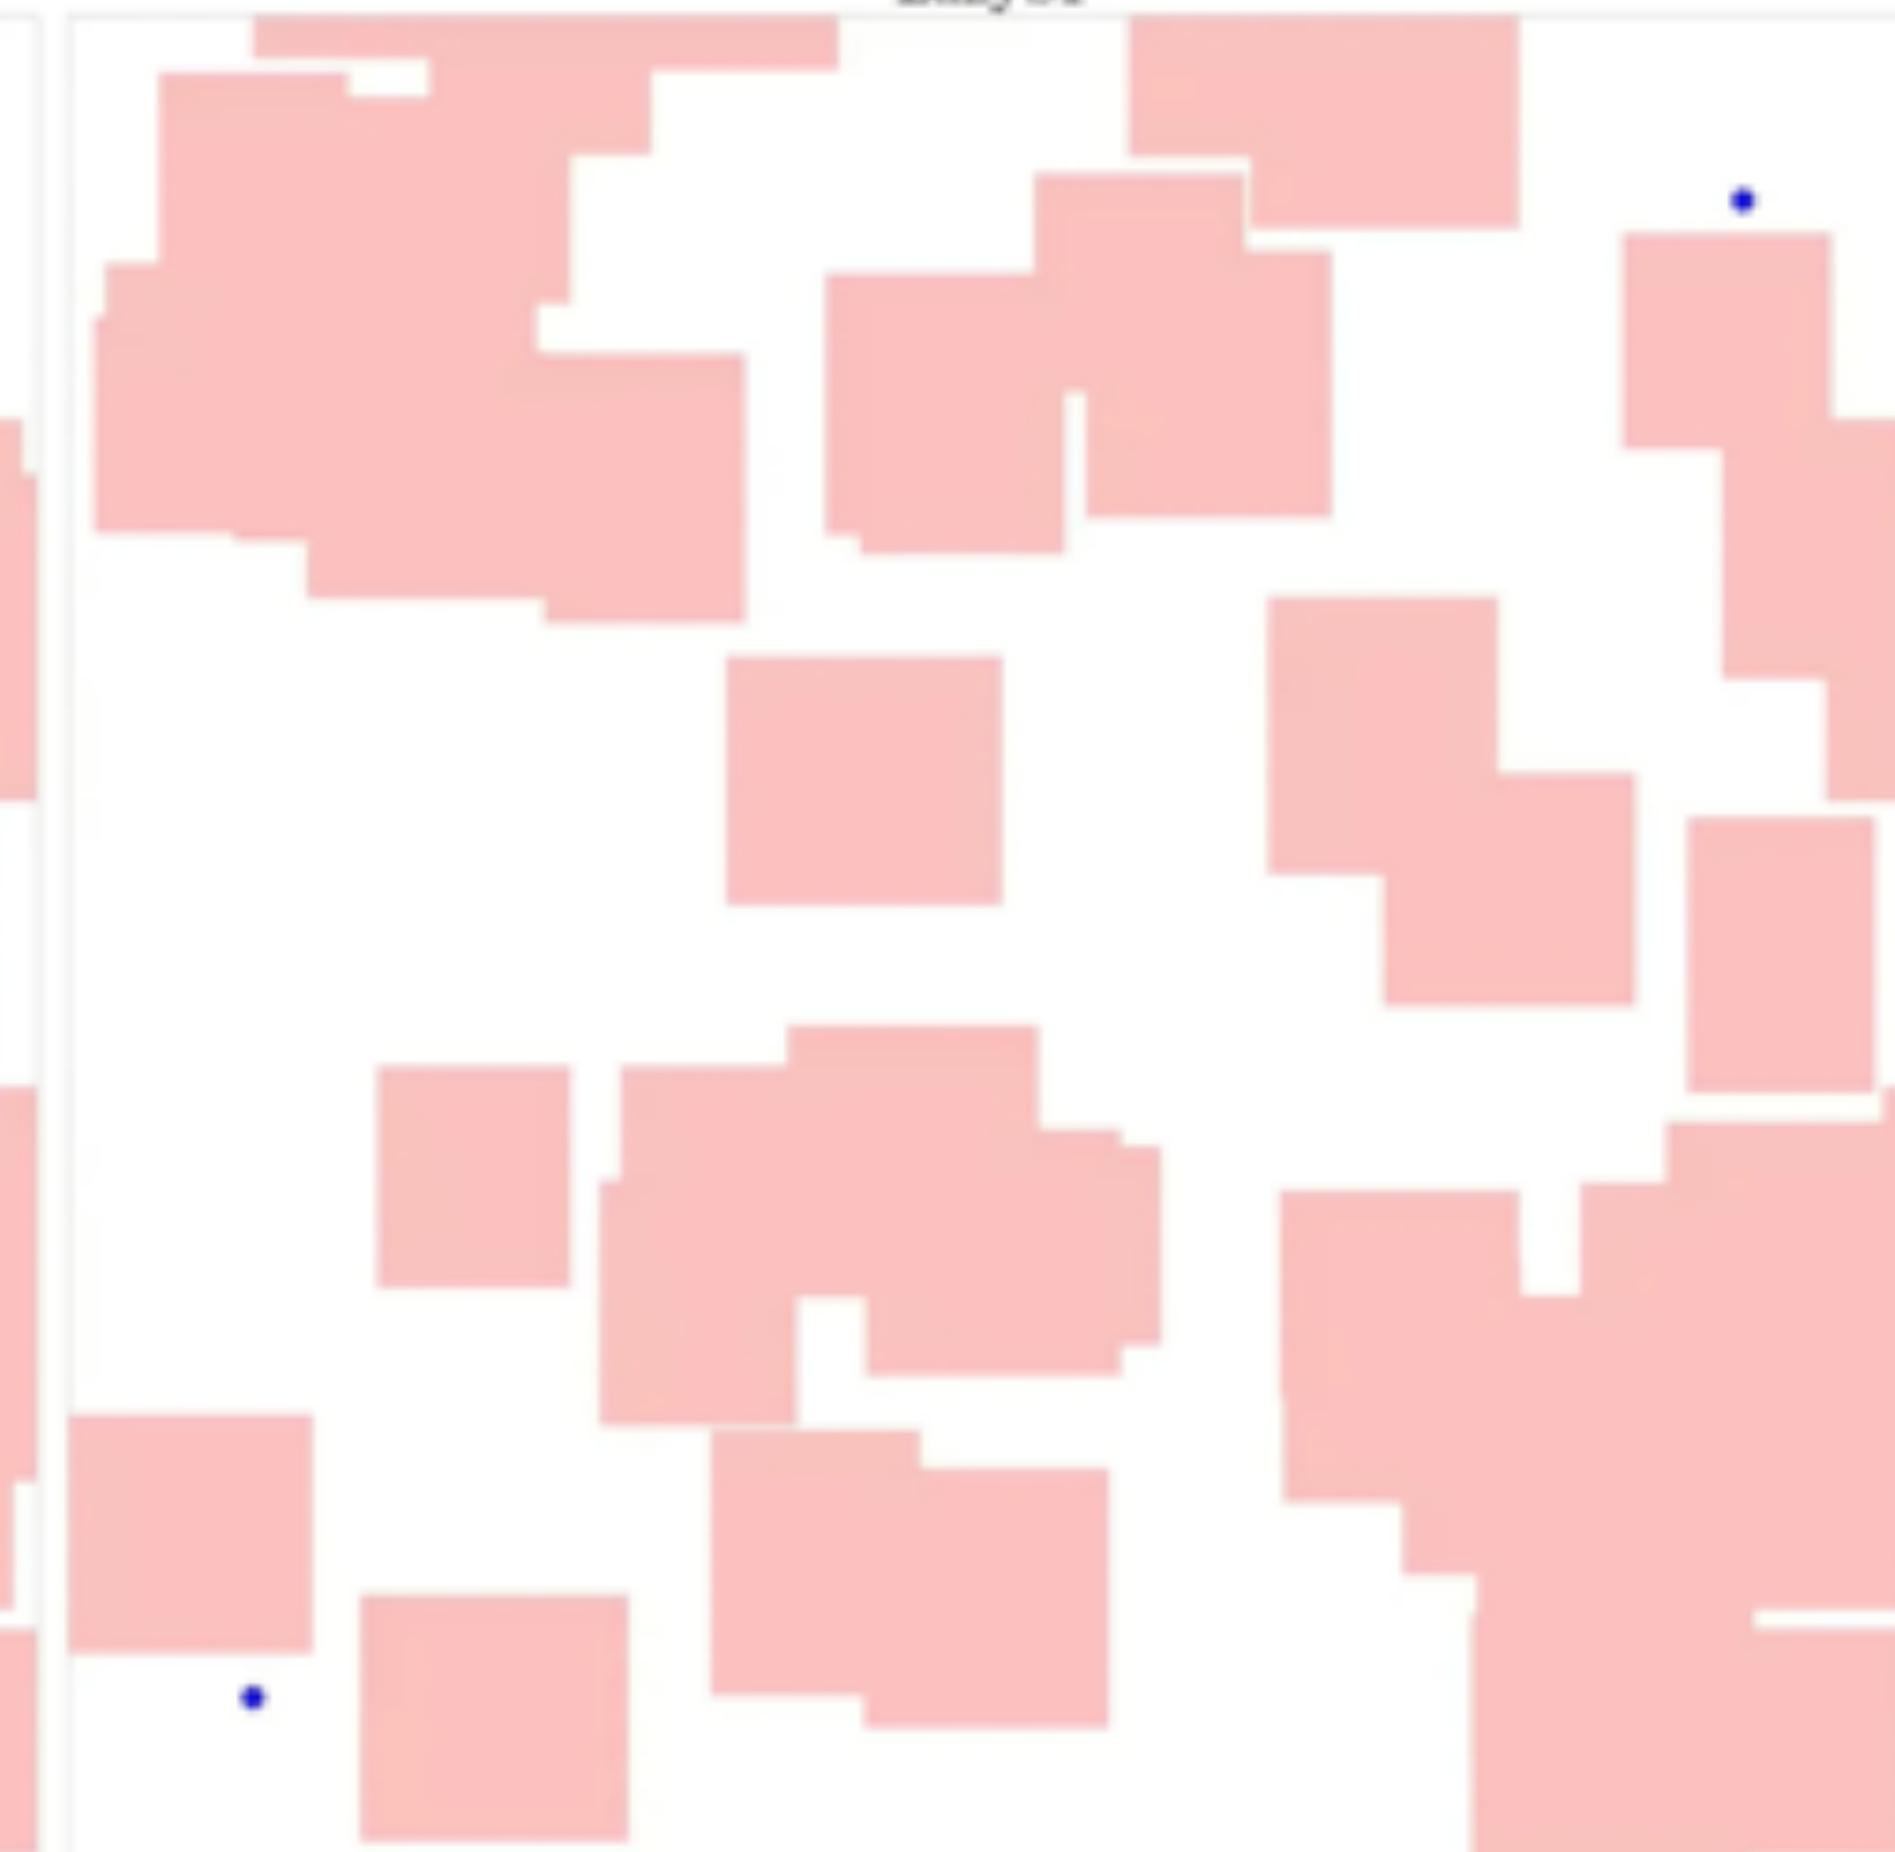
First Provably Edge-Optimal A\*-like Search Algorithm

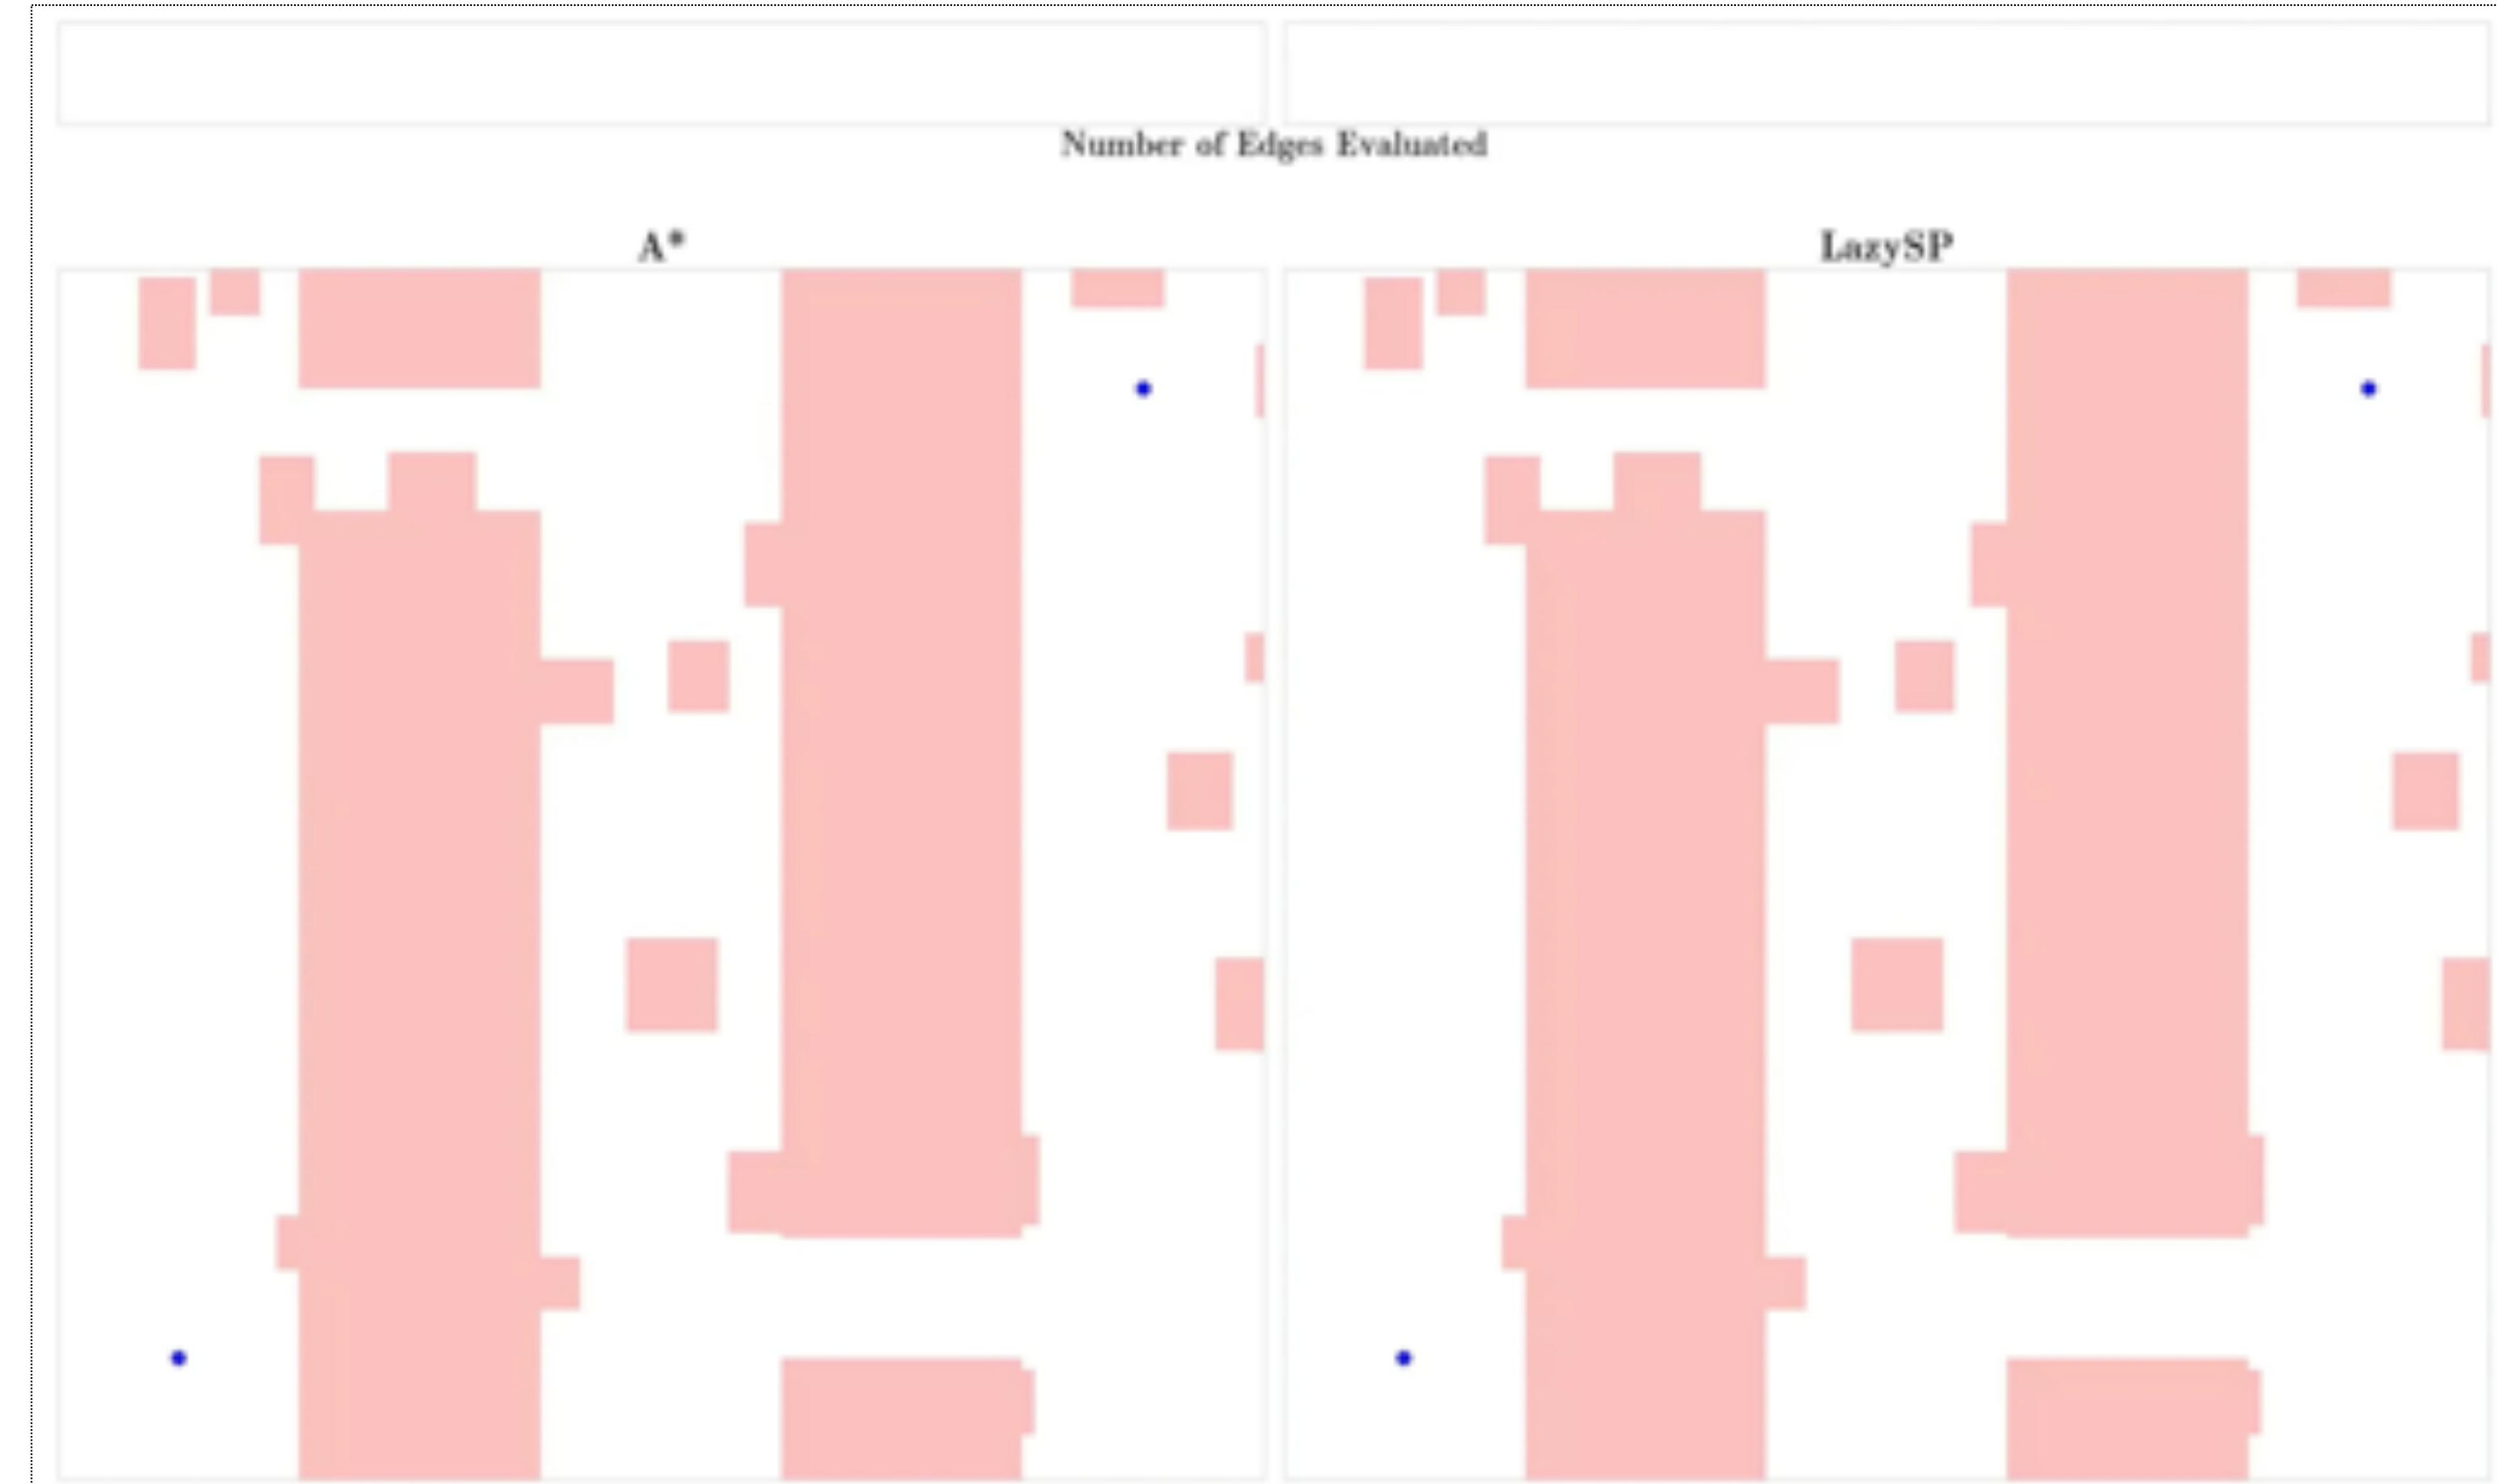
Number of Edges Evaluated

A\*



LazySP

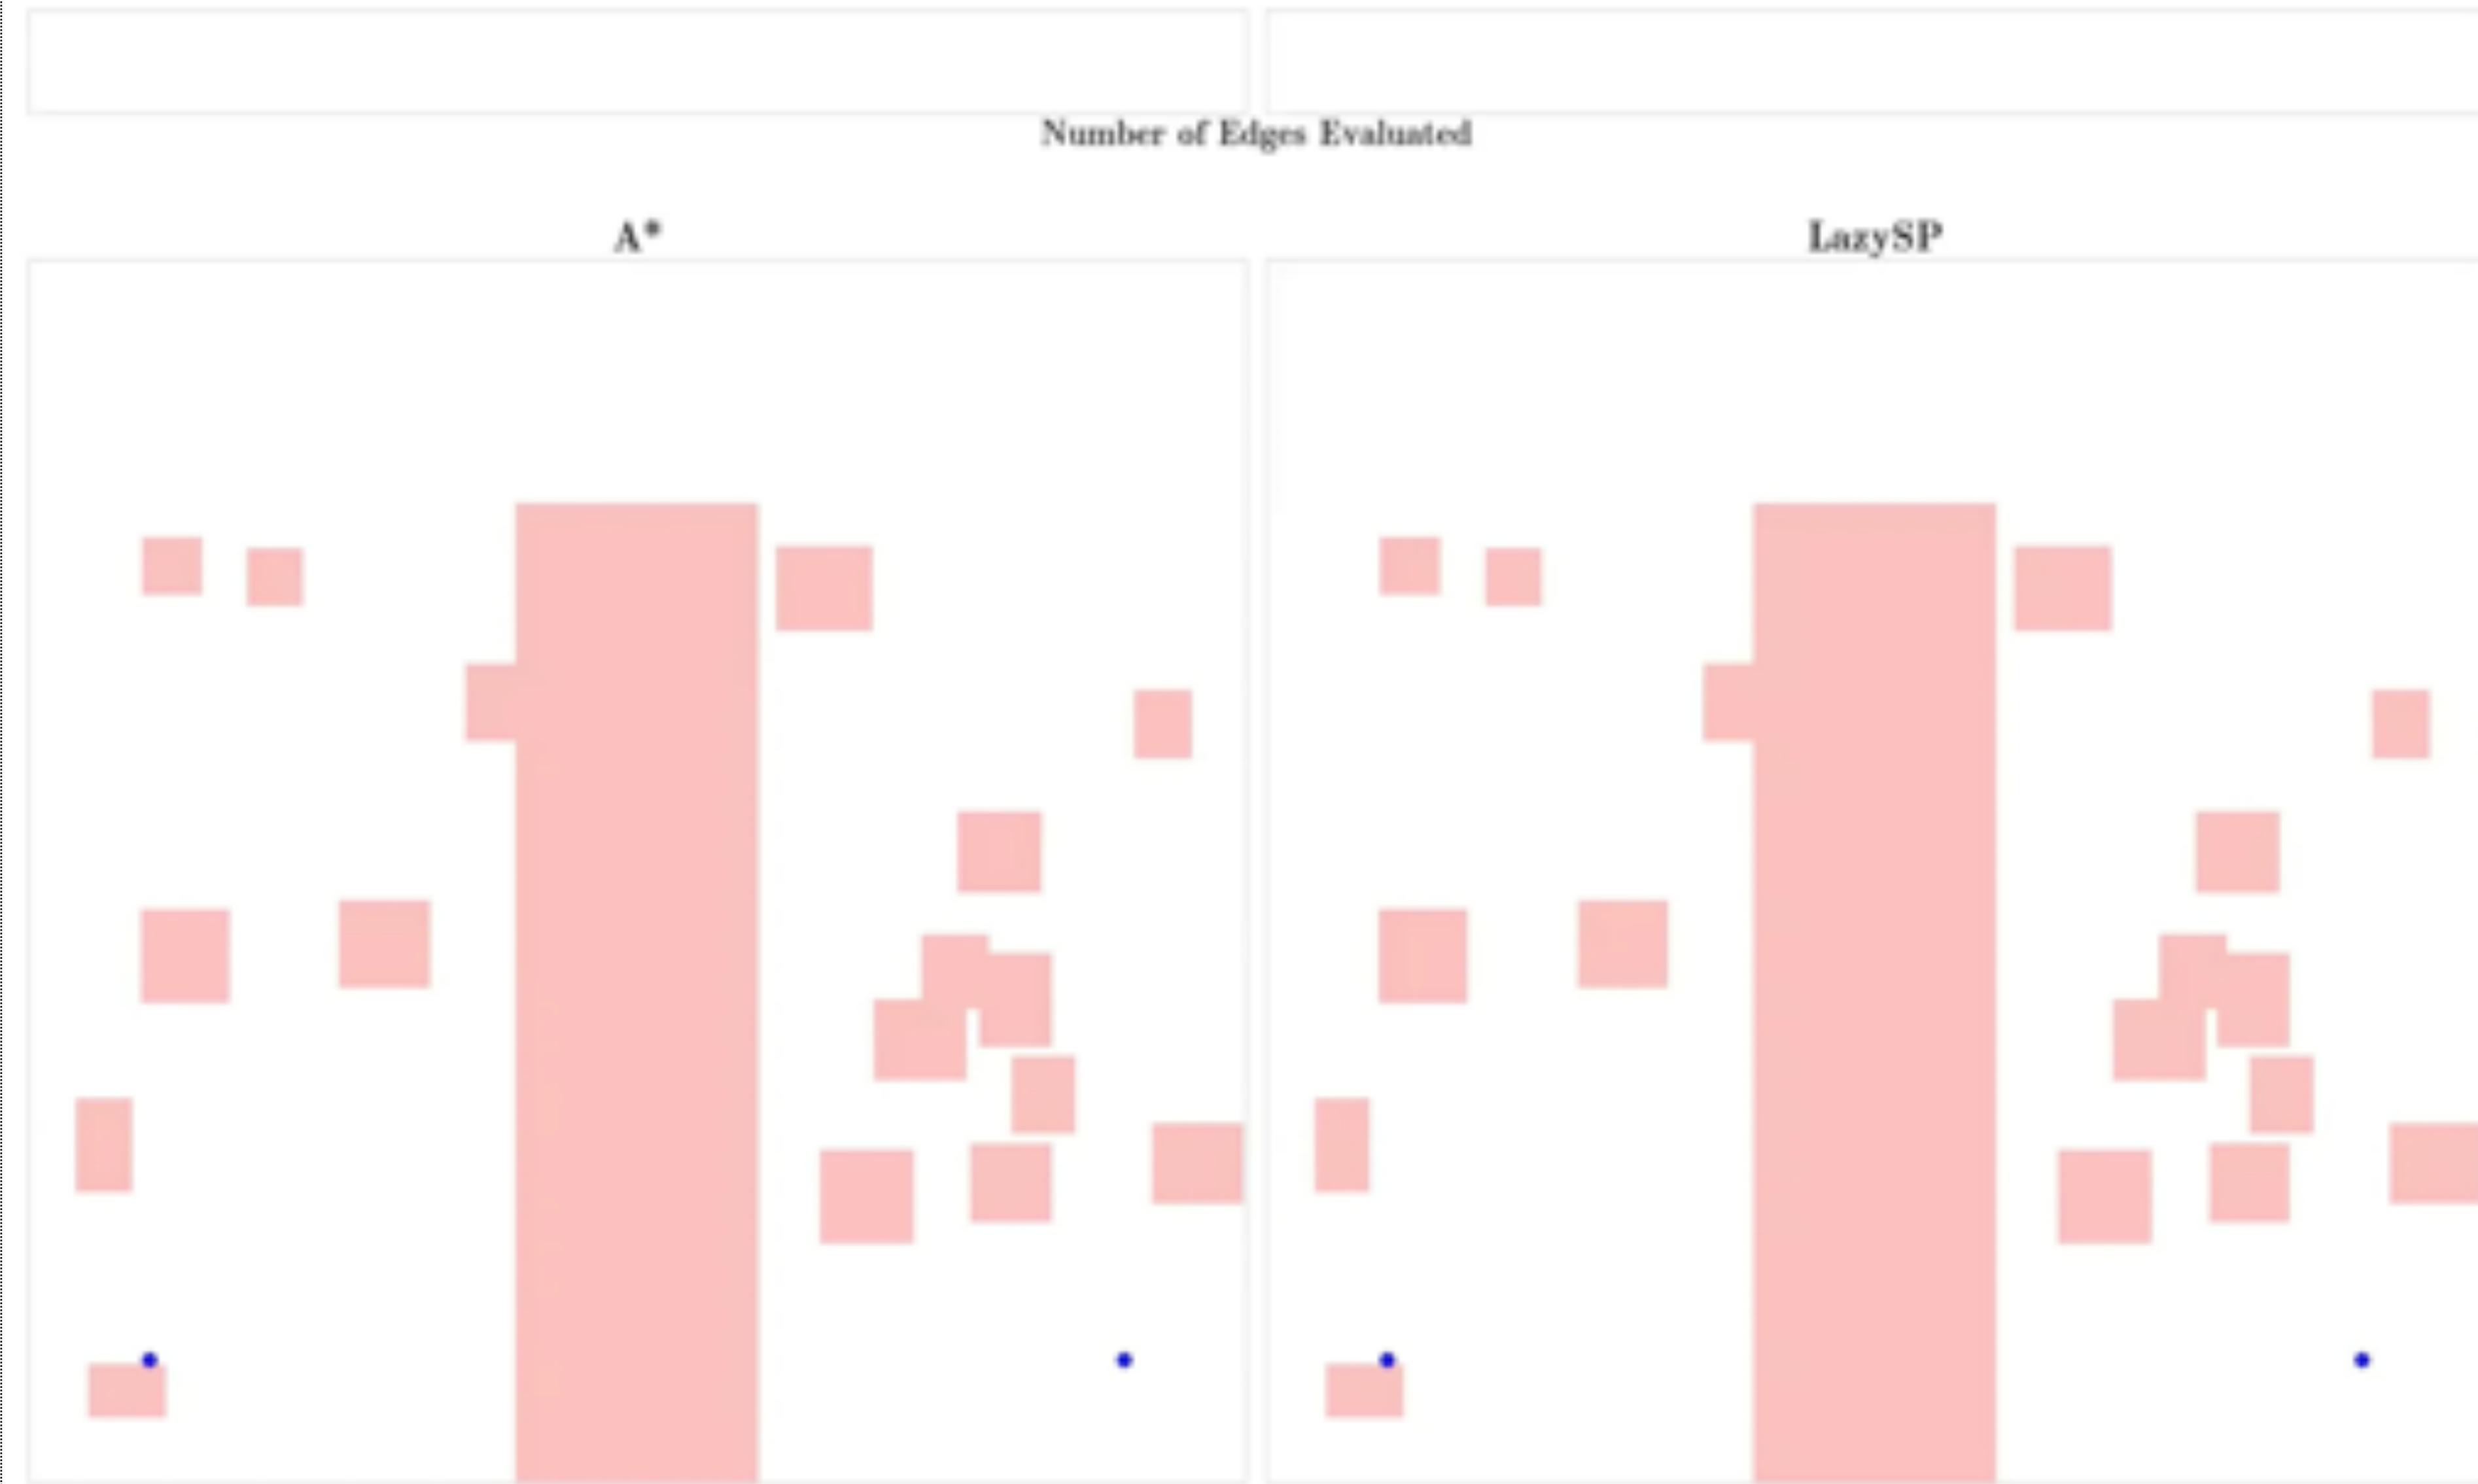




Number of Edges Evaluated

A\*

LazySP



Number of Edges Evaluated

A\*



LazySP



# LazySP

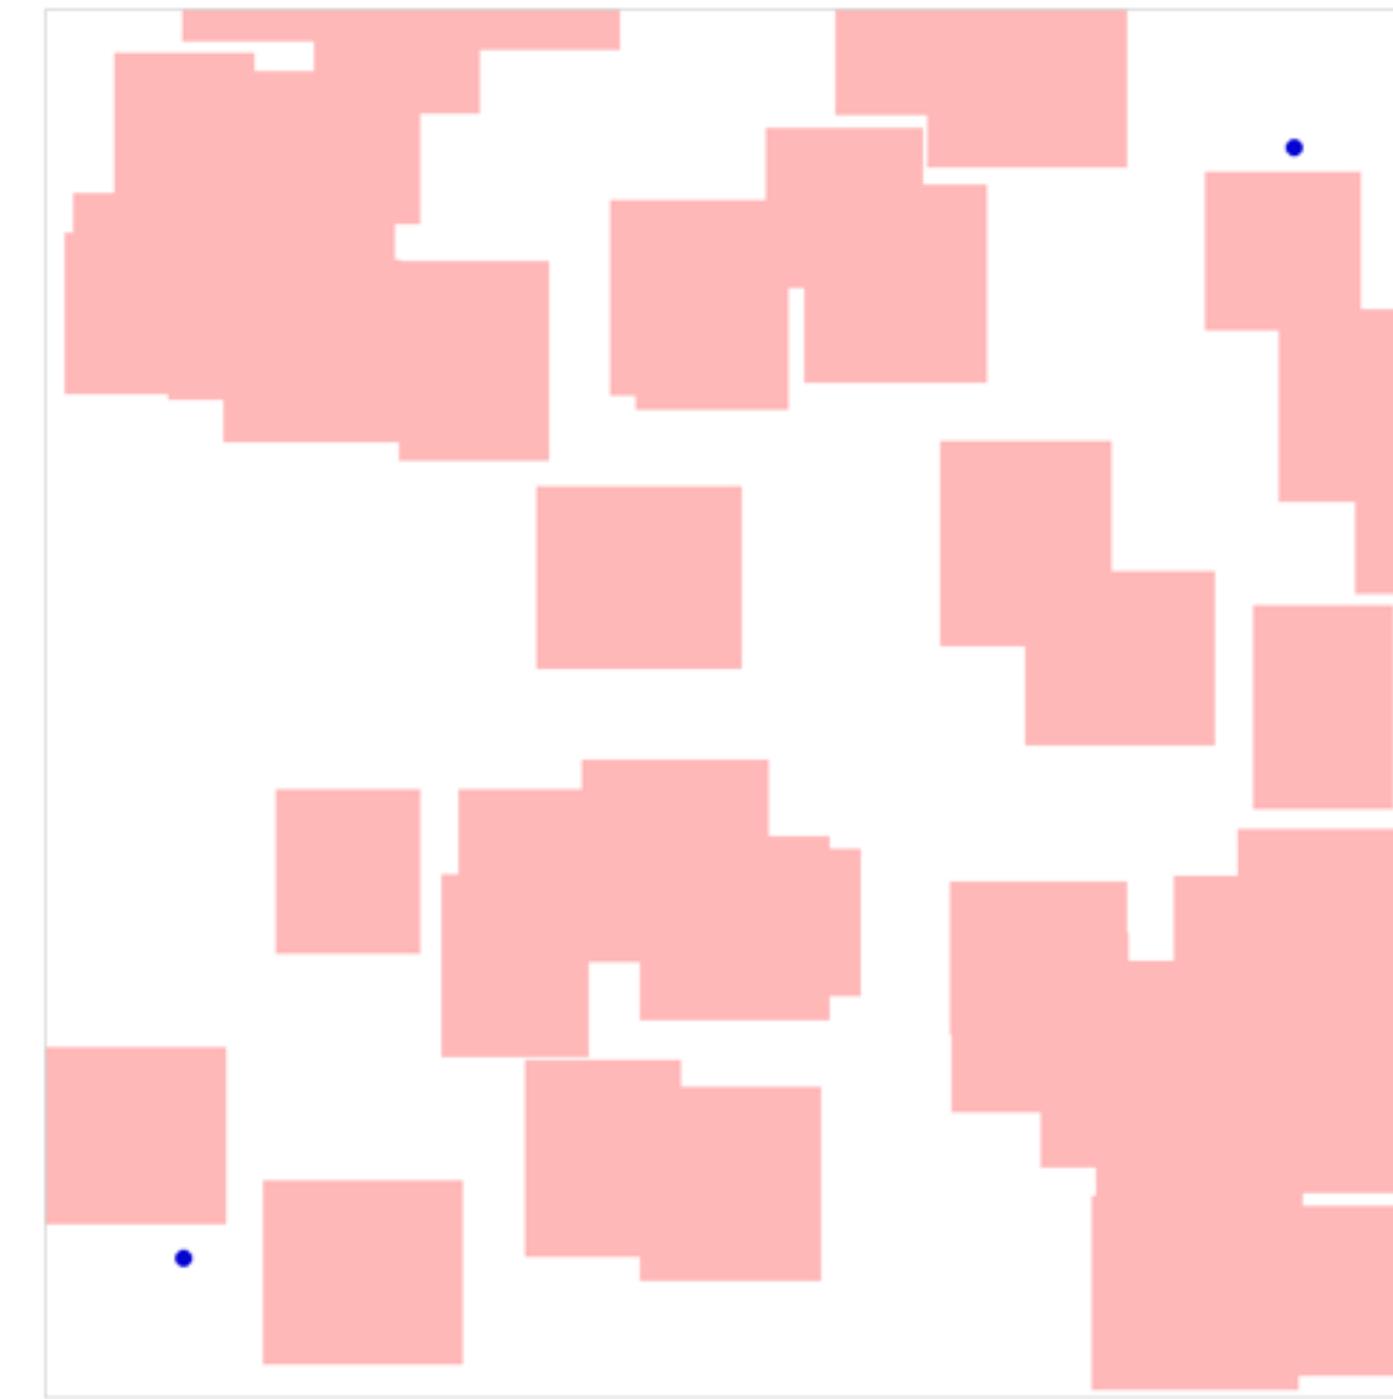
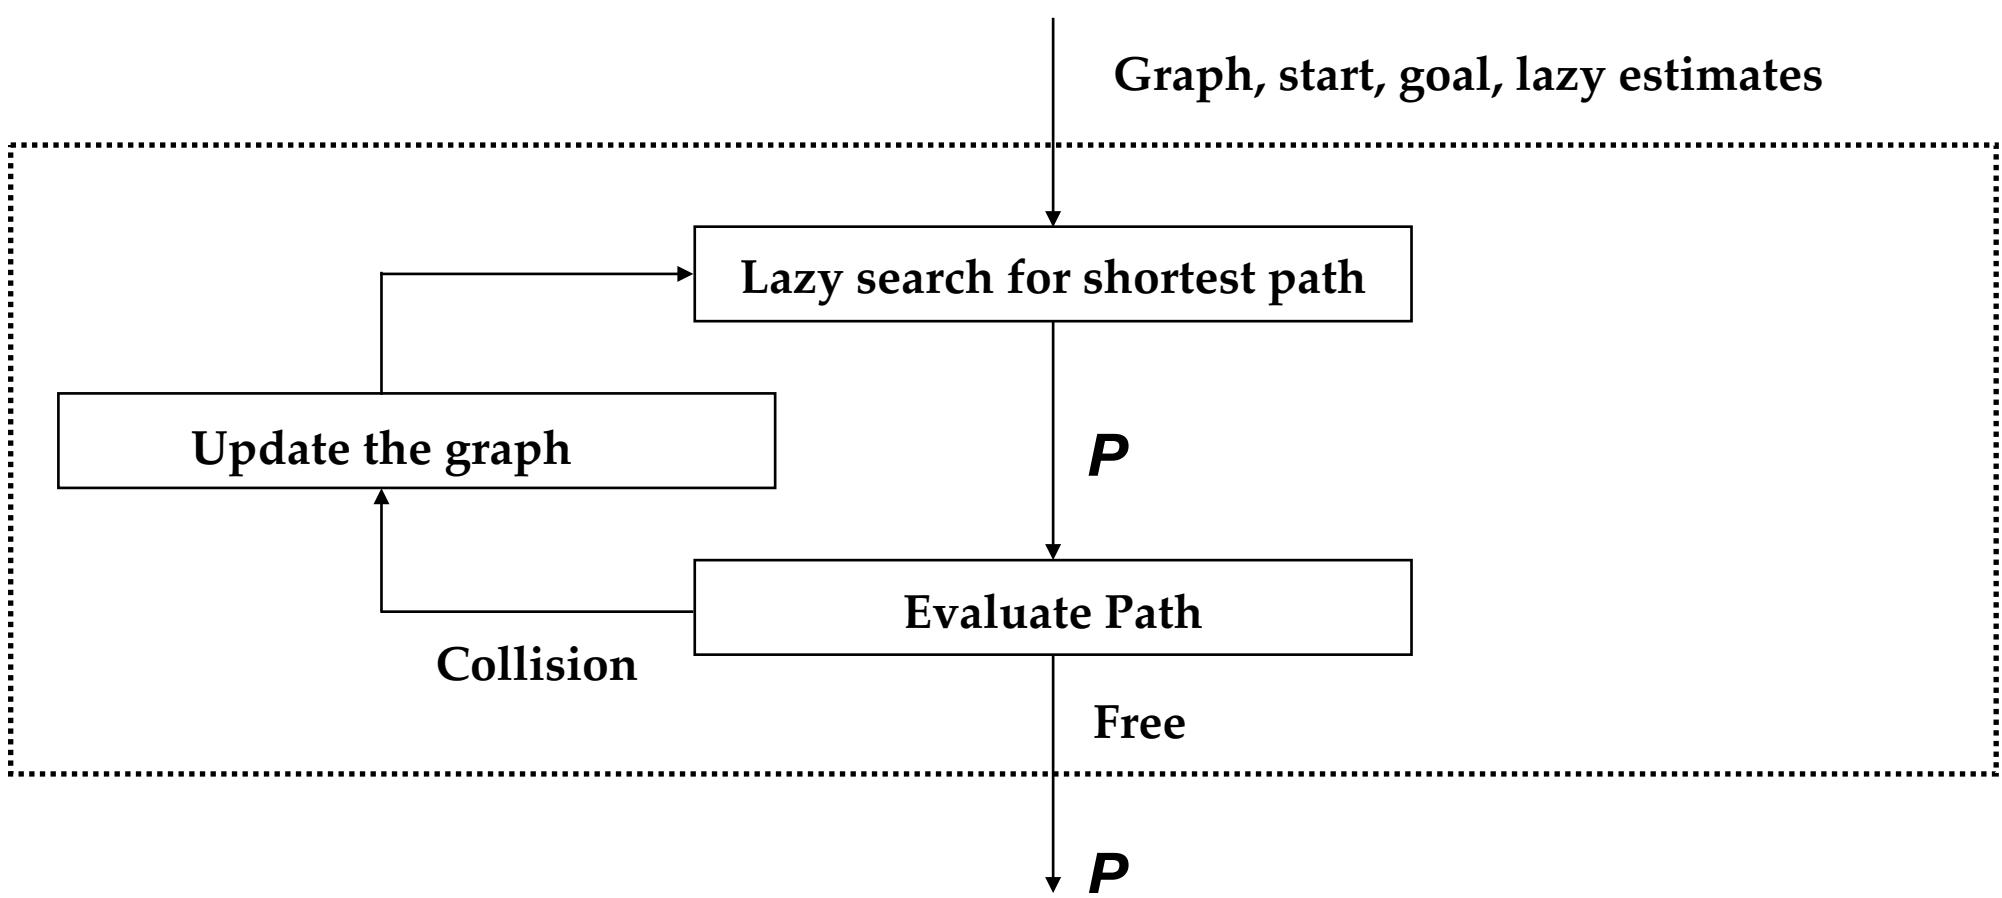
Greedy Best-first Search over Paths



*To find the shortest path,  
eliminate all shorter paths!*

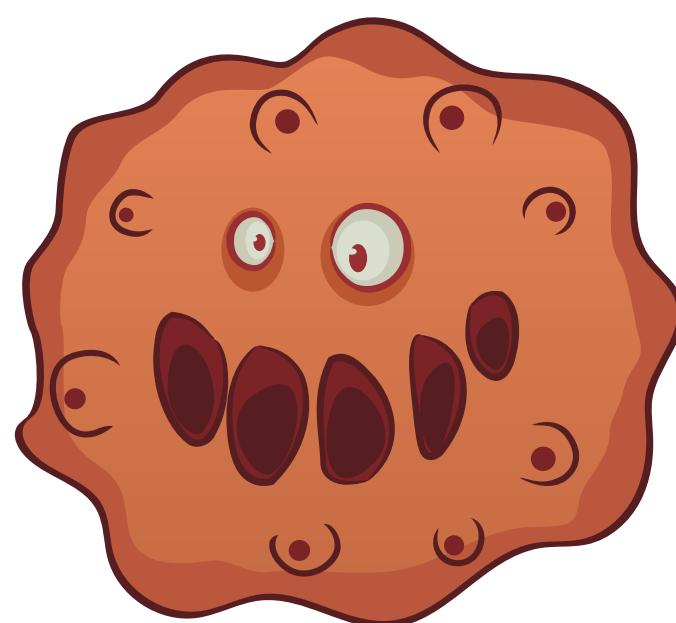
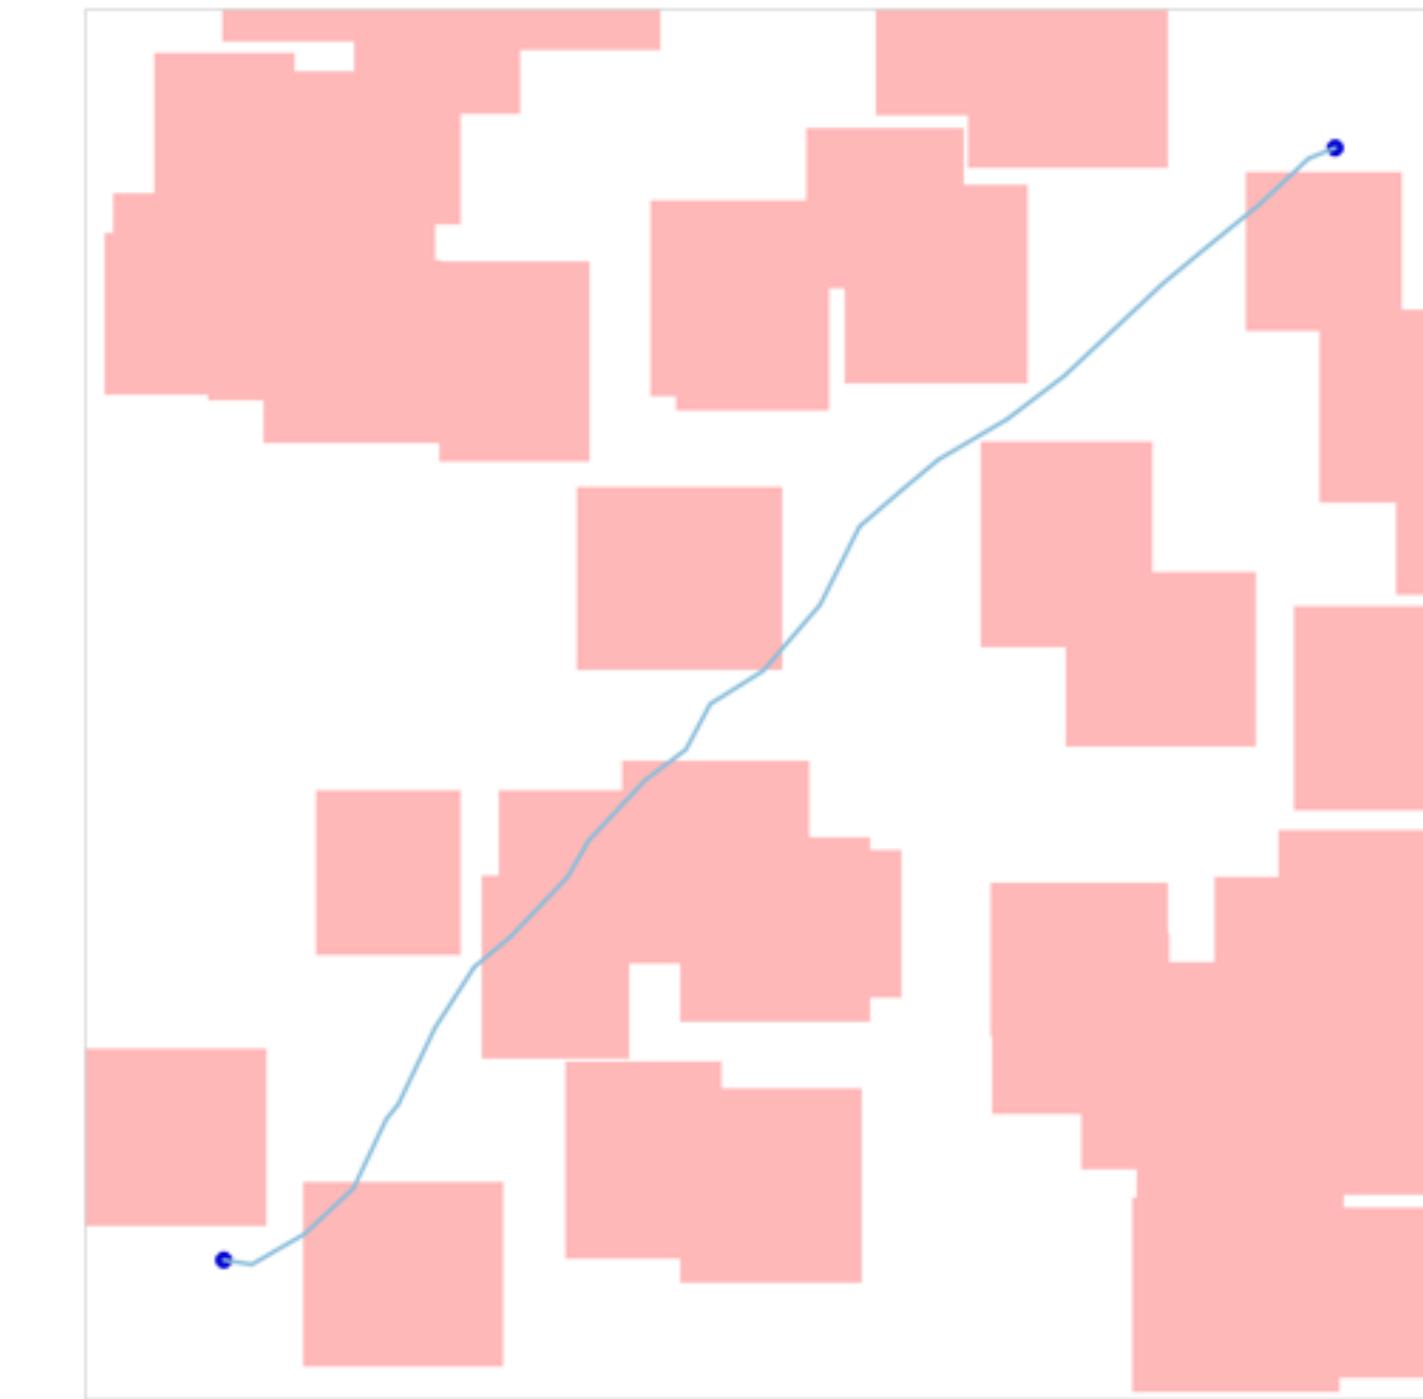
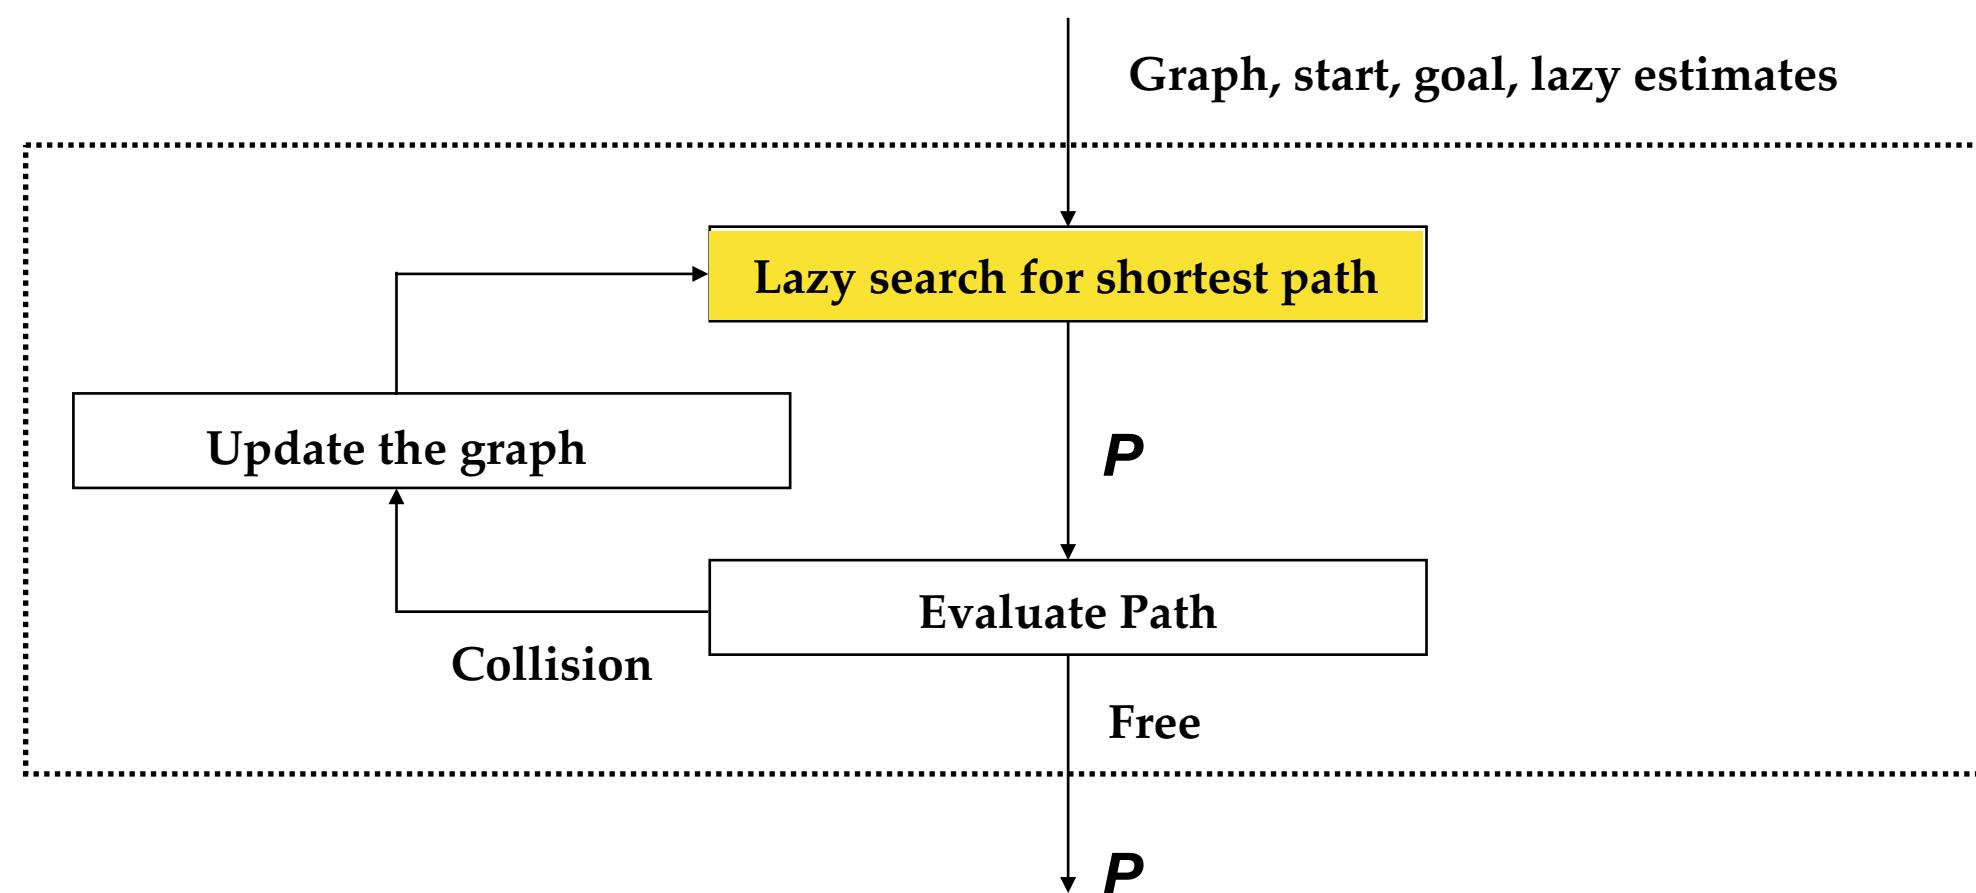
# LazySP

## OFU on Steroids!



# LazySP

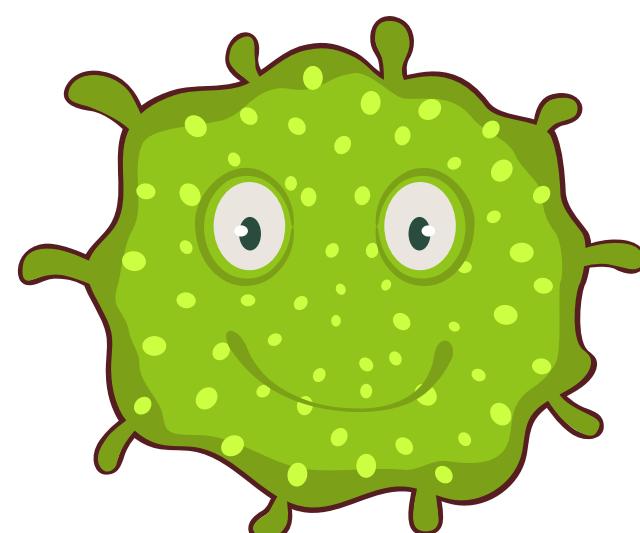
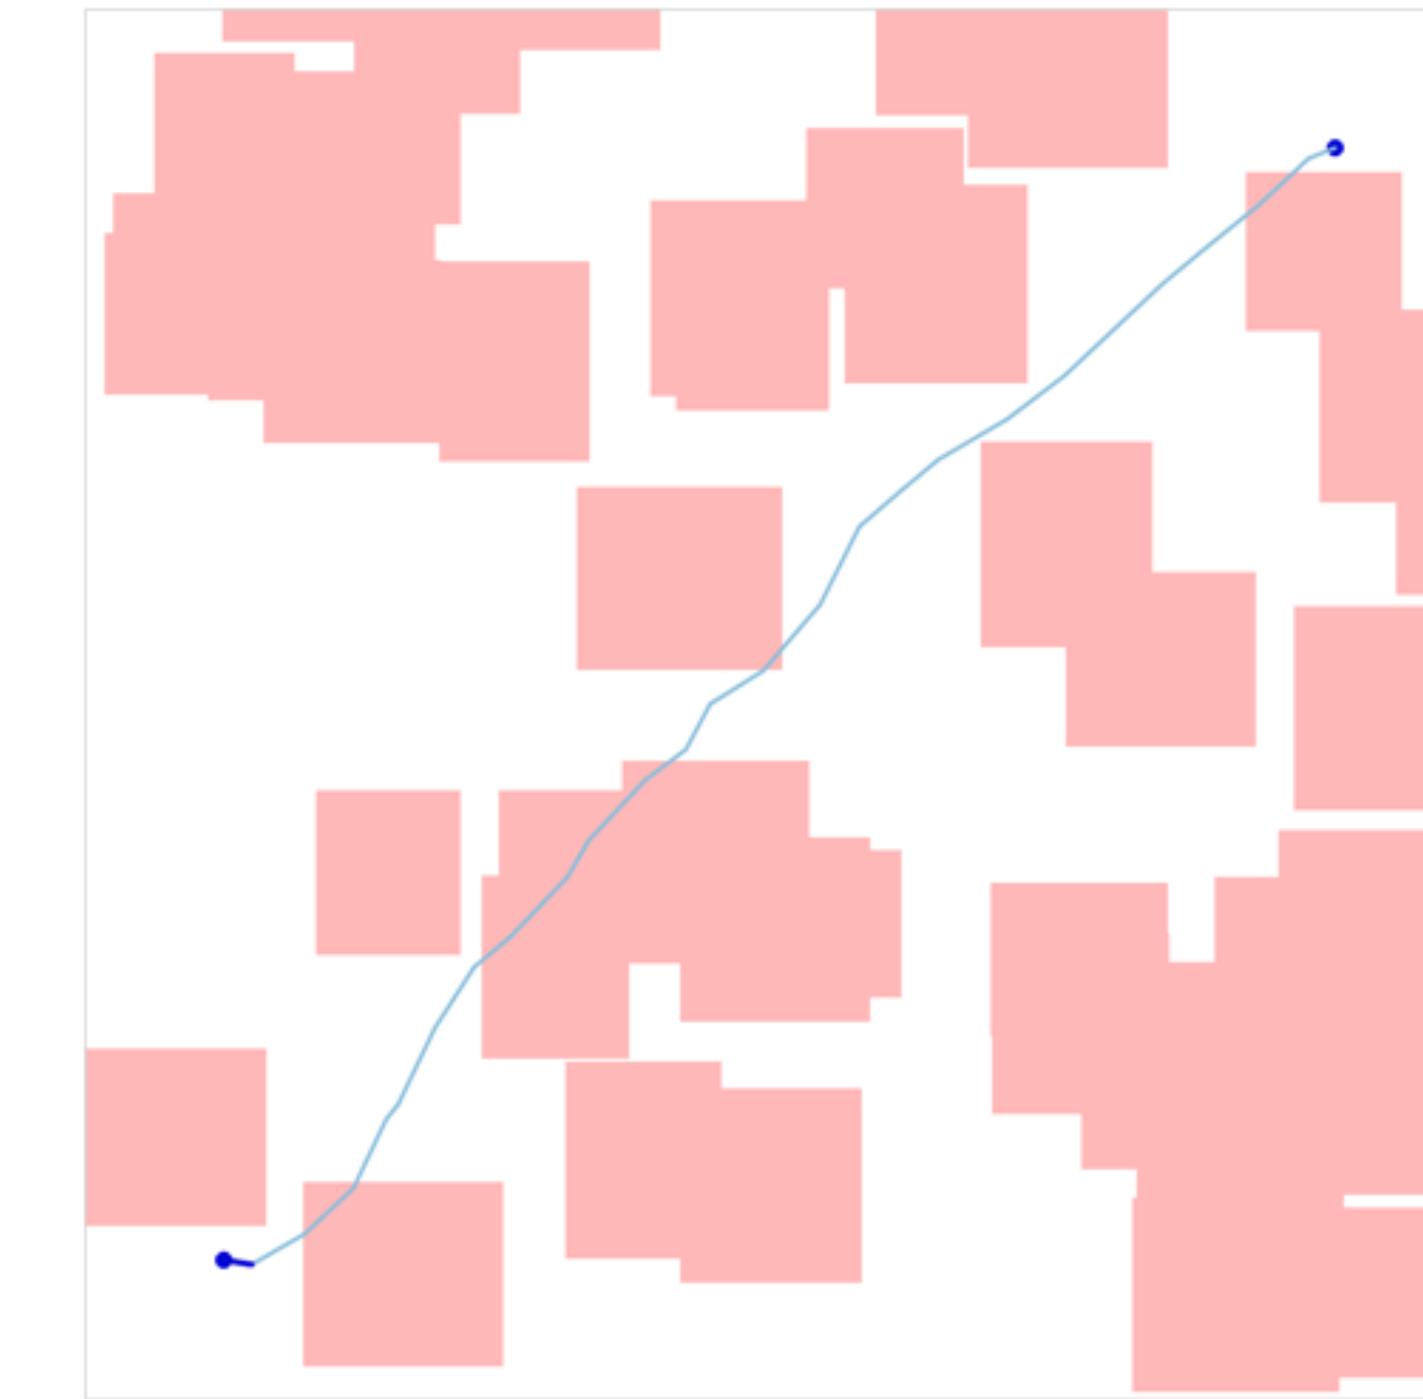
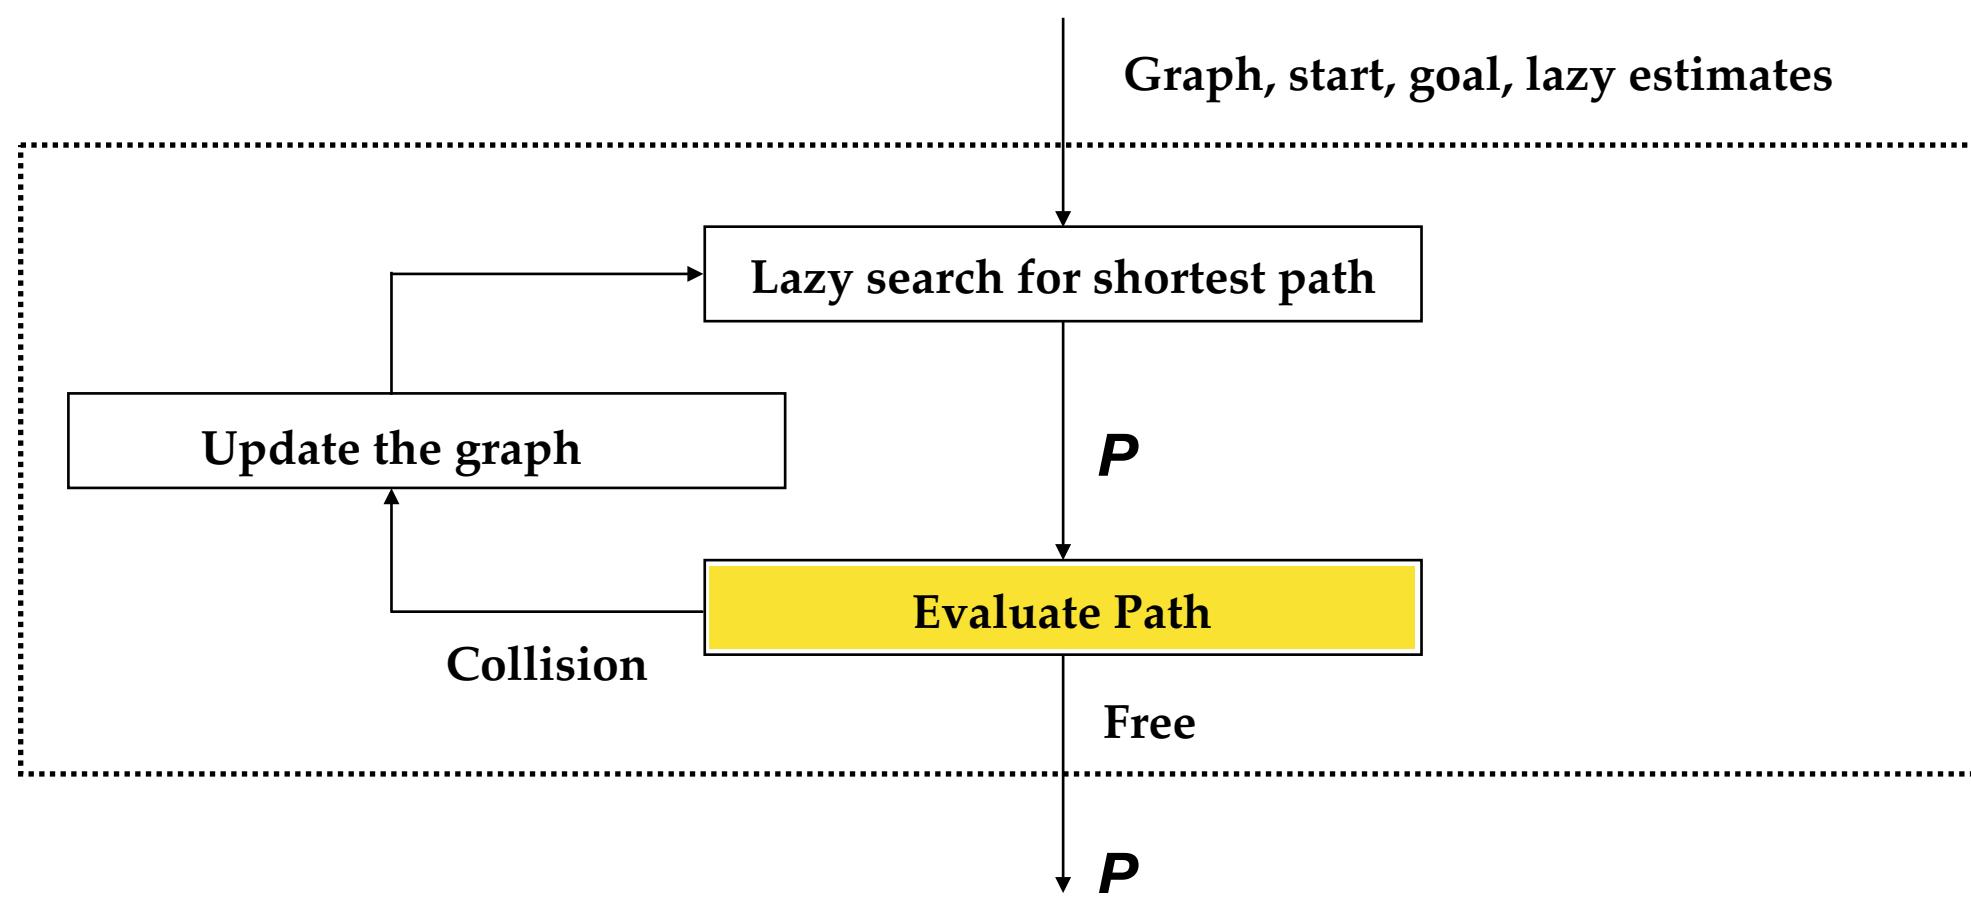
OFU on Steroids!



Send out the Ghost Amoebas

# LazySP

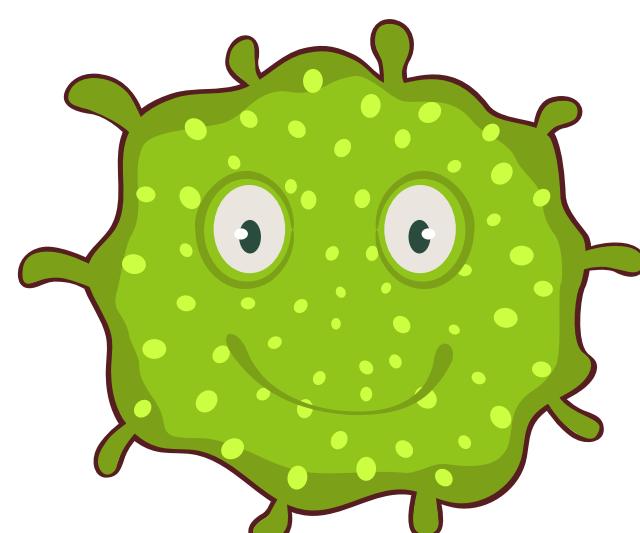
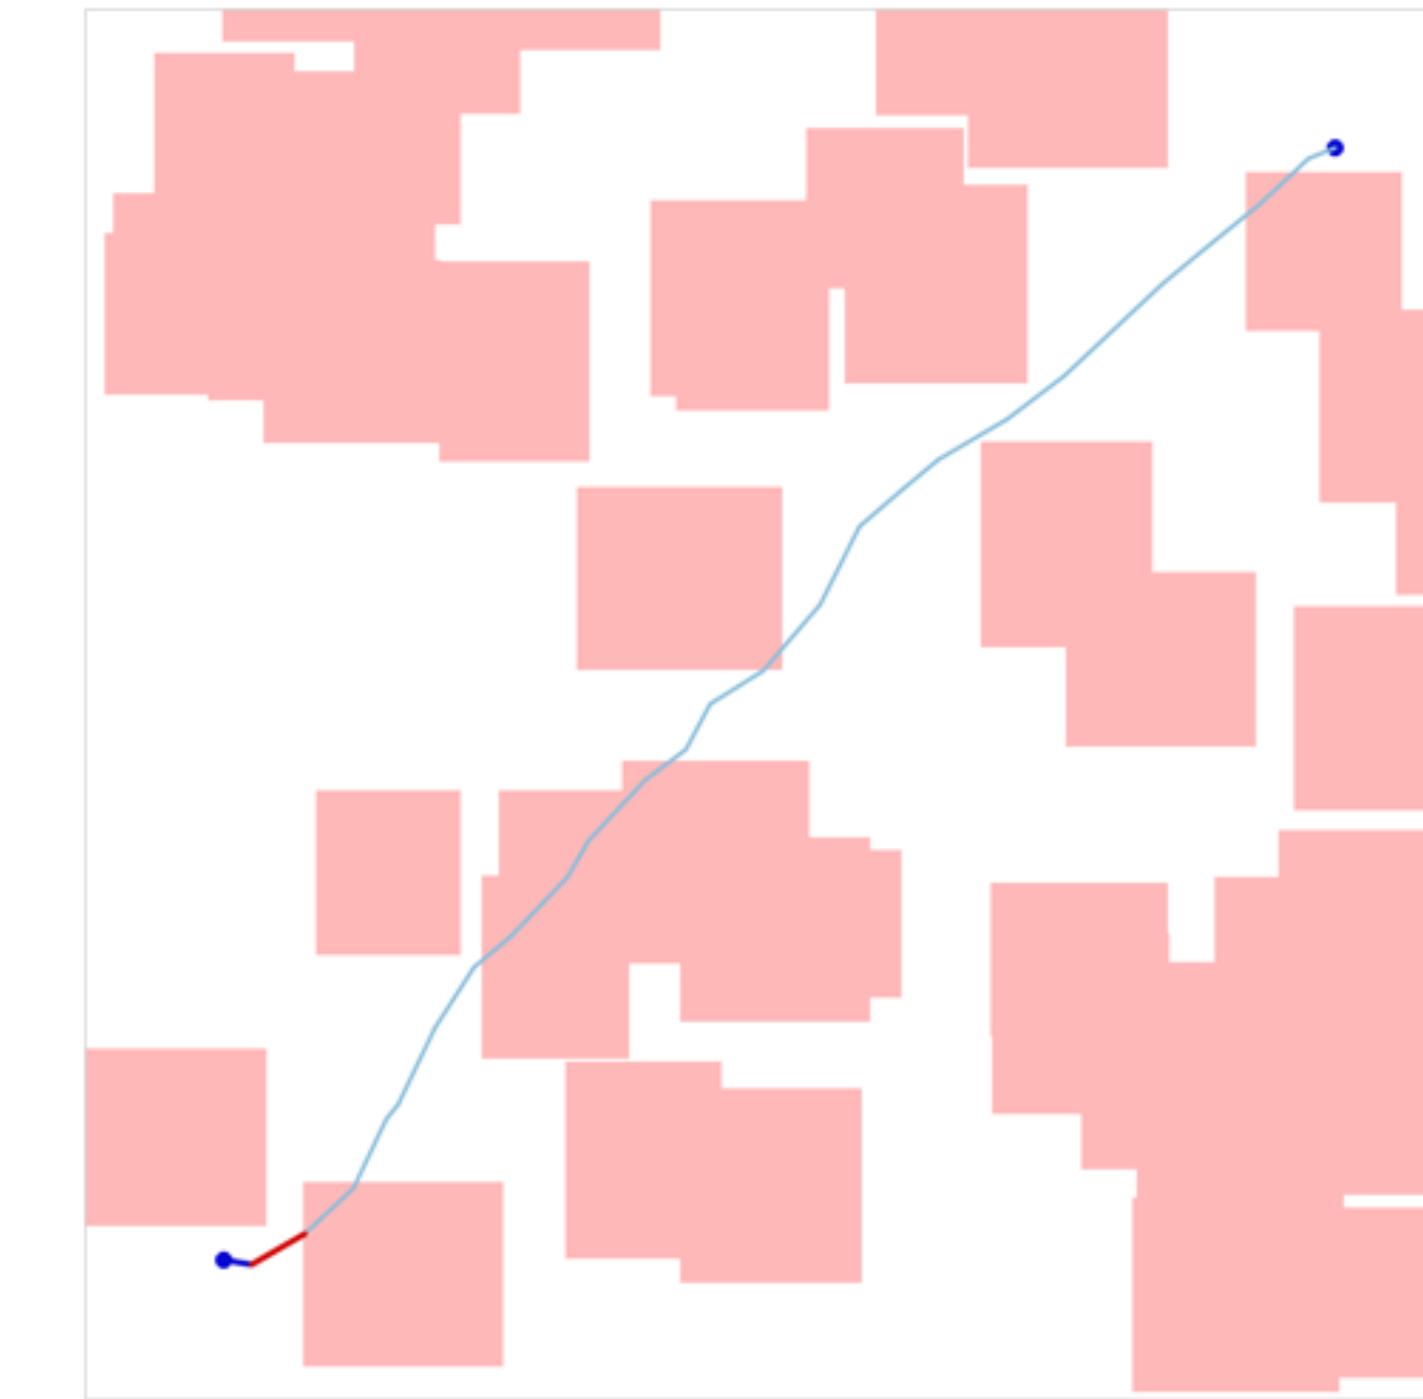
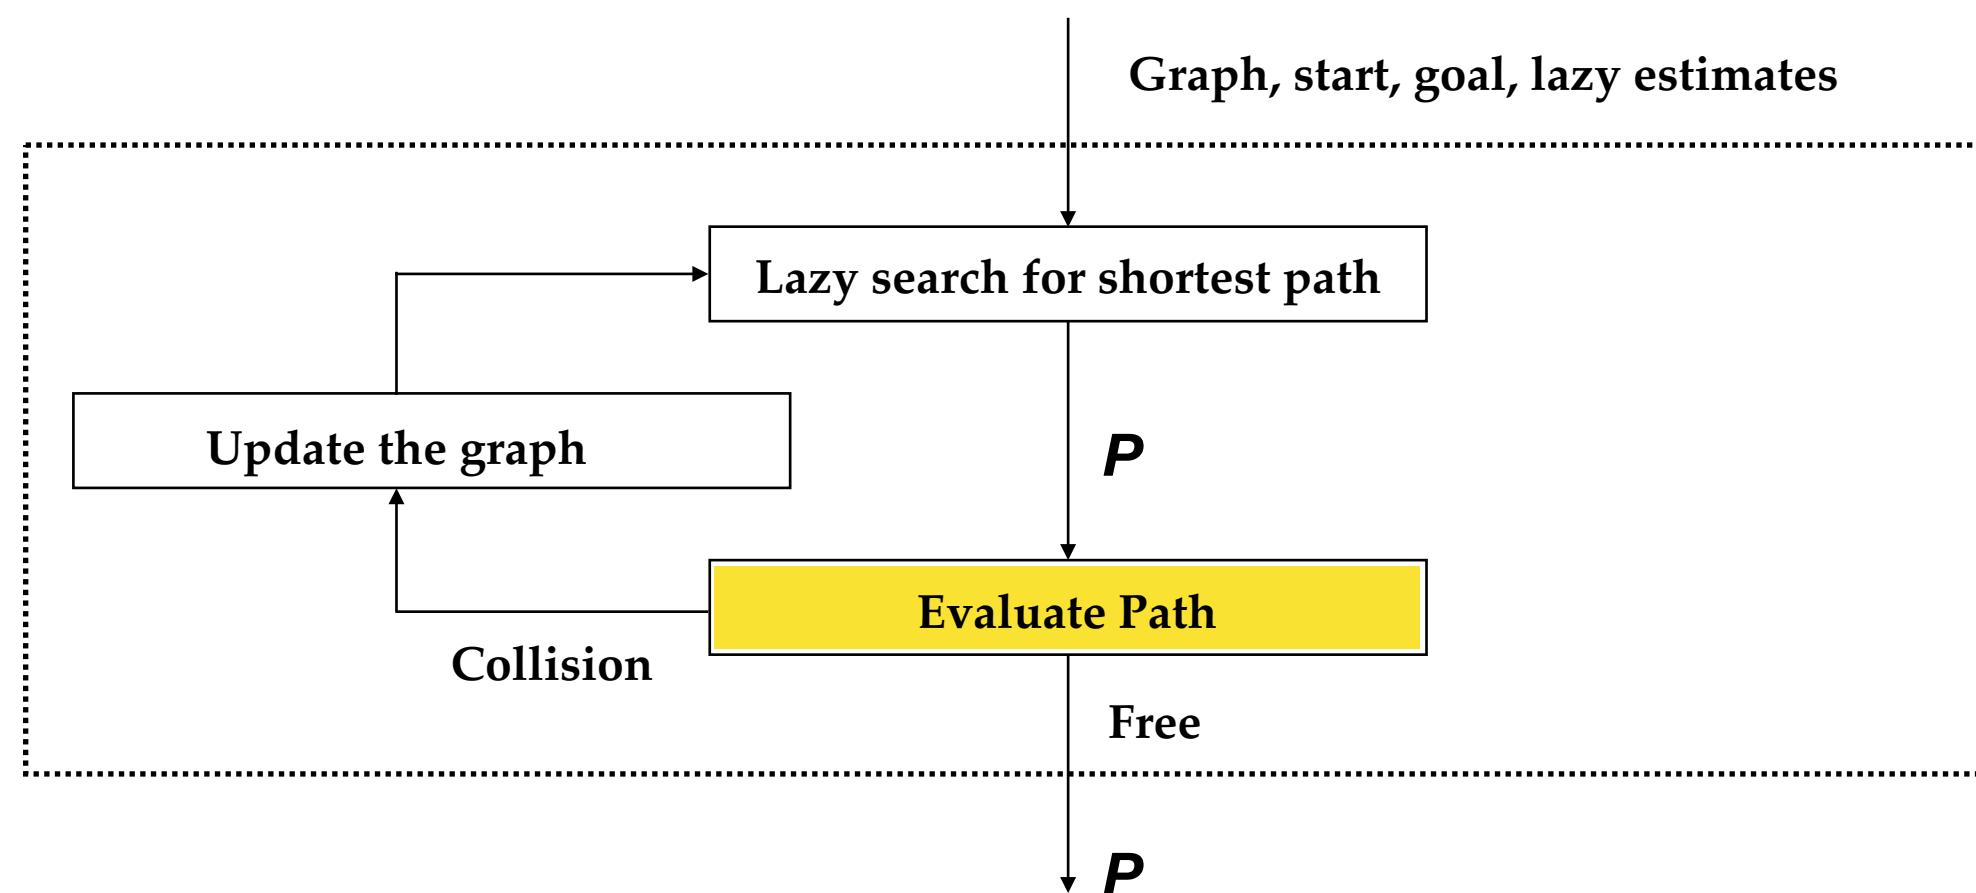
OFU on Steroids!



Only Slime Known Shortest Paths

# LazySP

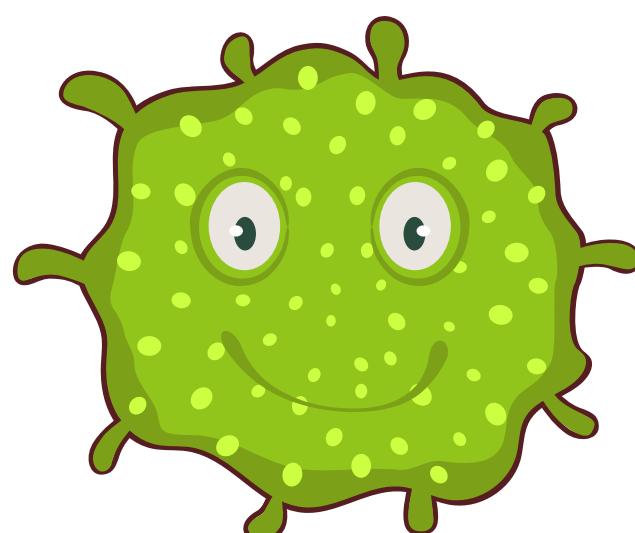
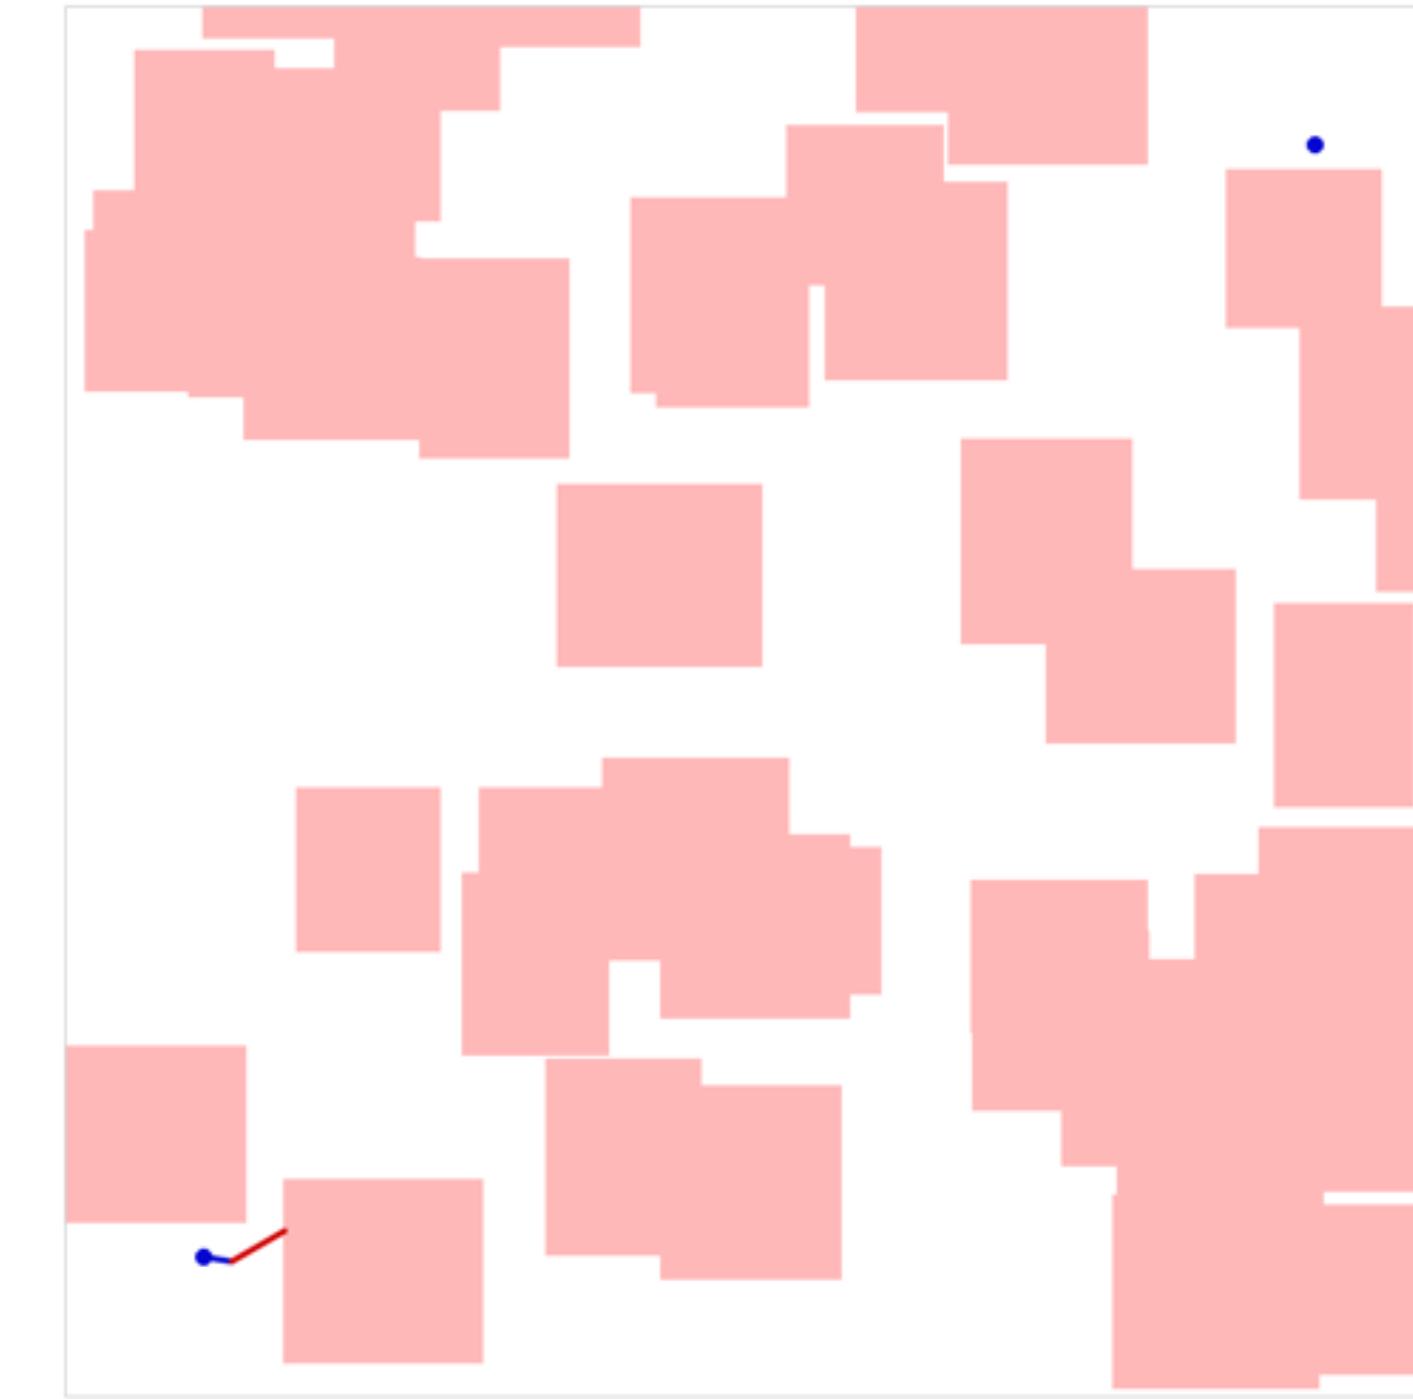
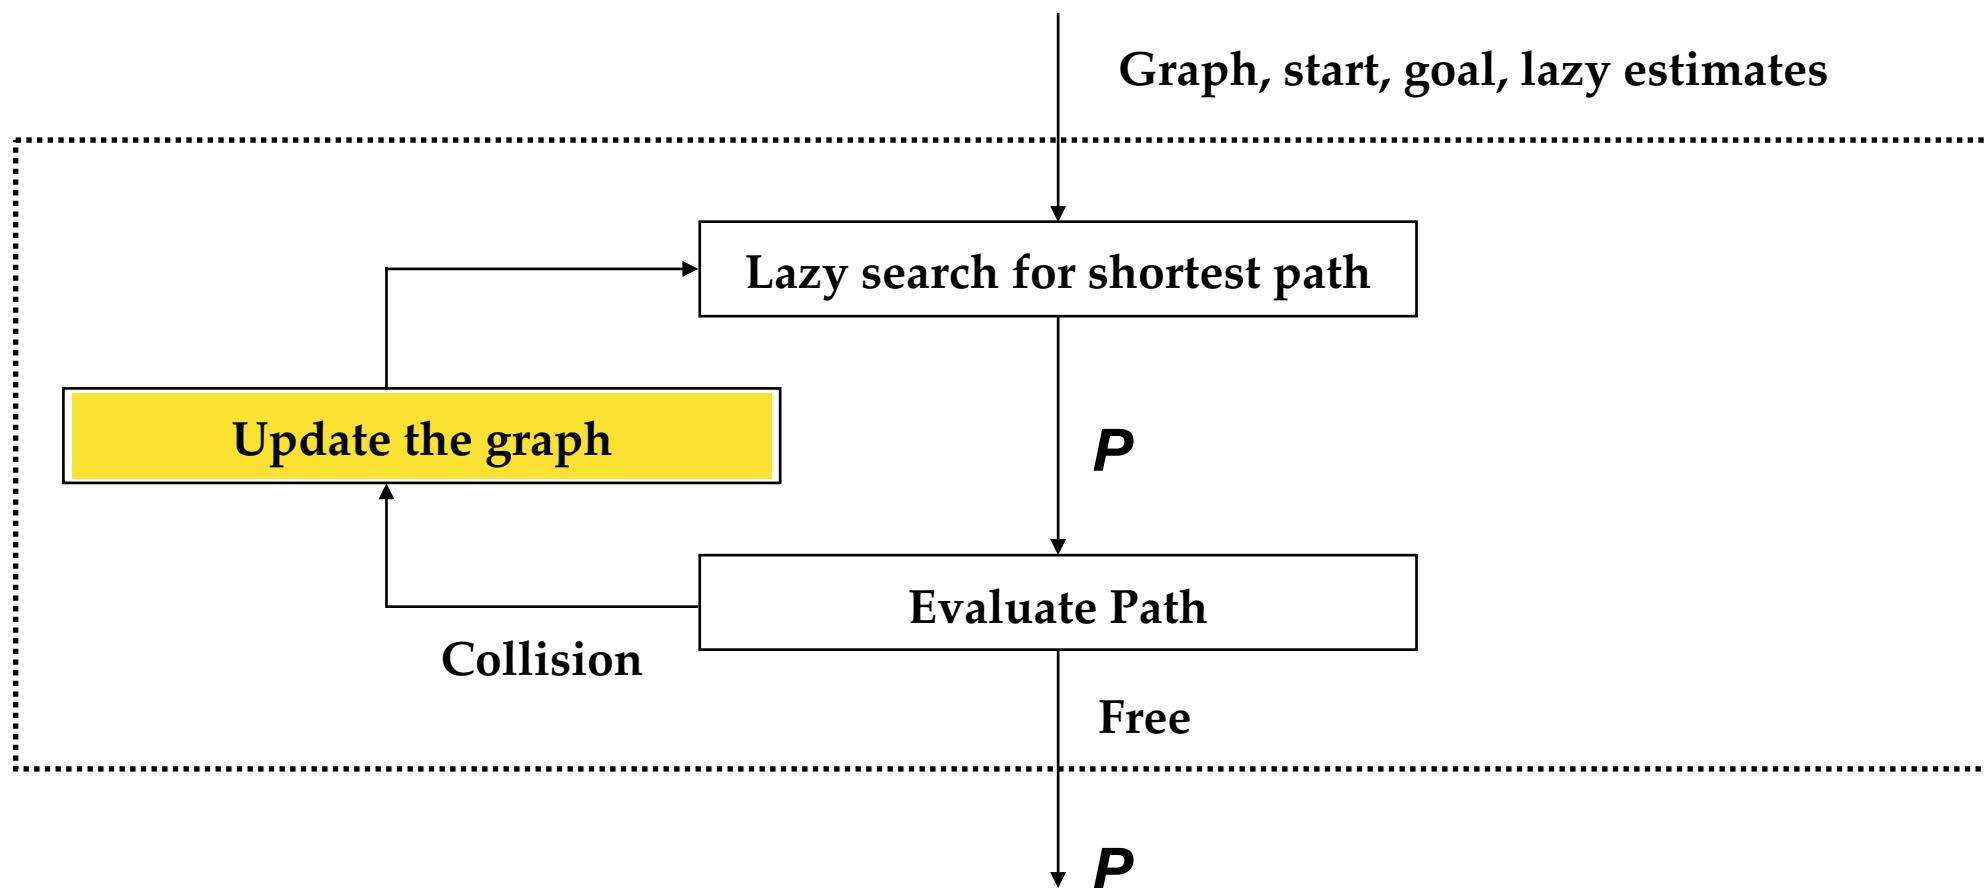
OFU on Steroids!



Only Slime Known Shortest Paths

# LazySP

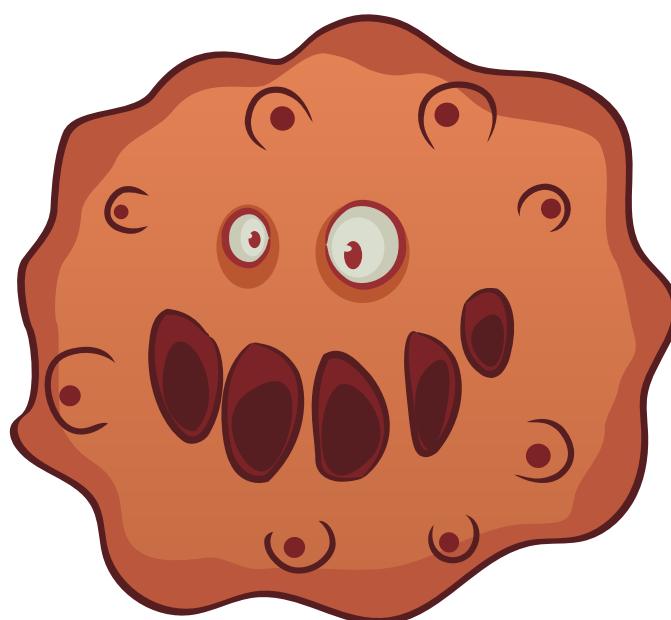
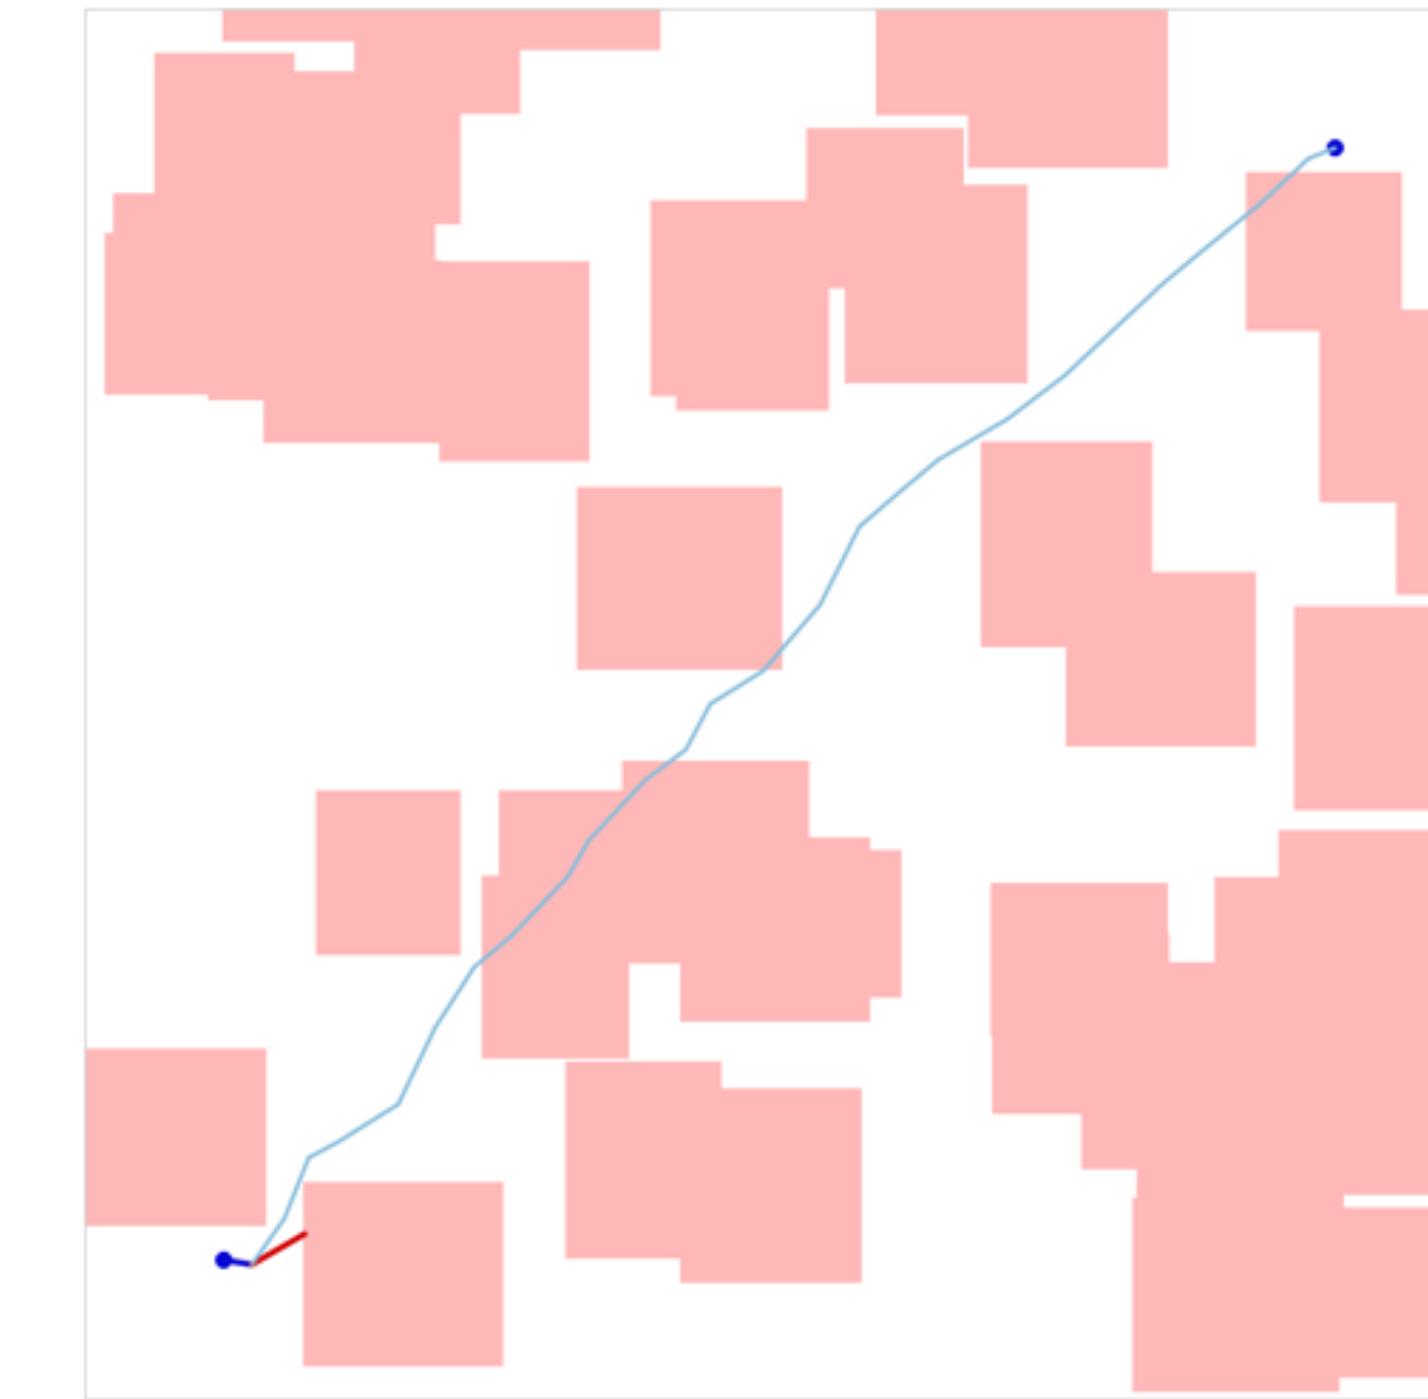
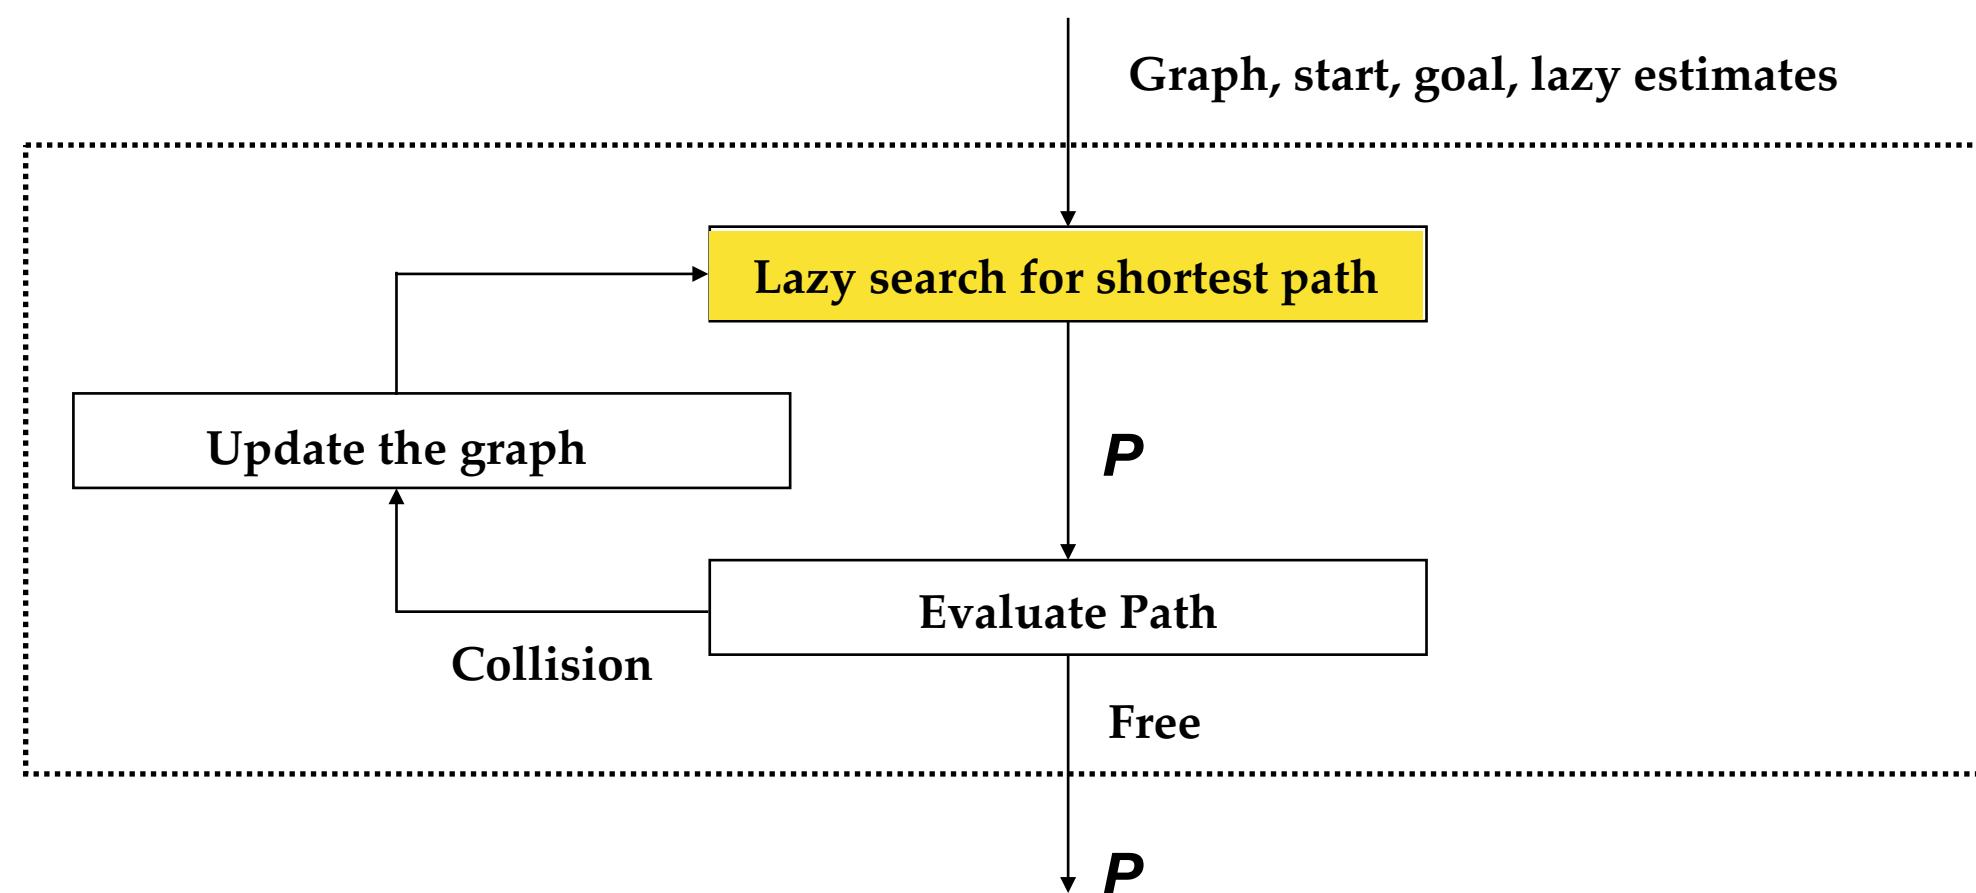
OFU on Steroids!



Only Slime Known Shortest Paths

# LazySP

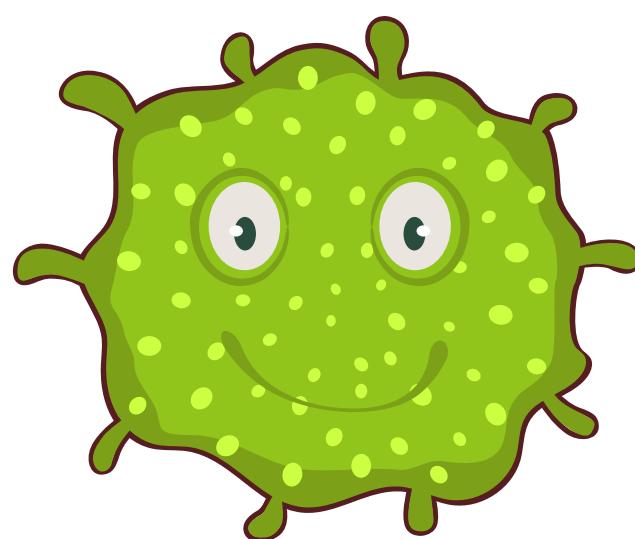
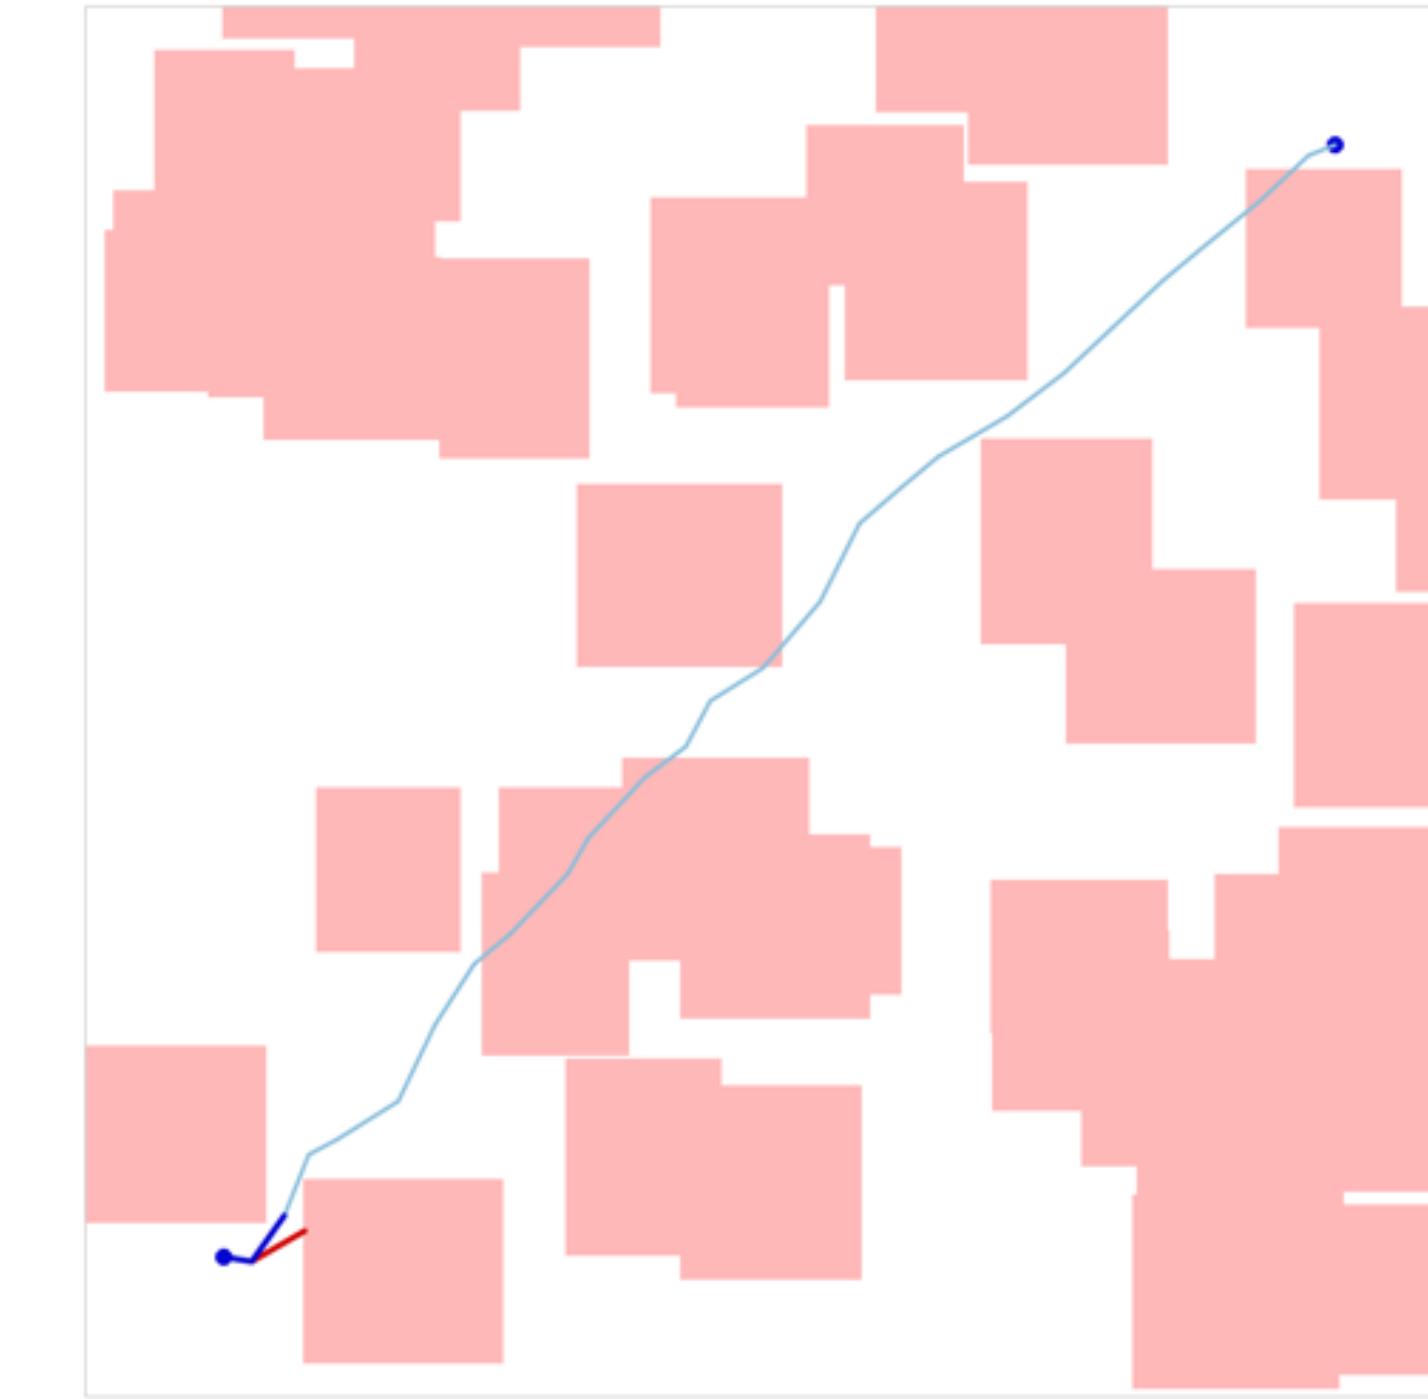
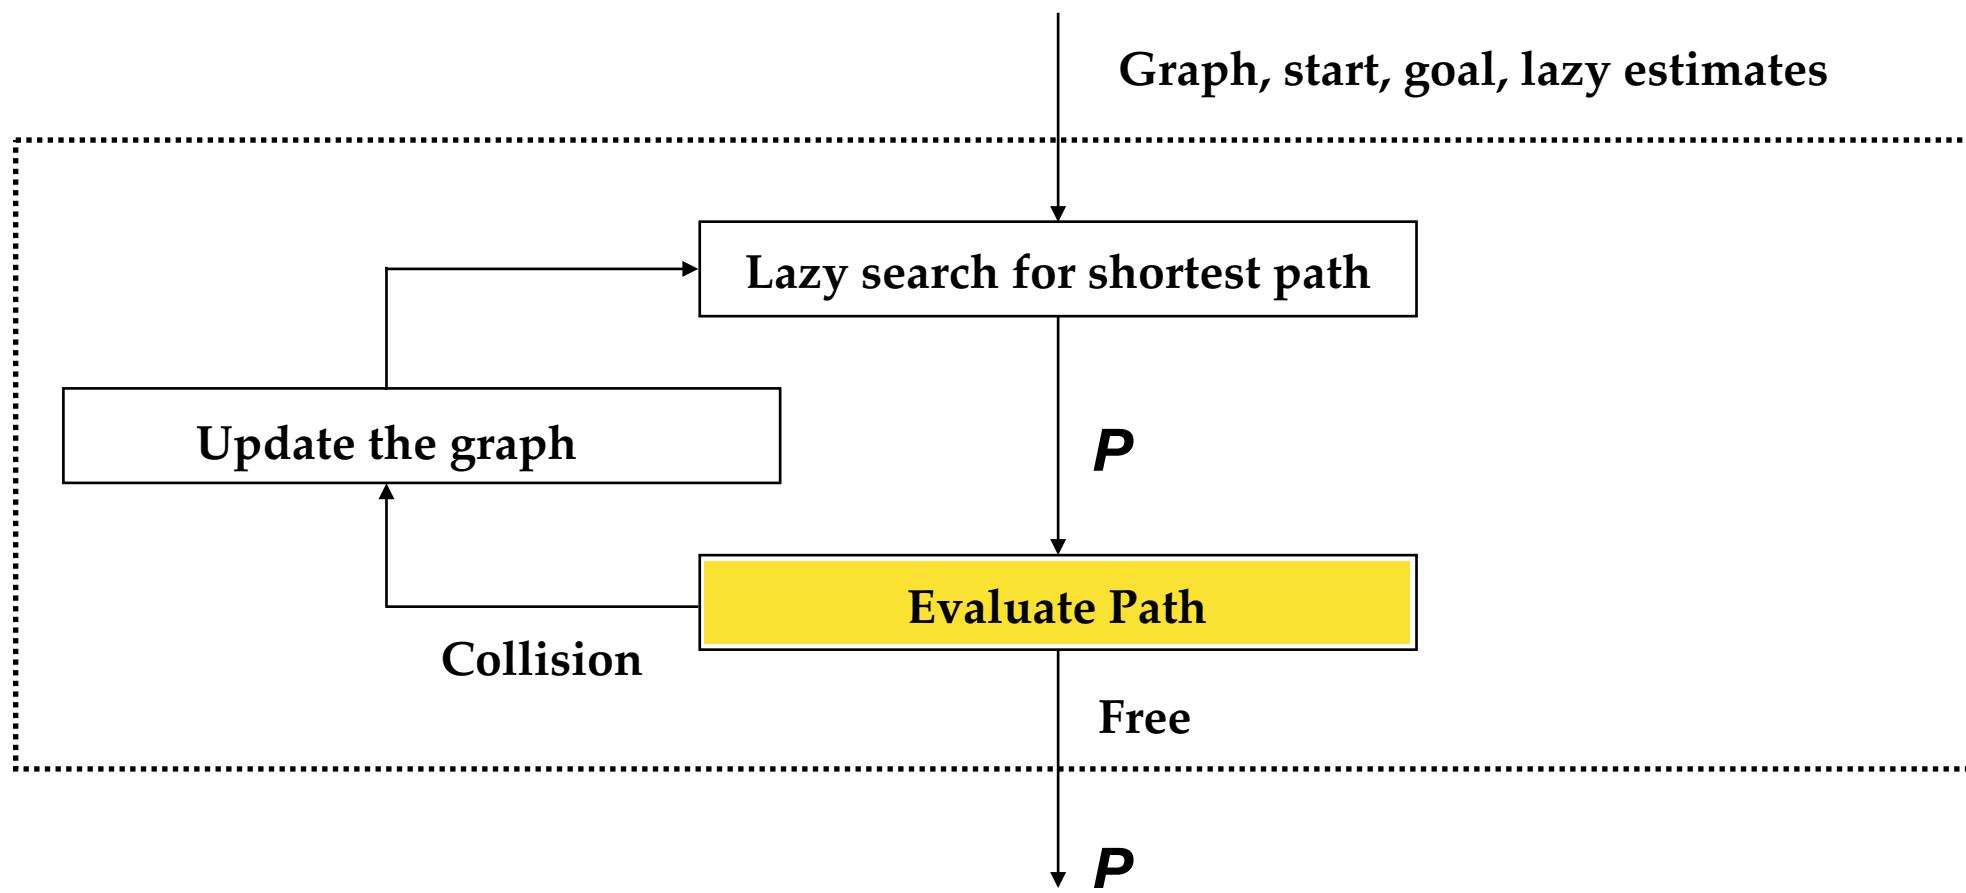
OFU on Steroids!



Send out the Ghost Amoebas

# LazySP

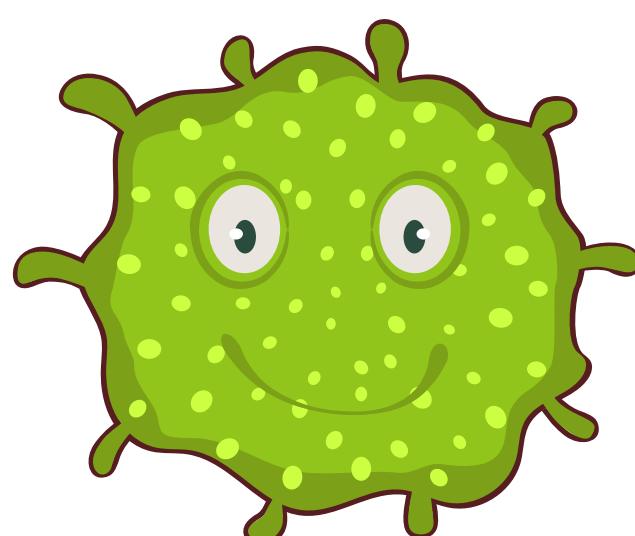
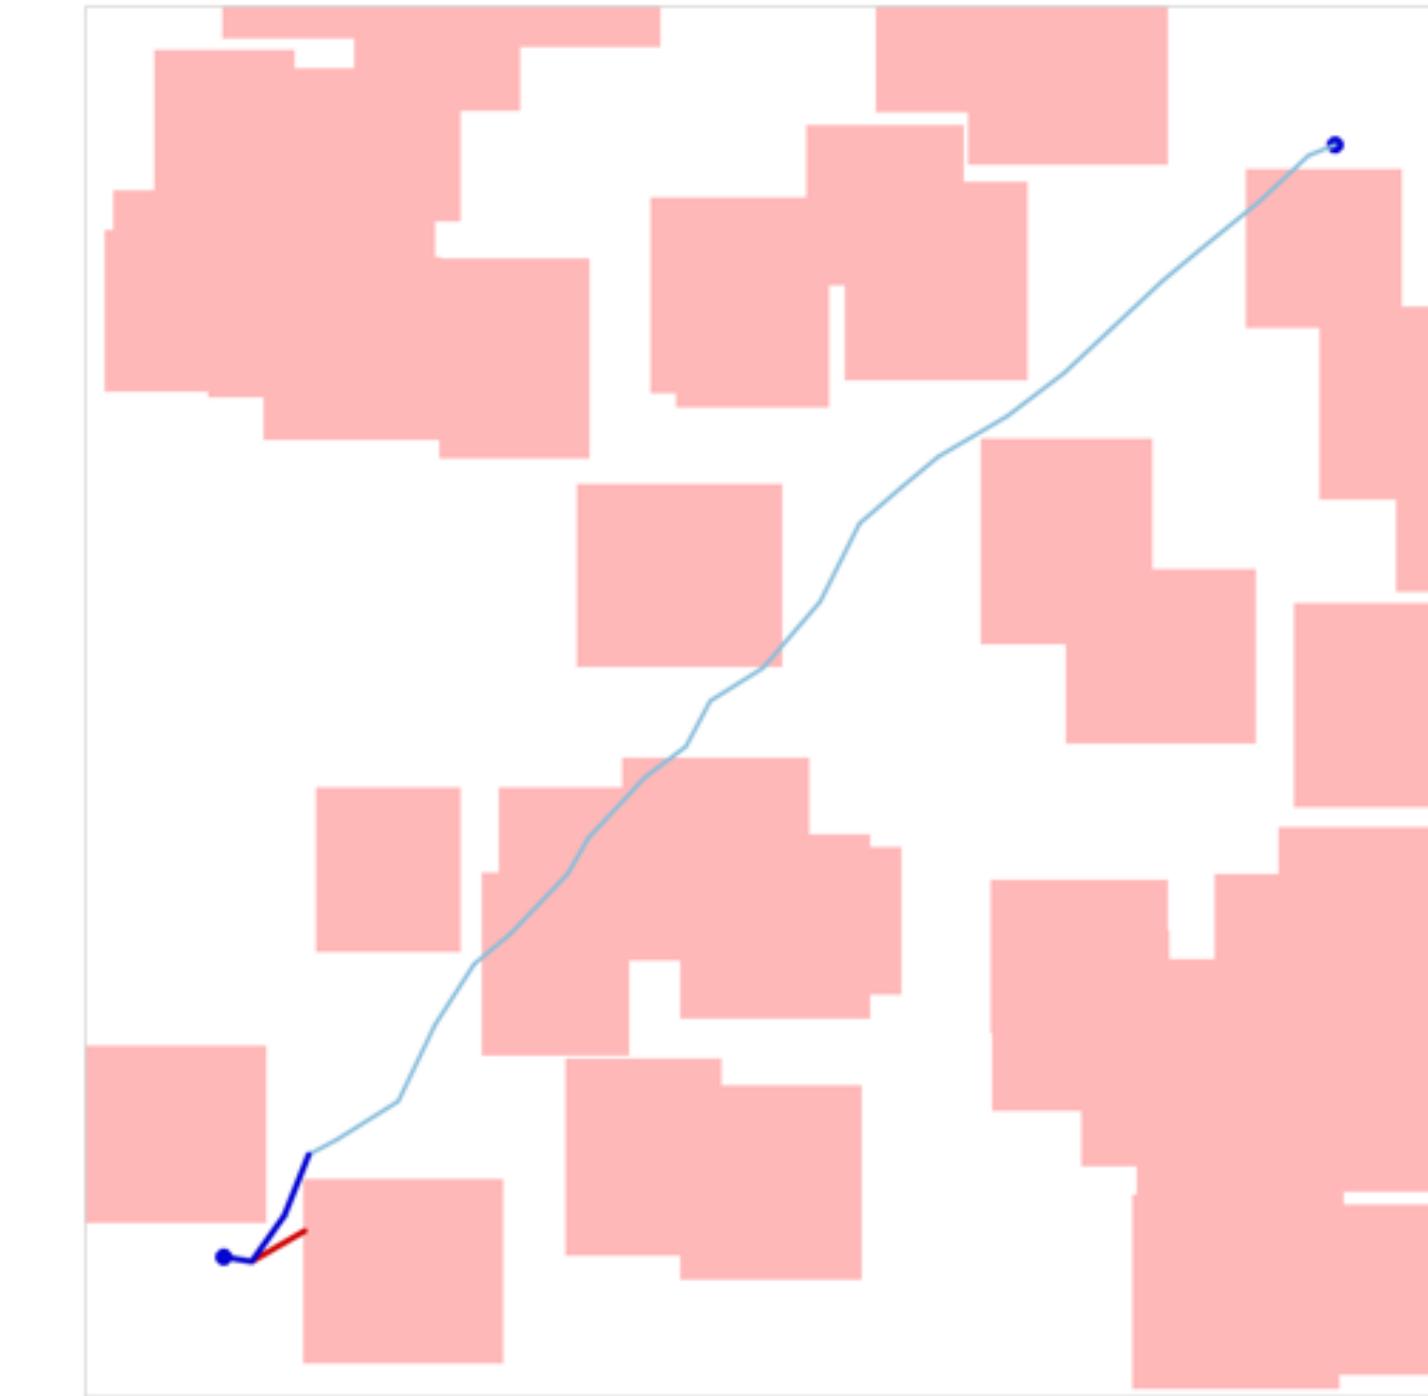
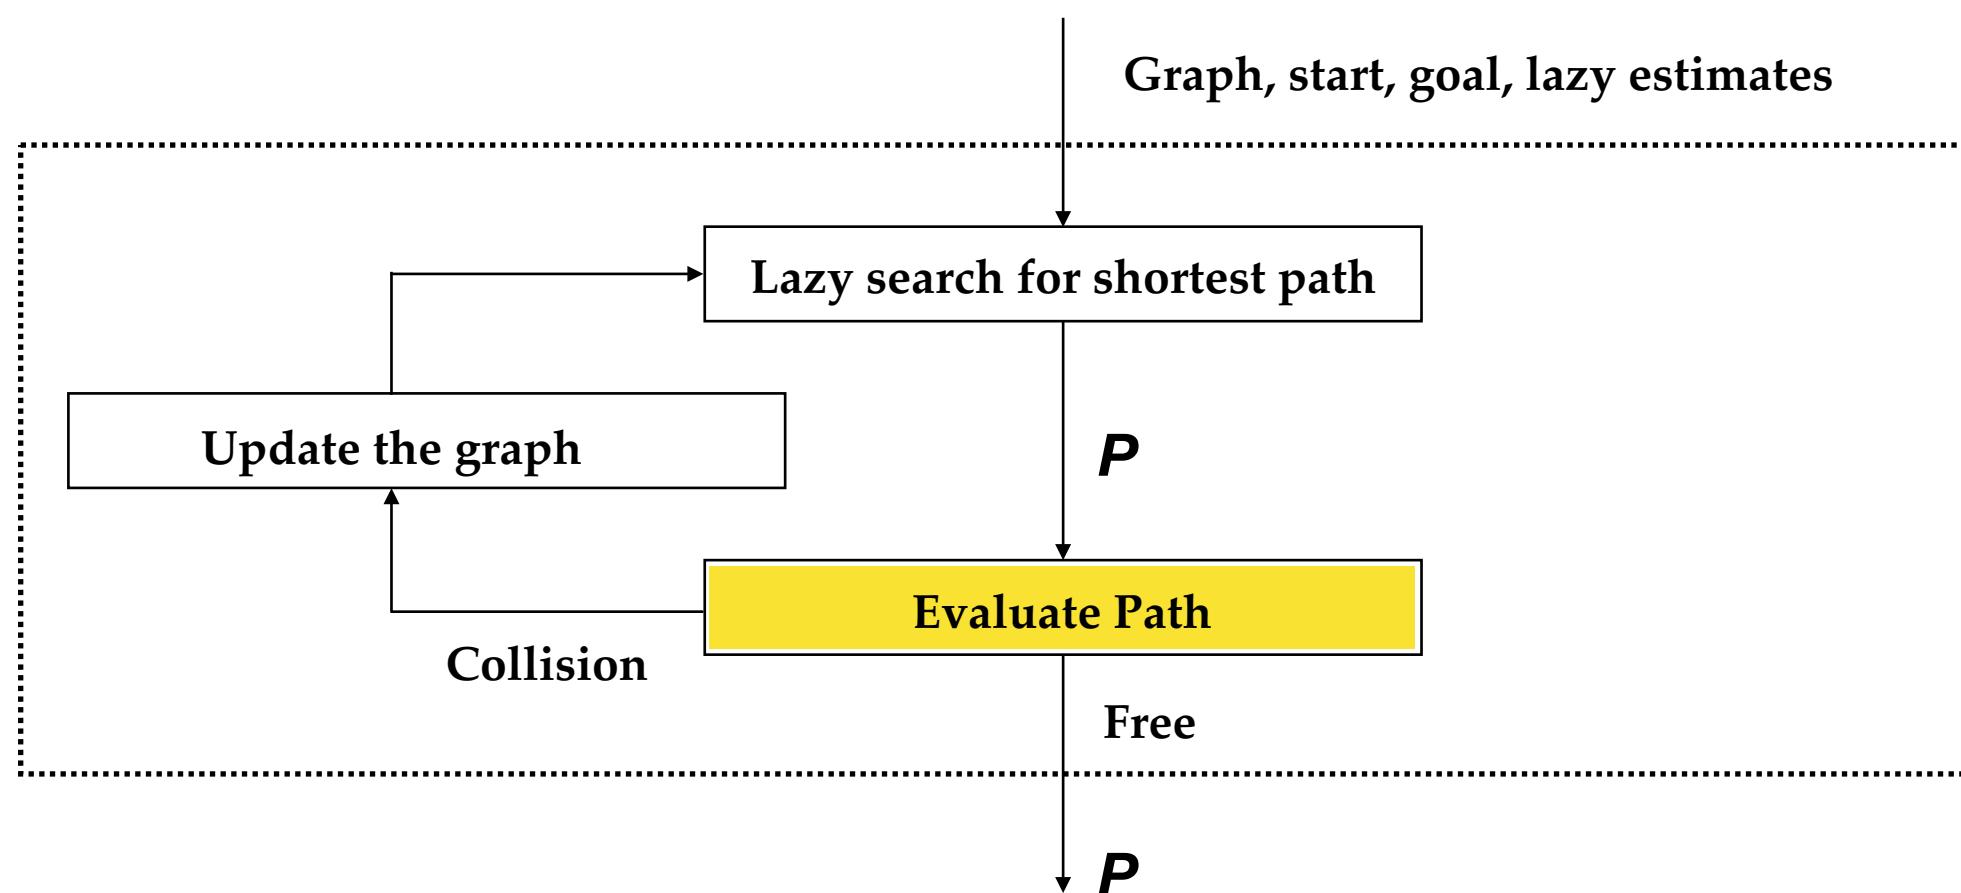
OFU on Steroids!



Only Slime Known Shortest Paths

# LazySP

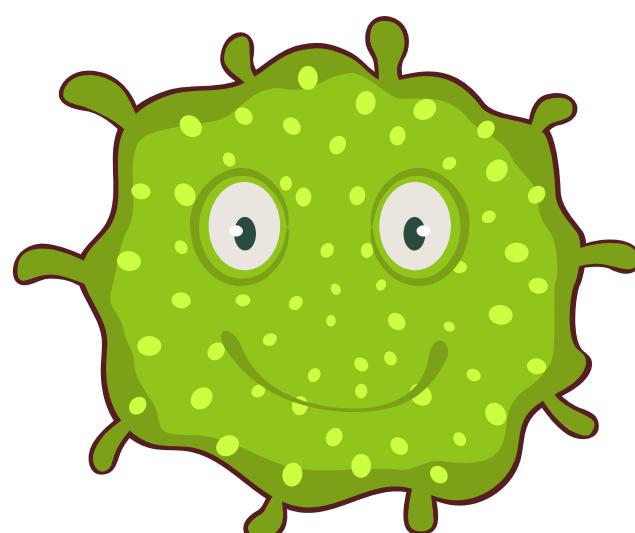
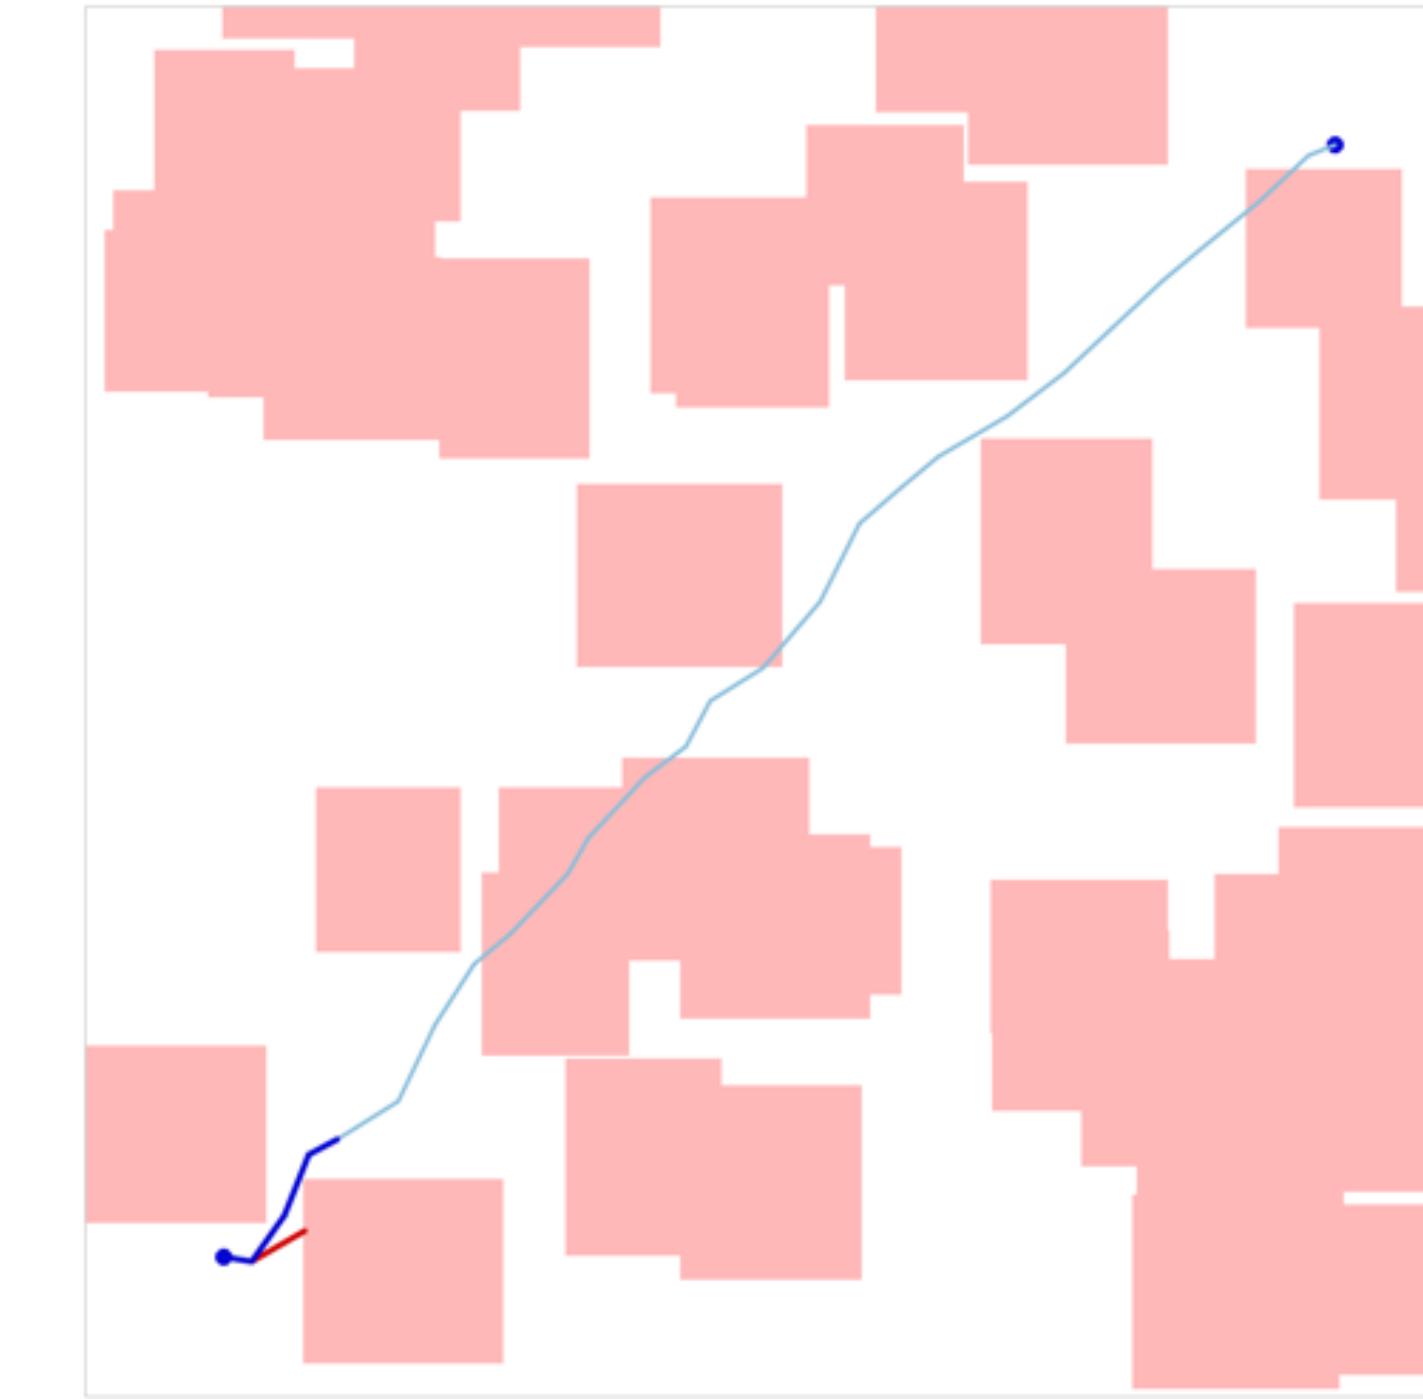
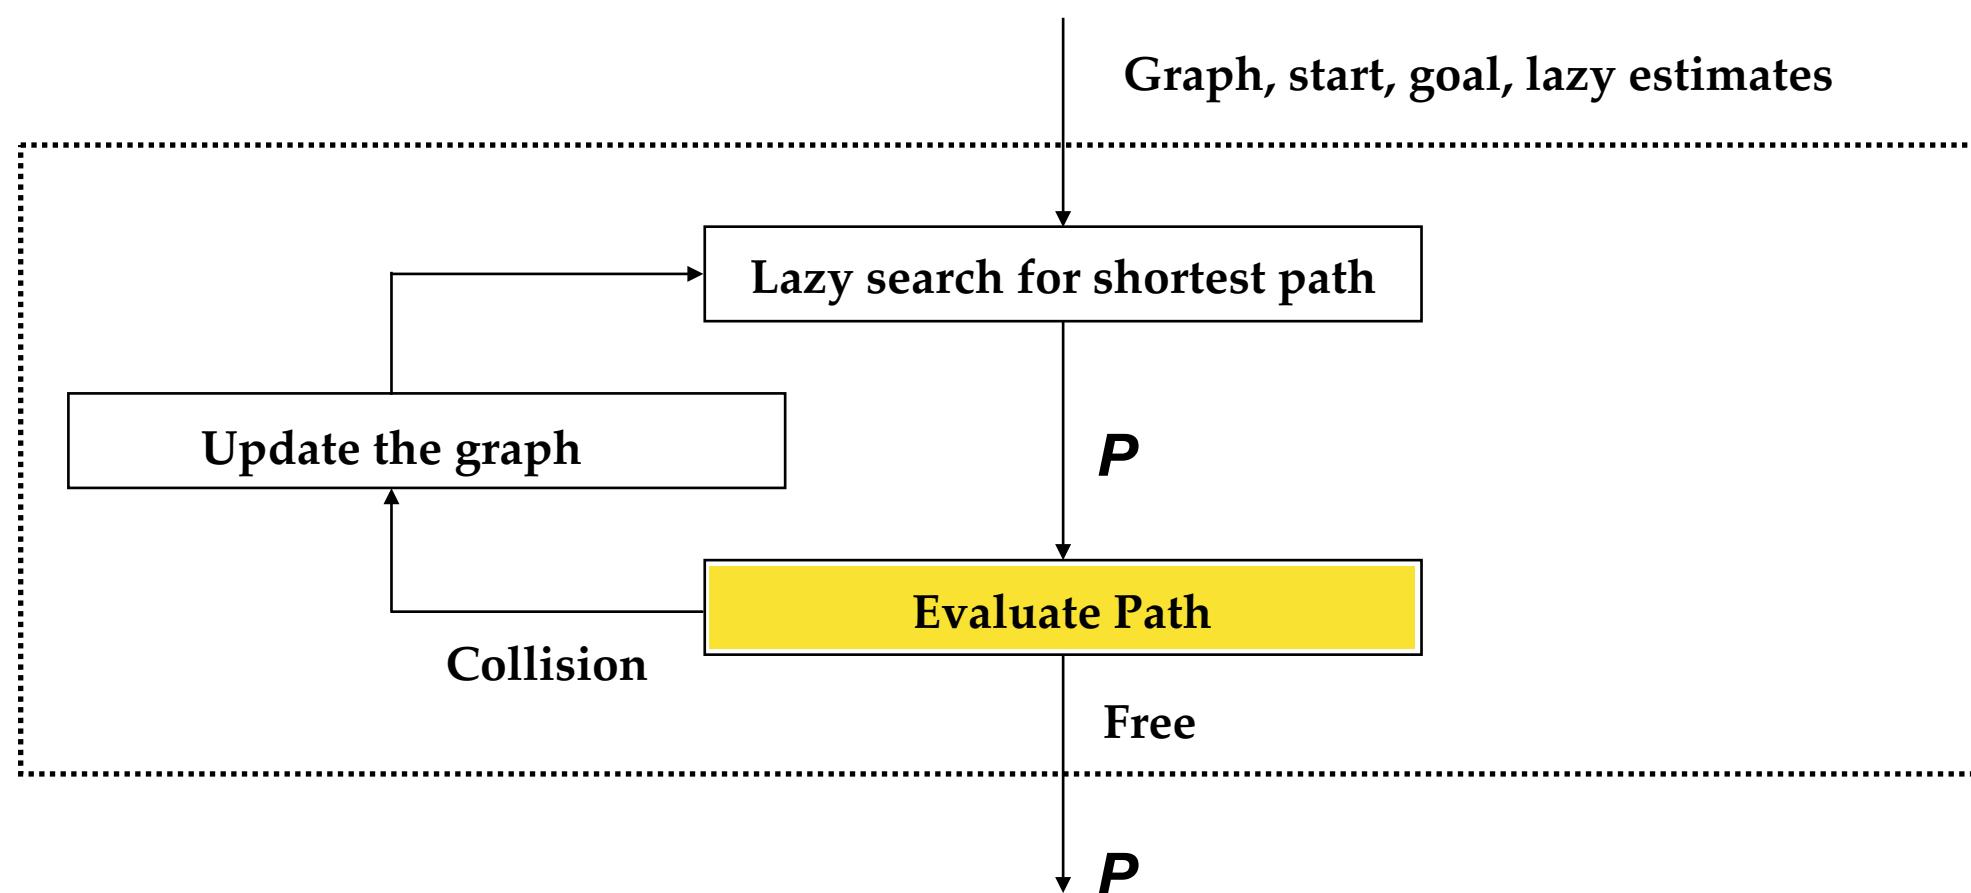
OFU on Steroids!



Only Slime Known Shortest Paths

# LazySP

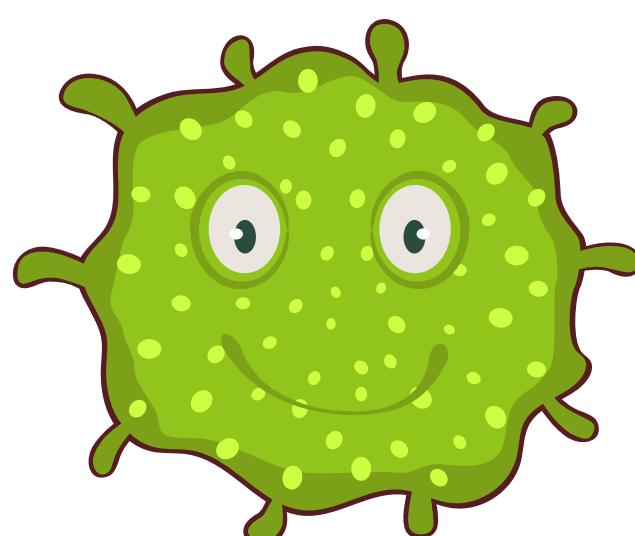
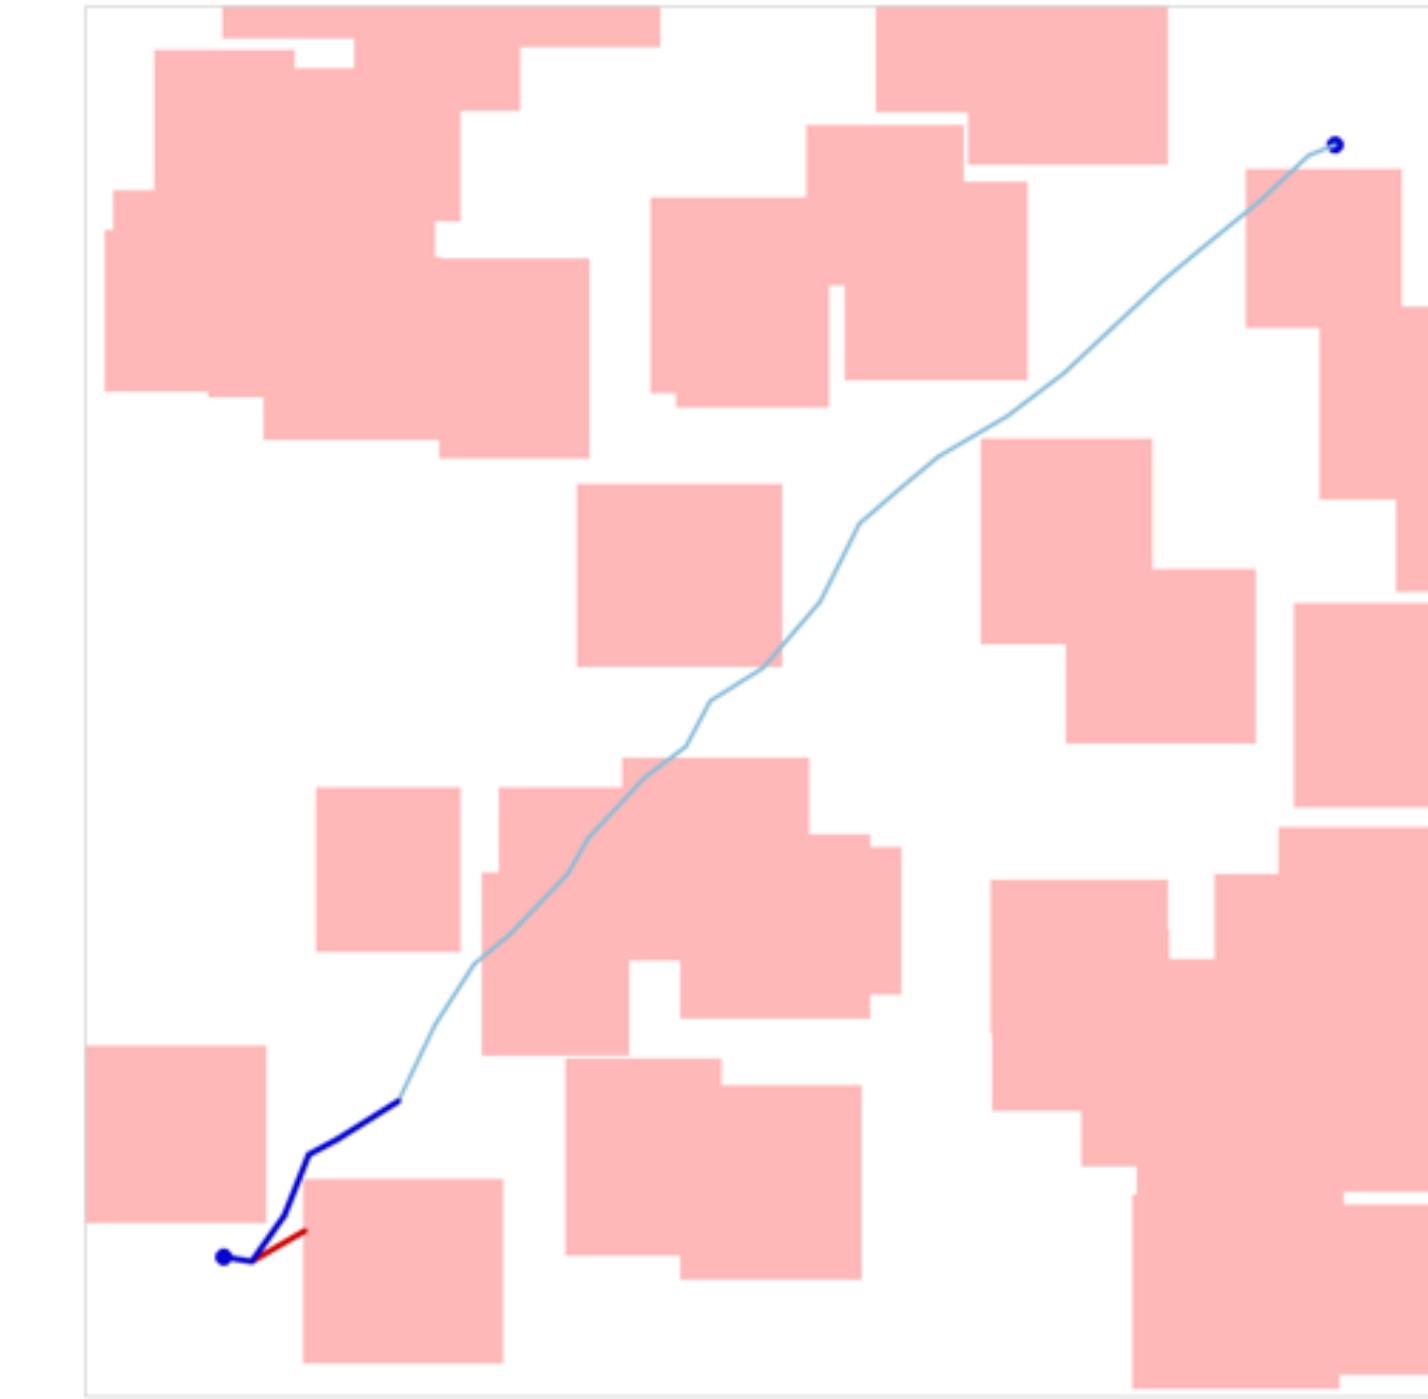
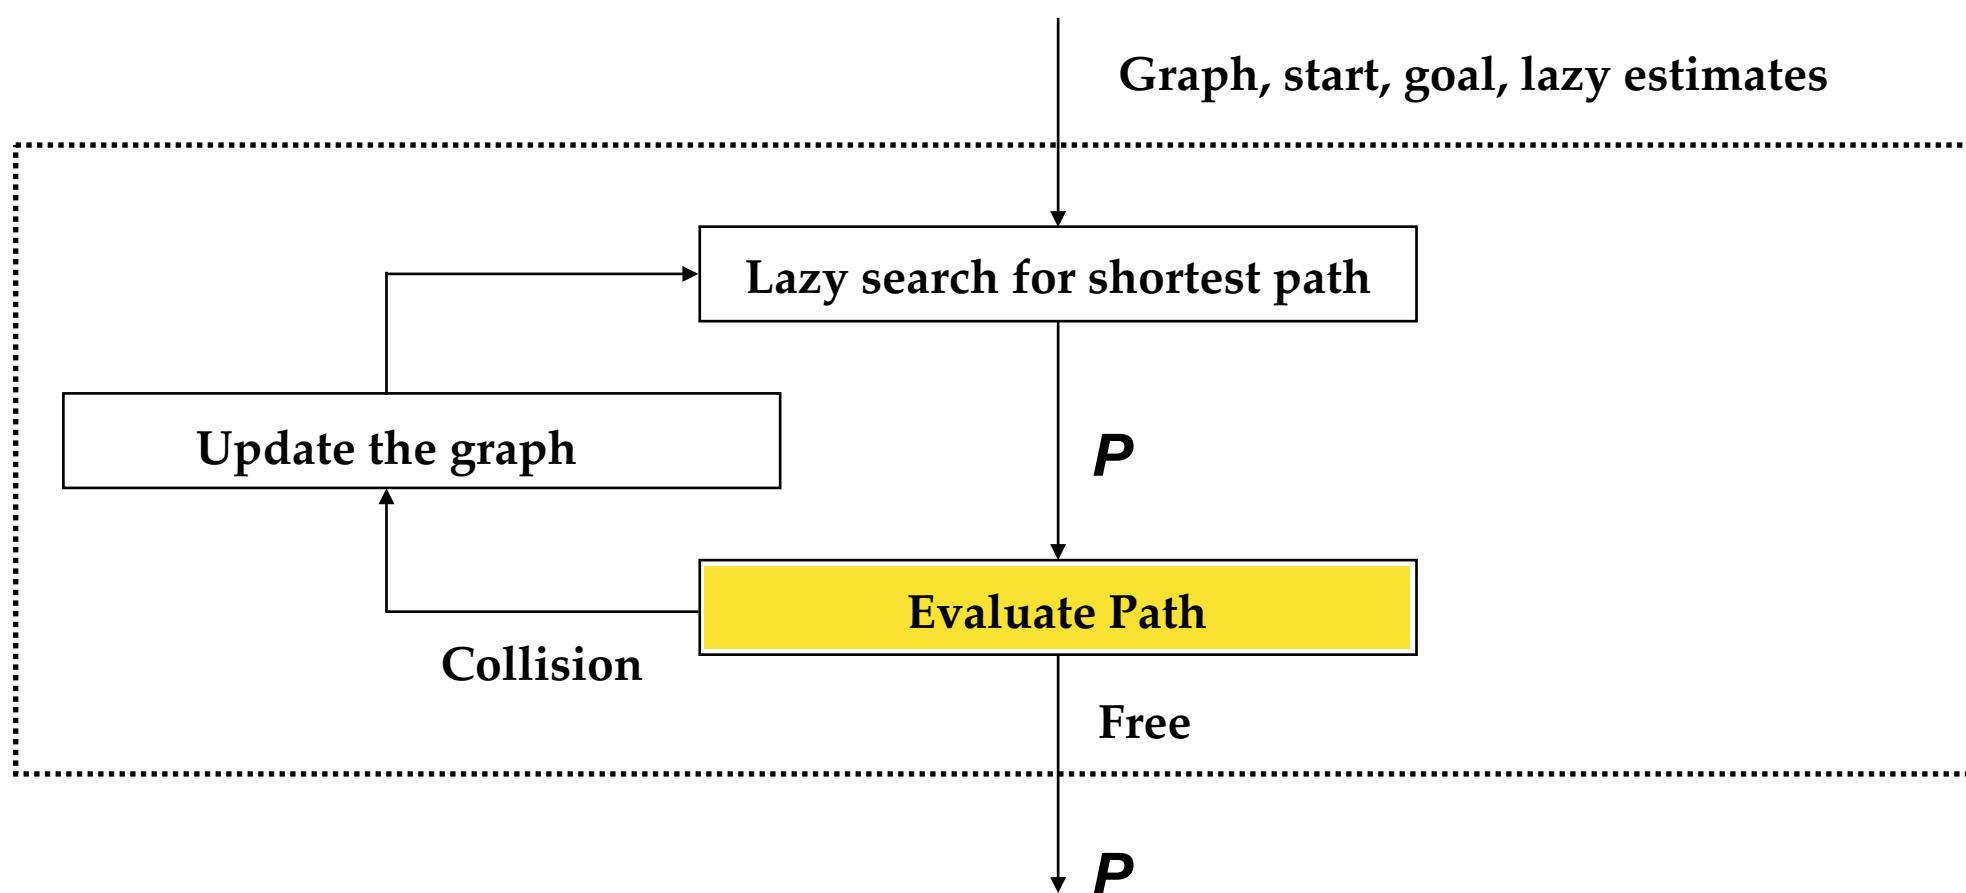
OFU on Steroids!



Only Slime Known Shortest Paths

# LazySP

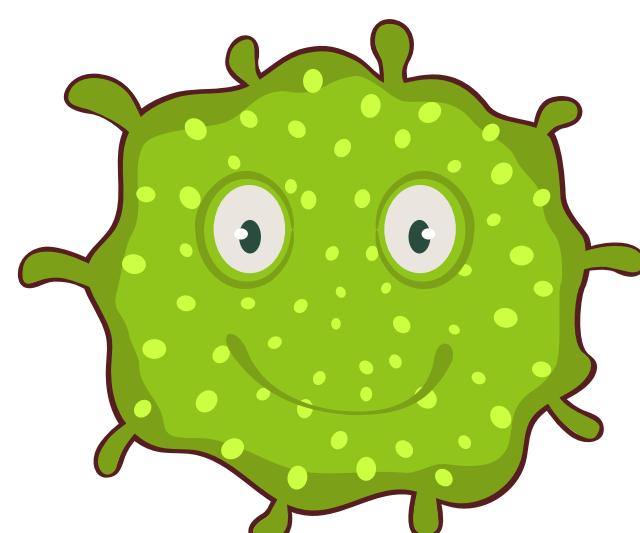
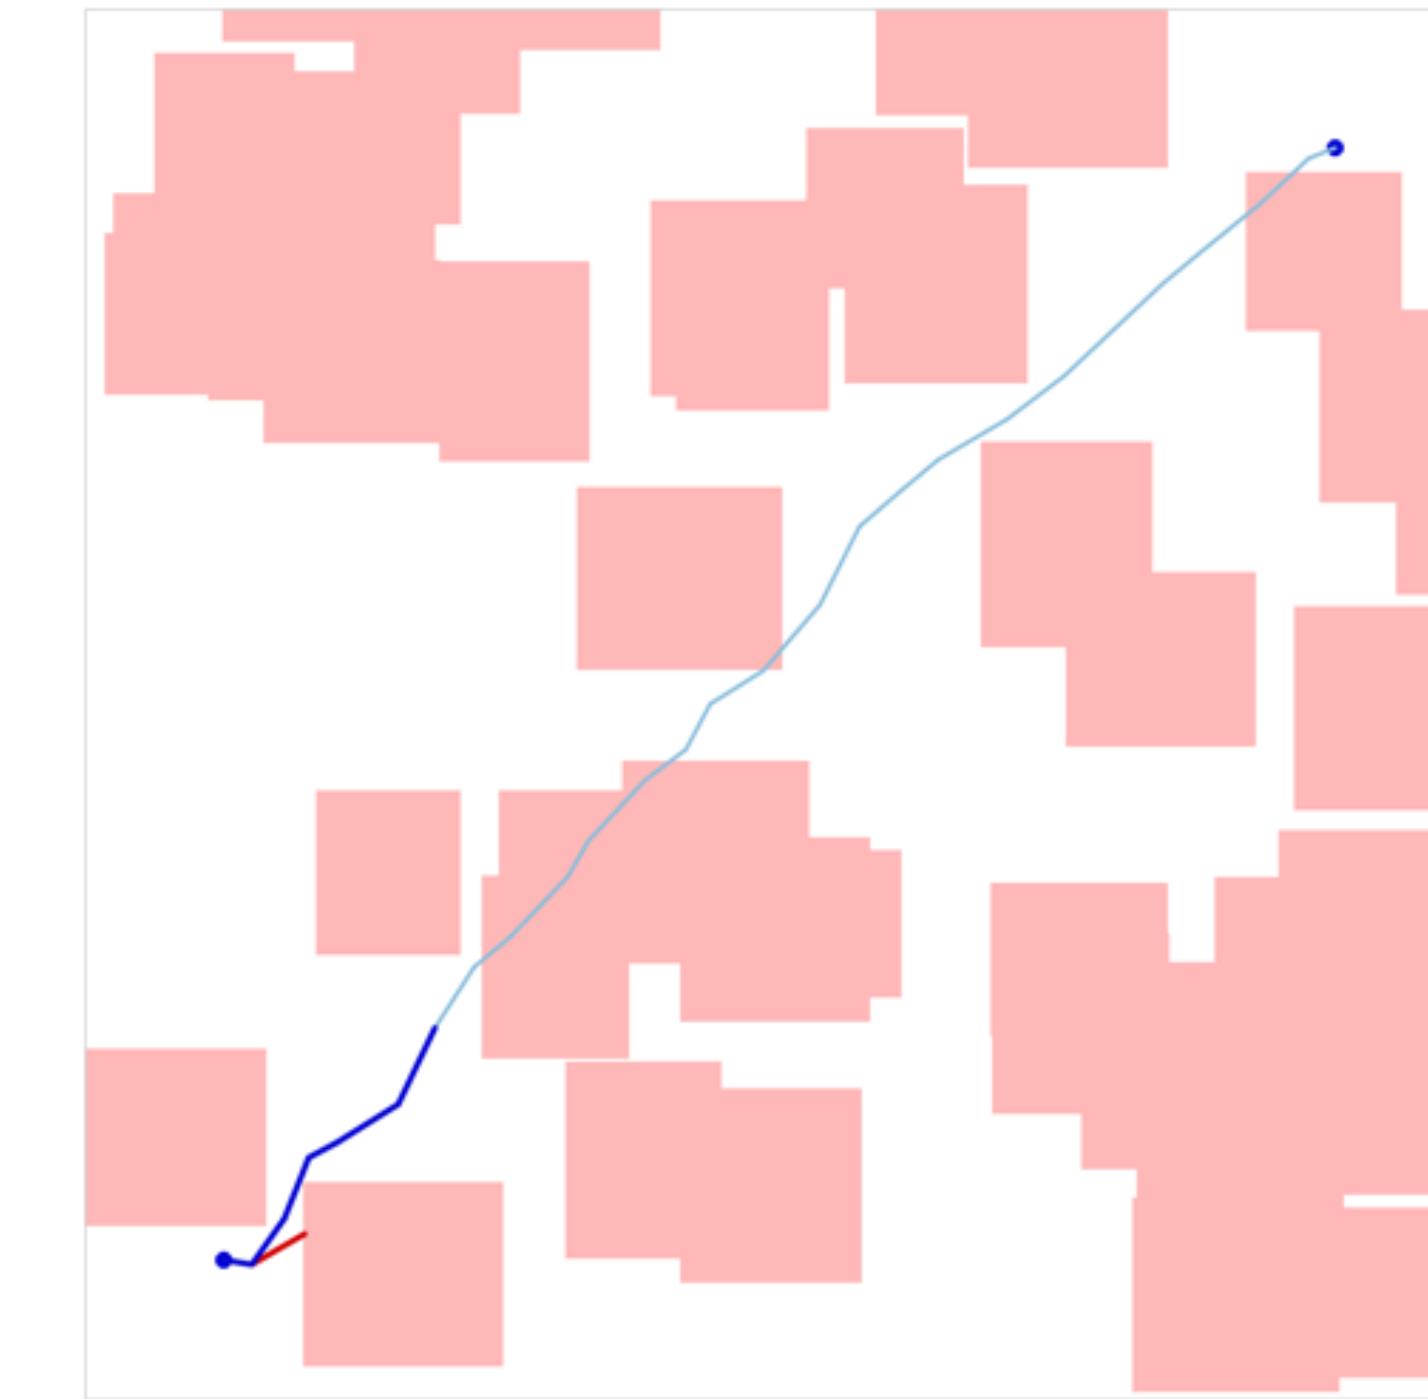
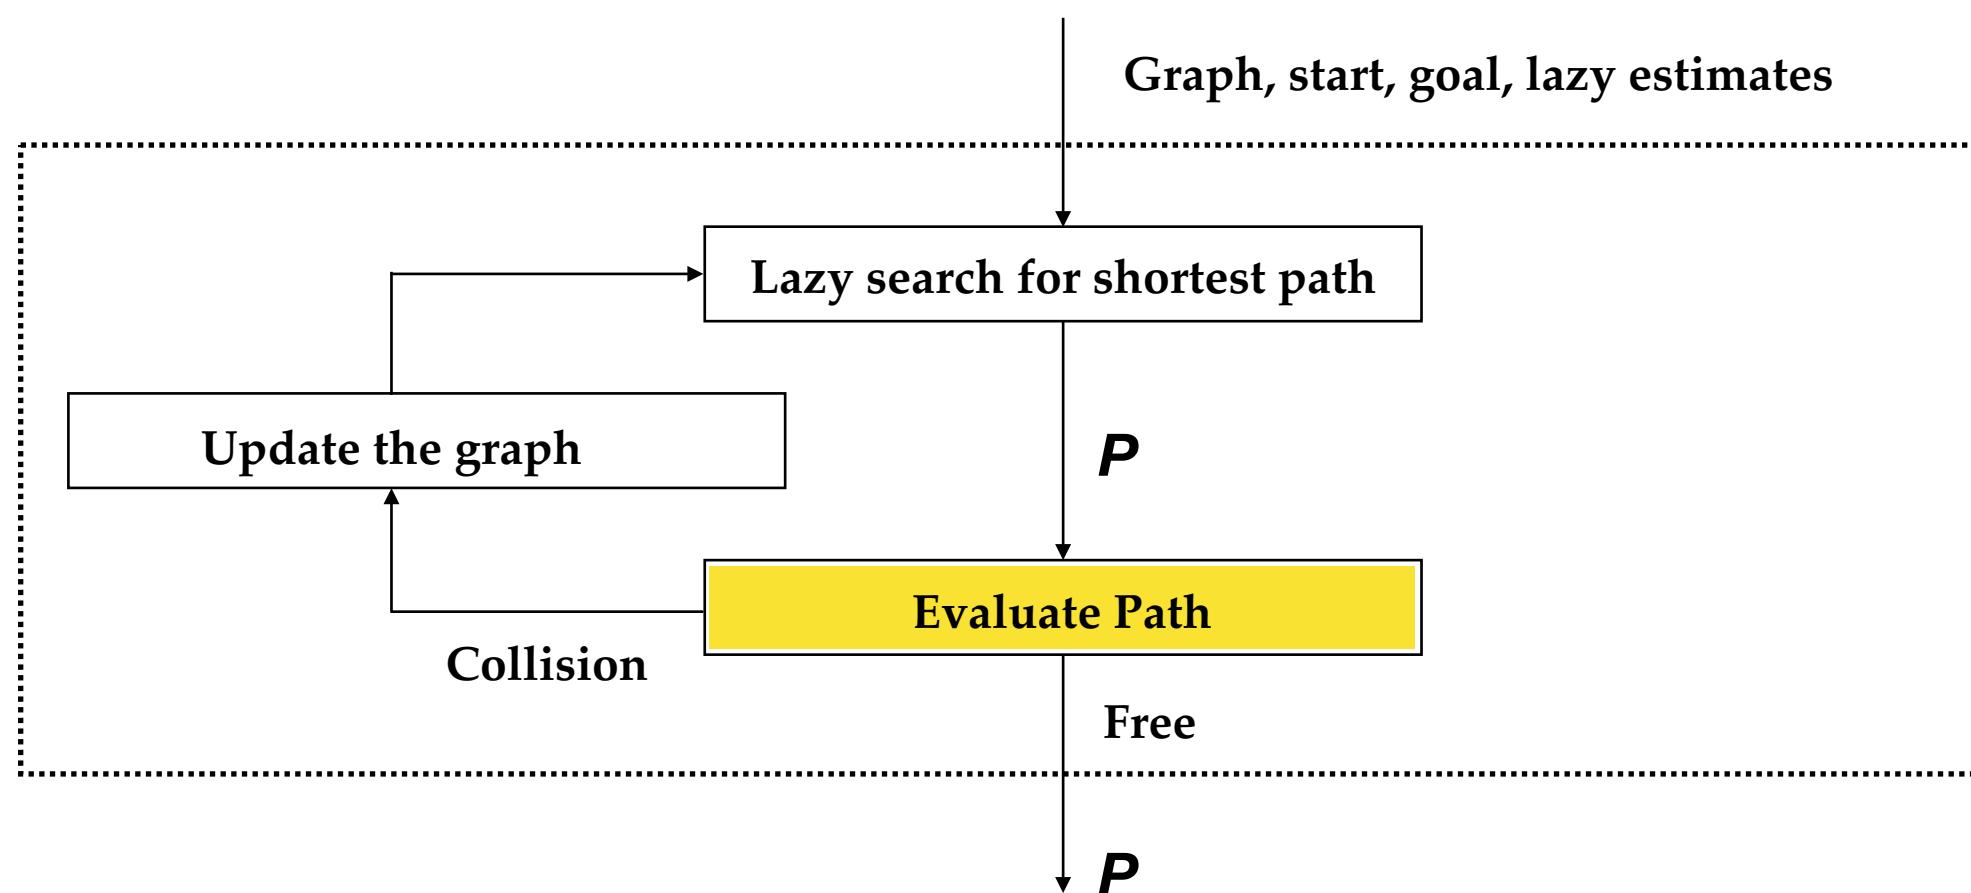
OFU on Steroids!



Only Slime Known Shortest Paths

# LazySP

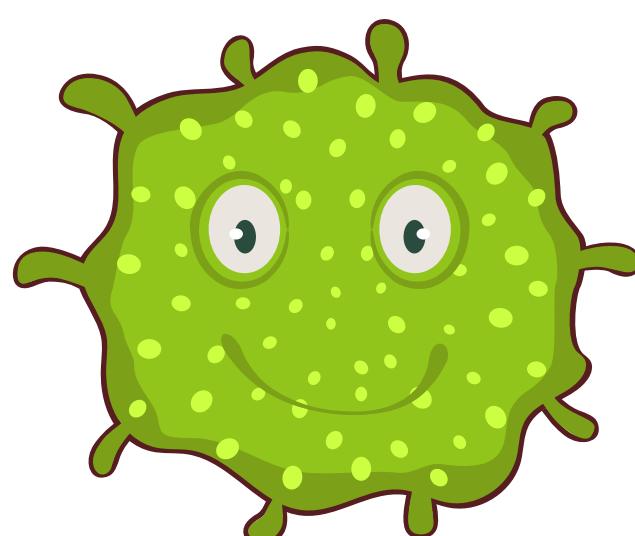
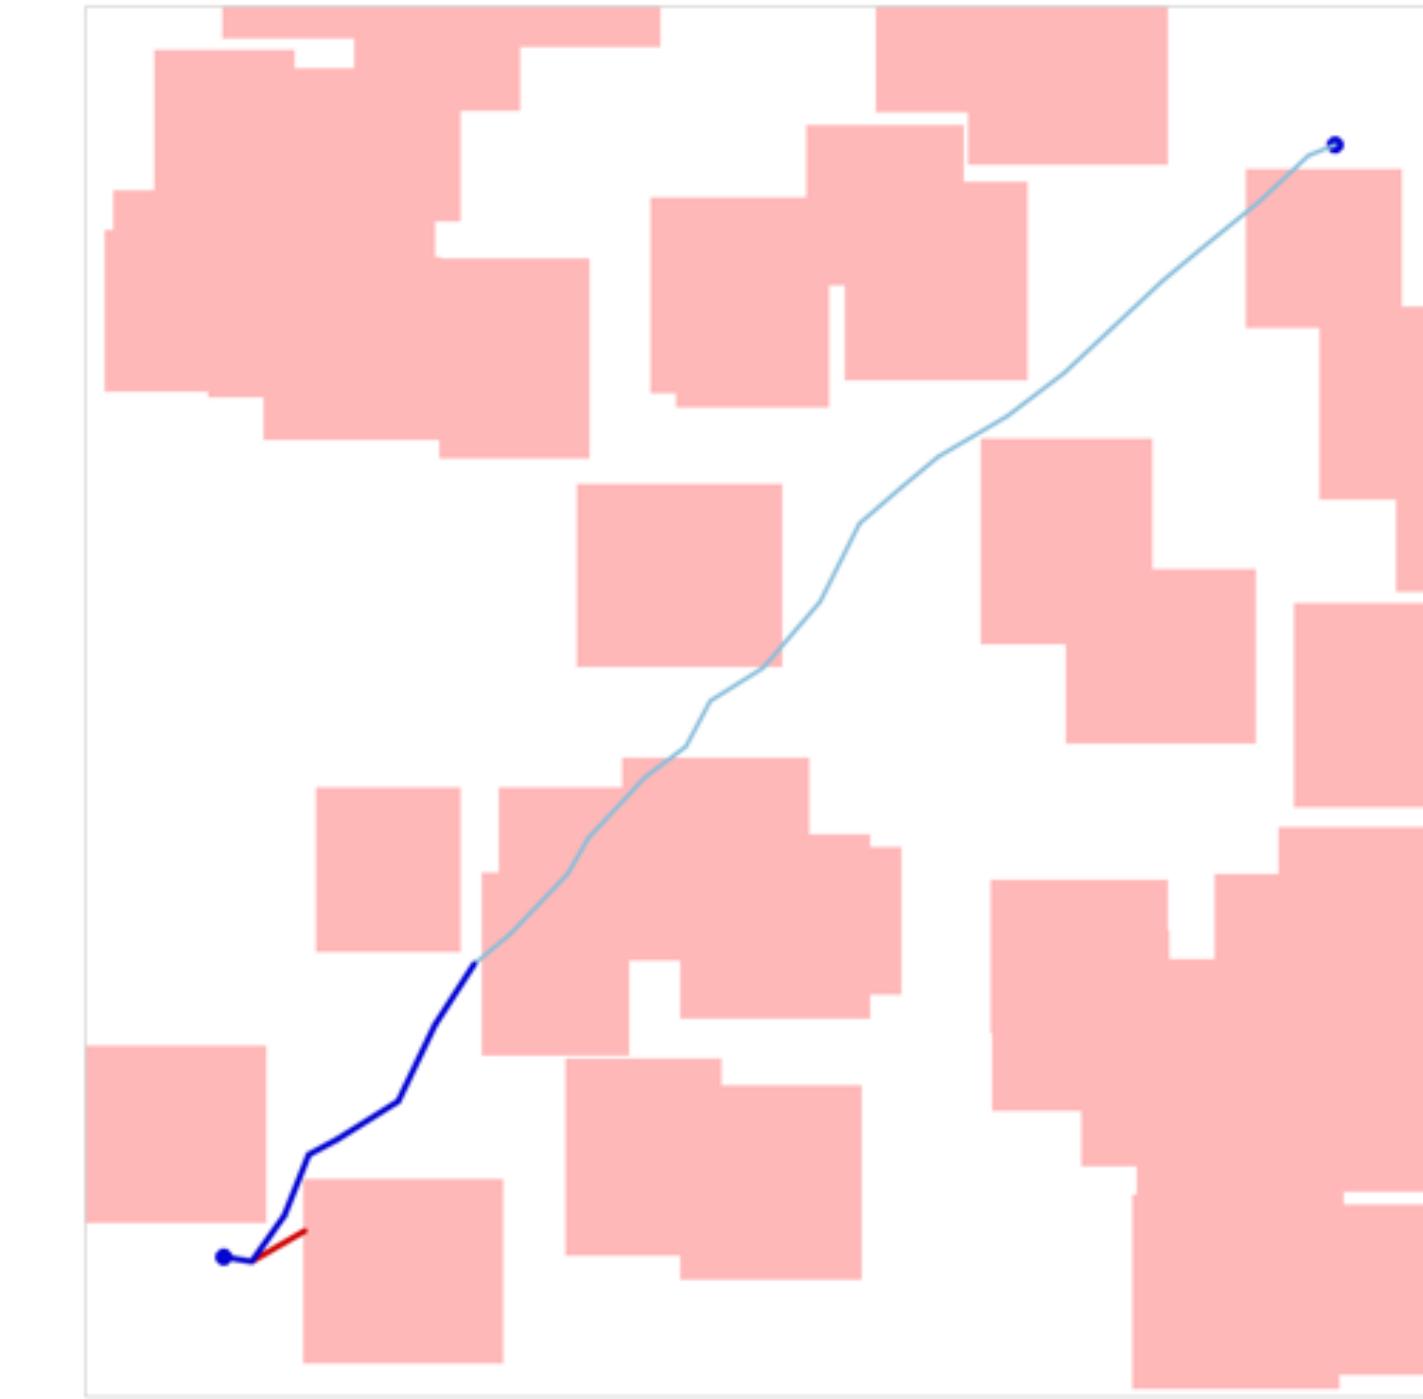
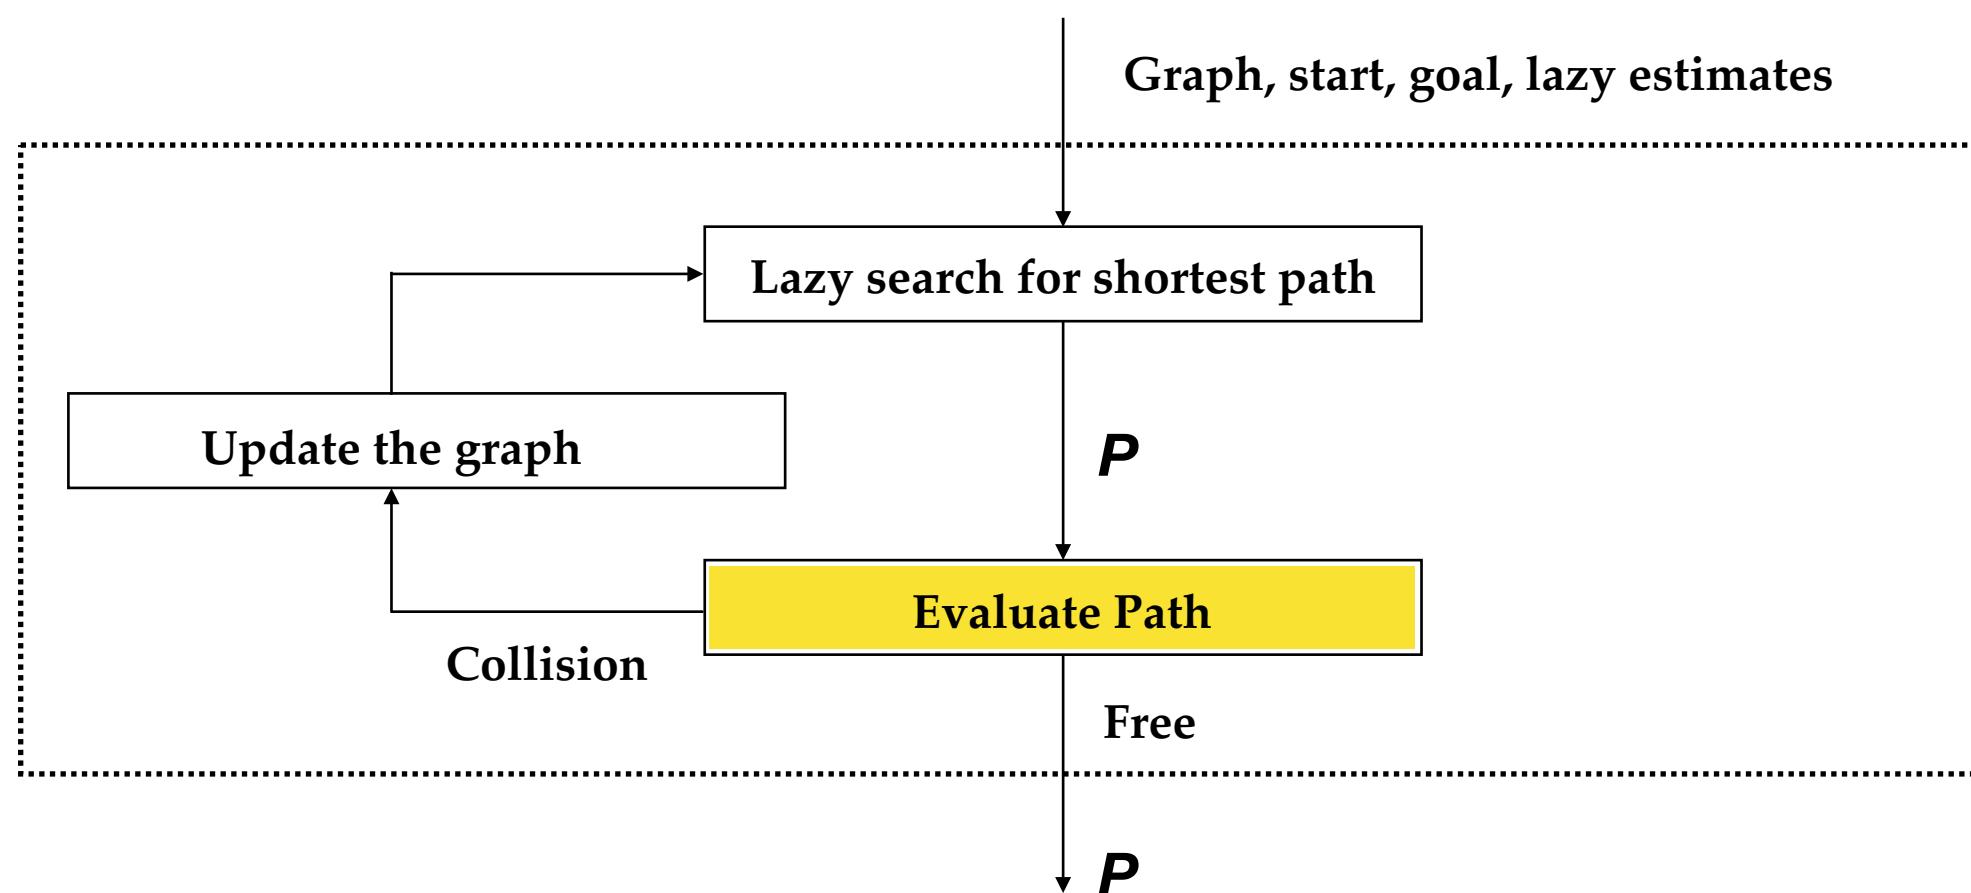
OFU on Steroids!



Only Slime Known Shortest Paths

# LazySP

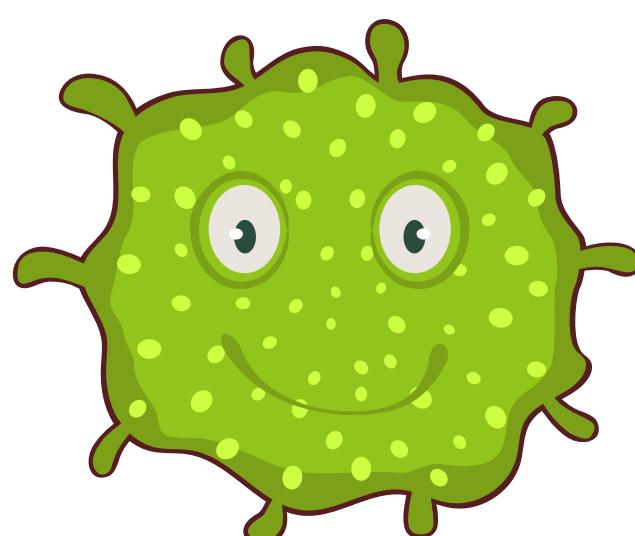
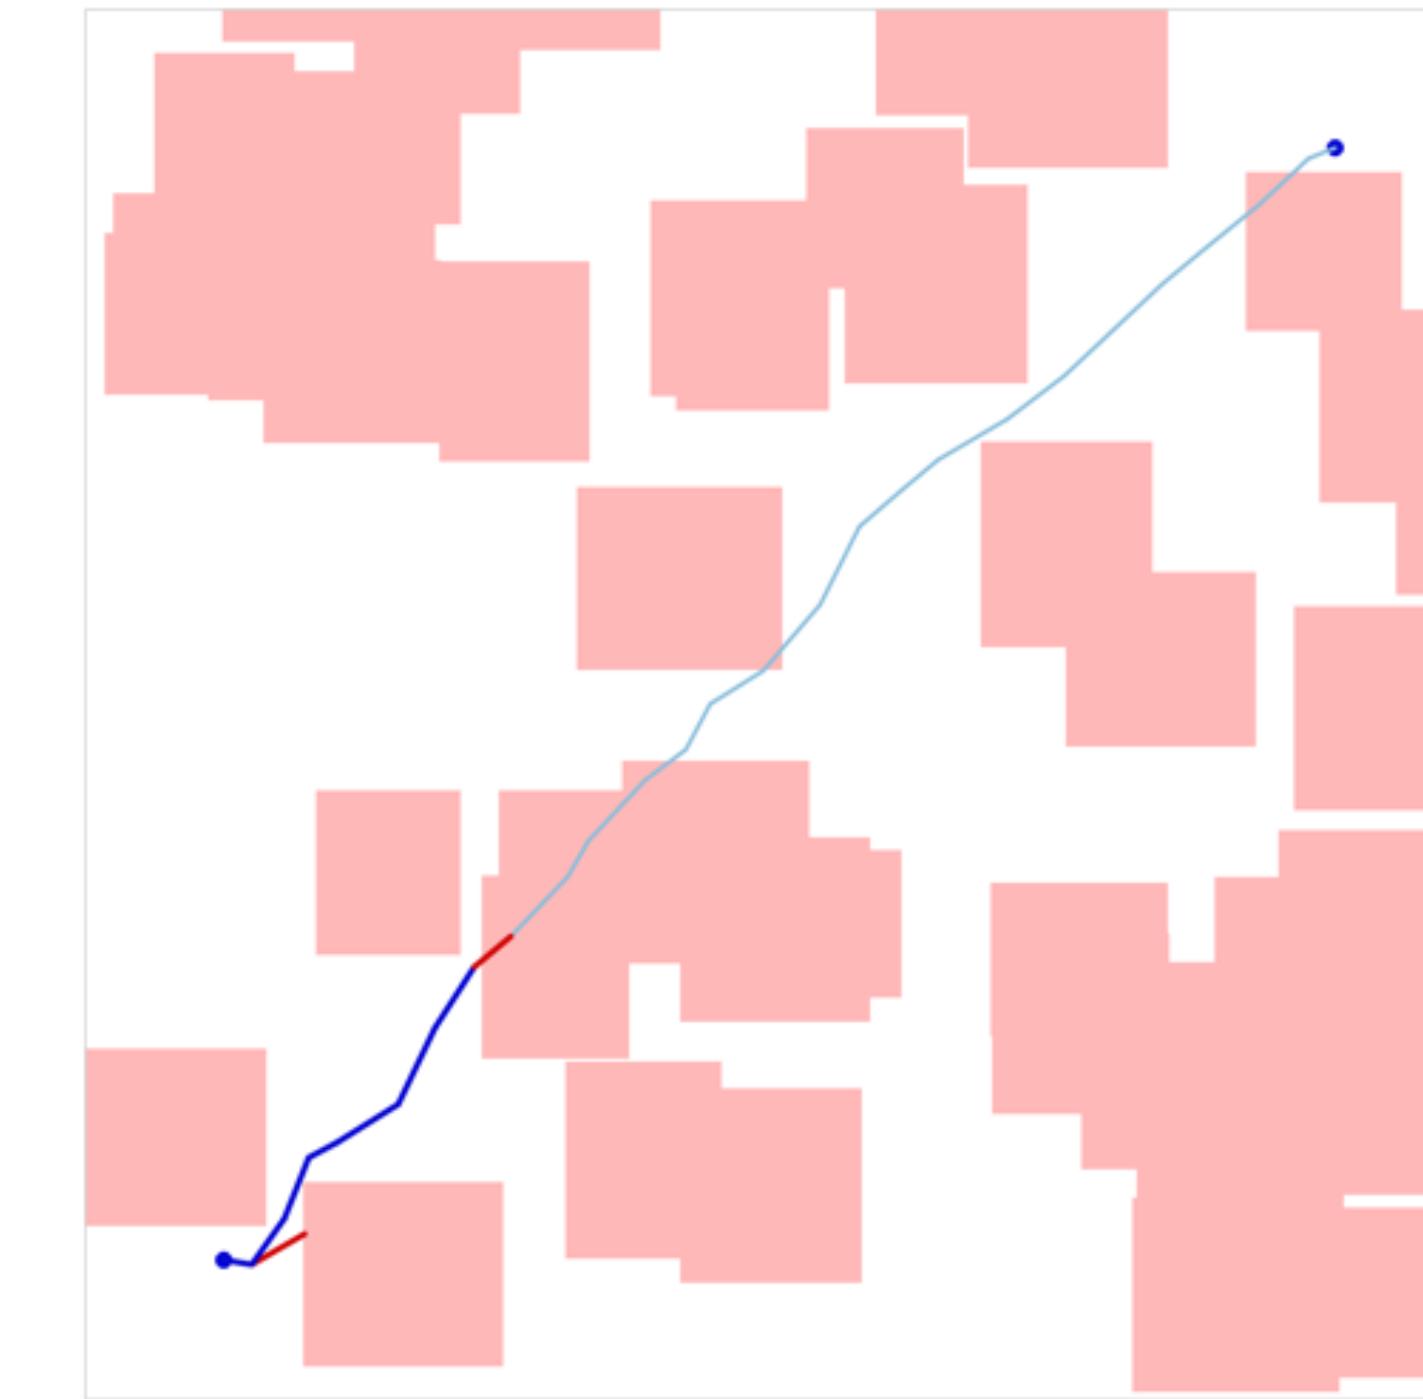
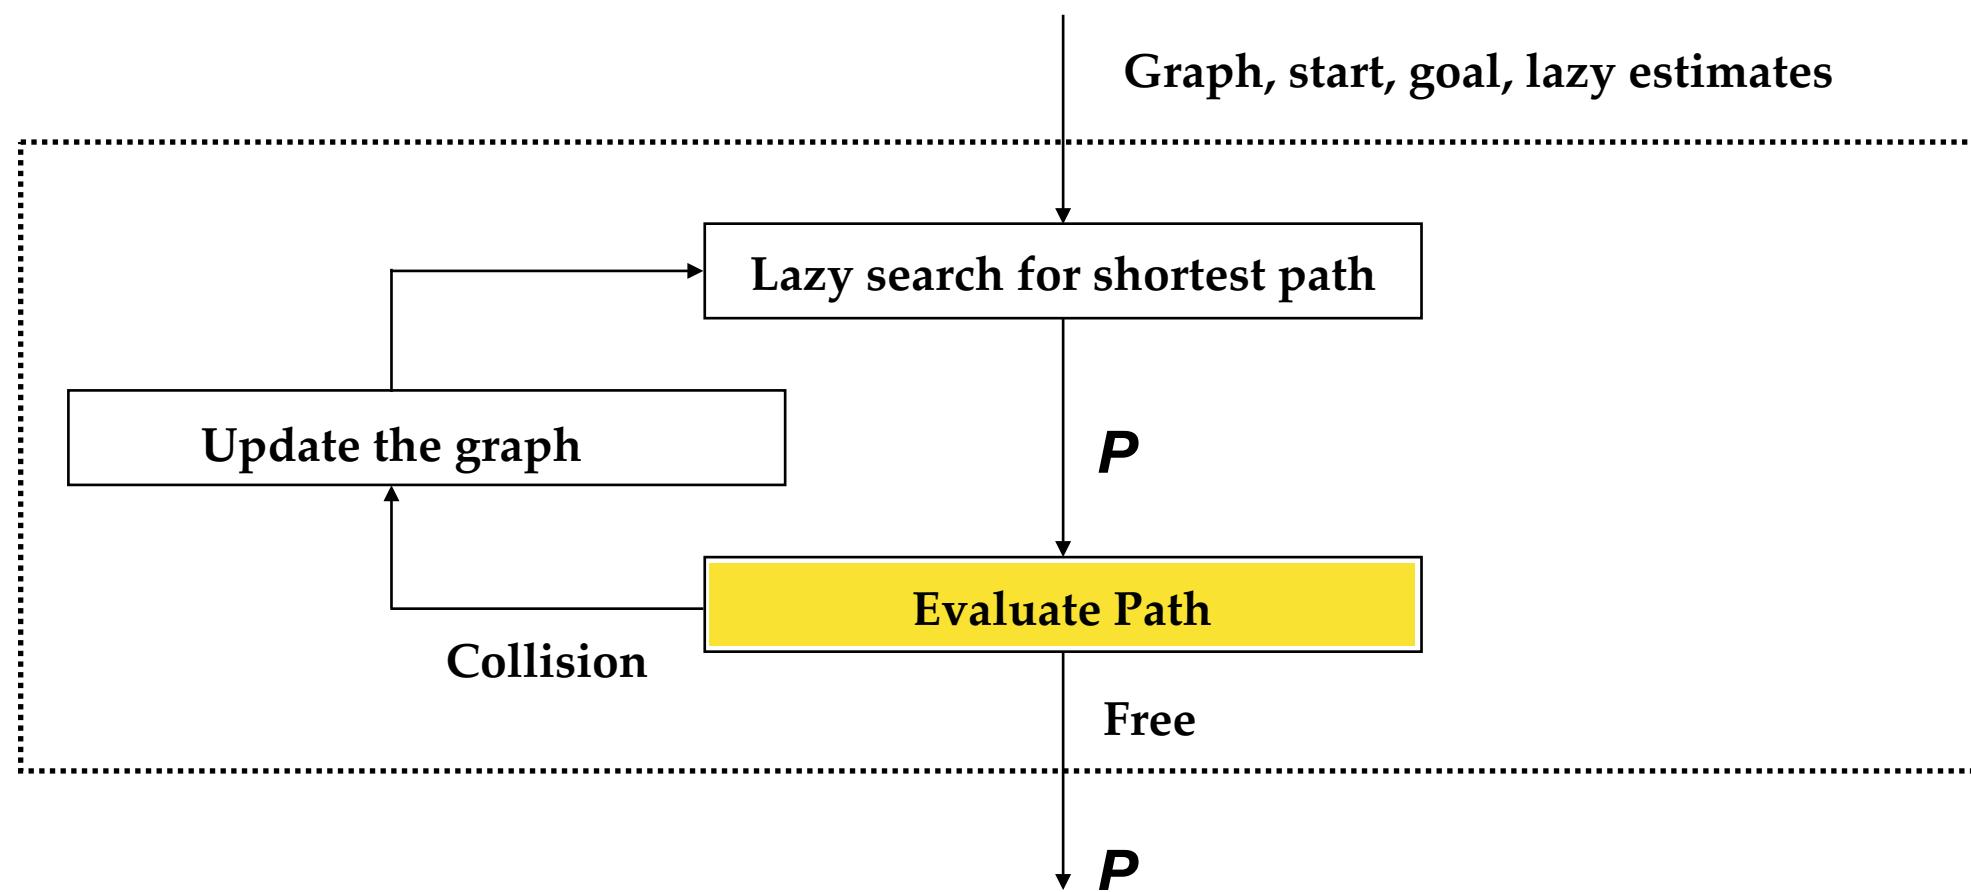
OFU on Steroids!



Only Slime Known Shortest Paths

# LazySP

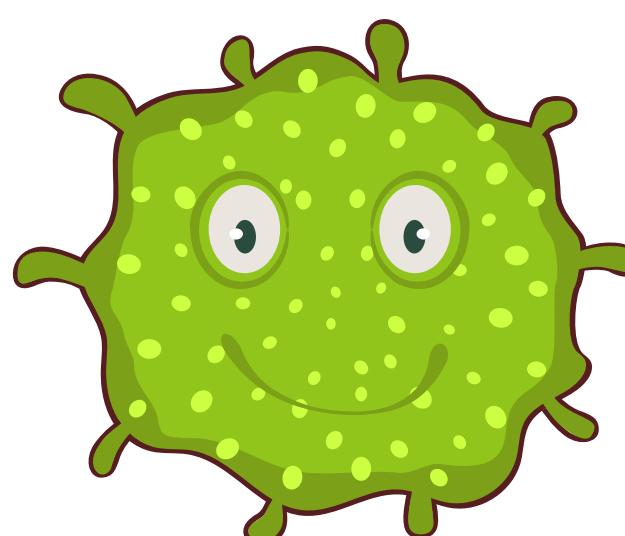
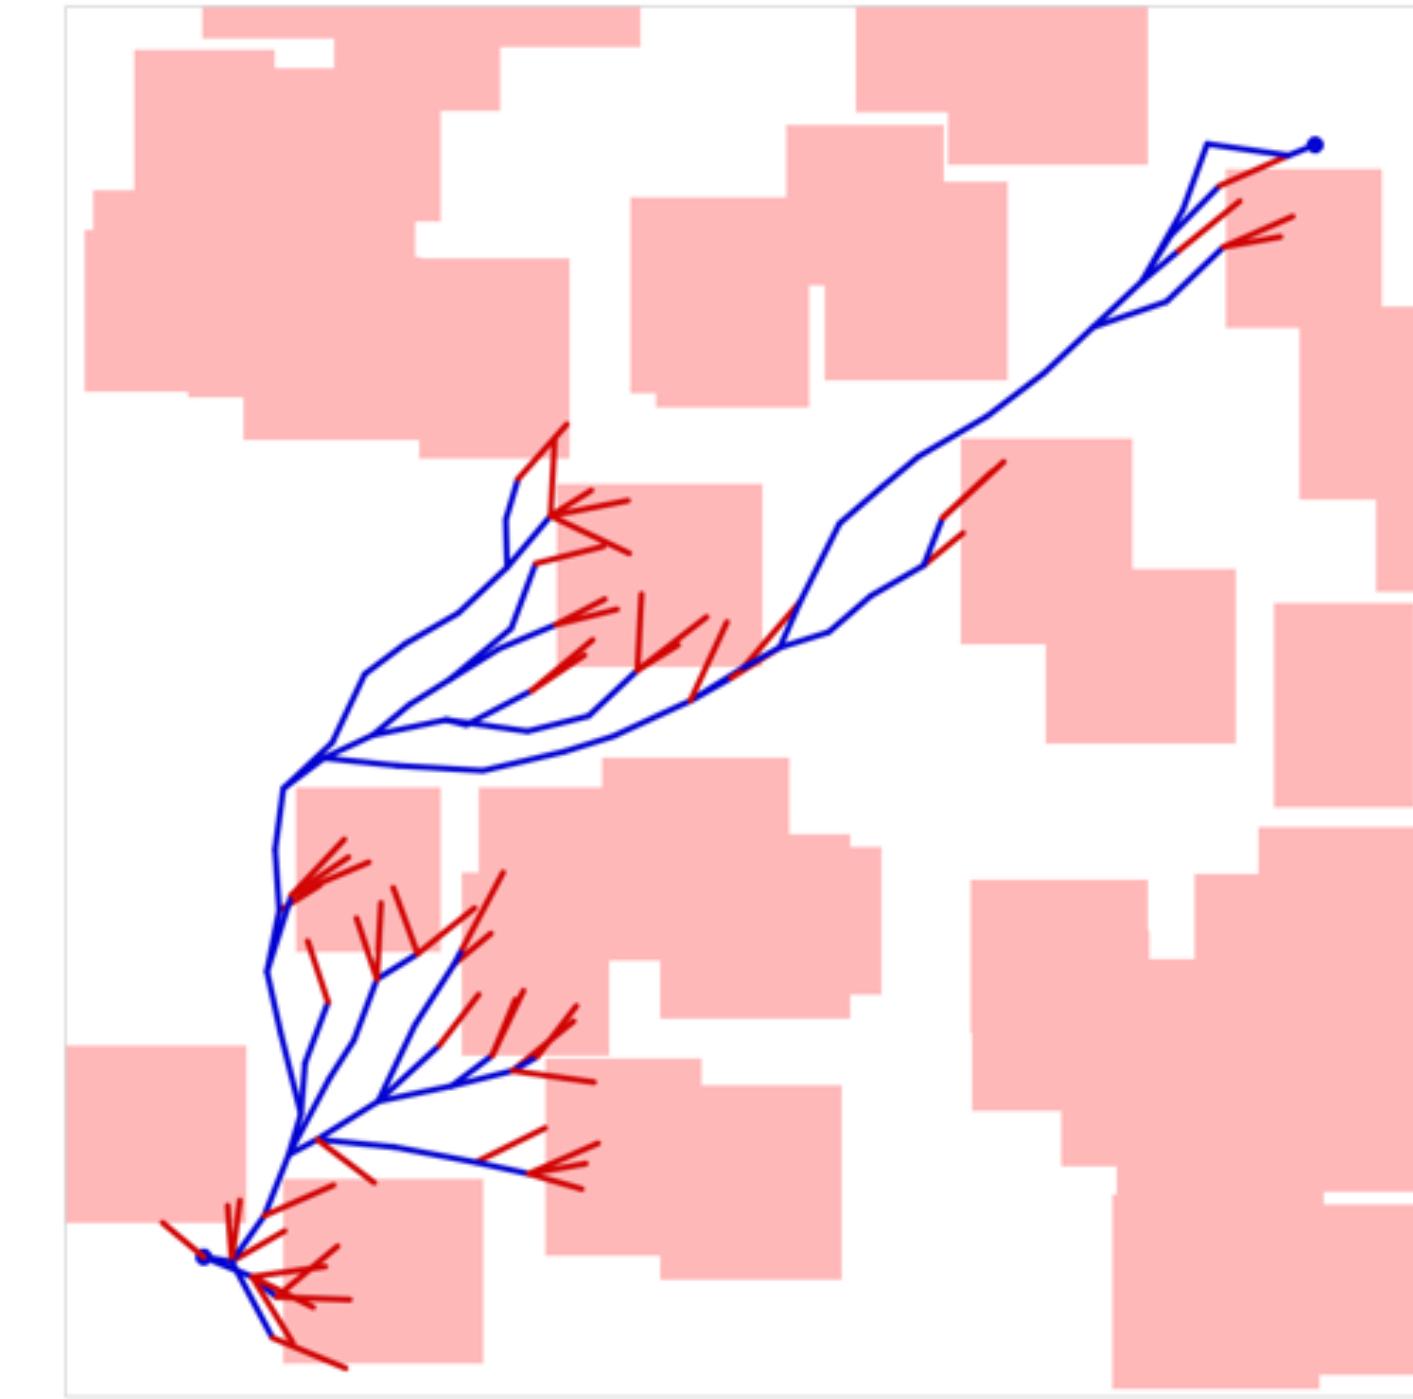
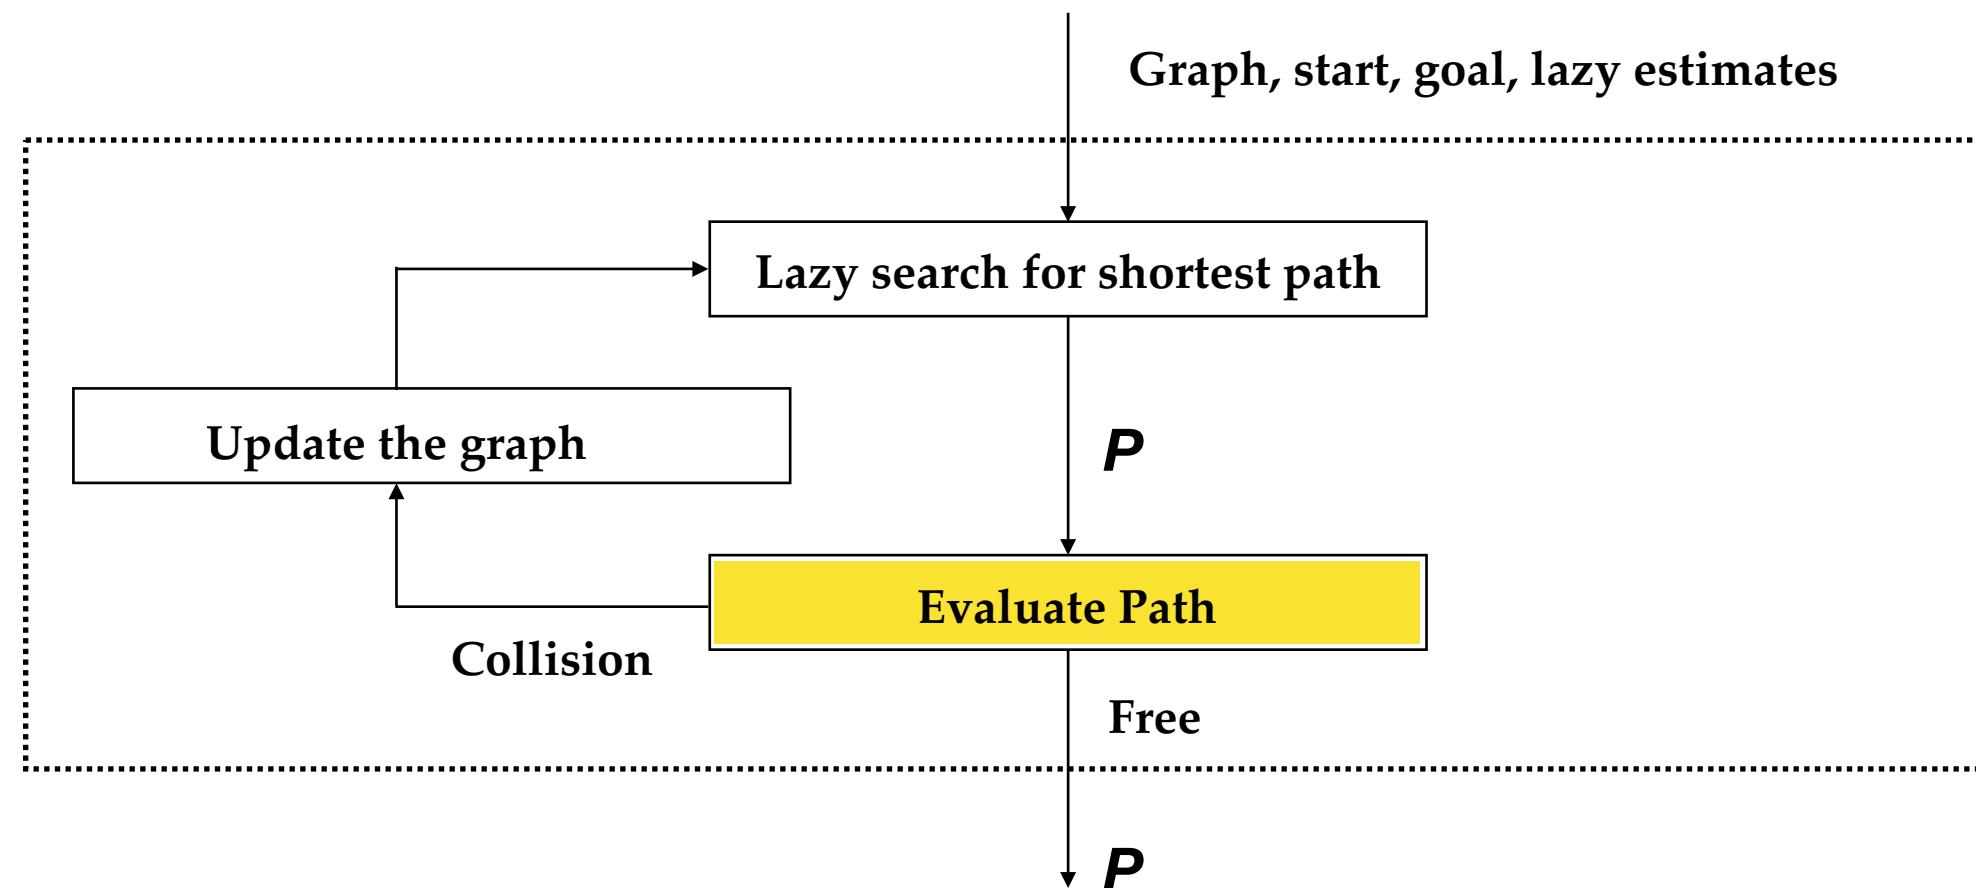
OFU on Steroids!



Only Slime Known Shortest Paths

# LazySP

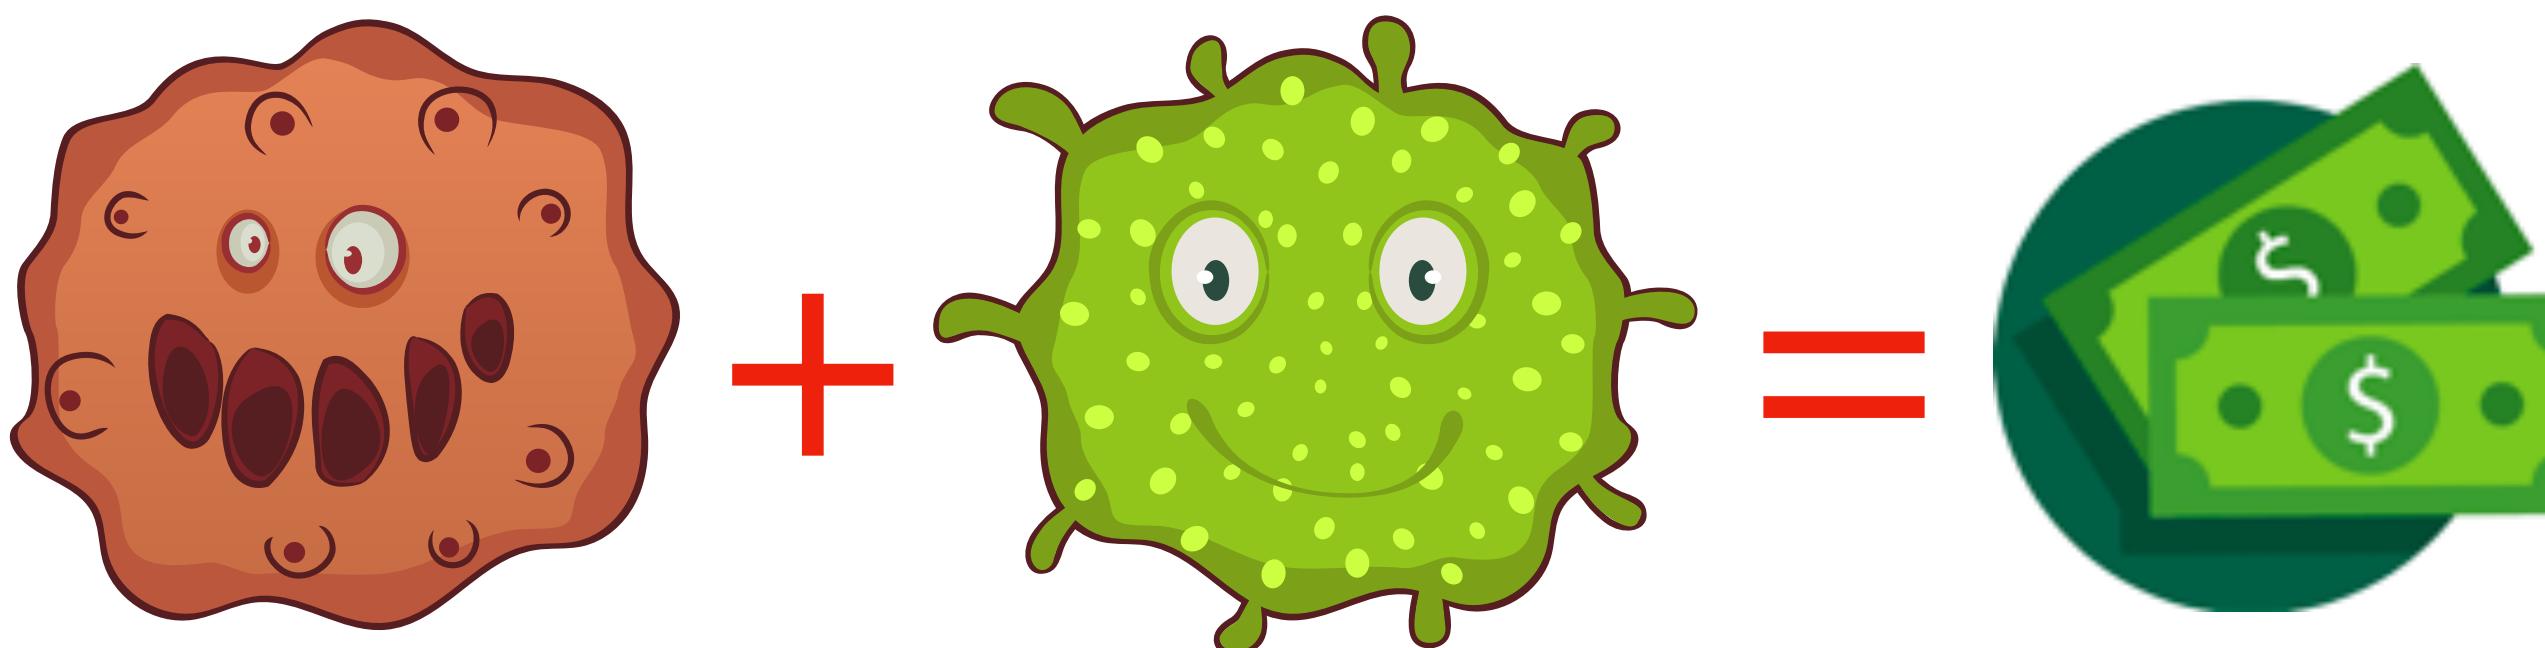
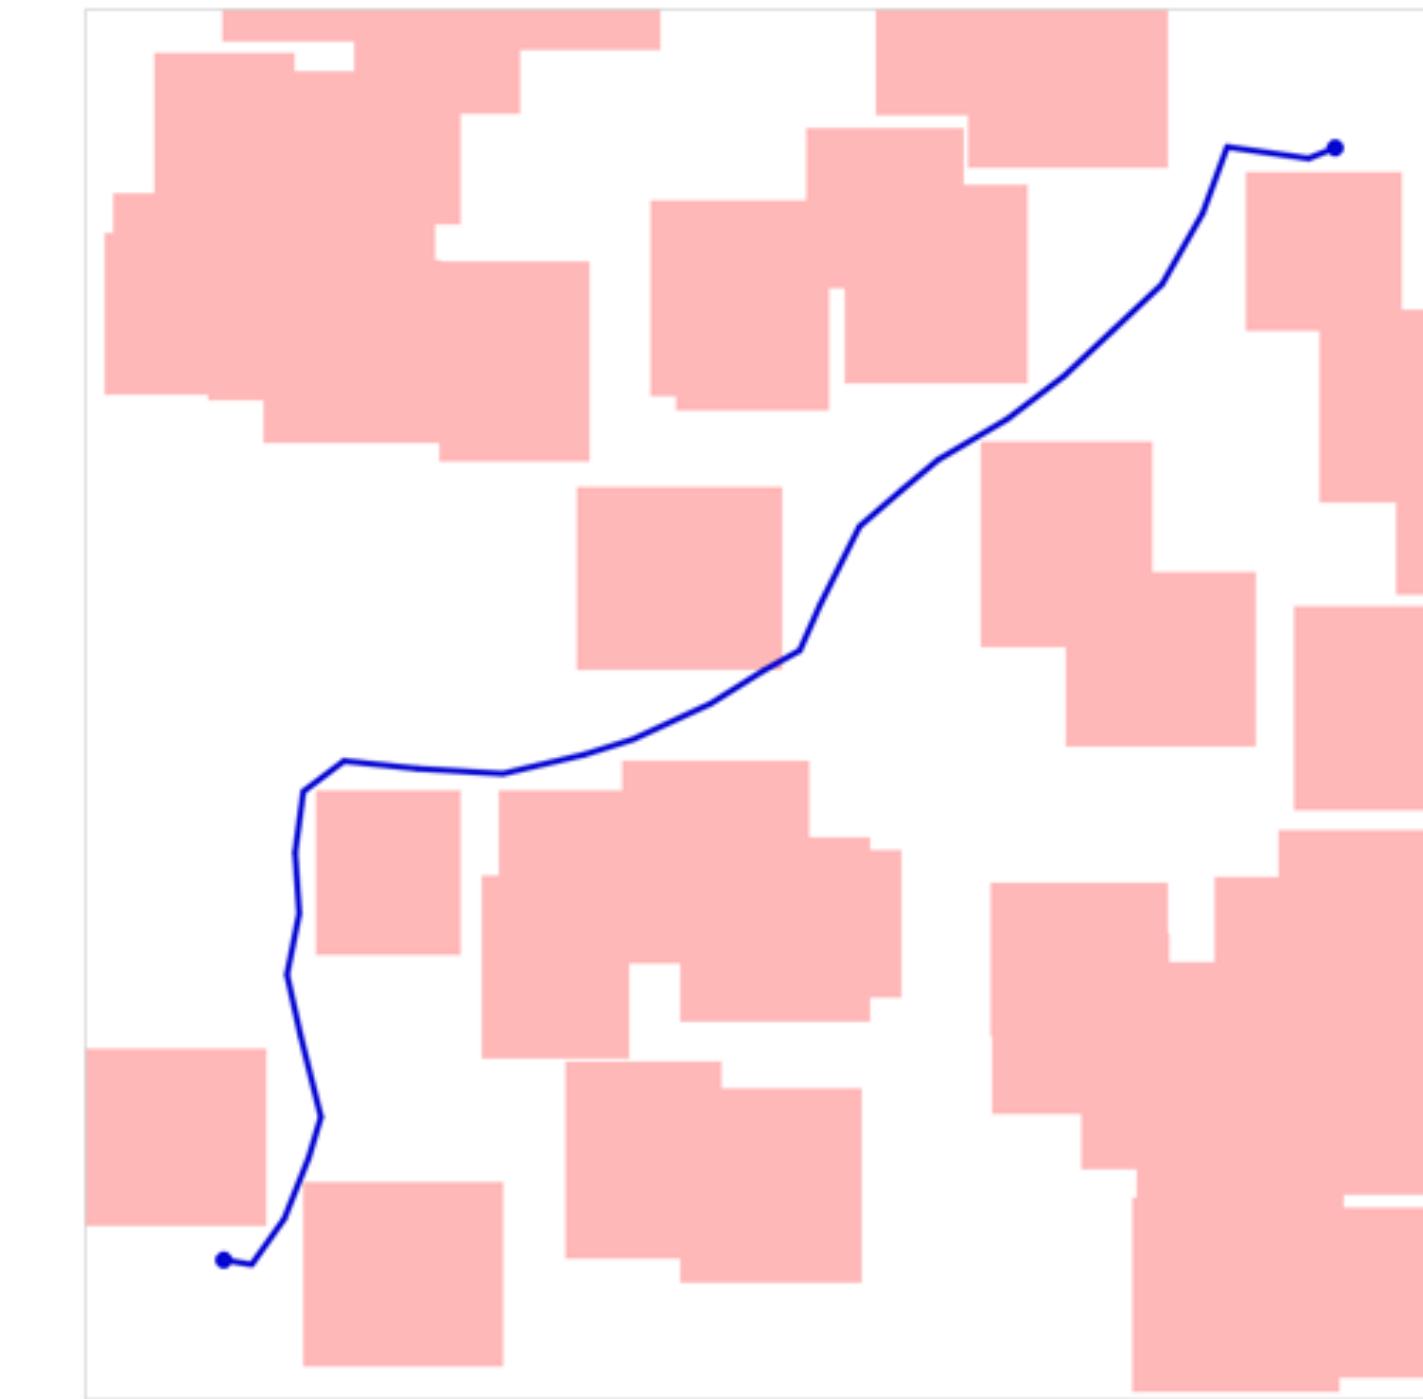
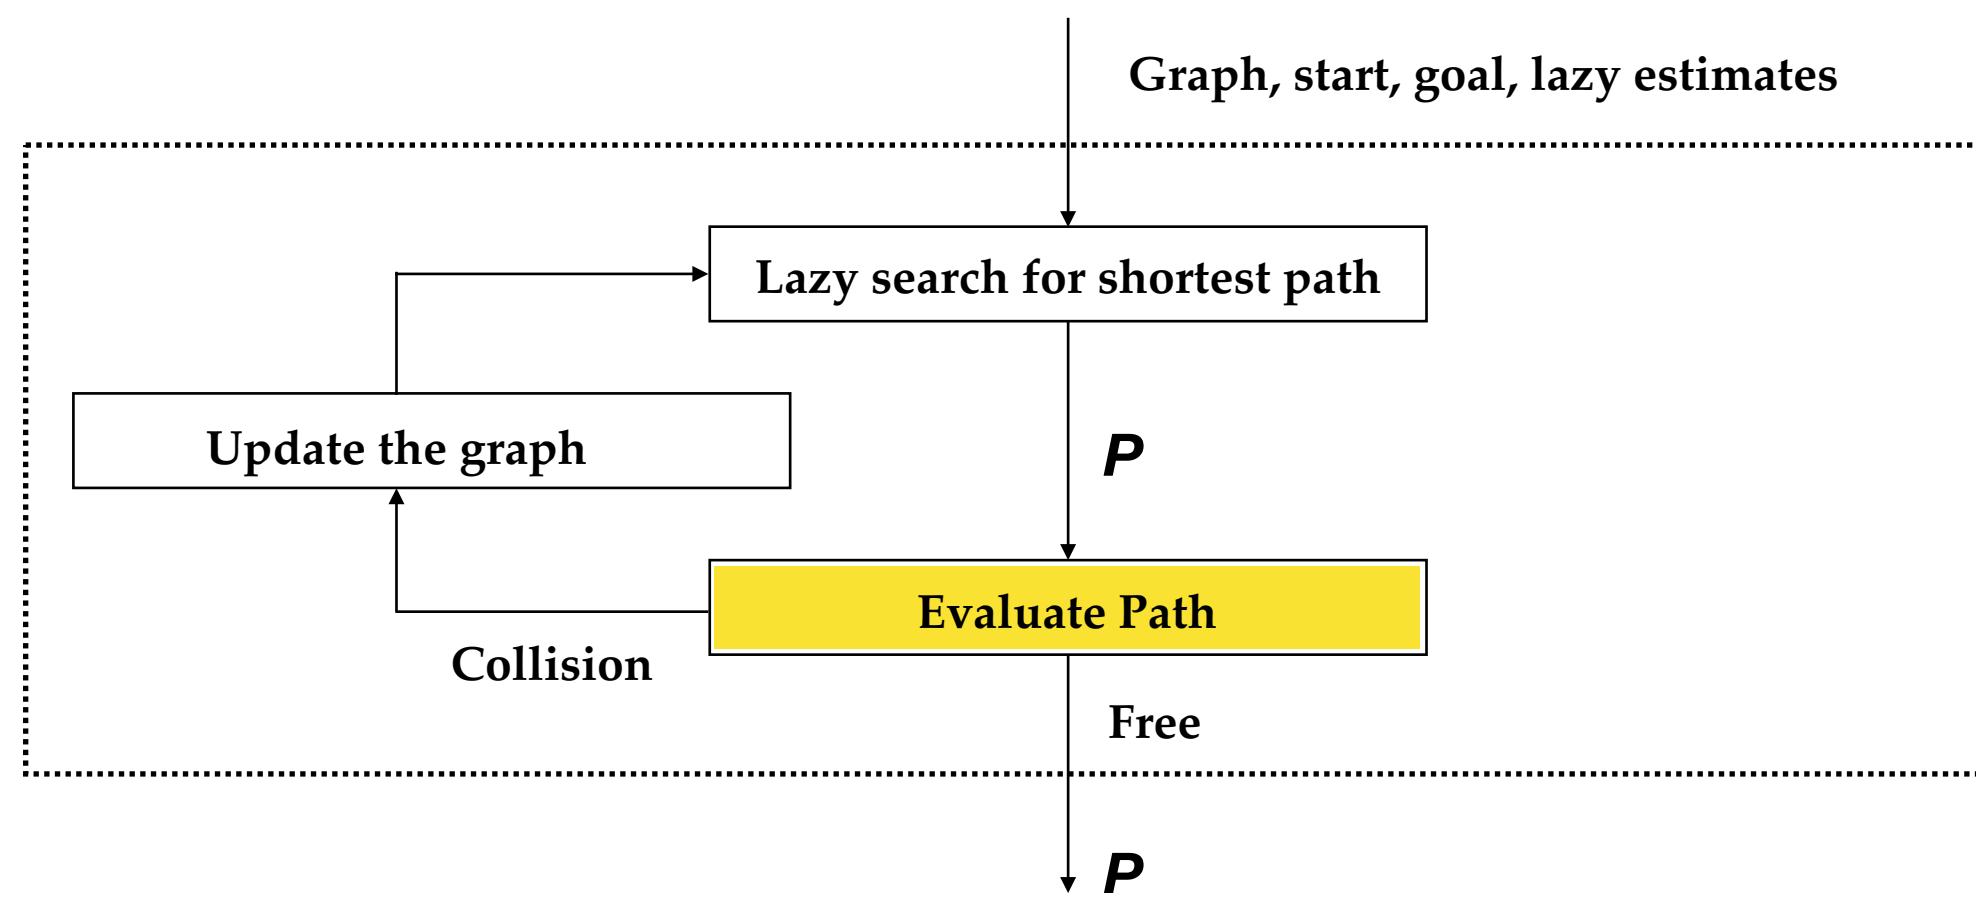
OFU on Steroids!



Optimal Slime!

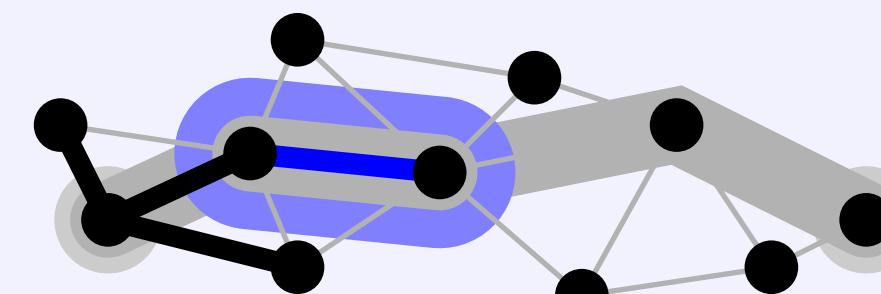
# LazySP

OFU on Steroids!

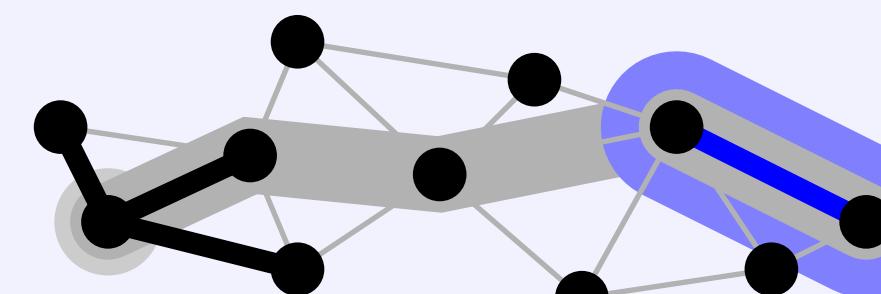


# Edge Selectors

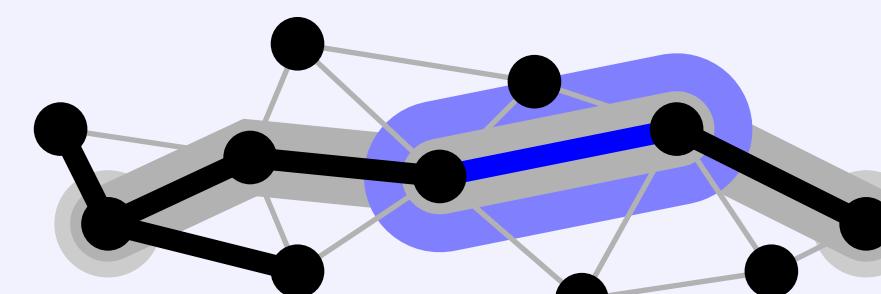
**Forward**  
(first unevaluated edge)



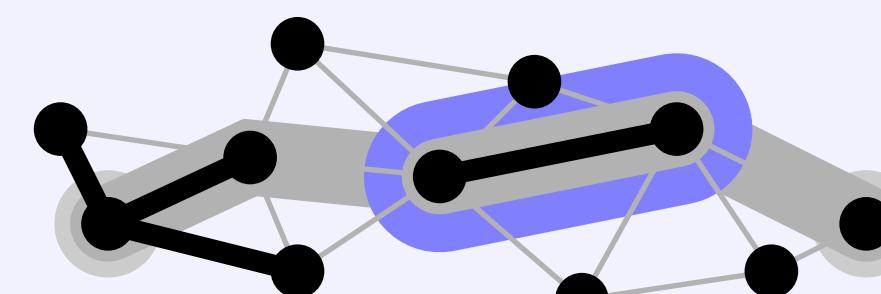
**Reverse**  
(last unevaluated edge)



**Alternate**  
(alternate Forward and Reverse)



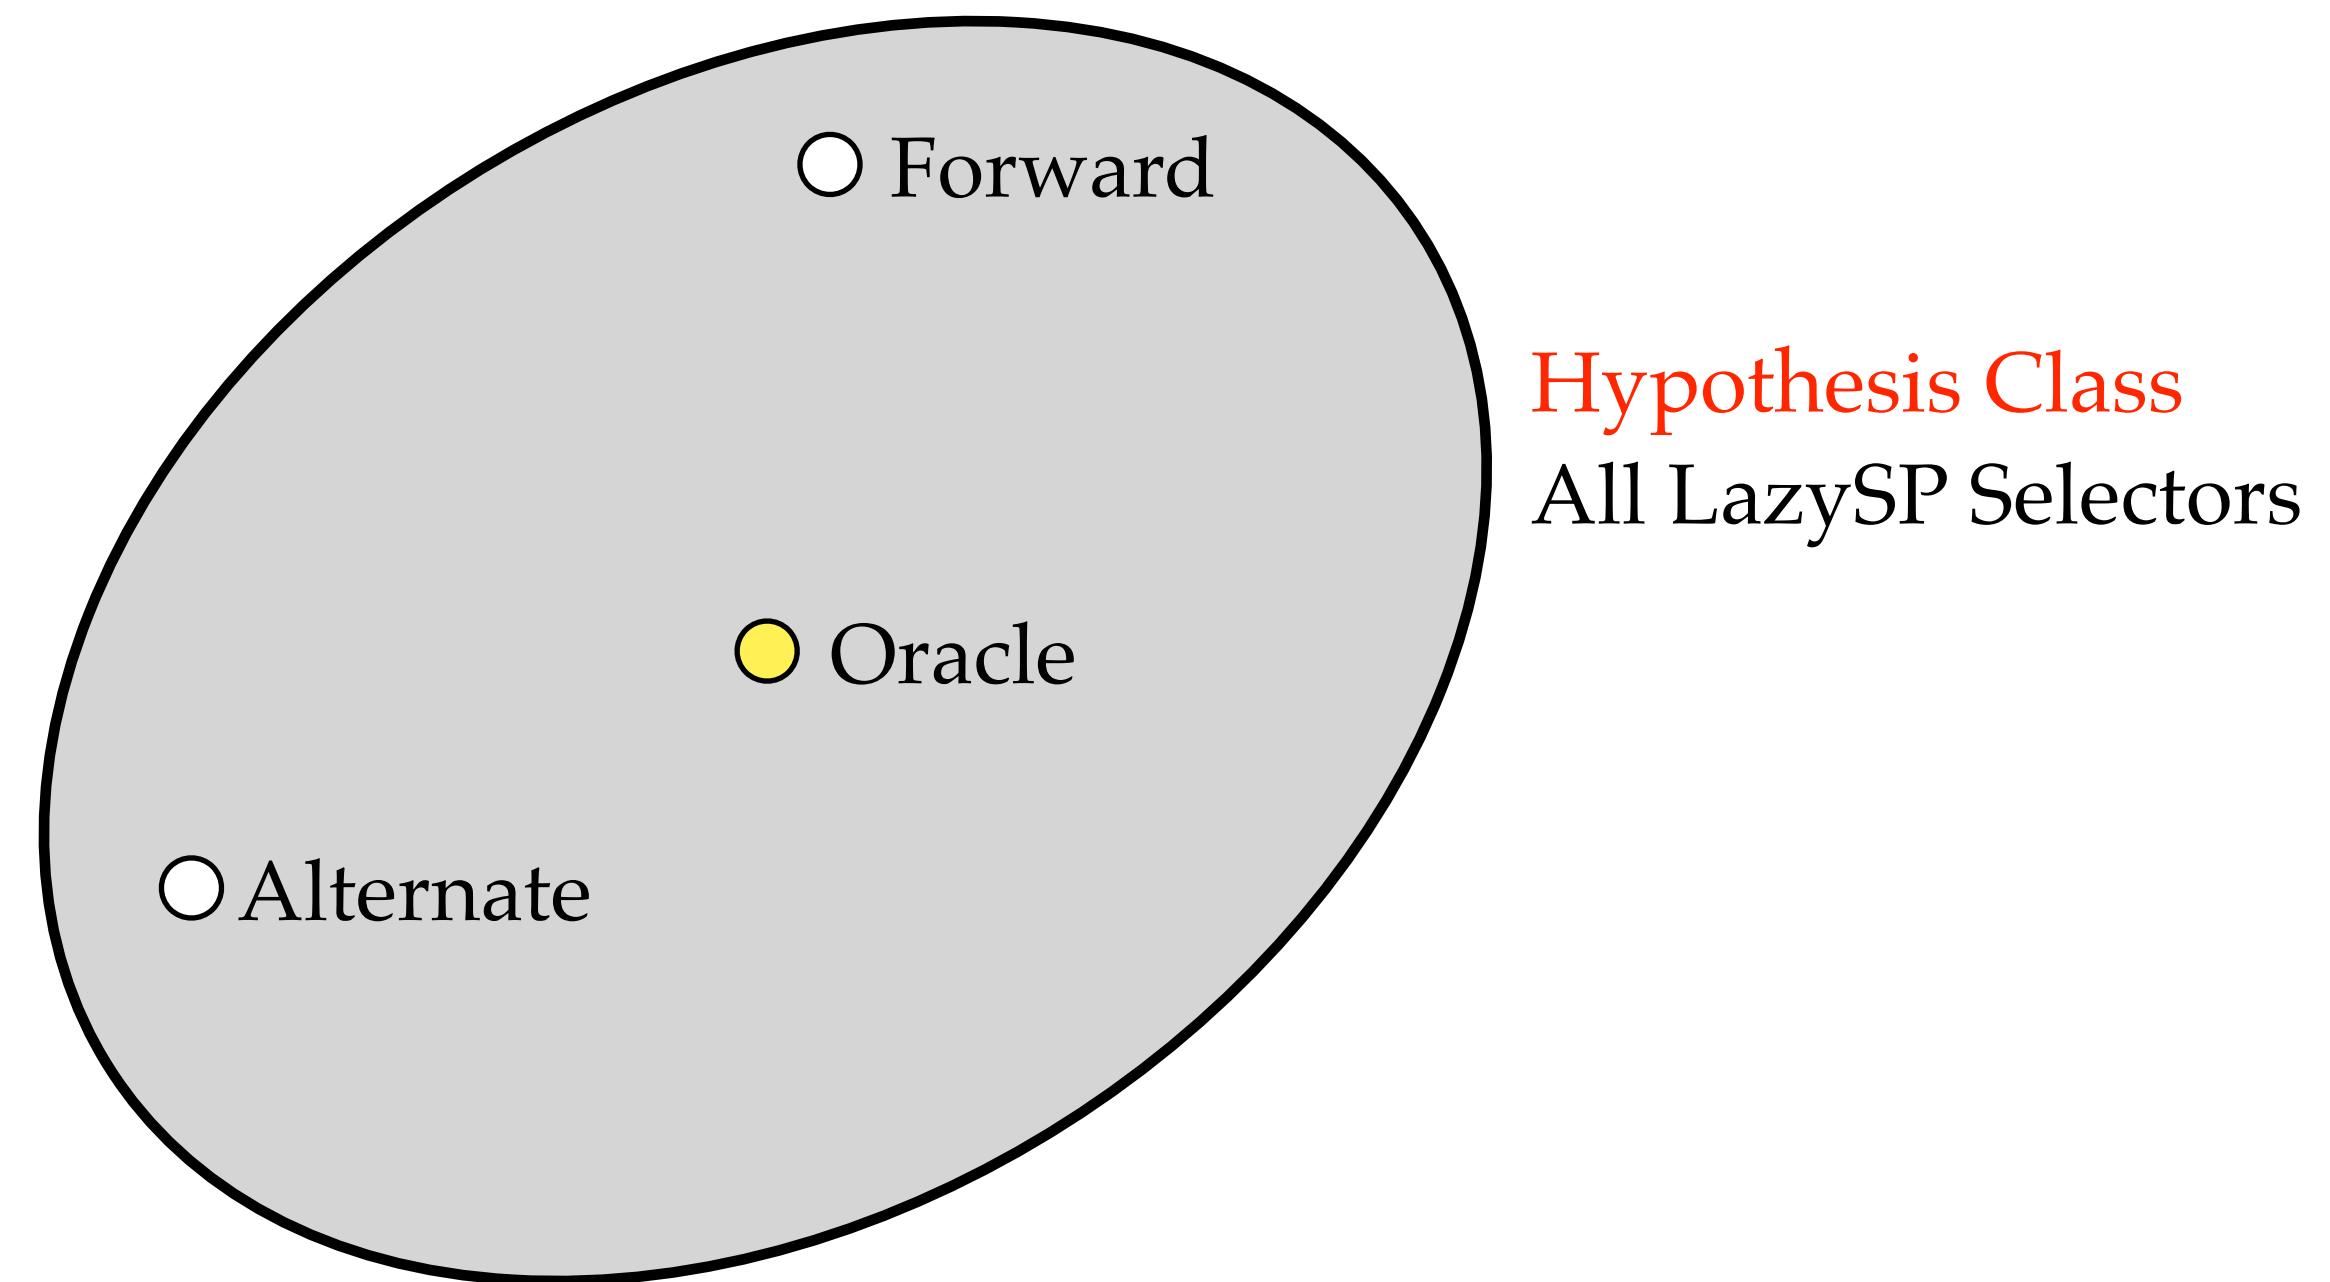
**Bisect**  
(furthest from an unevaluated edge)



# The Realizability Assumption

Can we Learn to  
Imitate the Oracle?

*Leveraging experience in lazy search, Bhardwaj et al., RSS 2019.*



## The Oracle is a LazySP Selector!

*The Provable Virtue of Laziness in Motion Planning, Hagtalab et al.,  
ICAPS 2018.*

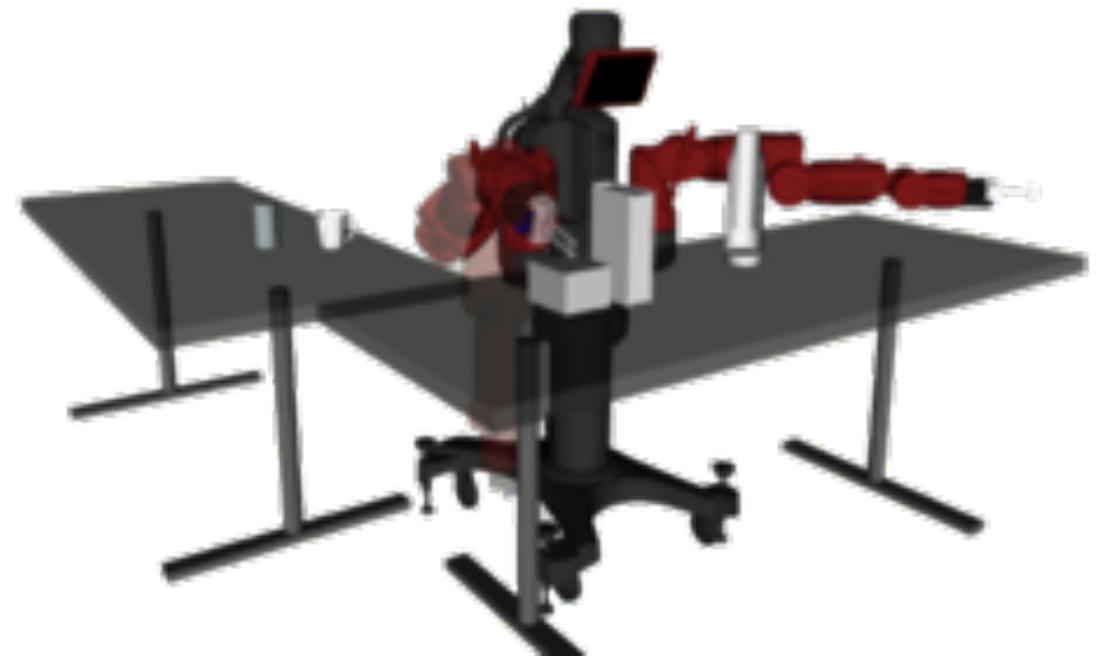
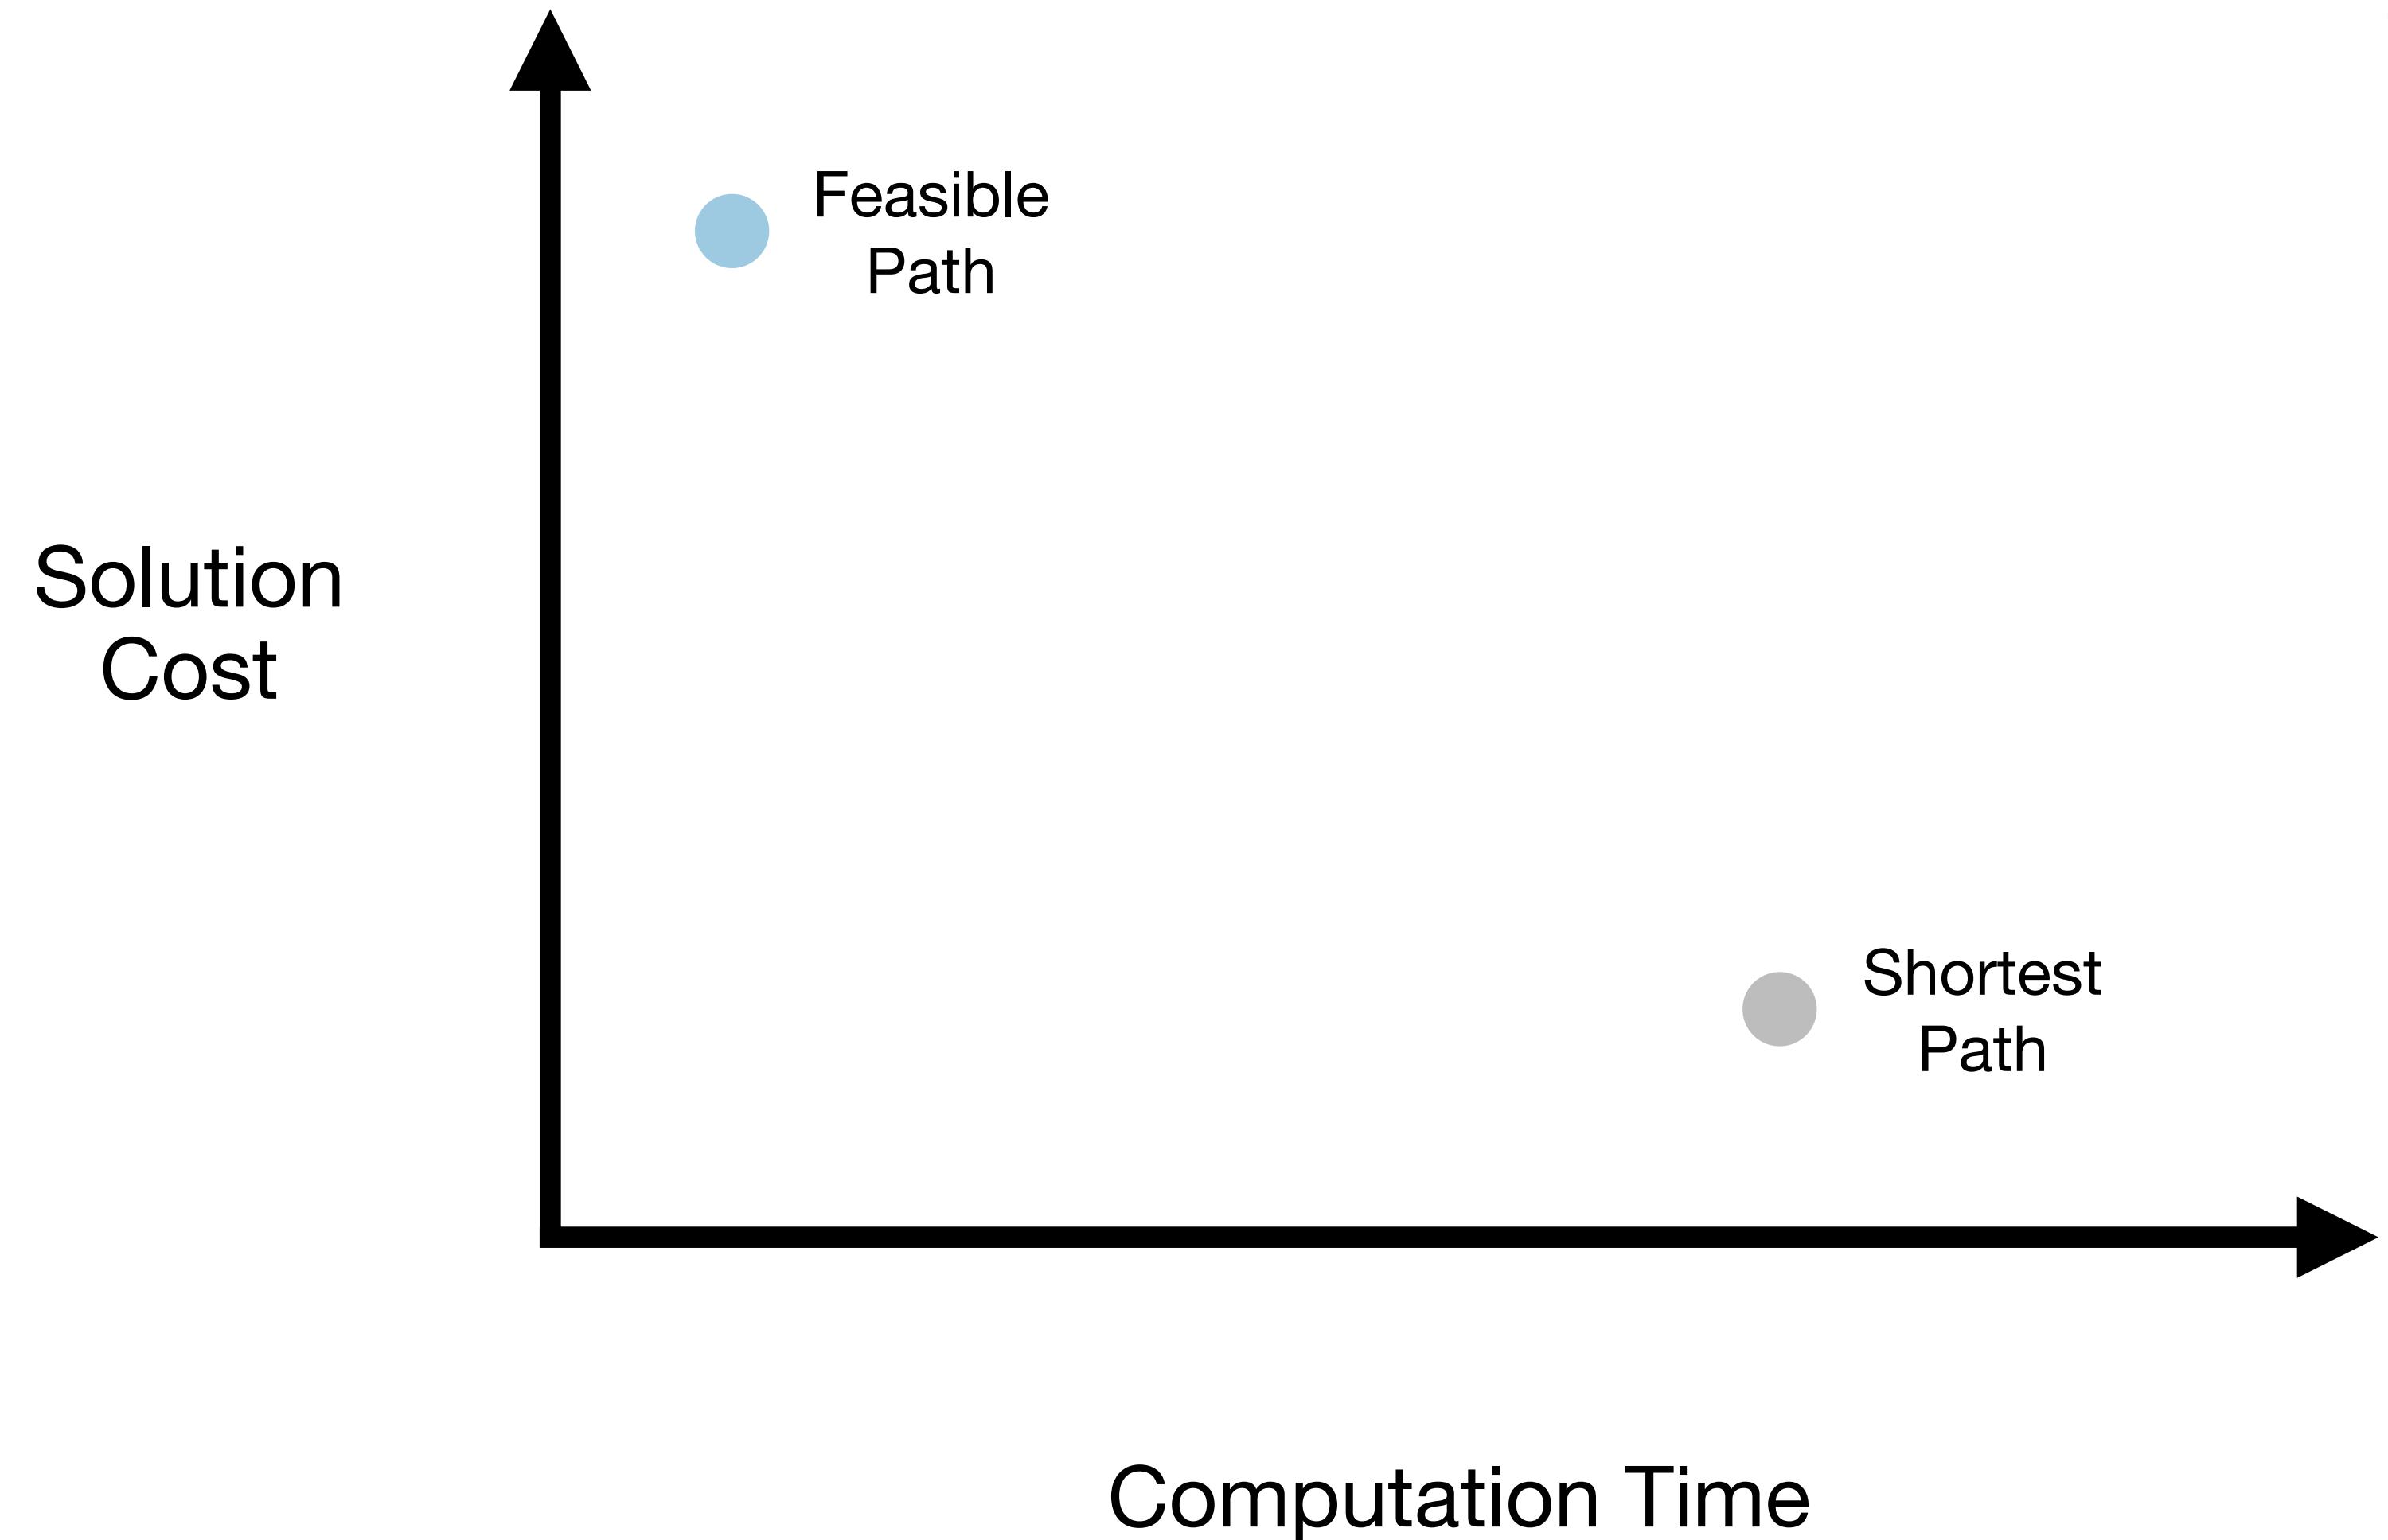
Is there a Search Algorithm  
that **Minimizes**  
the Number of Edge Evaluations?

# LazySP

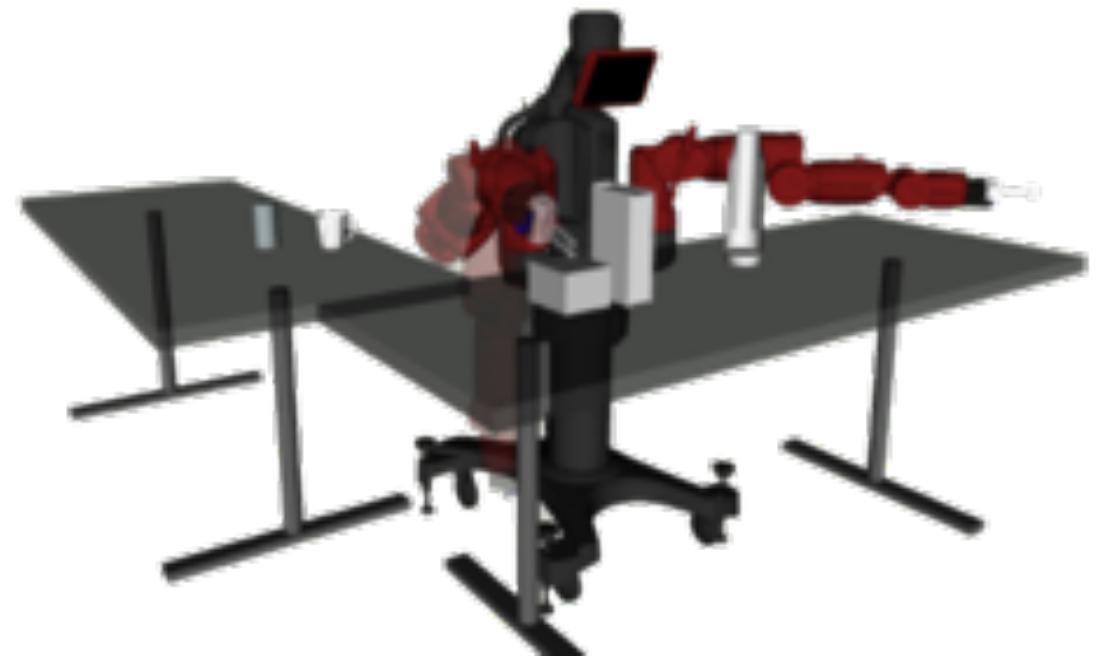
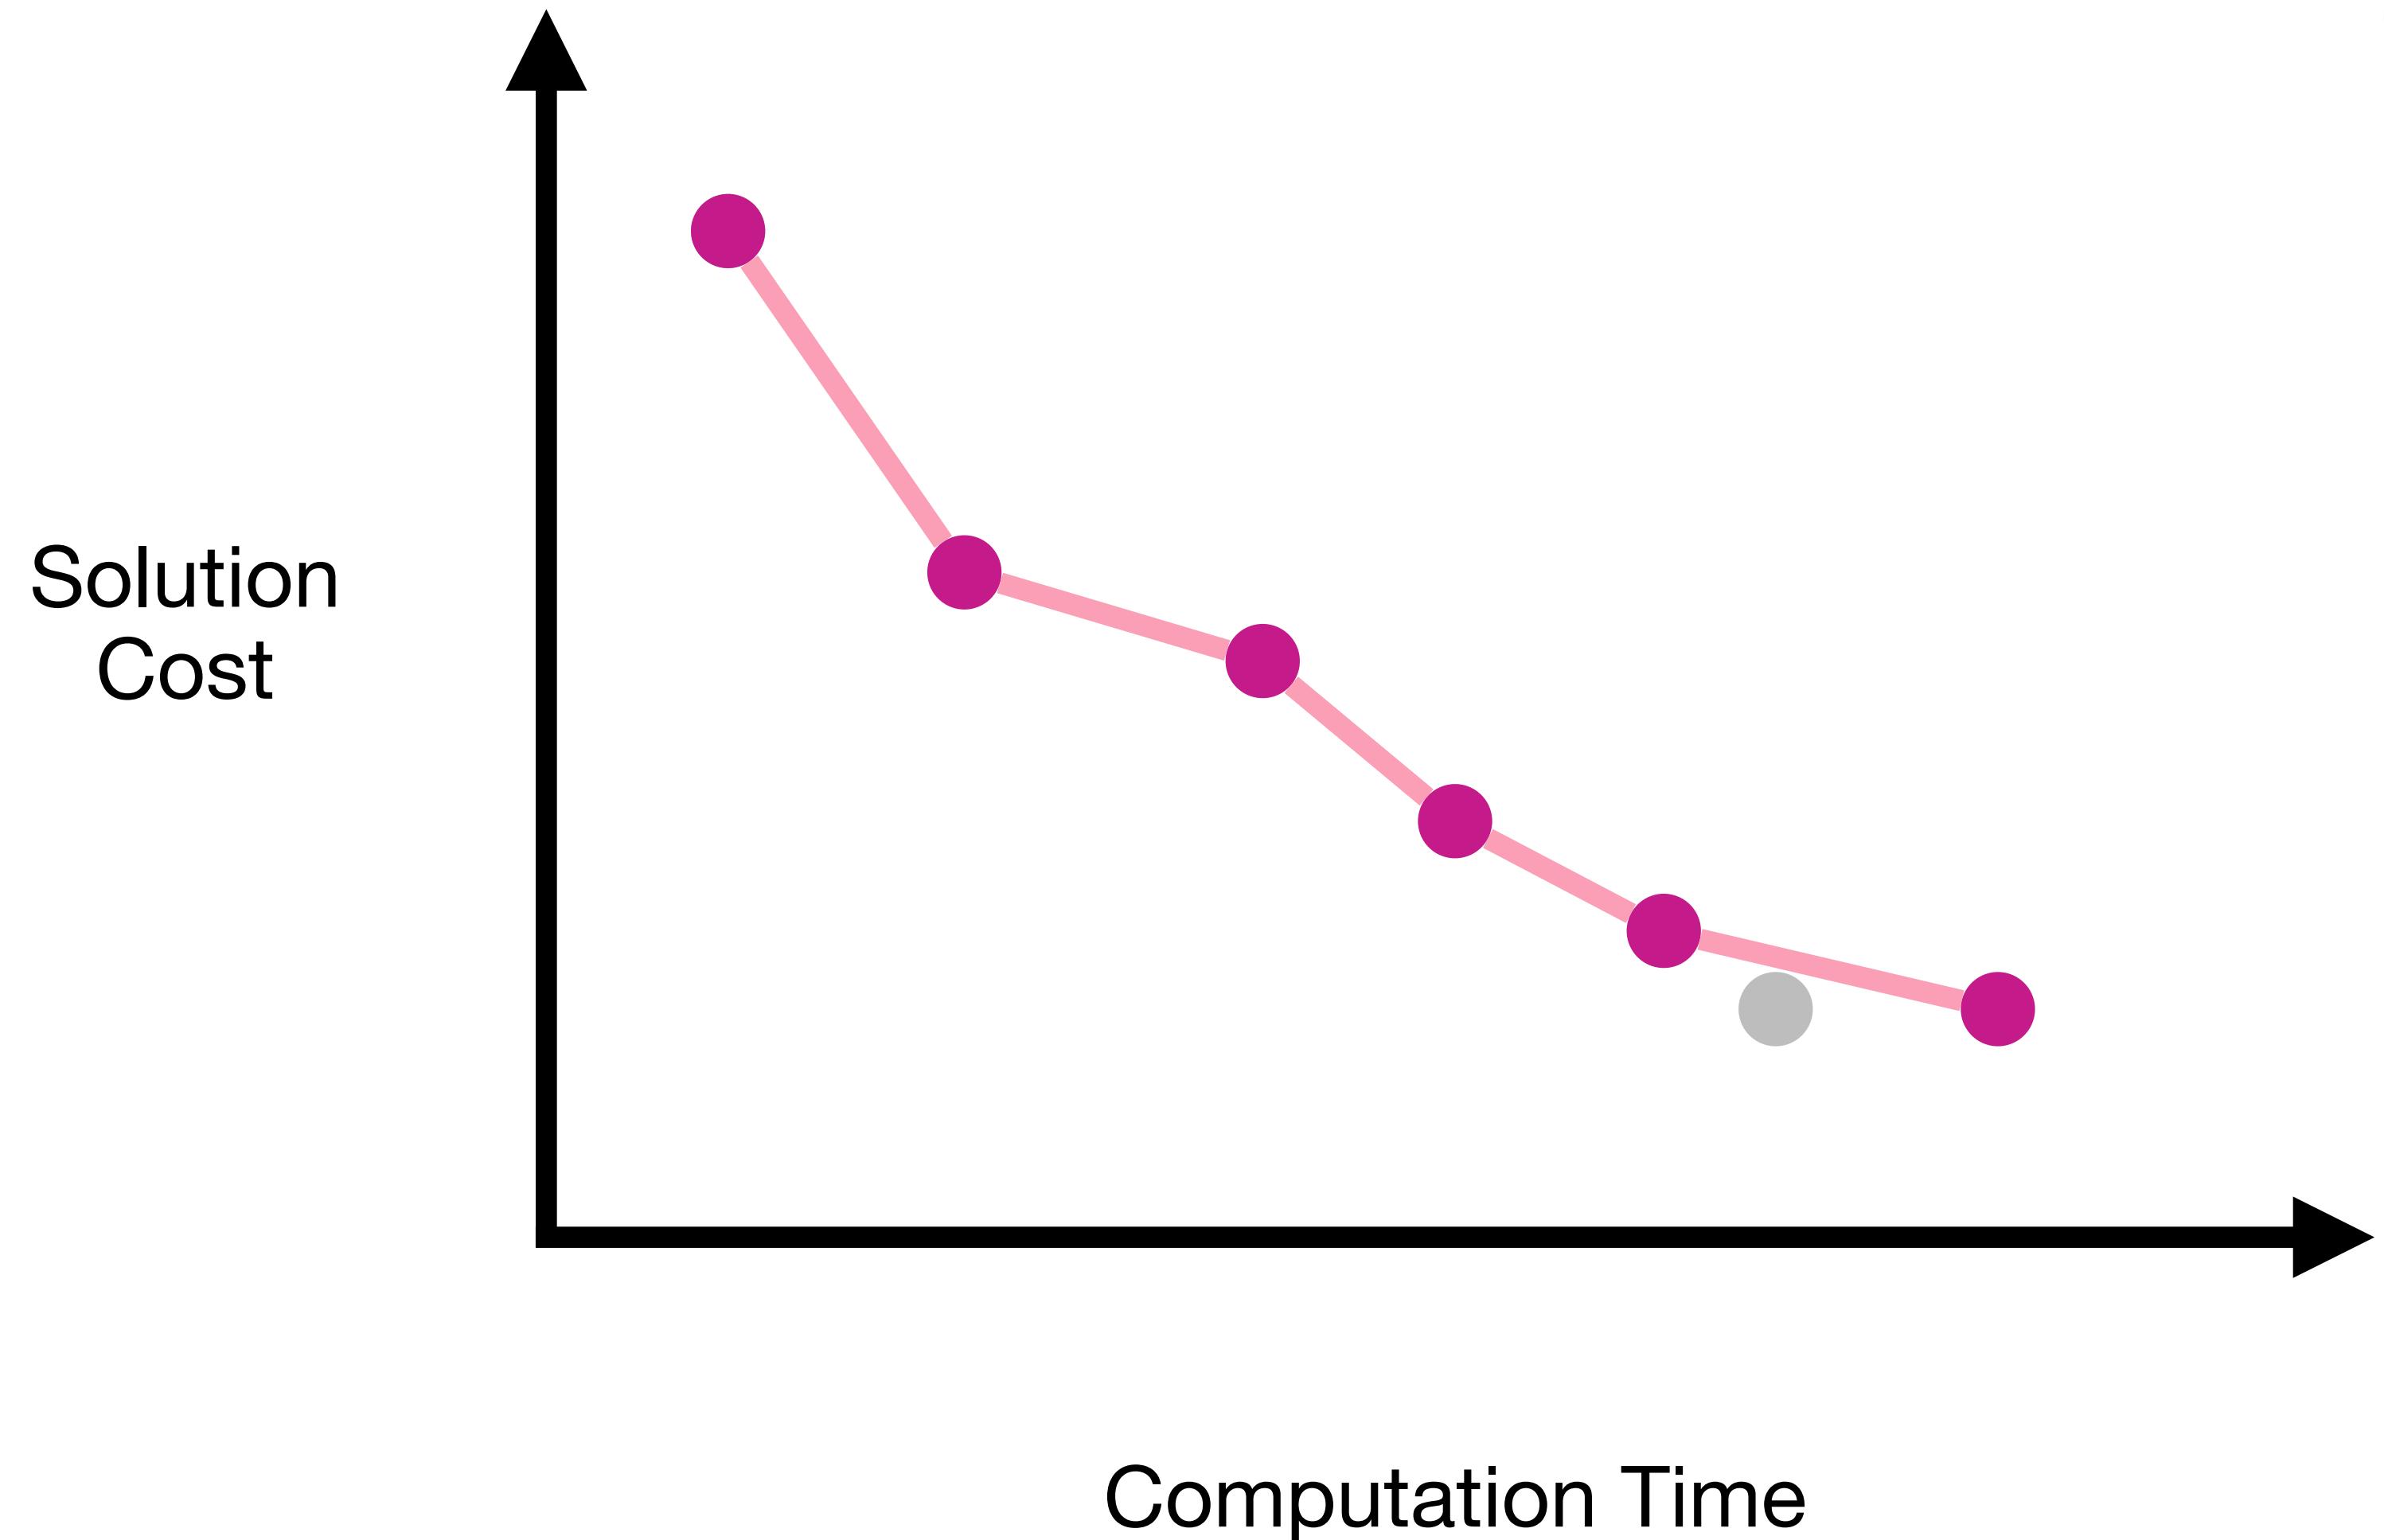
ICAPS 2018 [Best Conference Paper Award Winner]

First Provably Edge-Optimal A\*-like Search Algorithm

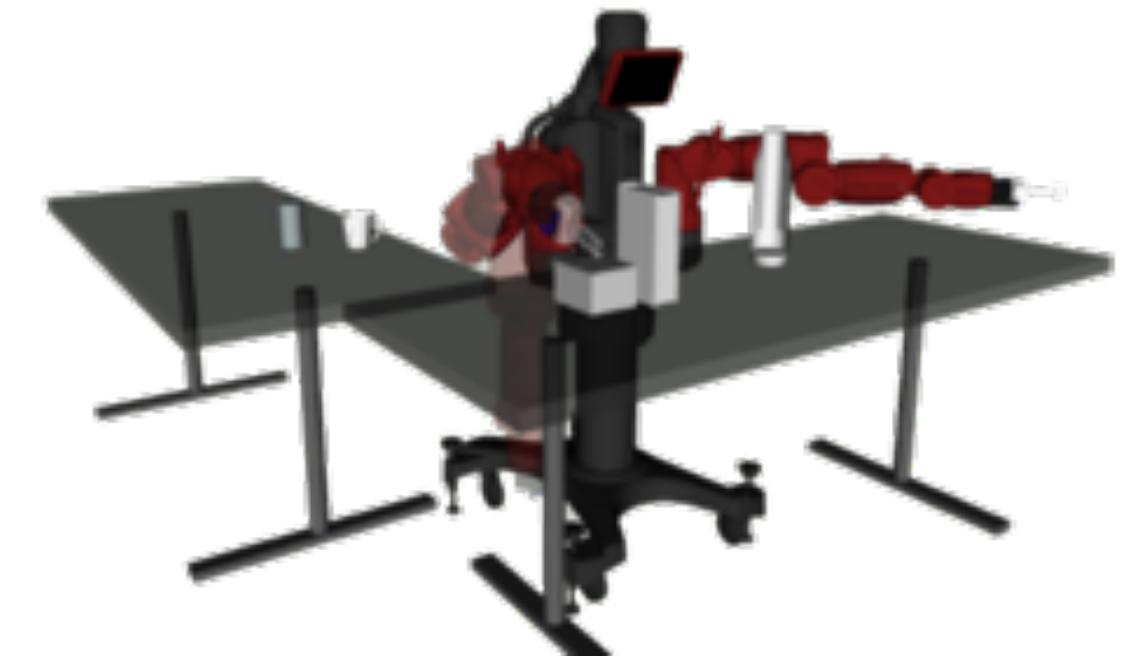
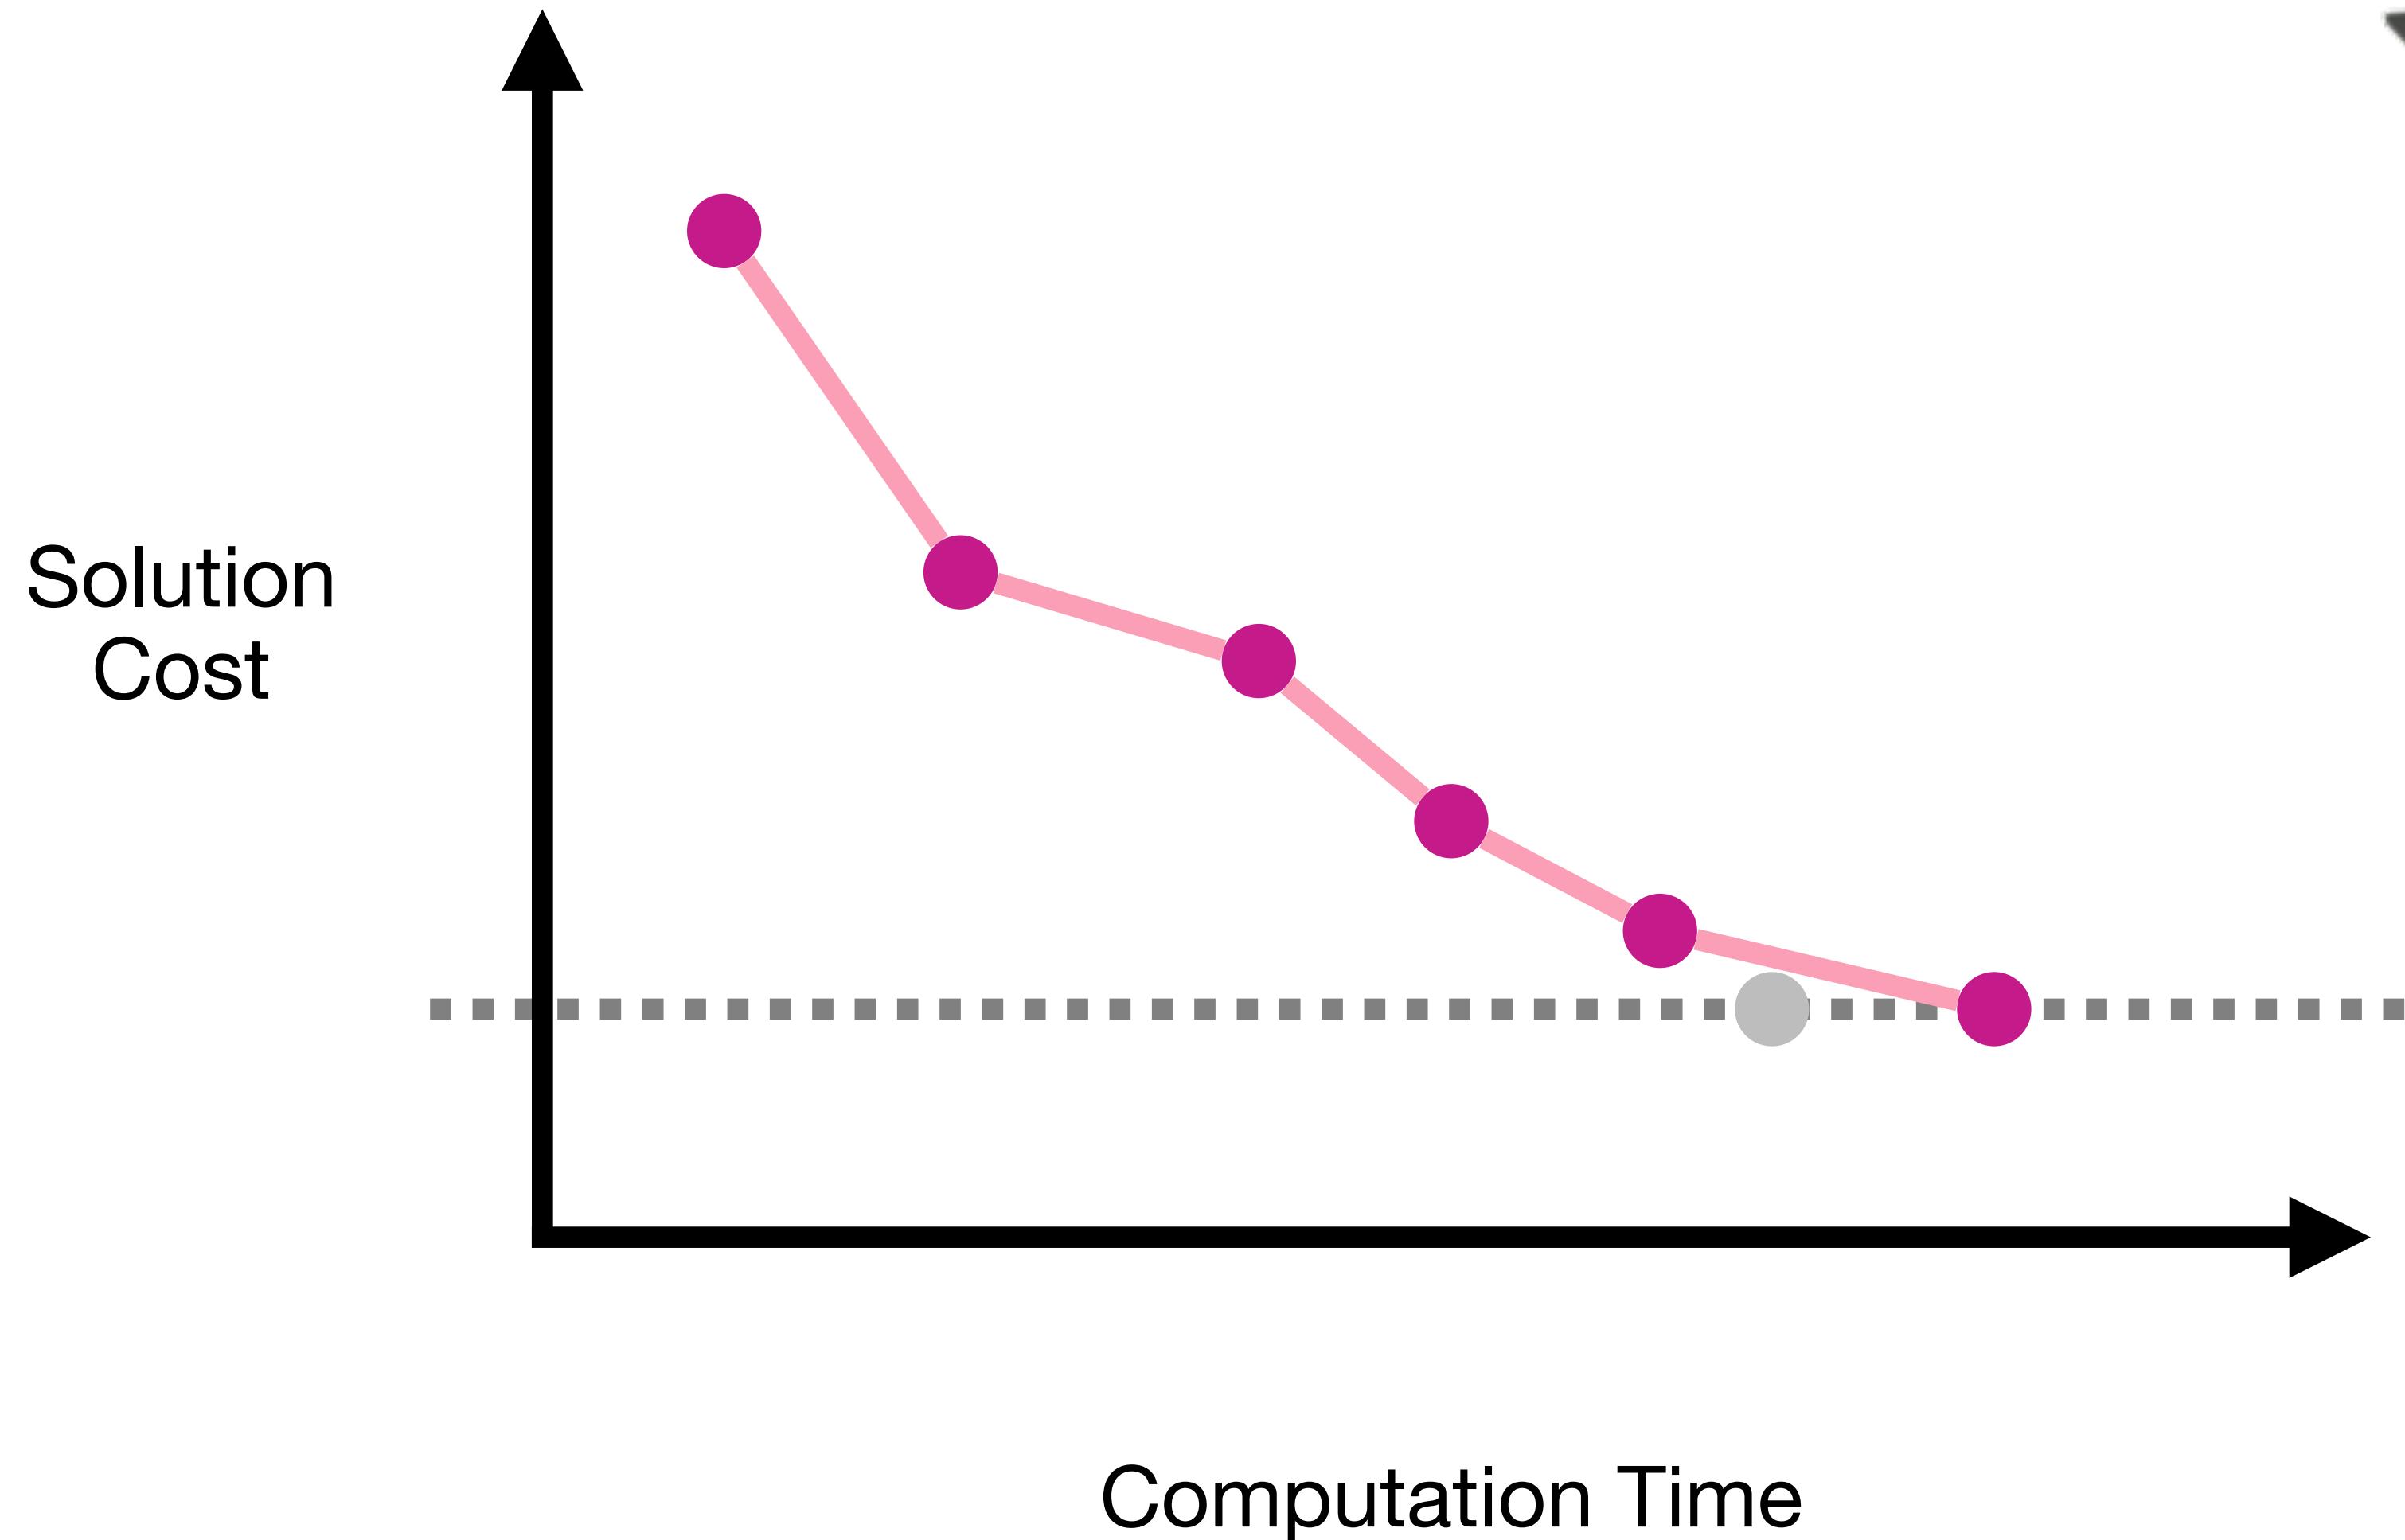
# Anytime Motion Planning



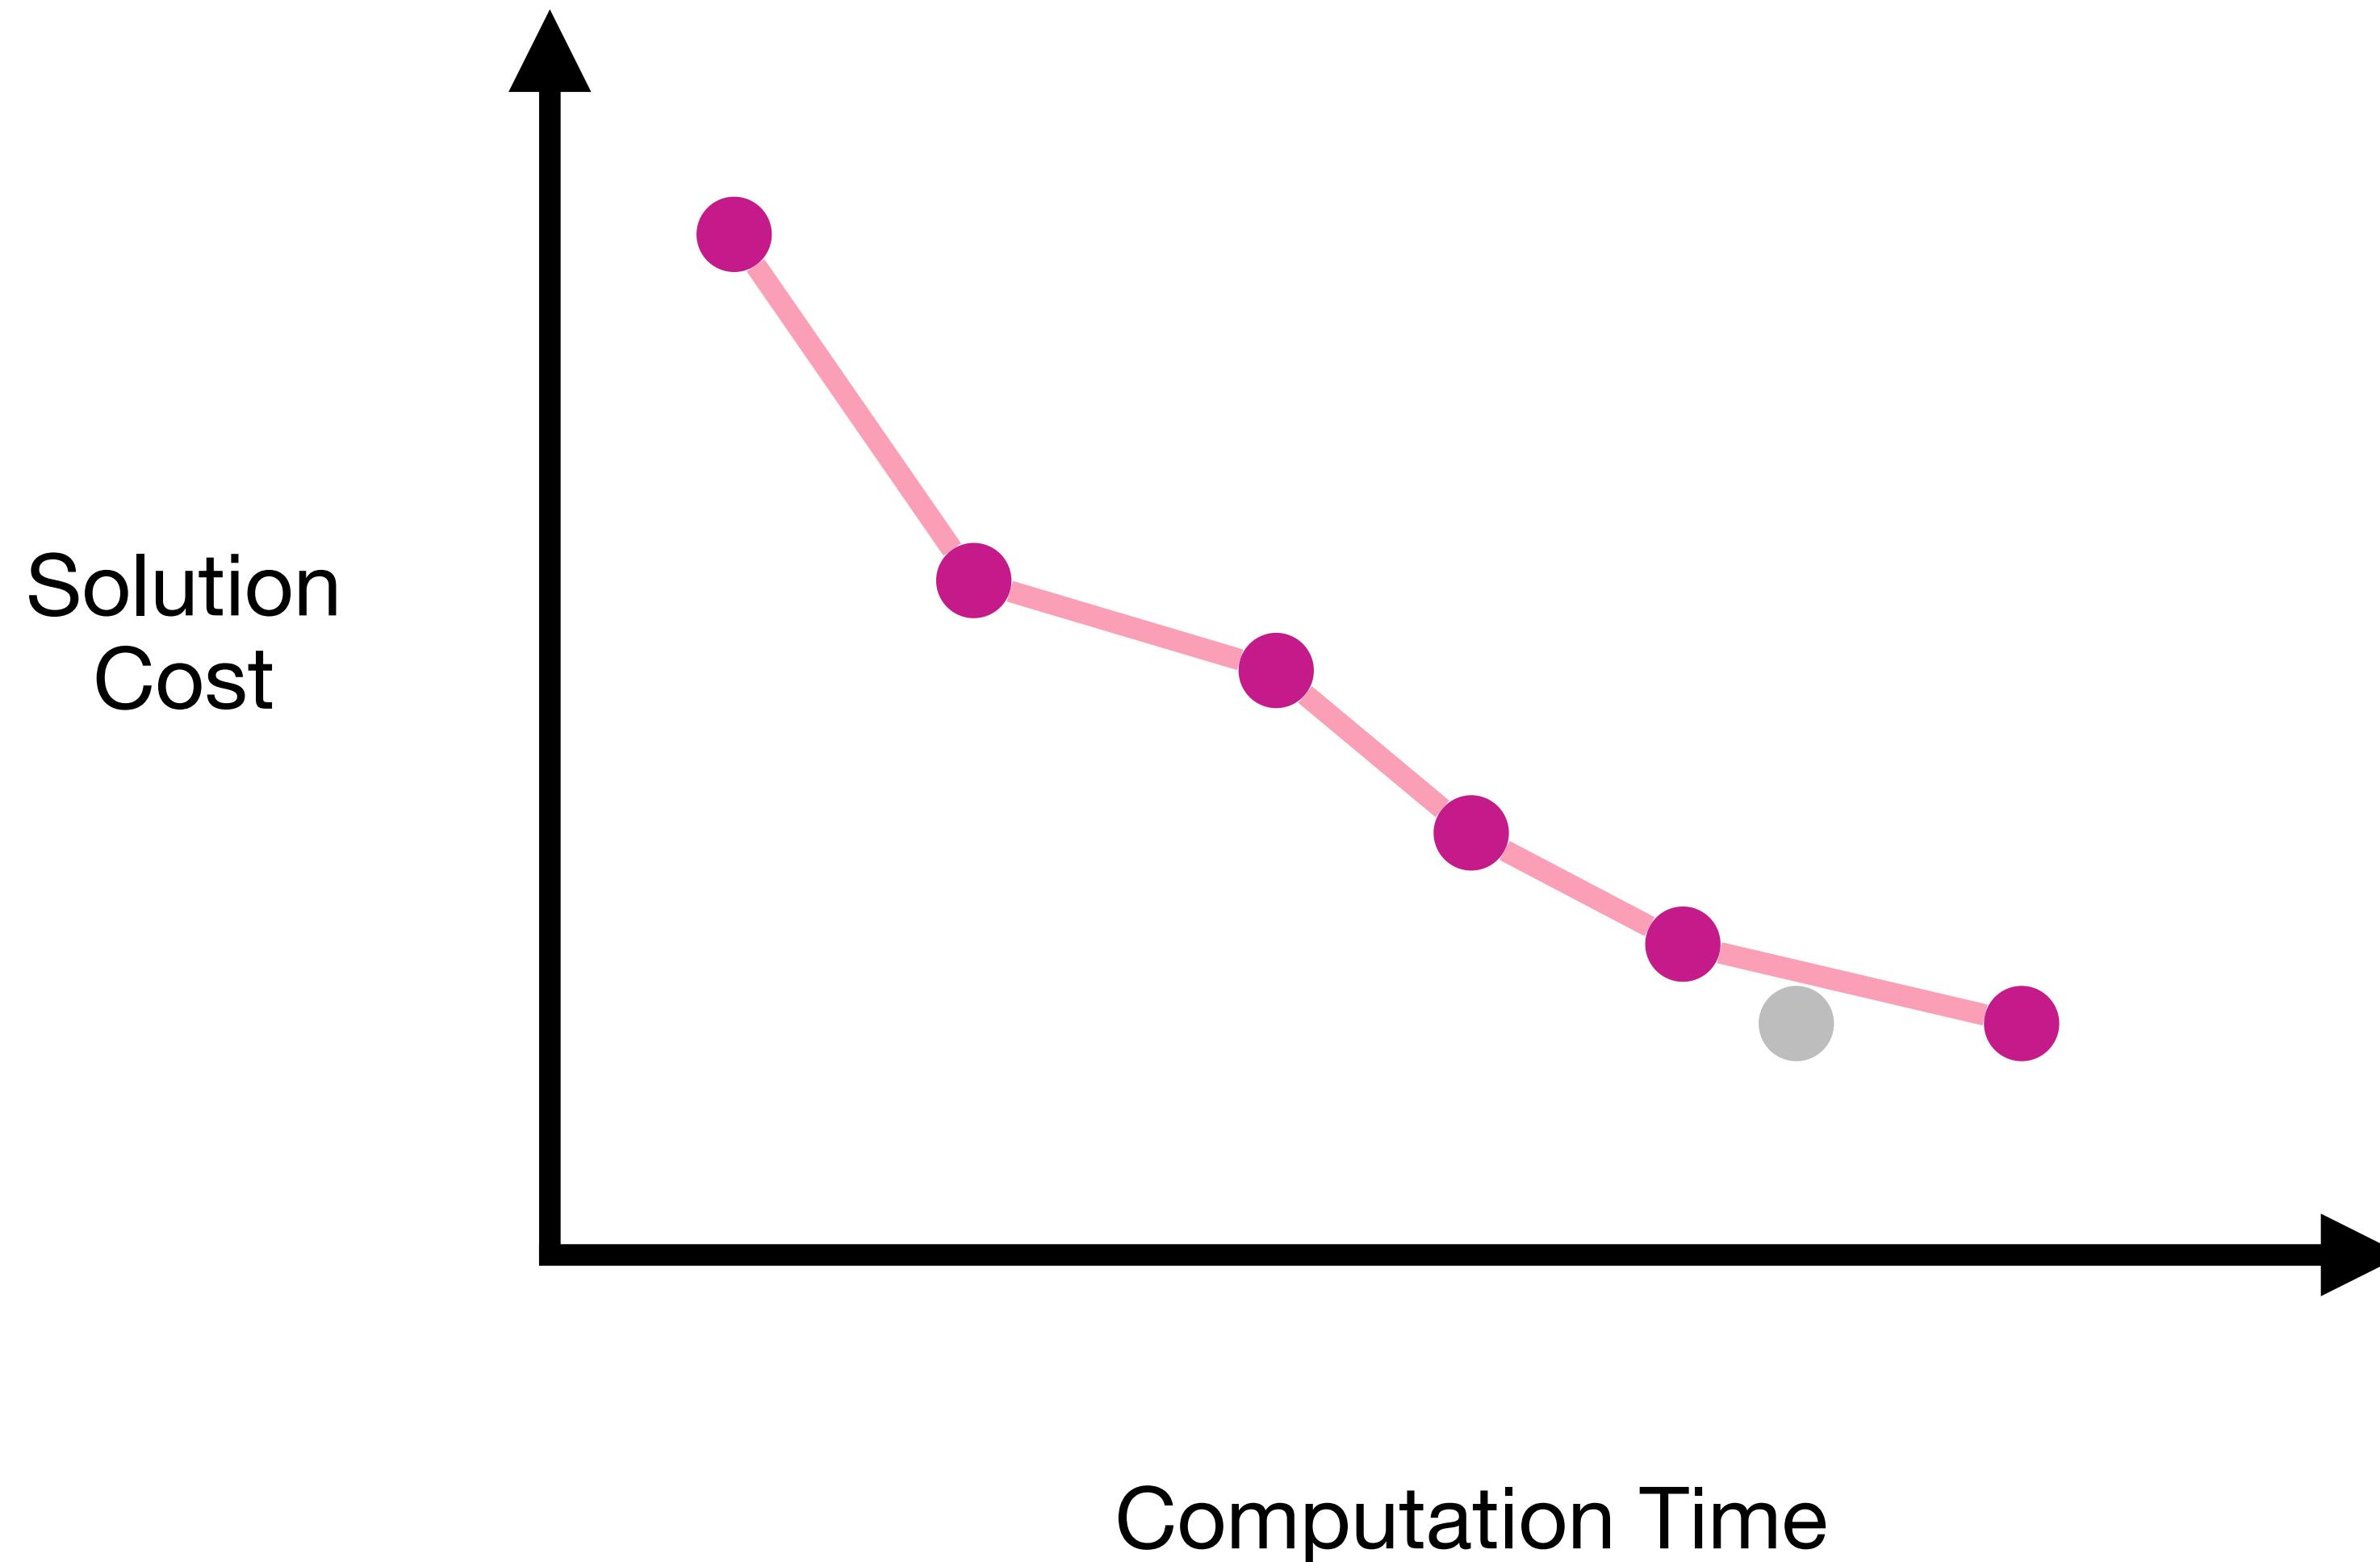
# Anytime Motion Planning



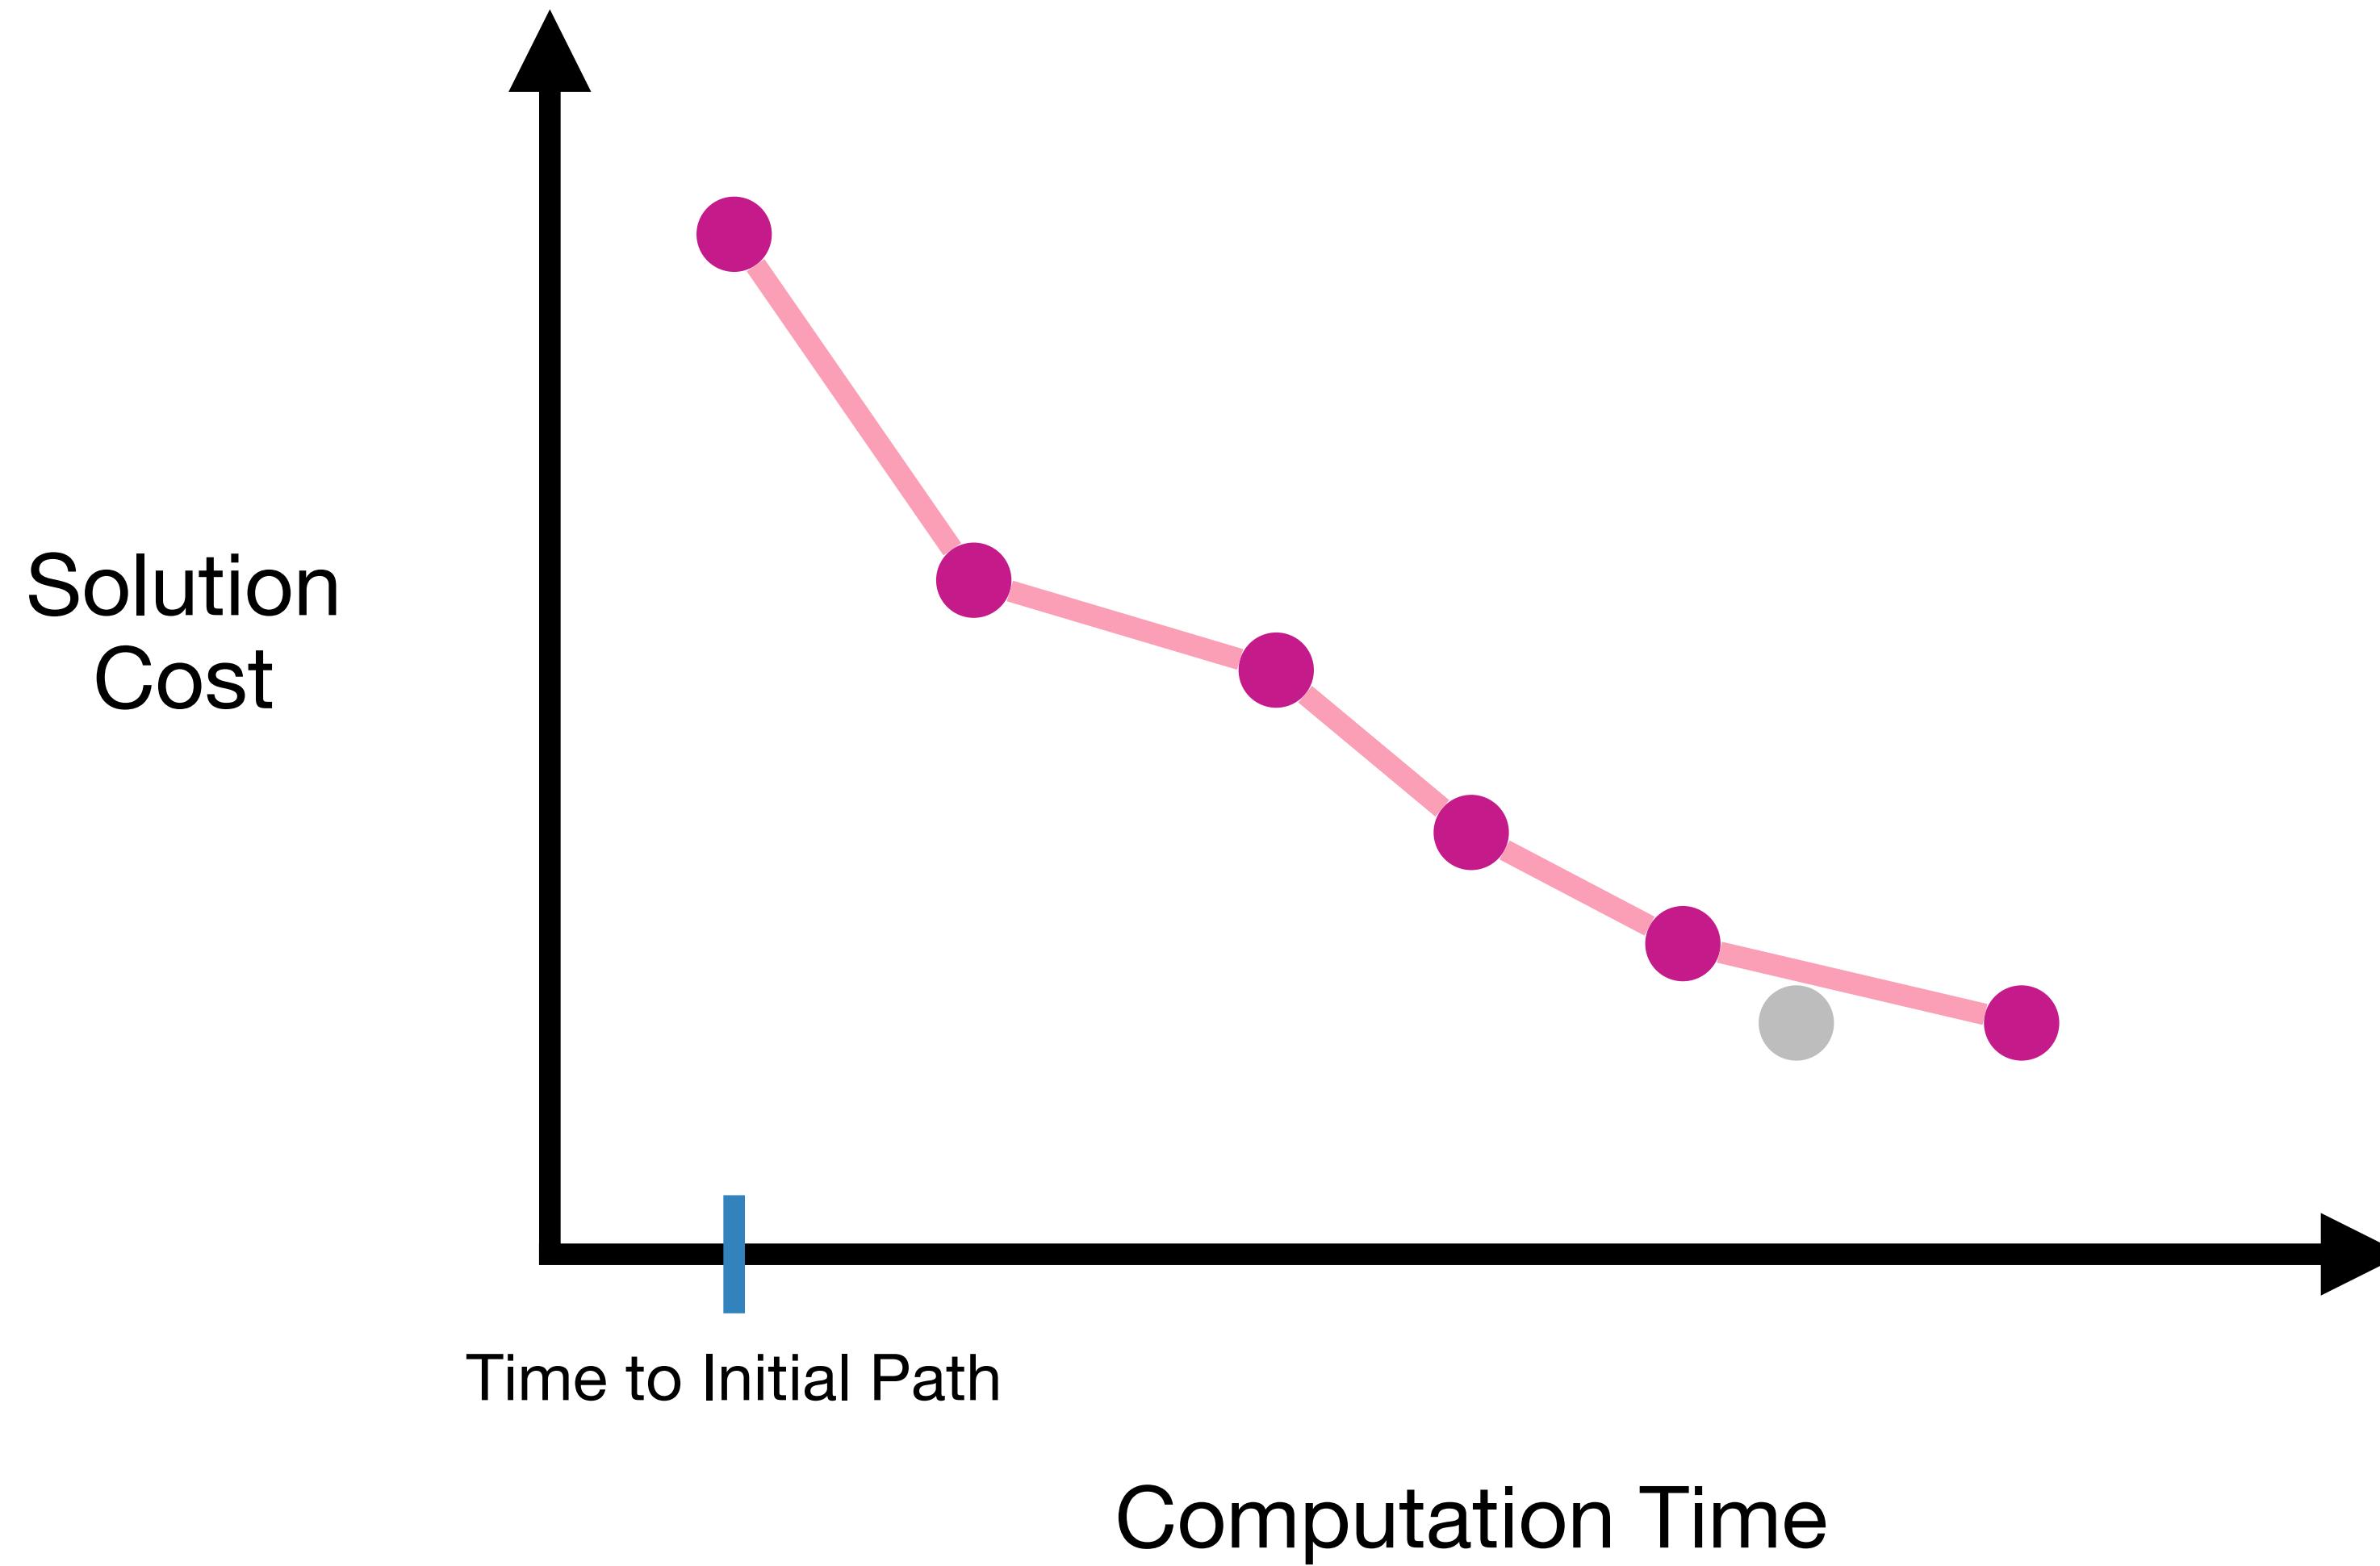
# Will it converge to the shortest path?



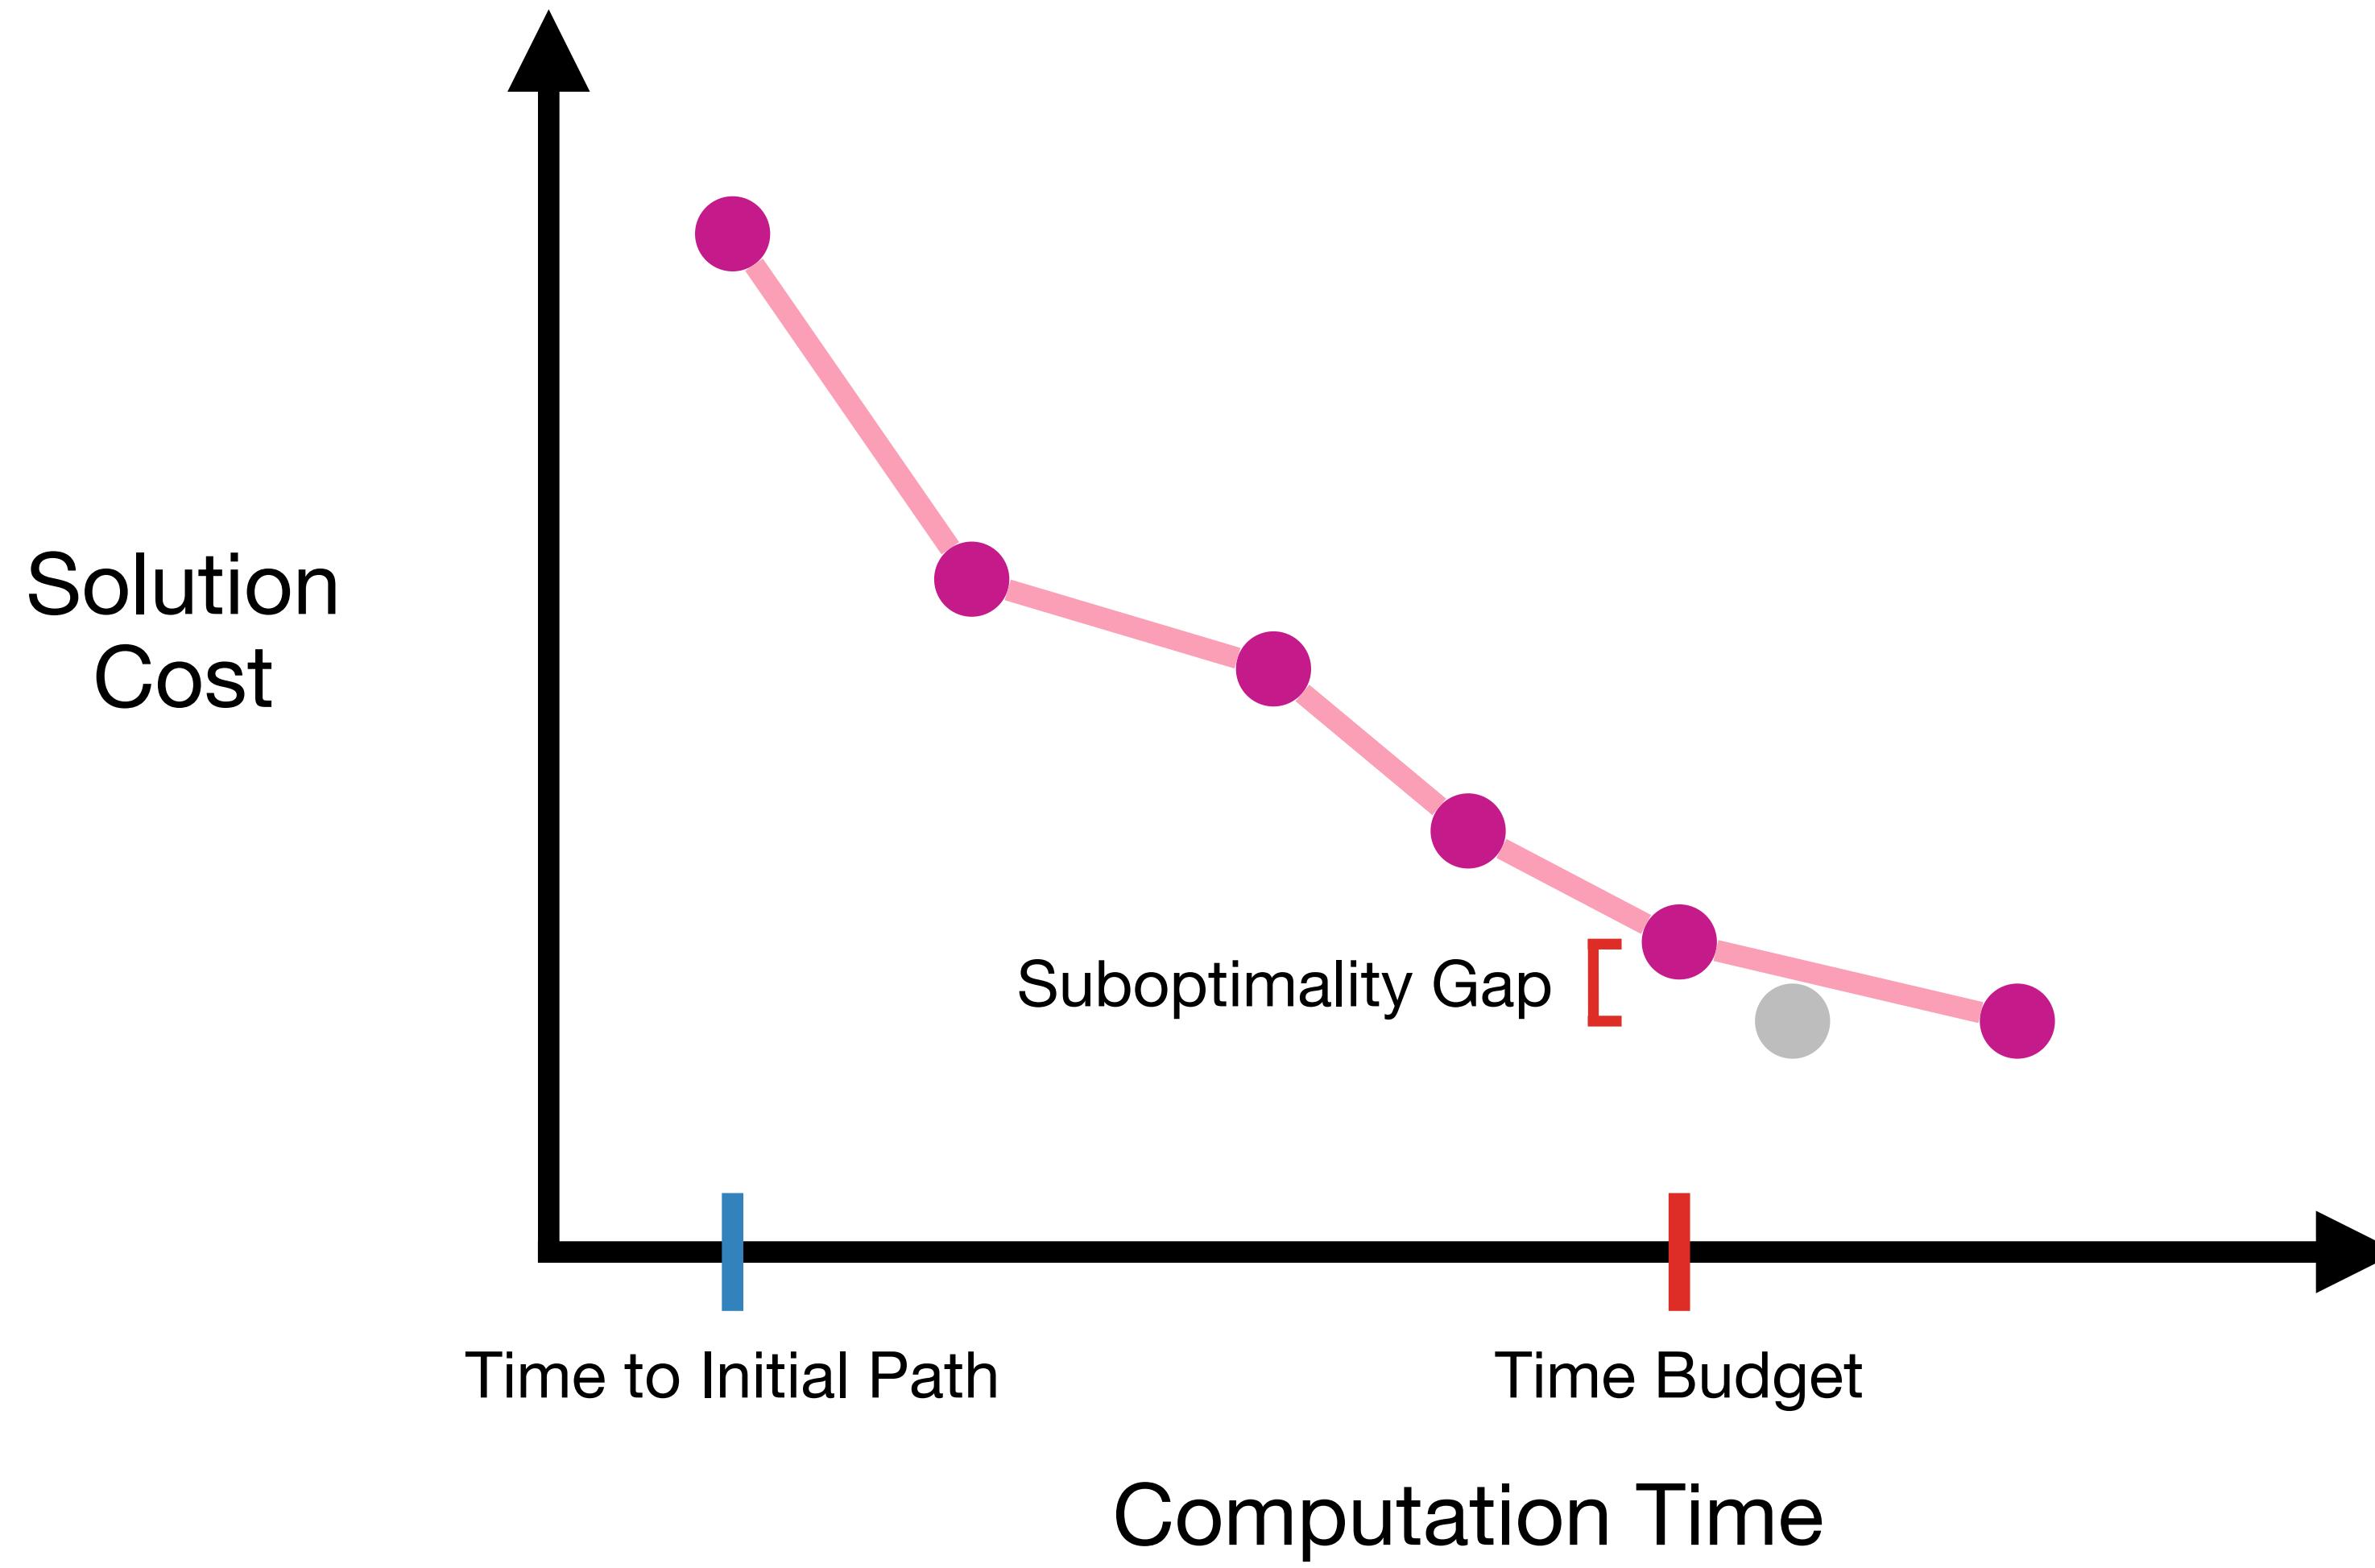
# Beyond Asymptotic Optimality



# Beyond Asymptotic Optimality



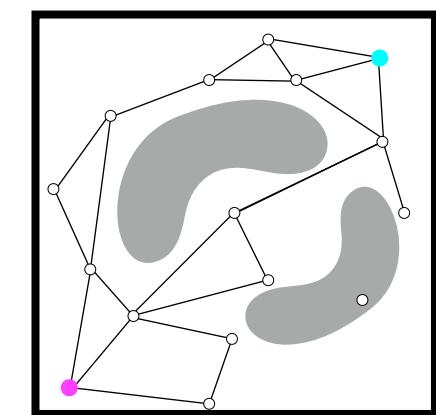
# Beyond Asymptotic Optimality



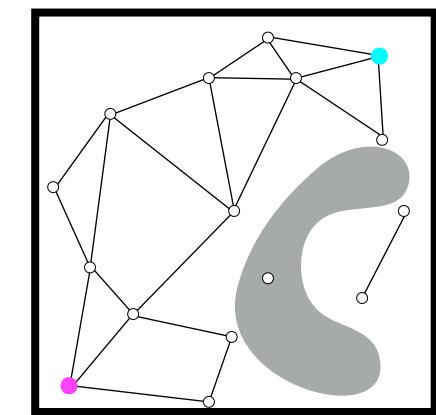
We formalize anytime search as  
Bayesian Reinforcement Learning

# Bayesian Anytime Motion Planning

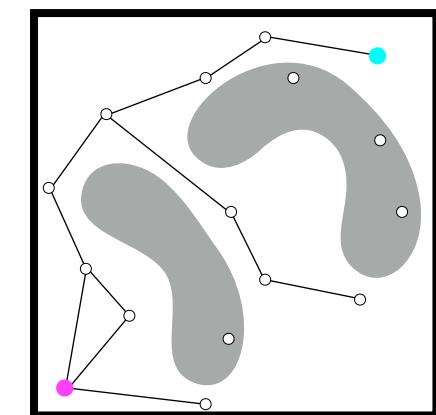
- Evaluating edges uncovers shorter paths  $\xi_t$ 
  - **Anytime Objective:** cumulative path lengths  $\sum_t w(\xi_t)$
- Given prior on collision statuses  $P(\phi)$ 
  - **Bayesian Anytime Objective:**  $\mathbb{E}_{P(\phi)} \sum_t w(\xi_t)$
- Bayesian planning algorithm uses edge evaluation history  $\psi_t$  to compute collision posterior  $P(\phi|\psi_t)$



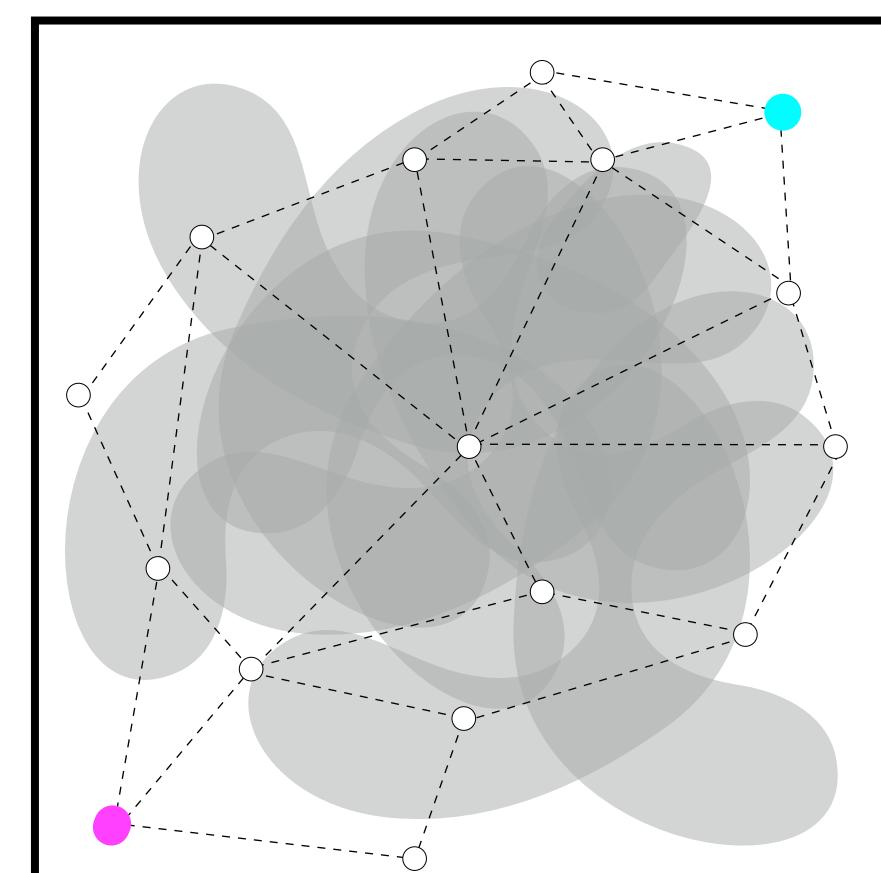
$\phi_1$



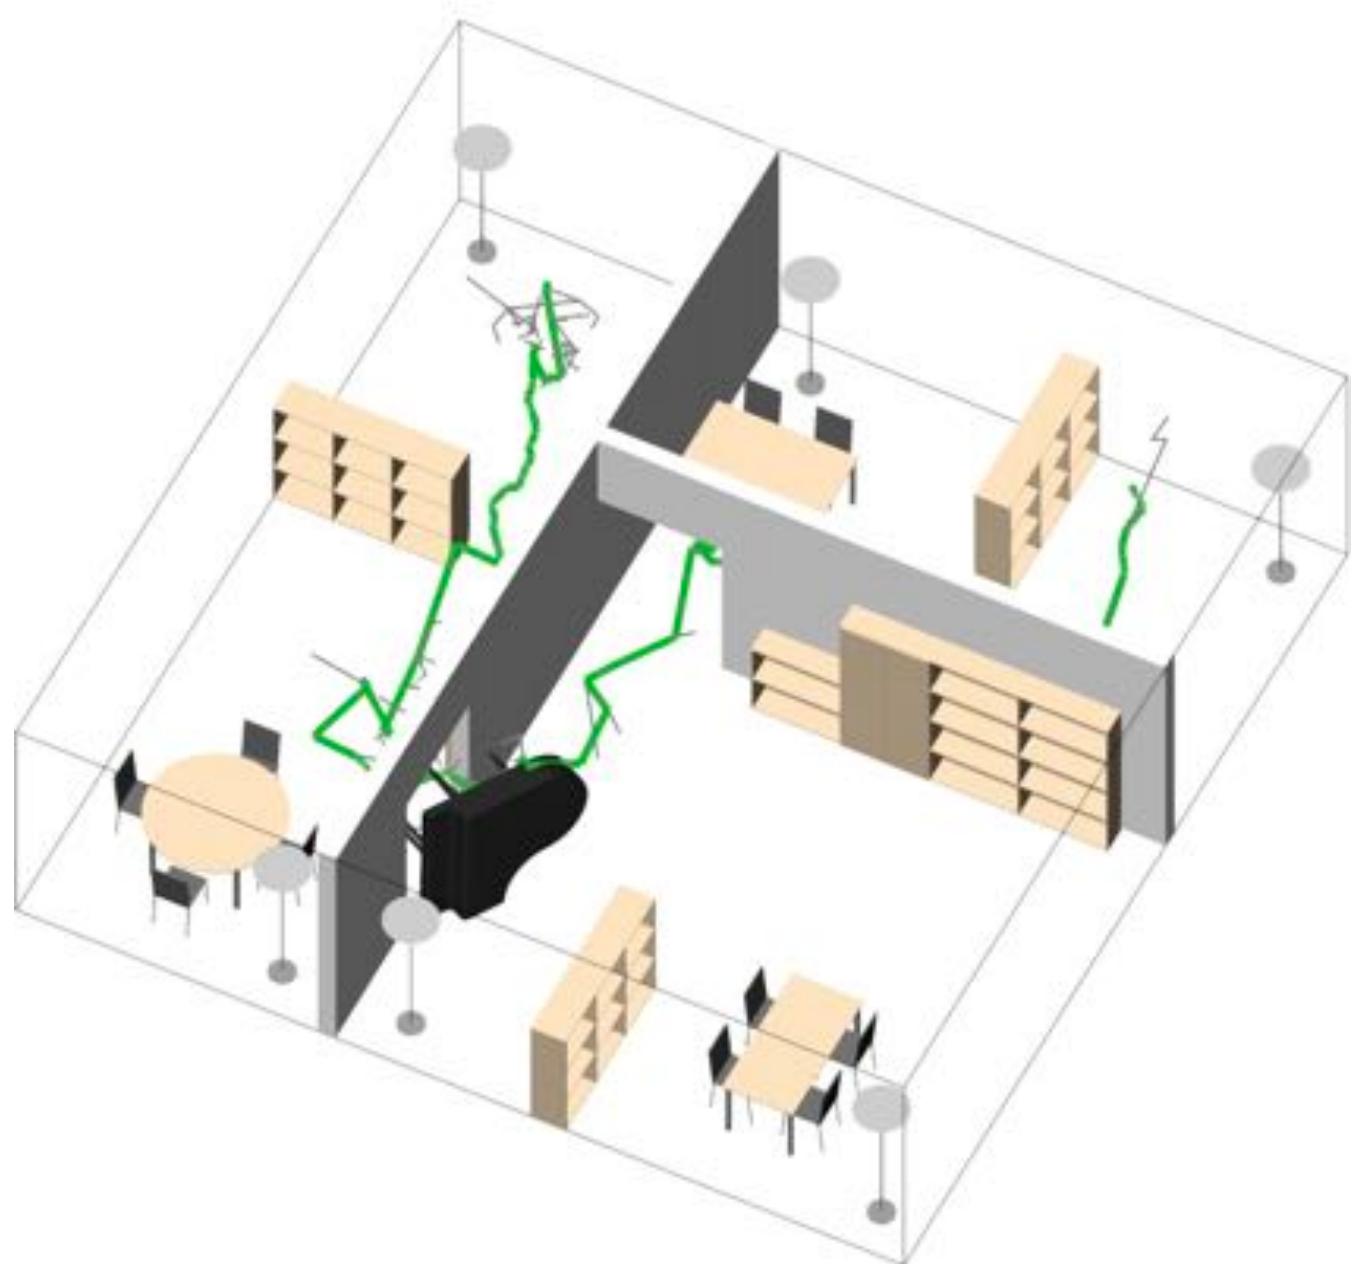
$\phi_2$



$\phi_3$



# The Experienced Piano Movers' Problem



New Piano.  
New House.  
Same Mover.

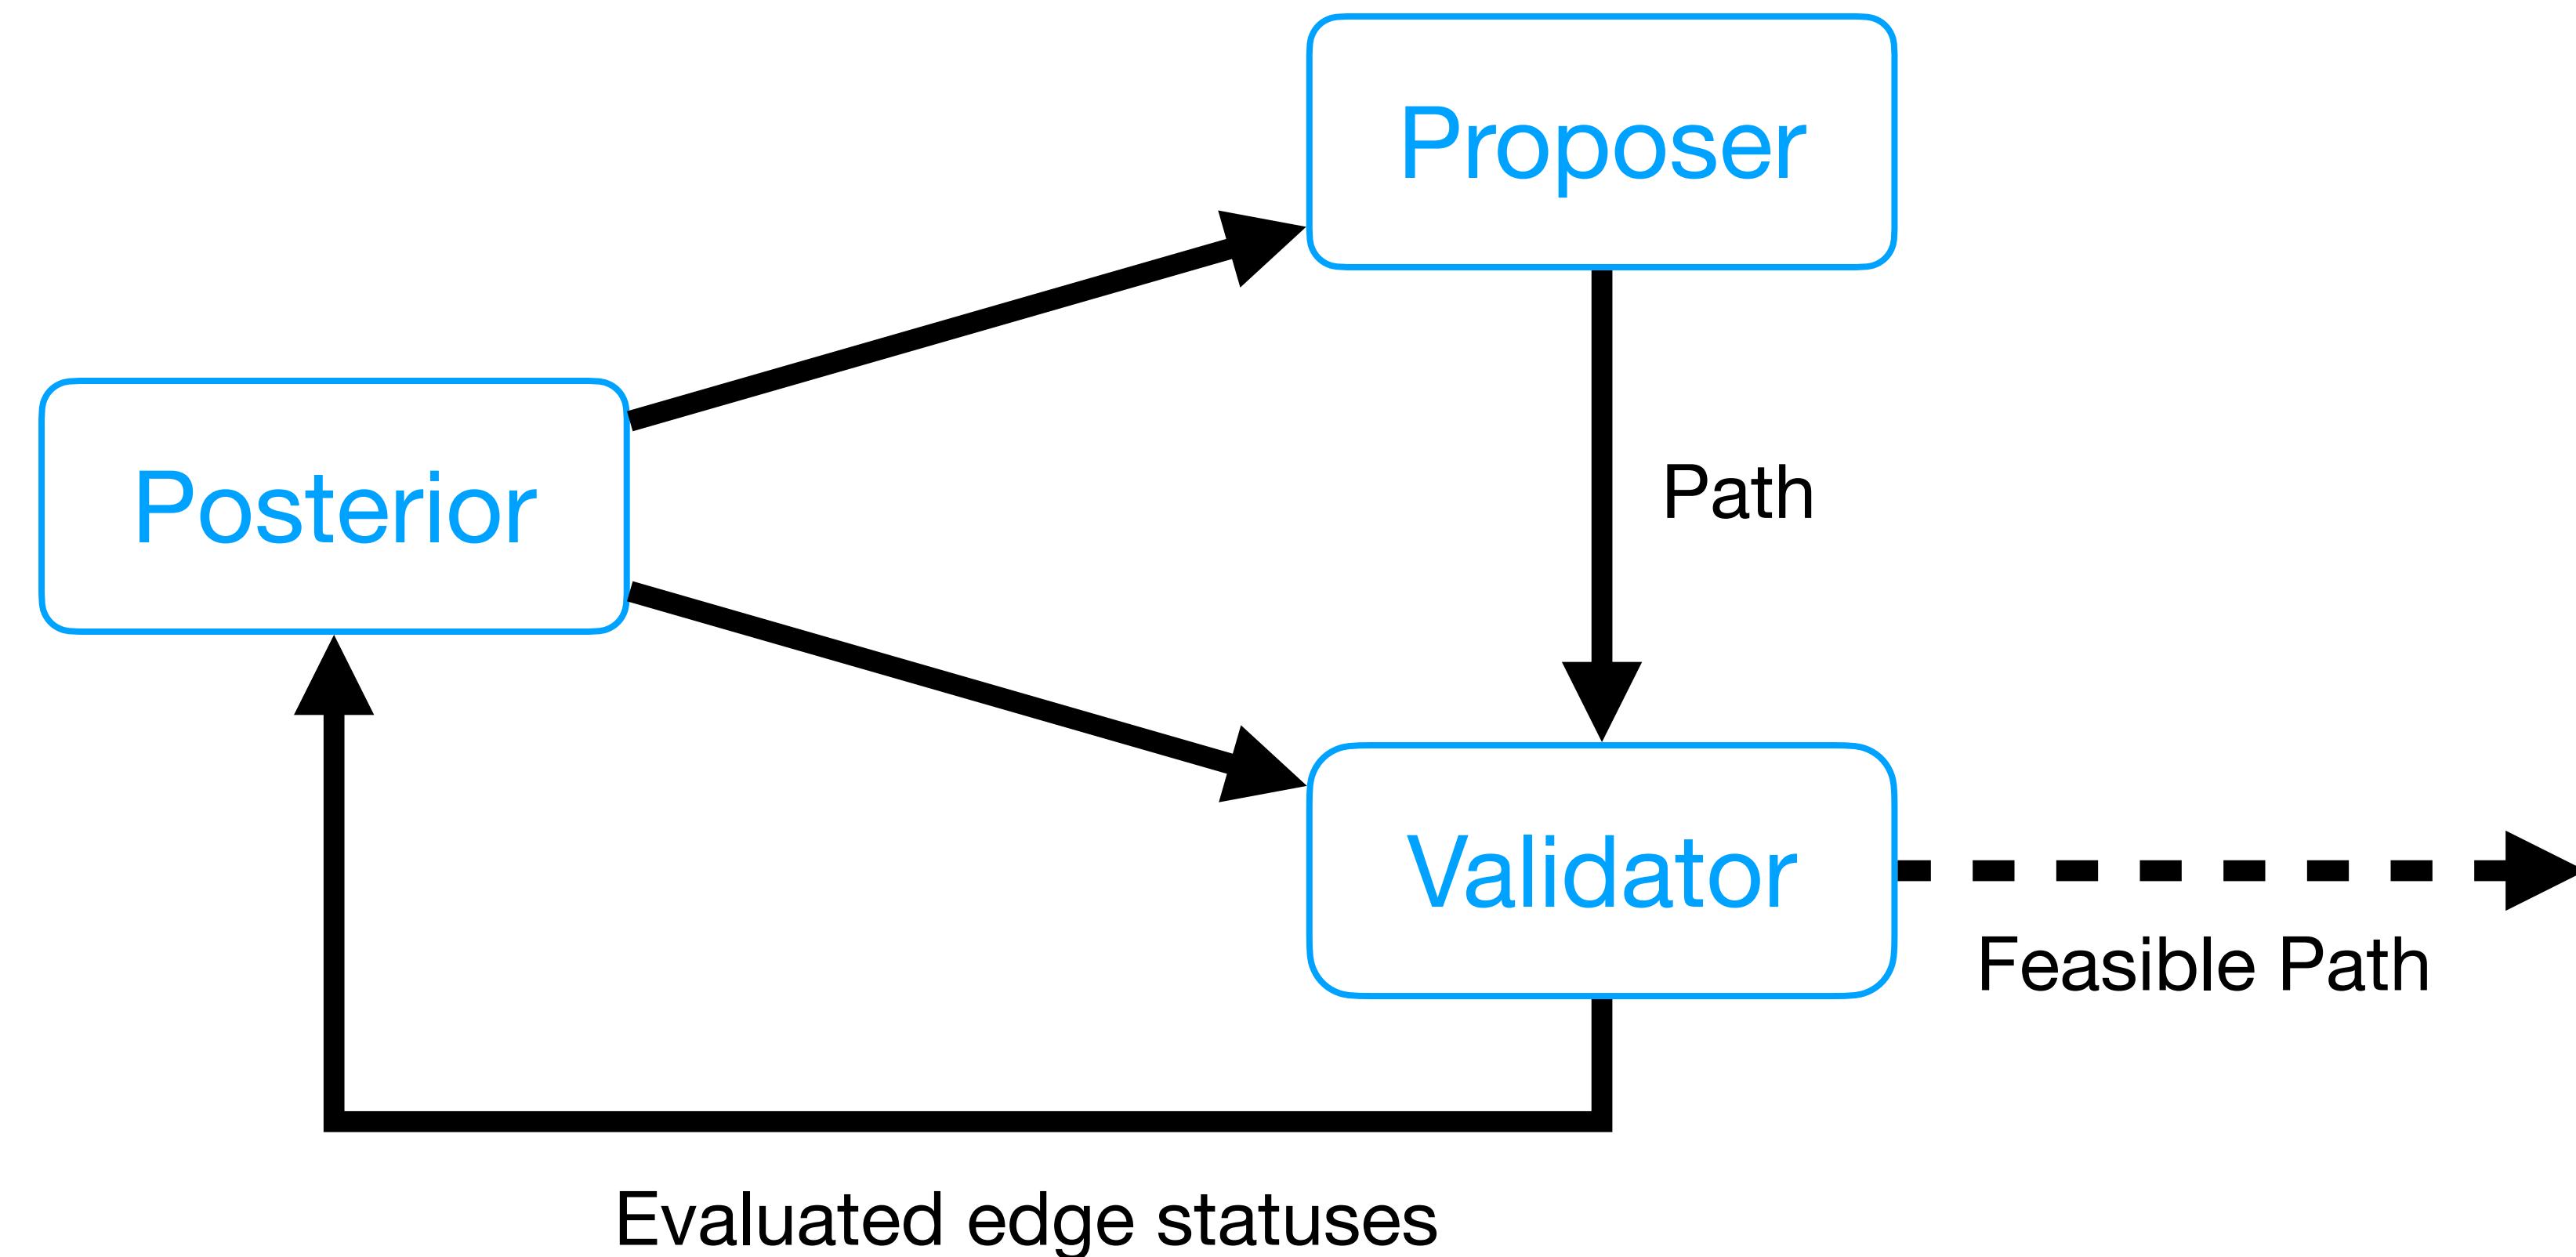
# Bayesian Anytime Motion Planning as Bayesian Reinforcement Learning

- Equivalence to episodic Bayesian RL [Osband et al, 2013]
- Infer unknown MDP  $\mathcal{M}$  through  $m$  repeated episodes

$$\mathbb{E}_{P(\mathcal{M})} [\text{REGRET}(m)] = \mathbb{E}_{P(\mathcal{M})} \sum_{k=1}^m [V(\mu^*) - V(\mu_k)] = \mathbb{E}_{P(\phi)} \sum_{k=1}^m [w(\xi_k) - w(\xi^*)]$$

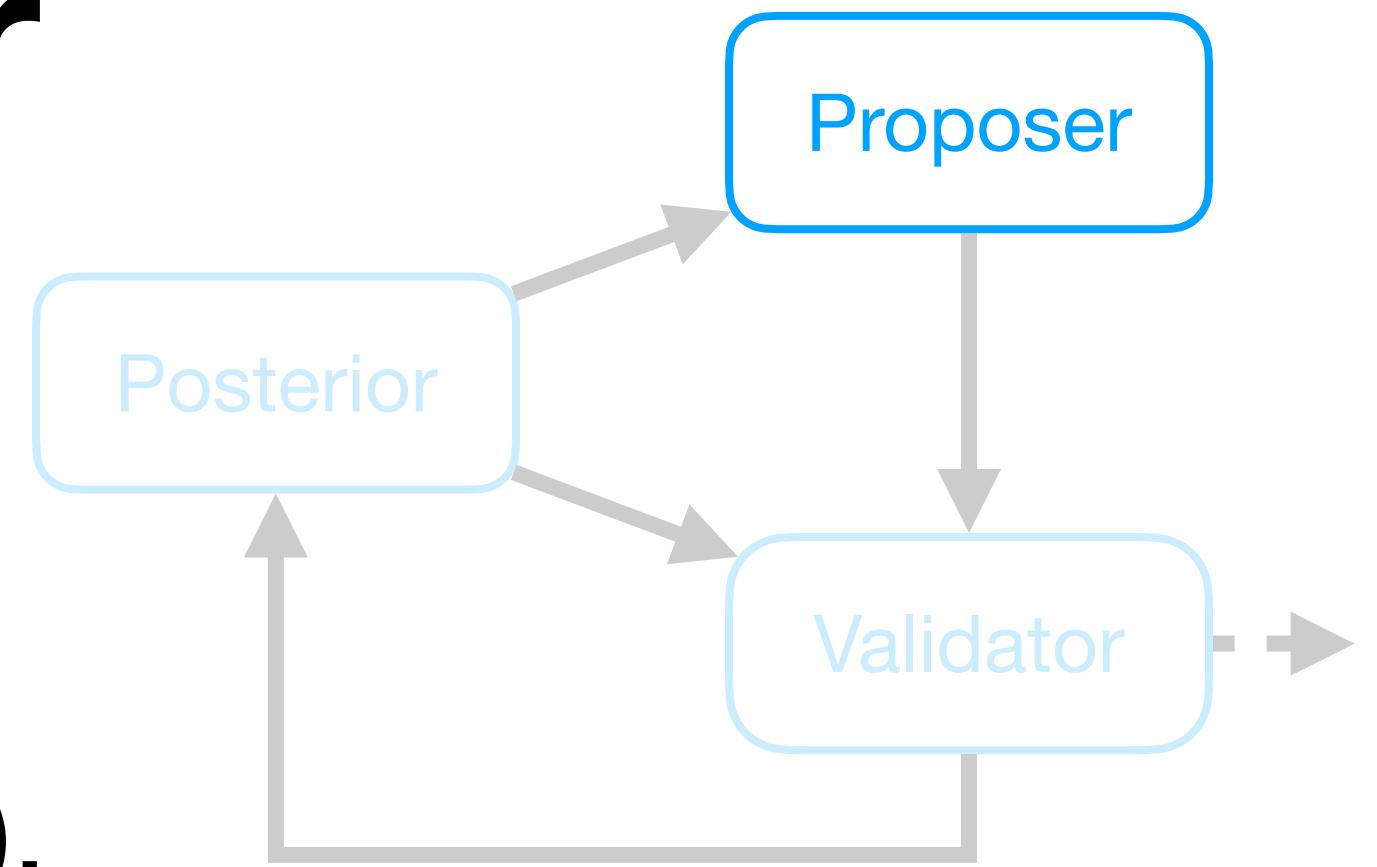
Minimizing Bayesian regret is **equivalent** to minimizing the Bayesian anytime planning objective!

# Experienced Lazy Path Search



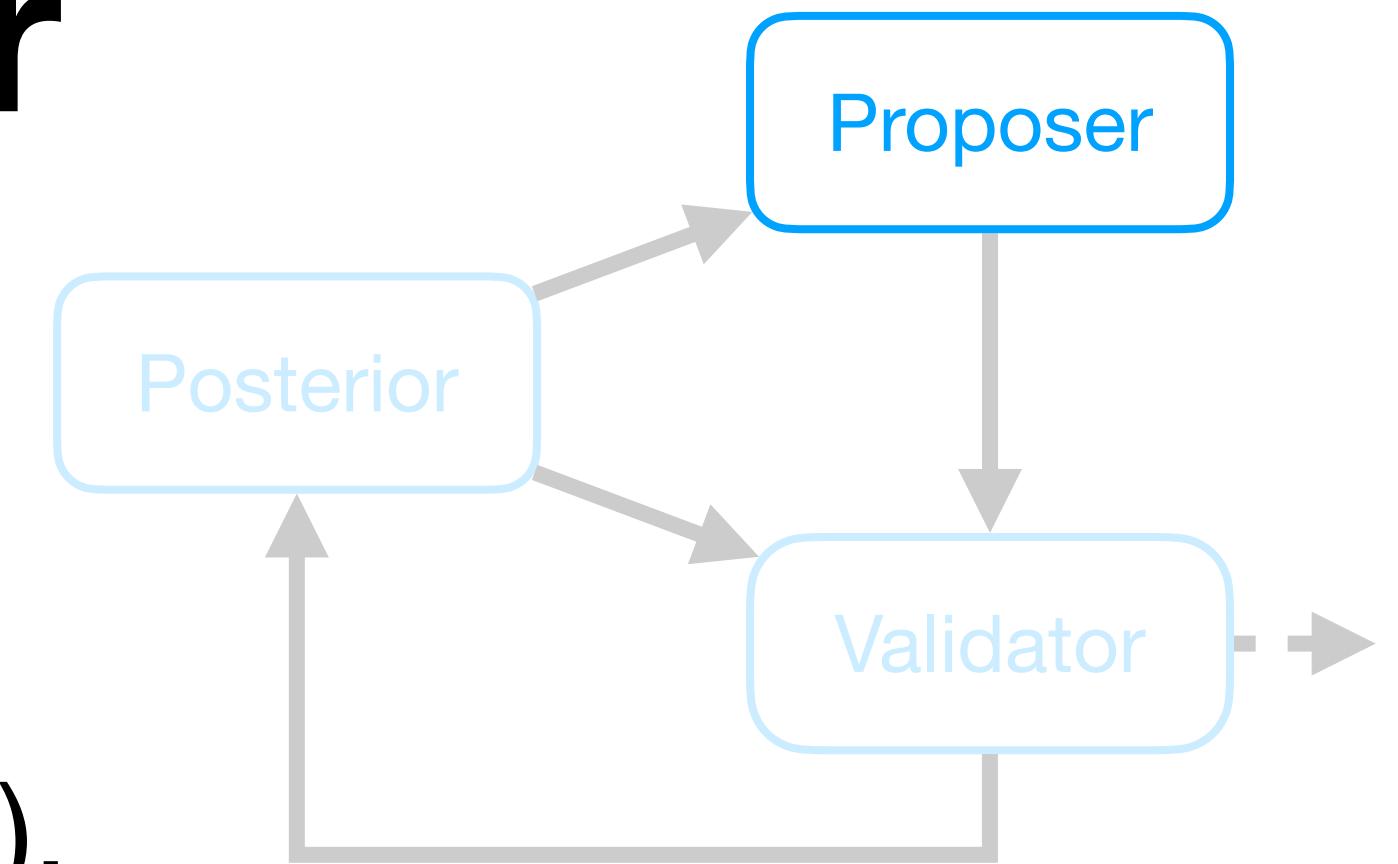
# The Posterior Sampling Proposer

- **Posterior Sampling for Motion Planning (PSMP):** propose paths according to probability they are optimal
  - Idea from multi-armed bandits (as Thompson sampling), Posterior Sampling for RL



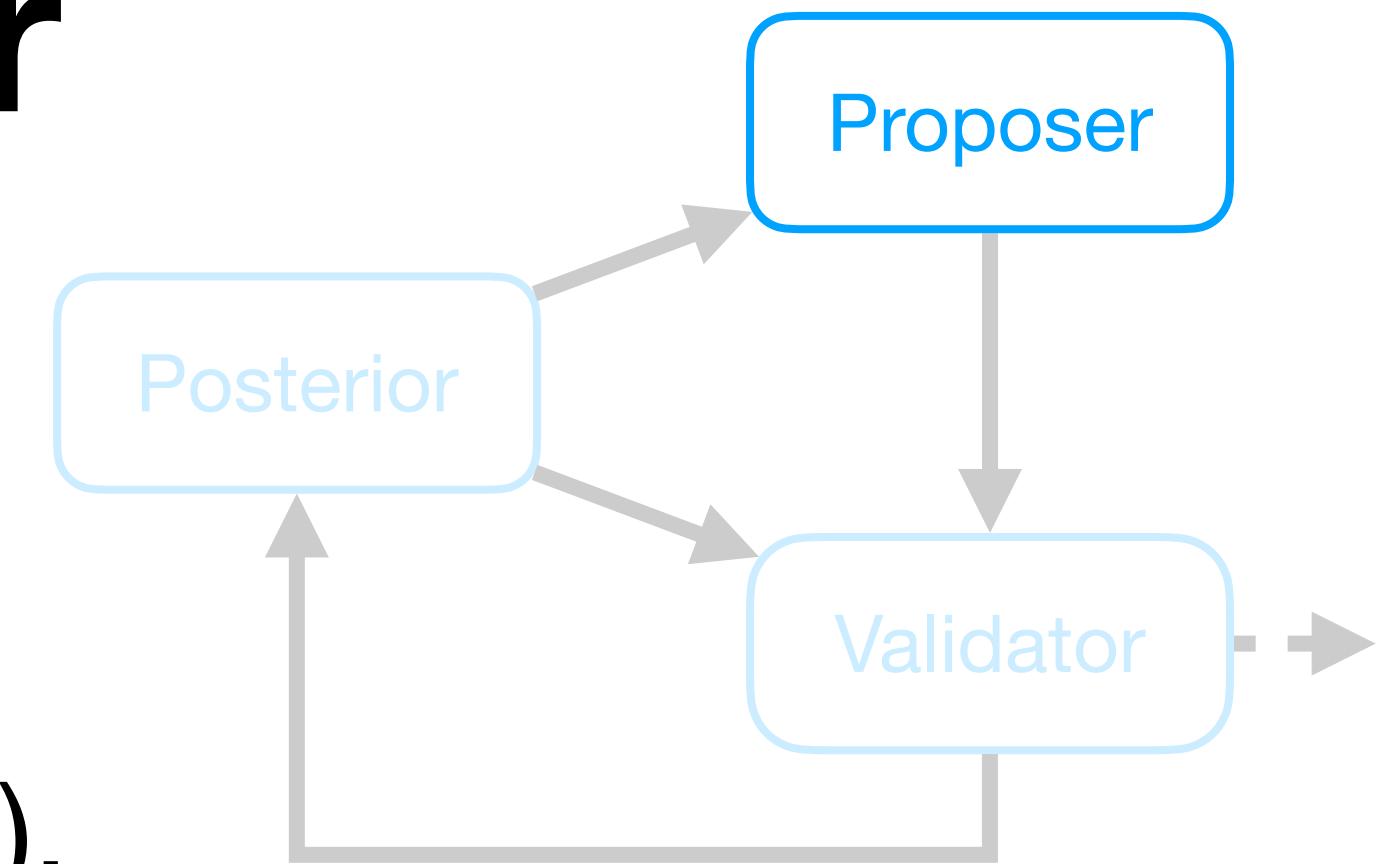
# The Posterior Sampling Proposer

- **Posterior Sampling for Motion Planning (PSMP):** propose paths according to probability they are optimal
  - Idea from multi-armed bandits (as Thompson sampling), Posterior Sampling for RL [Osband et al, 2013]
  - First anytime motion planning algorithm with Bayesian regret bounds
    - Analysis adapts [Osband et al, 2013] for deterministic MDPs
    - Bound of  $O(\tau \sqrt{|V||E|T \log(|V||E|T)})$  matches known lower bounds



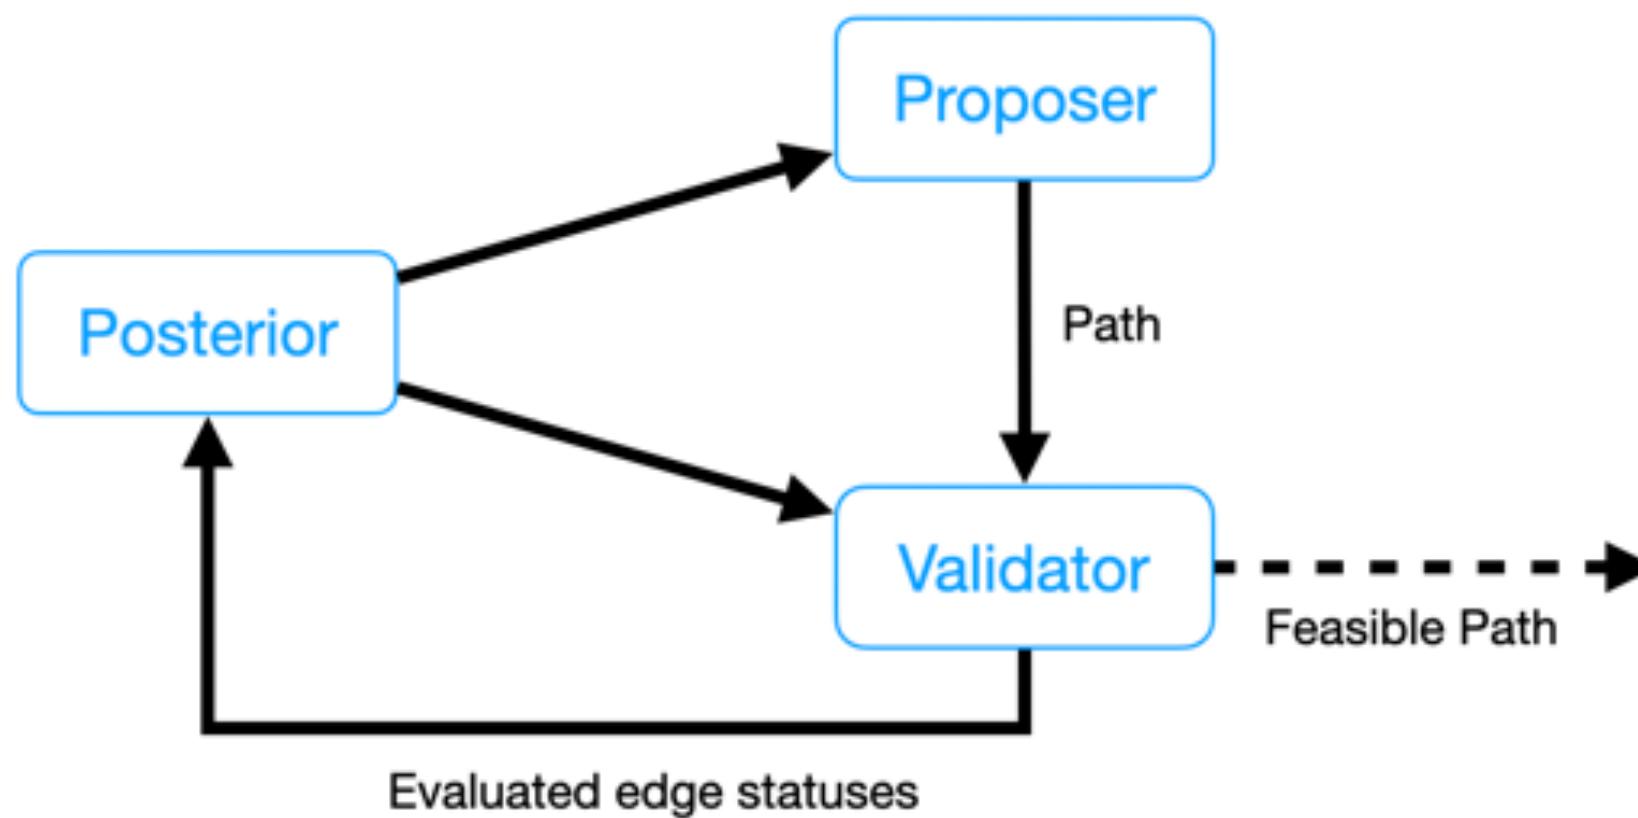
# The Posterior Sampling Proposer

- **Posterior Sampling for Motion Planning (PSMP):** propose paths according to probability they are optimal
  - Idea from multi-armed bandits (as Thompson sampling), Posterior Sampling for RL [Osband et al, 2013]
  - First anytime motion planning algorithm with Bayesian regret bounds
    - Analysis adapts [Osband et al, 2013] for deterministic MDPs
    - Bound of  $O(\tau \sqrt{|V||E|T \log(|V||E|T)})$  matches known lower bounds
    - Solves one shortest path problem per proposal



# But Whatever Happened to Optimism?!

Optimism in the  
Face of Uncertainty (OFU)



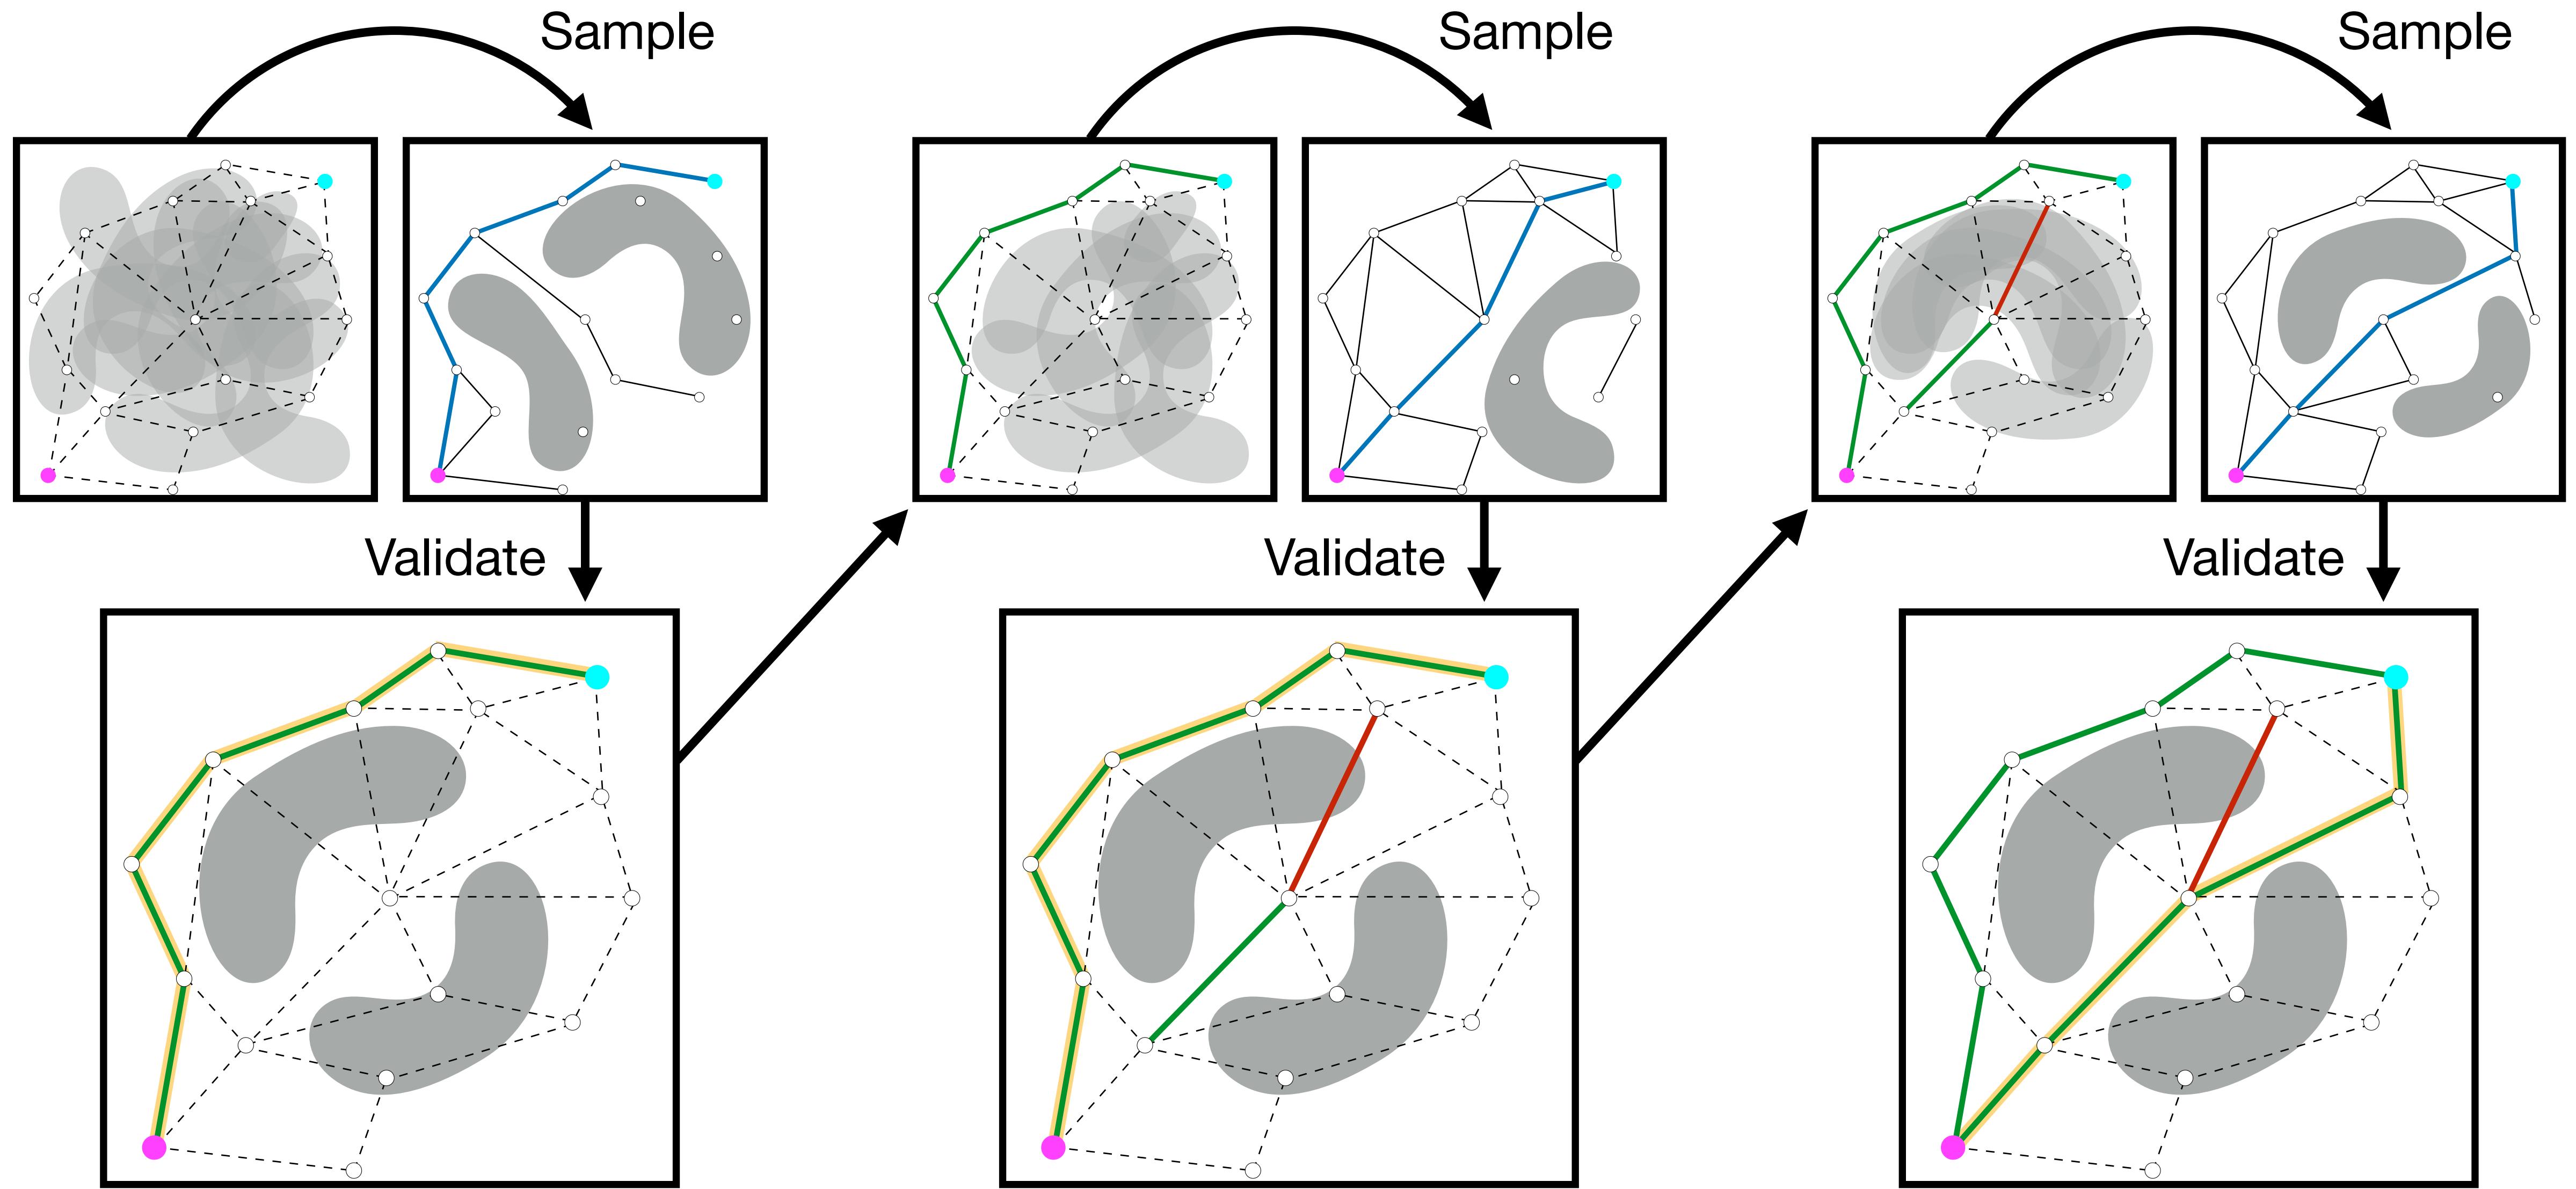
Bayesian  
Regret

Bayes  
Optimality

Shortest Path

Anytime Performance

Feasible Path

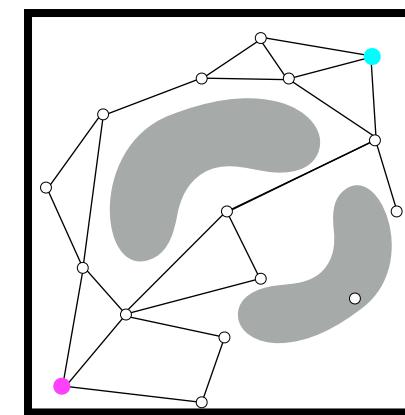


# Bayesian Anytime Motion Planning via Posterior Sampling

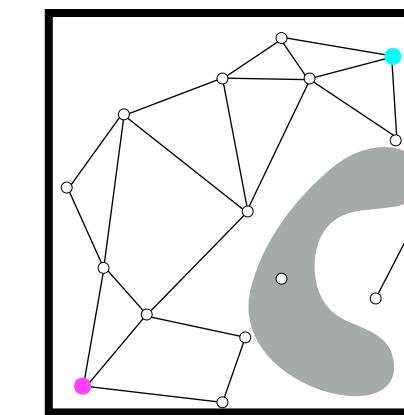
## Posterior Distributions

$$\mathbb{E}_{P(\phi)} \sum_t w(\xi_t)$$

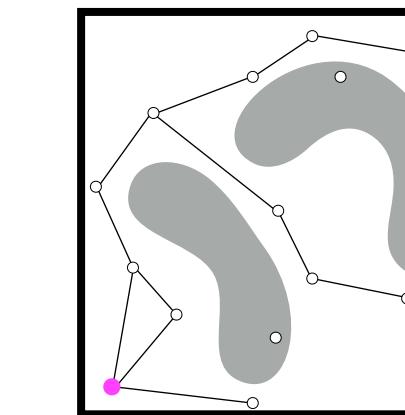
*When you have eliminated the impossible, whatever remains, however improbable, must be the truth.*



$\phi_1$

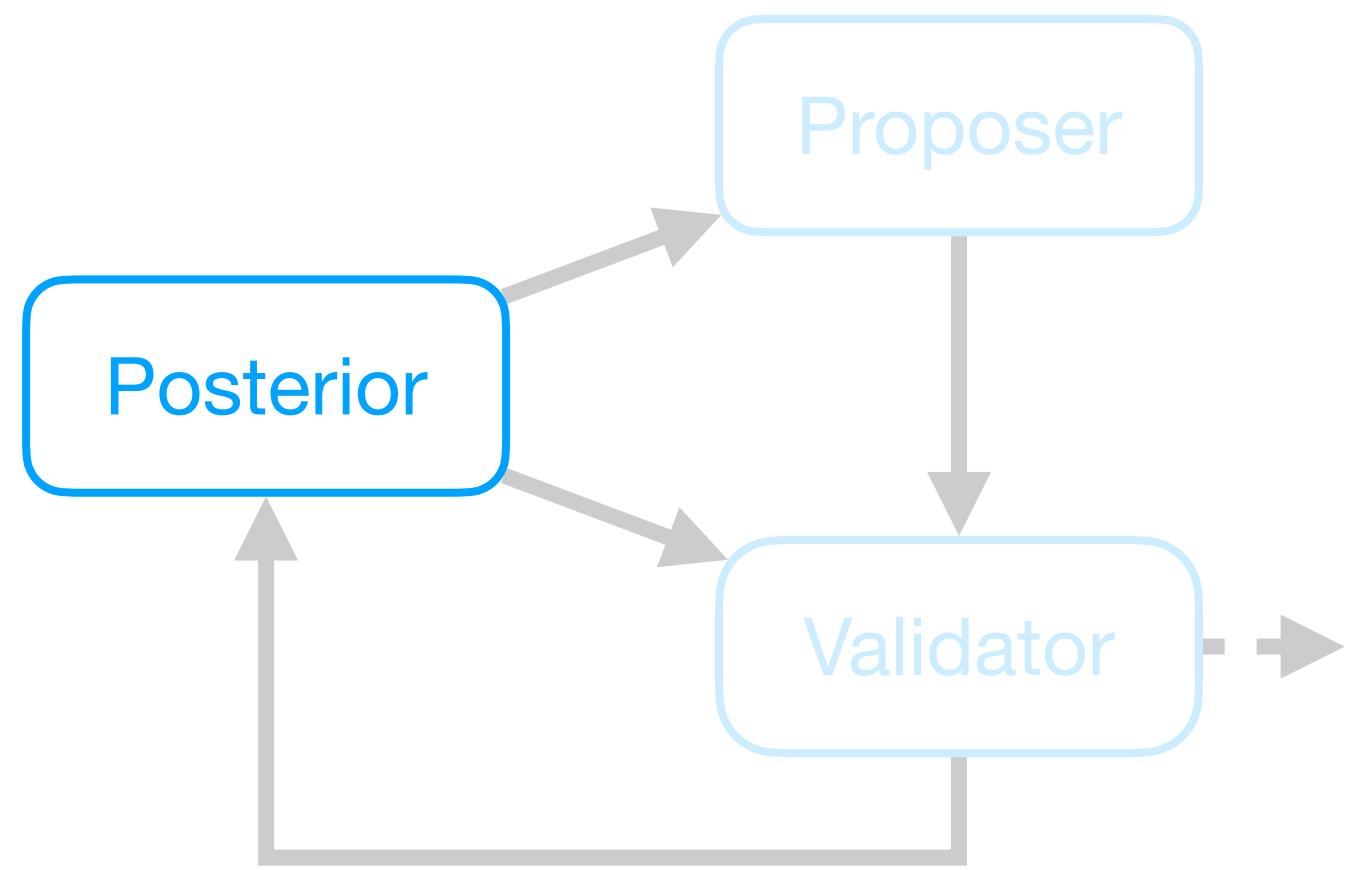


$\phi_2$



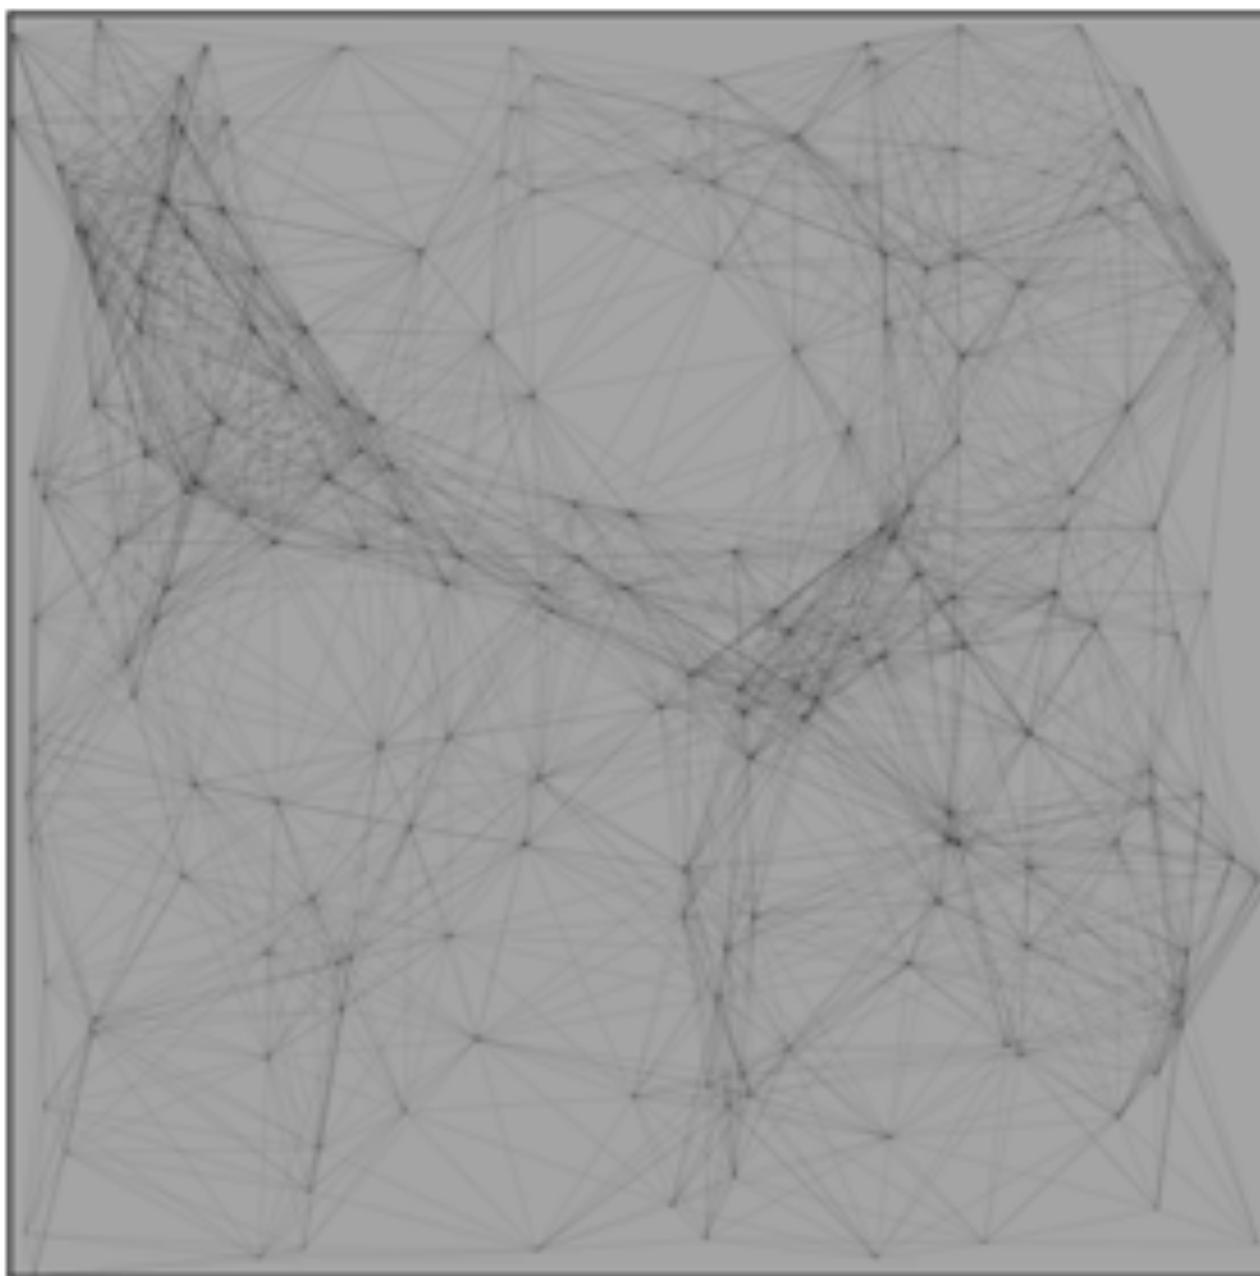
$\phi_3$

# Learning Collision Posteriors



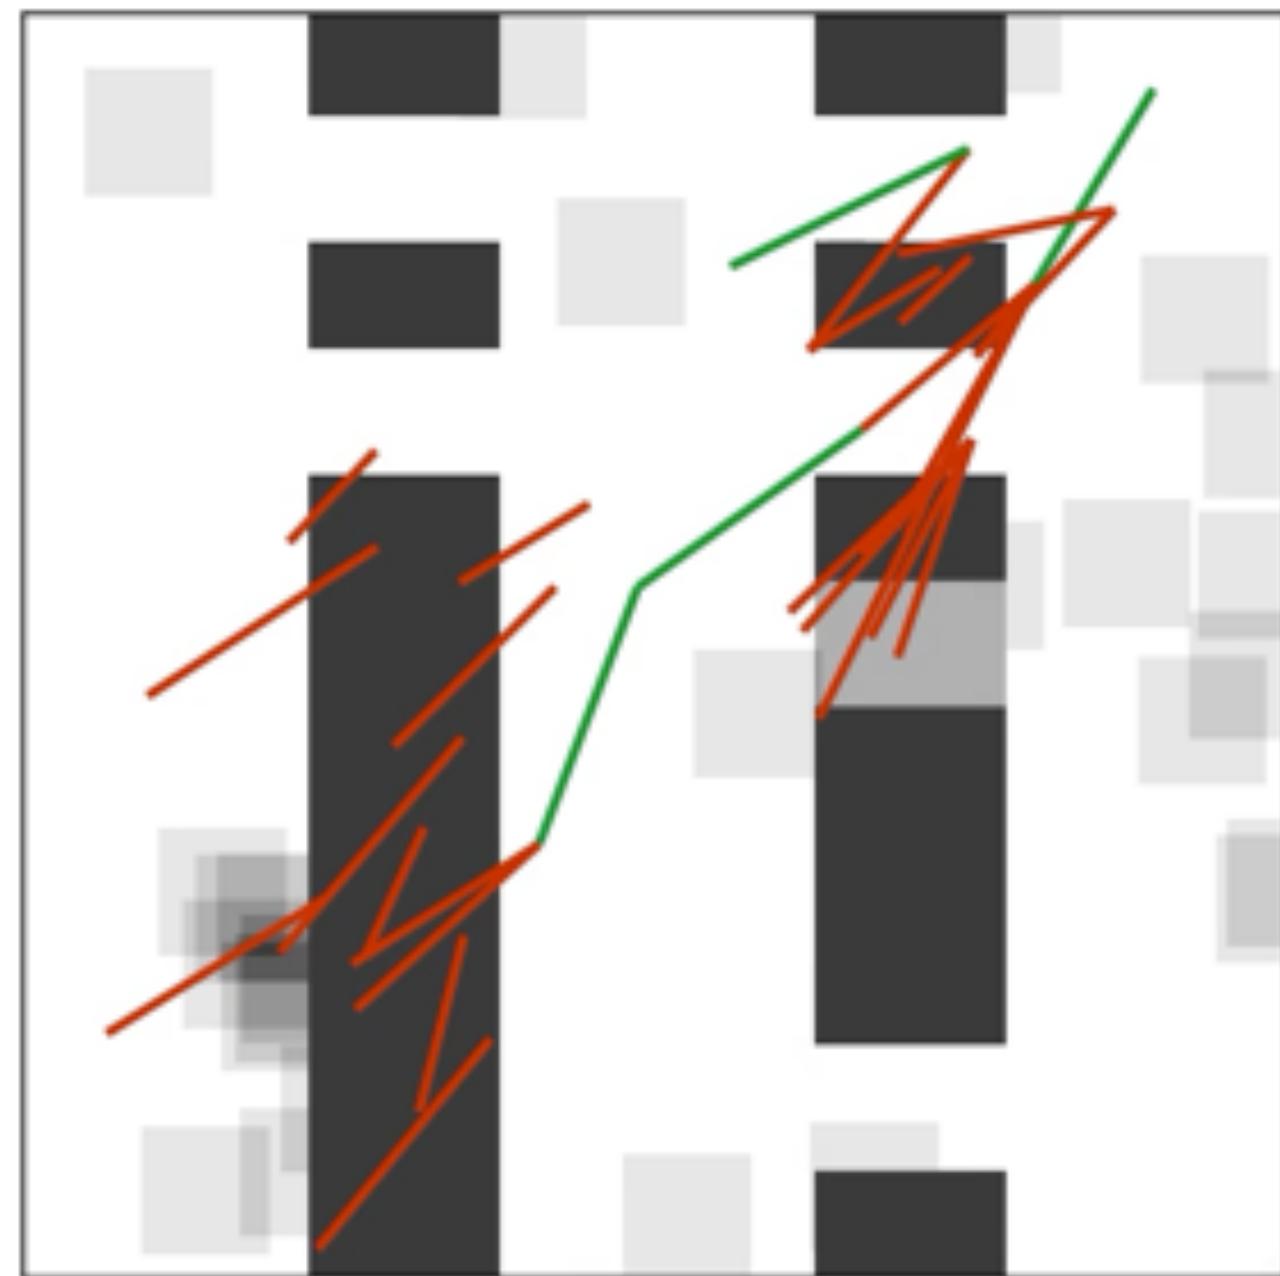
# Learning Collision Posteriors

Nearest Neighbor (NN)

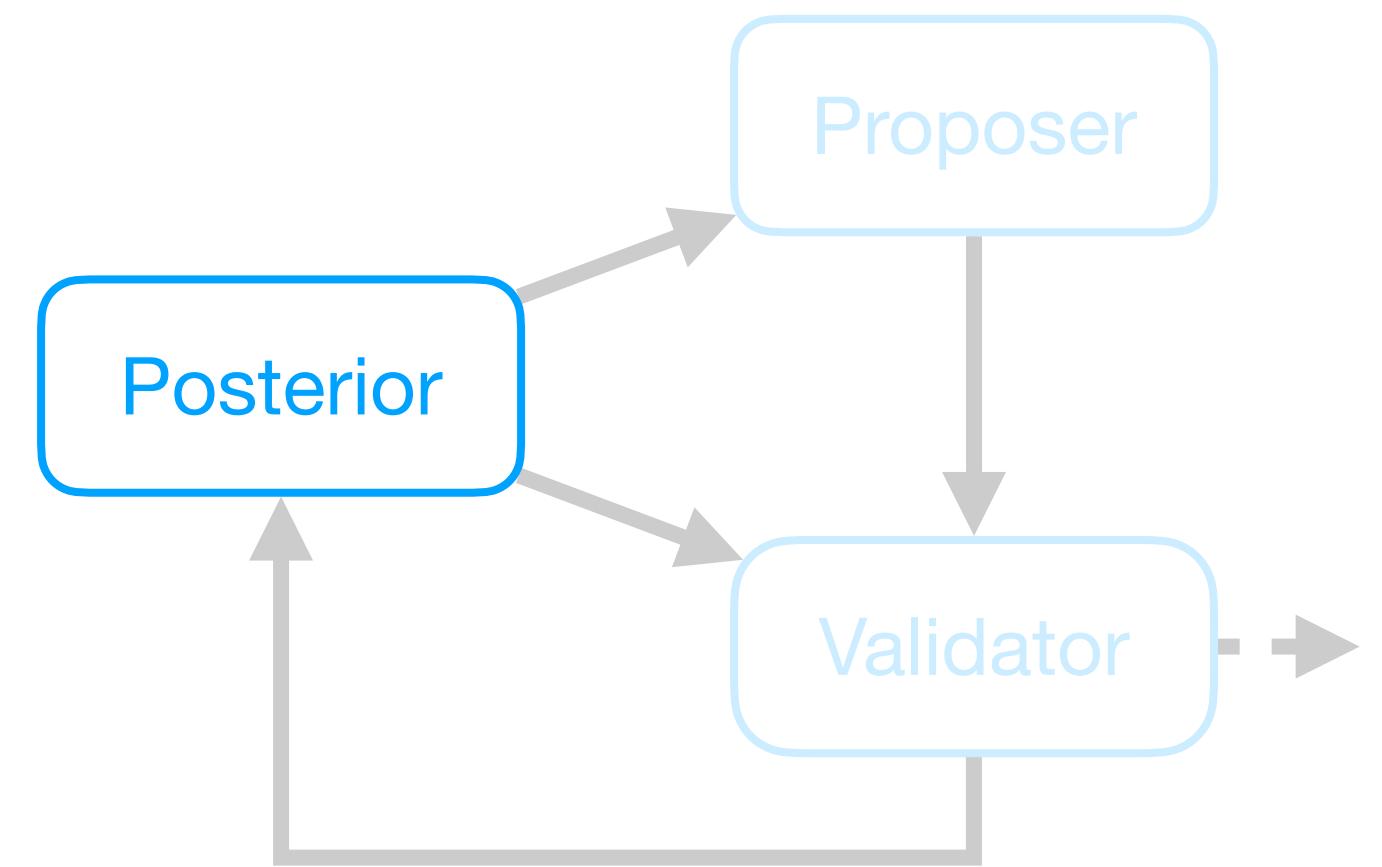


new environments with  
unknown structure

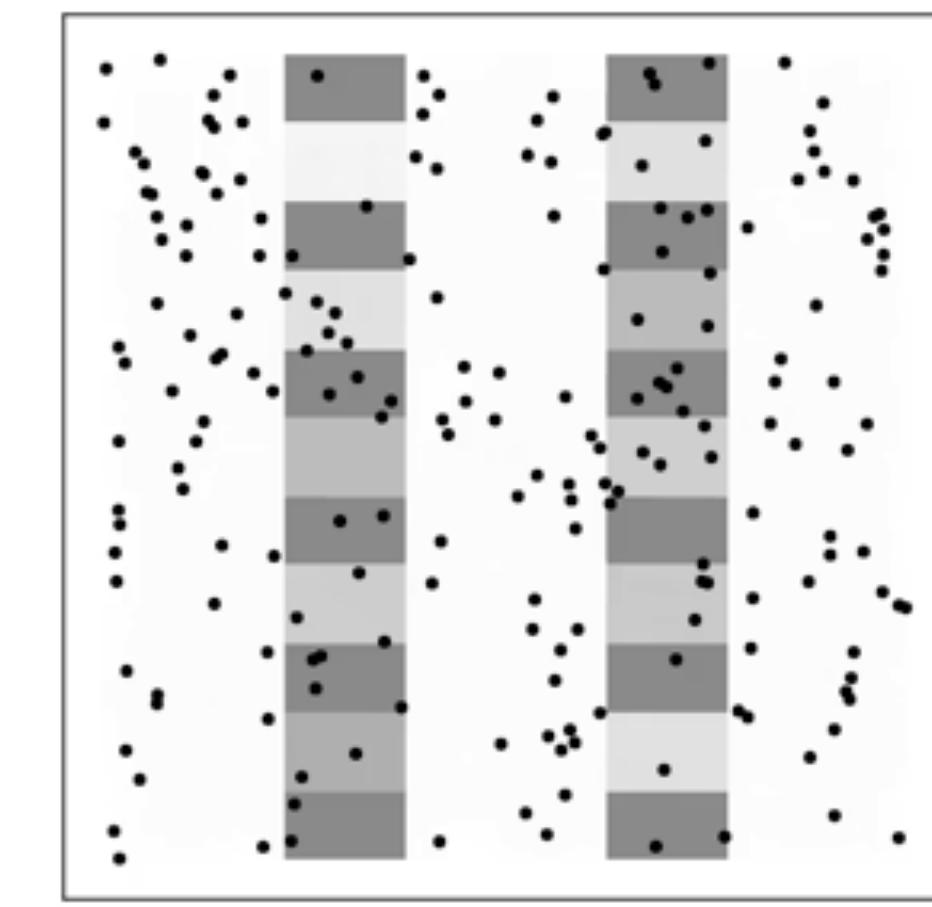
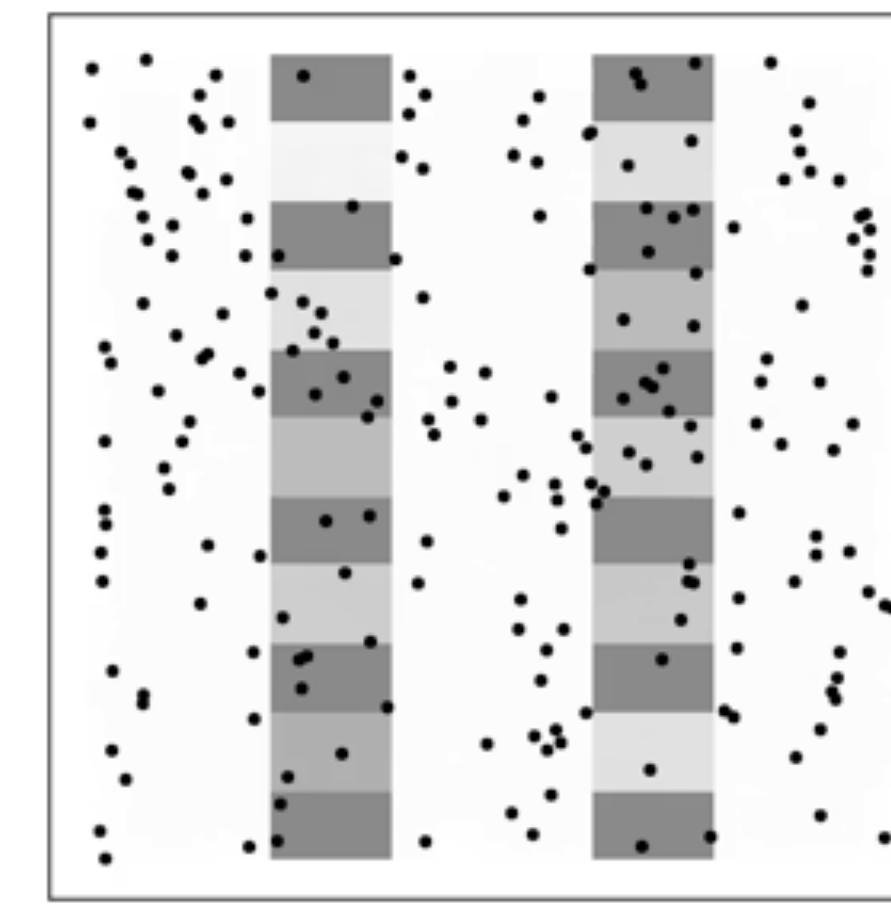
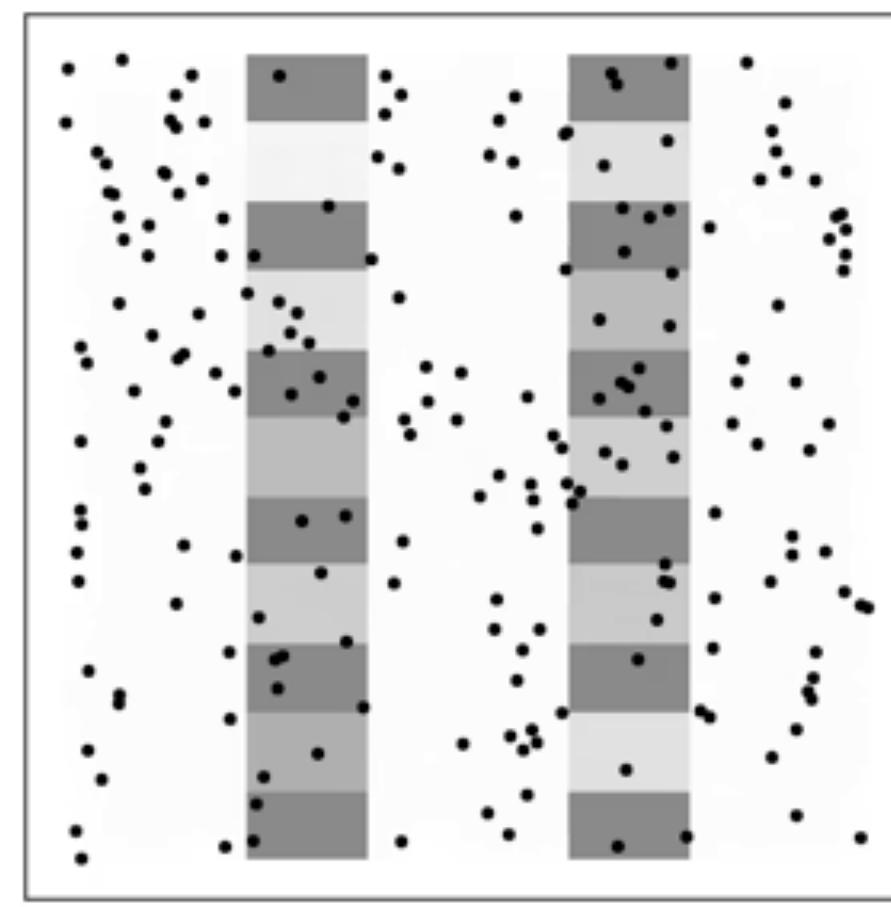
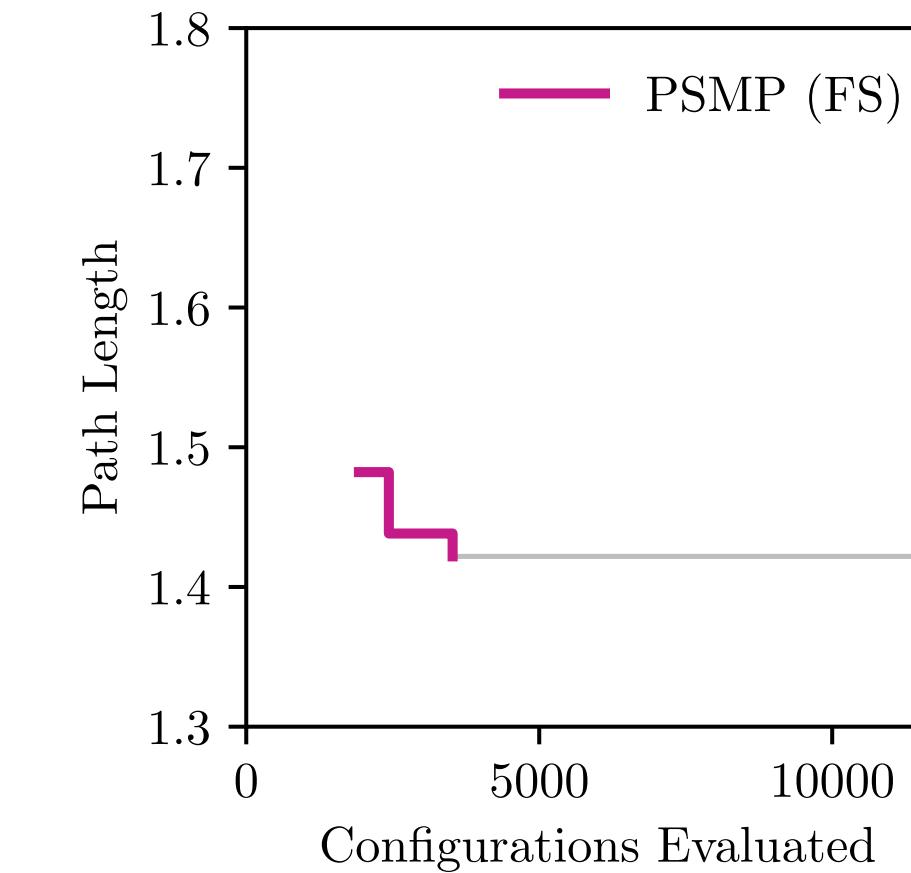
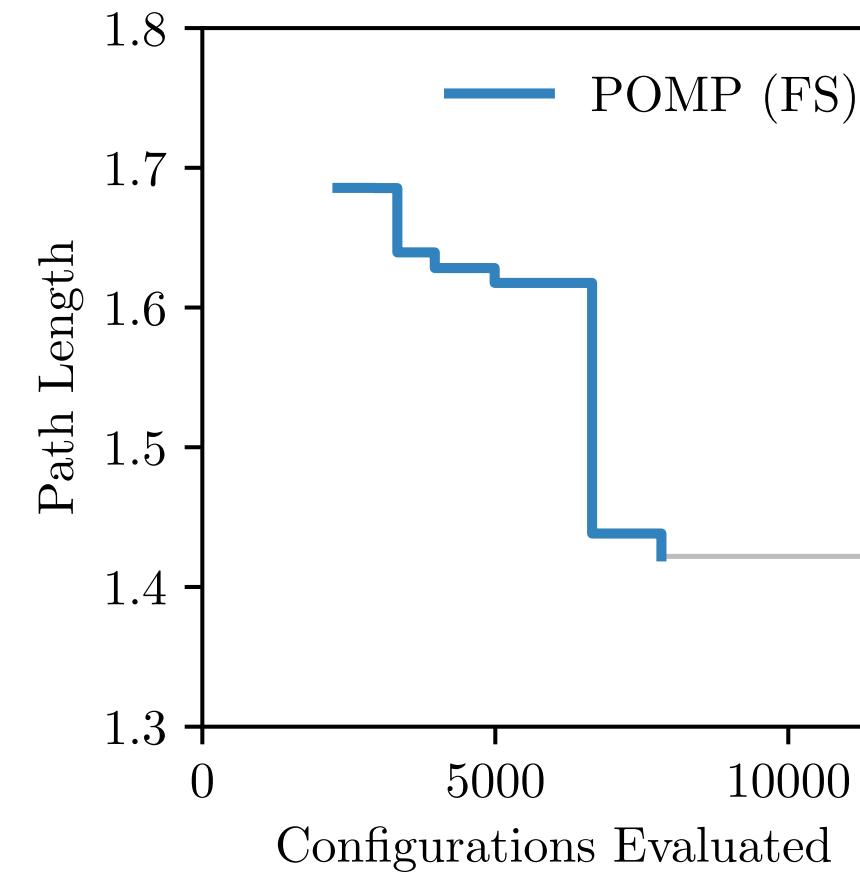
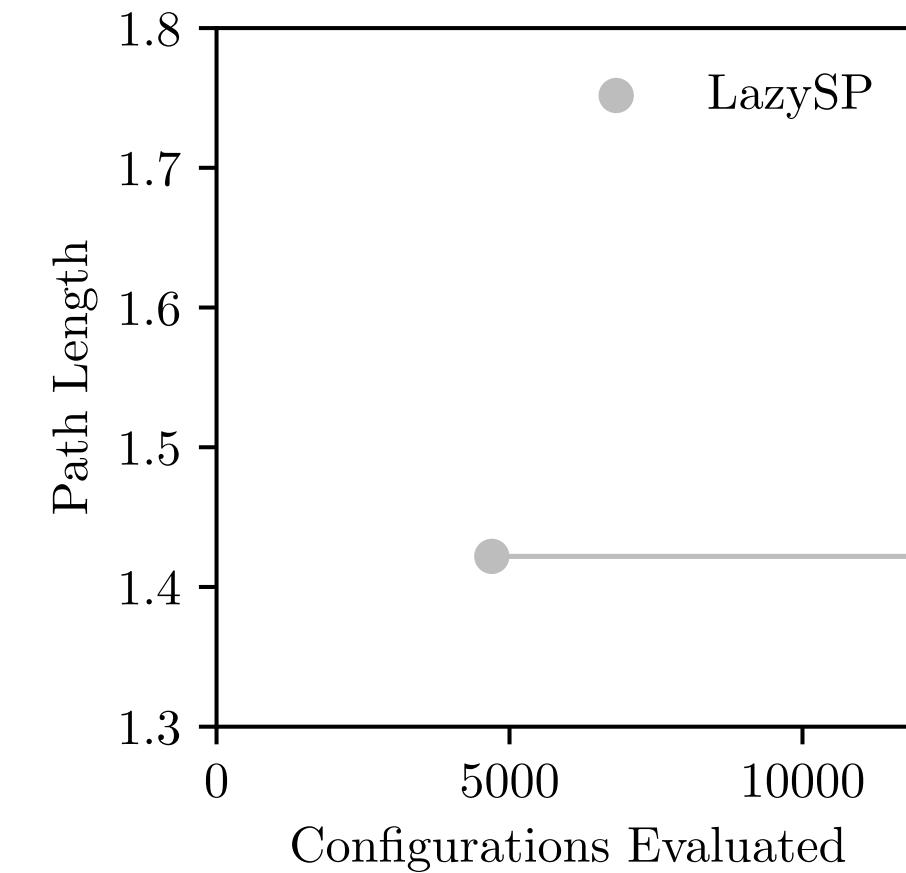
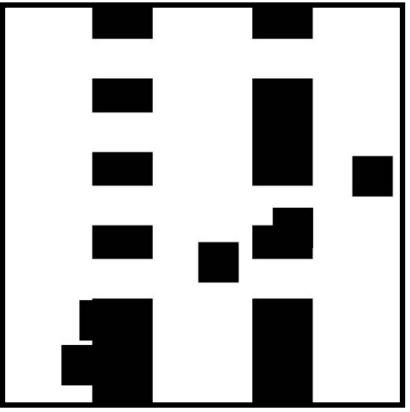
Finite Set (FS)

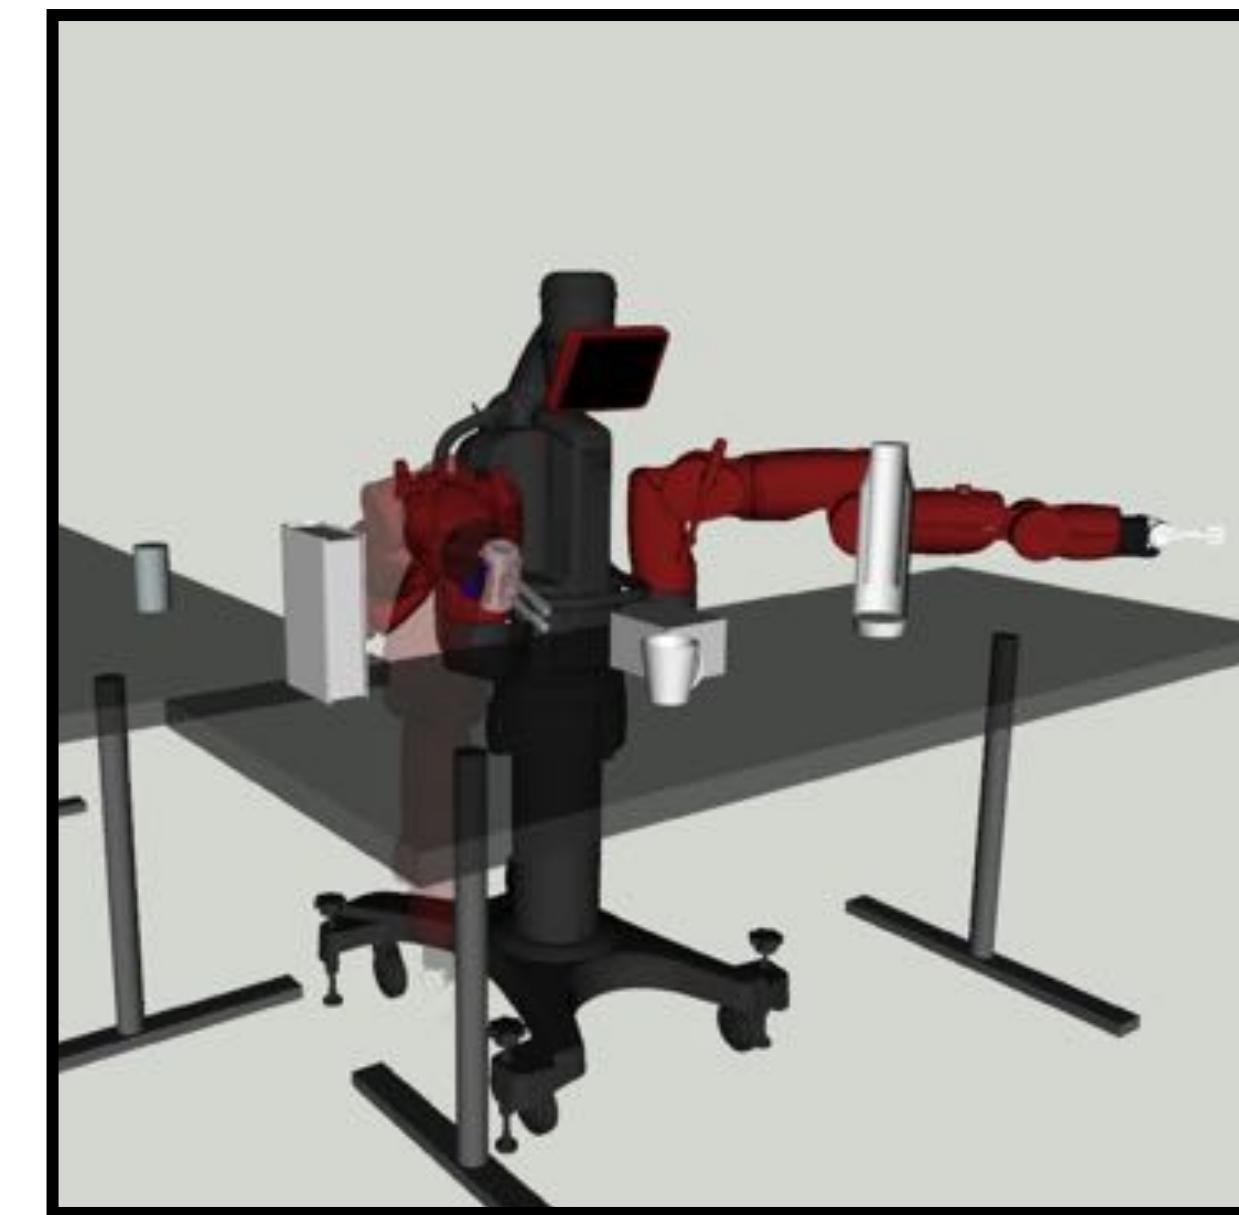
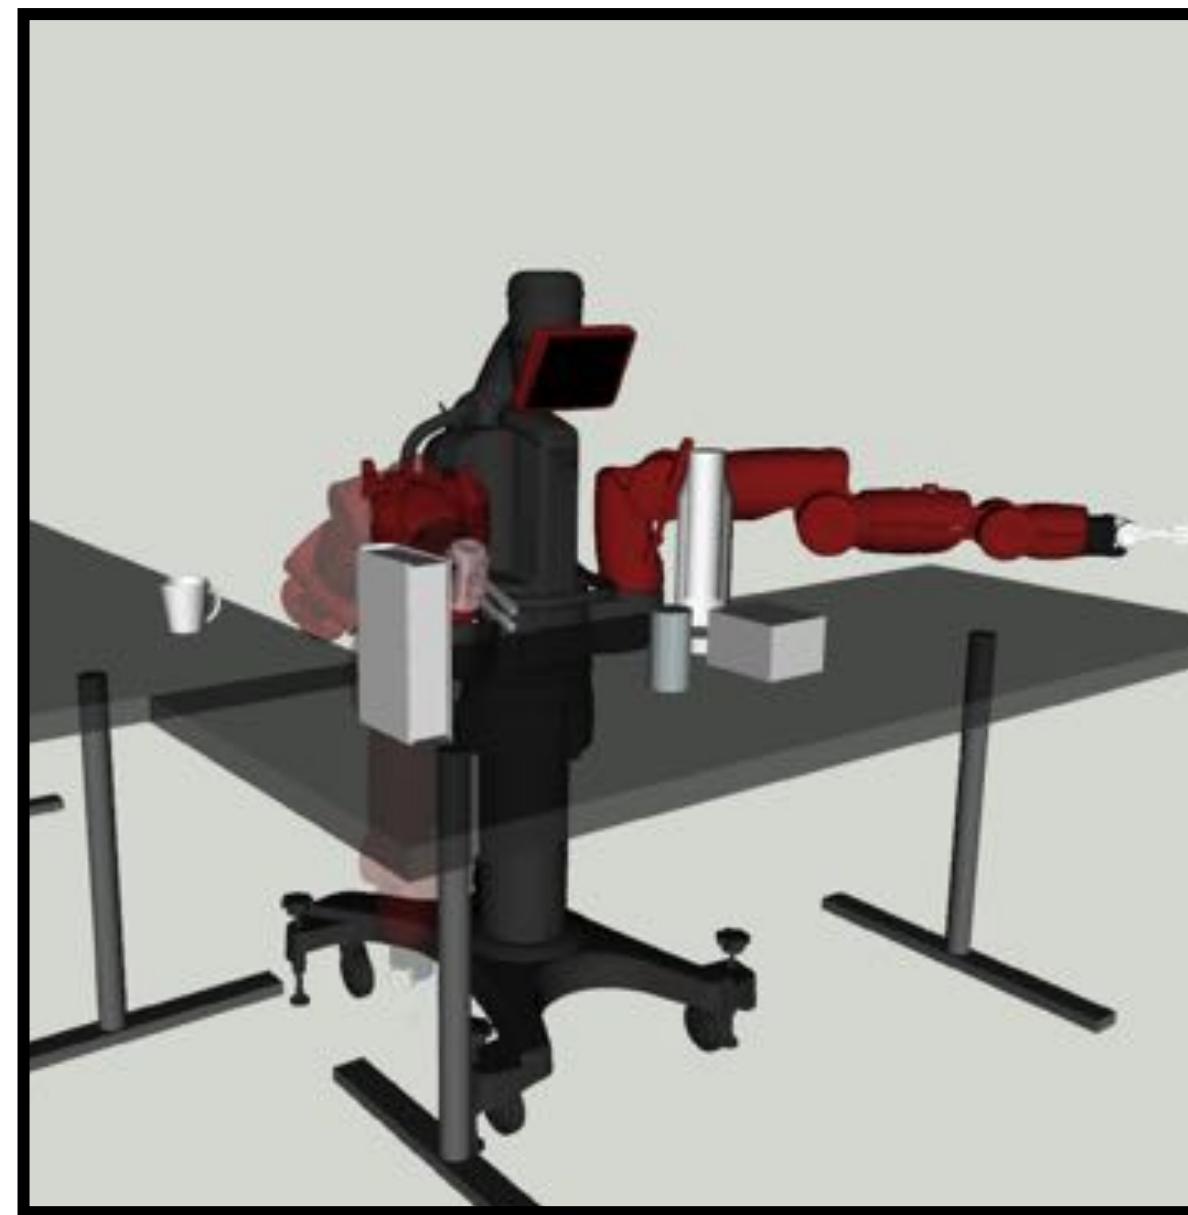


known structure from  
past experience

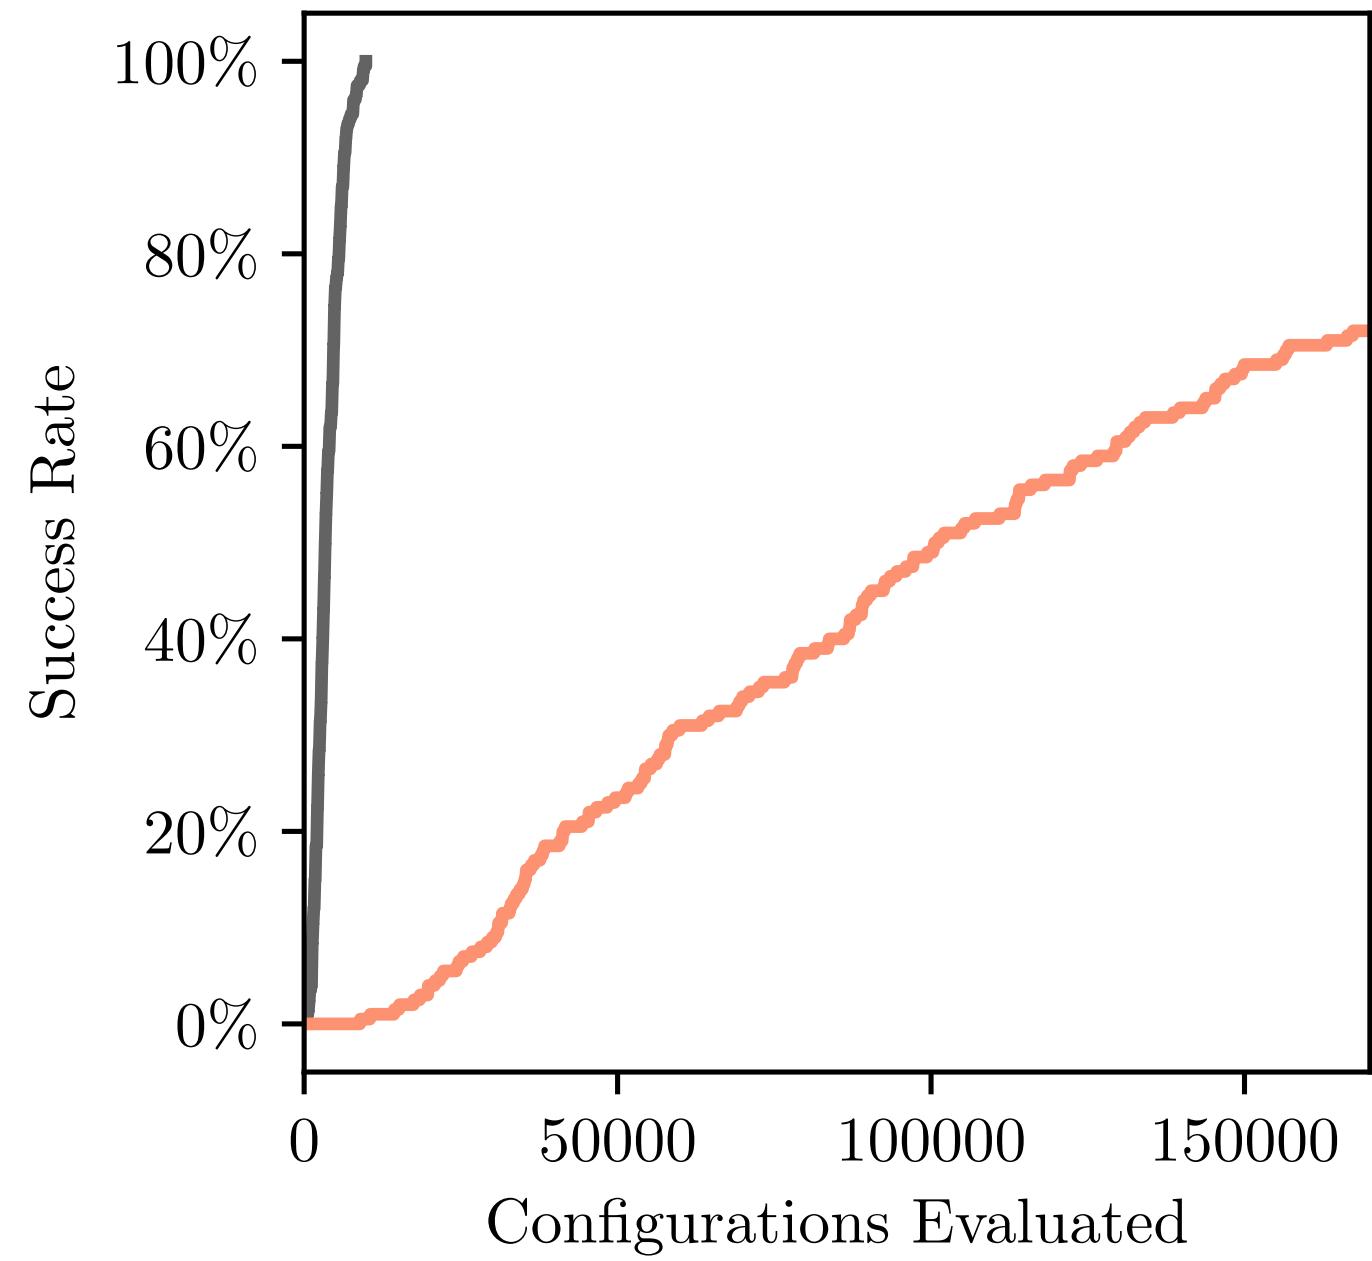
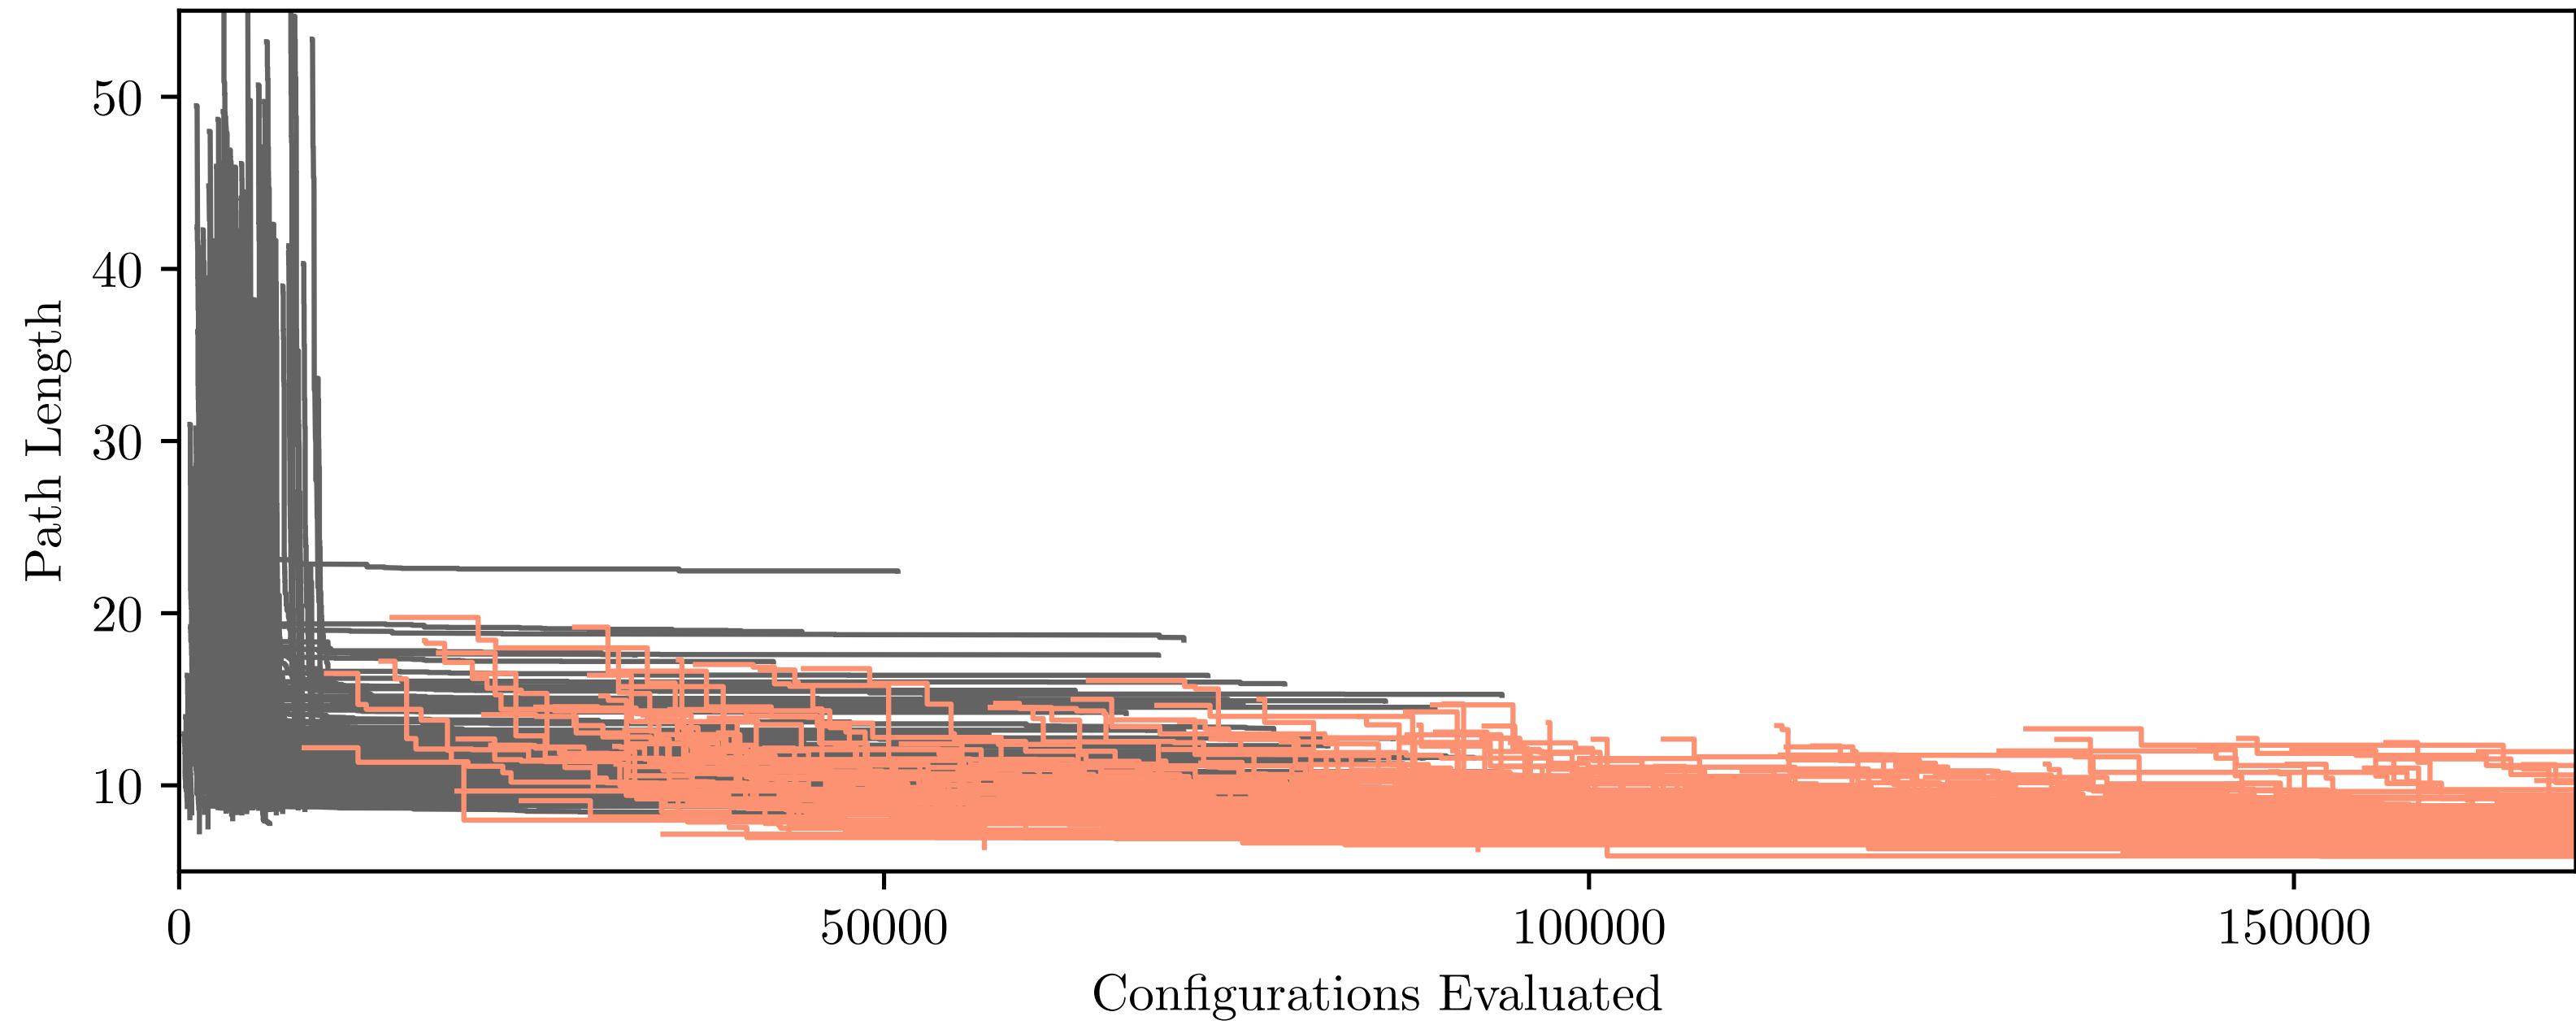


# Shorter paths in fewer collision checks





# RRT\* requires many collision checks

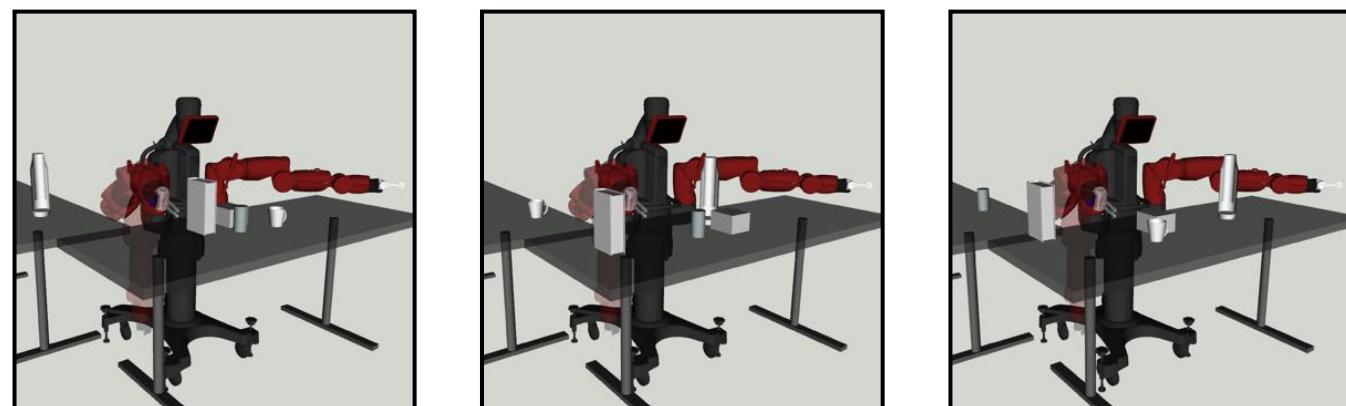


*RRT-Connect: An efficient approach to single-query path planning,*  
Kuffner and Lavalle, ICRA 2000.

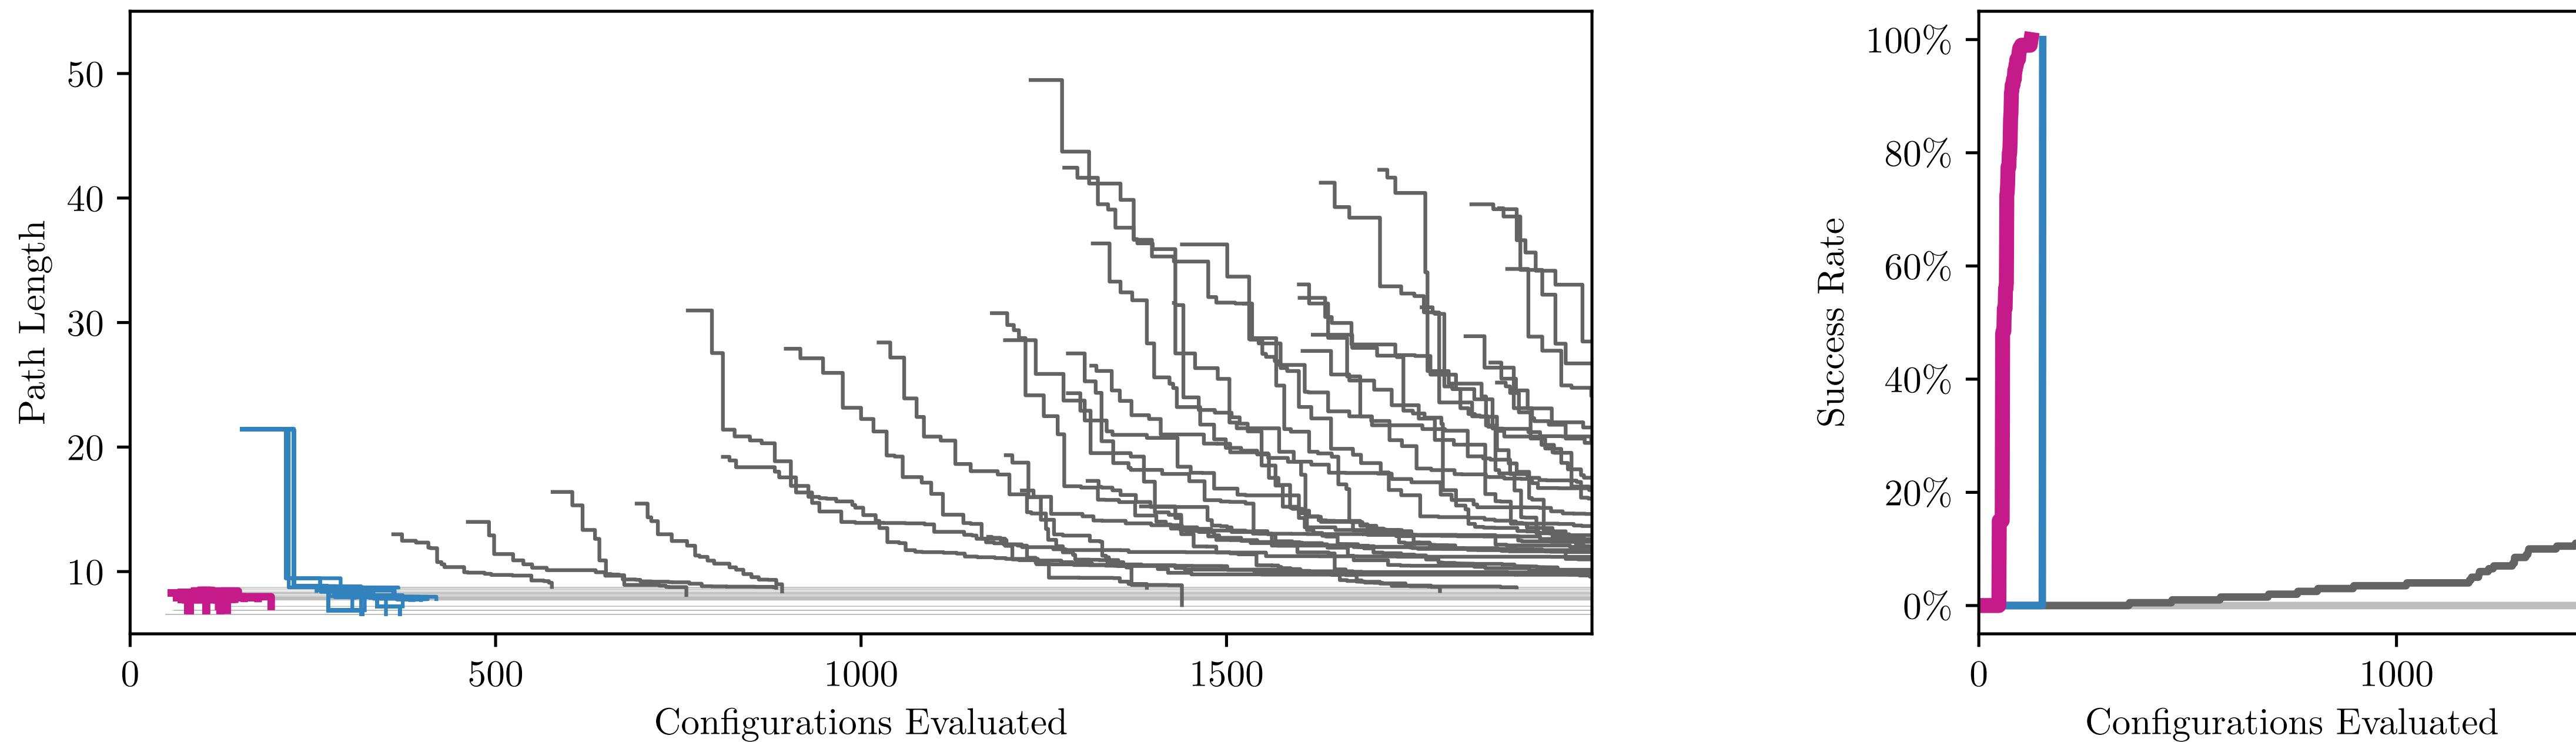
*Sampling-based Algorithms for Optimal Motion Planning, Karaman  
and Frazzoli, IJRR 2011.*

RRT\*

RRT + PS



# Outperforms common anytime heuristics



■ POMP (FS)

■ PSMP (FS)

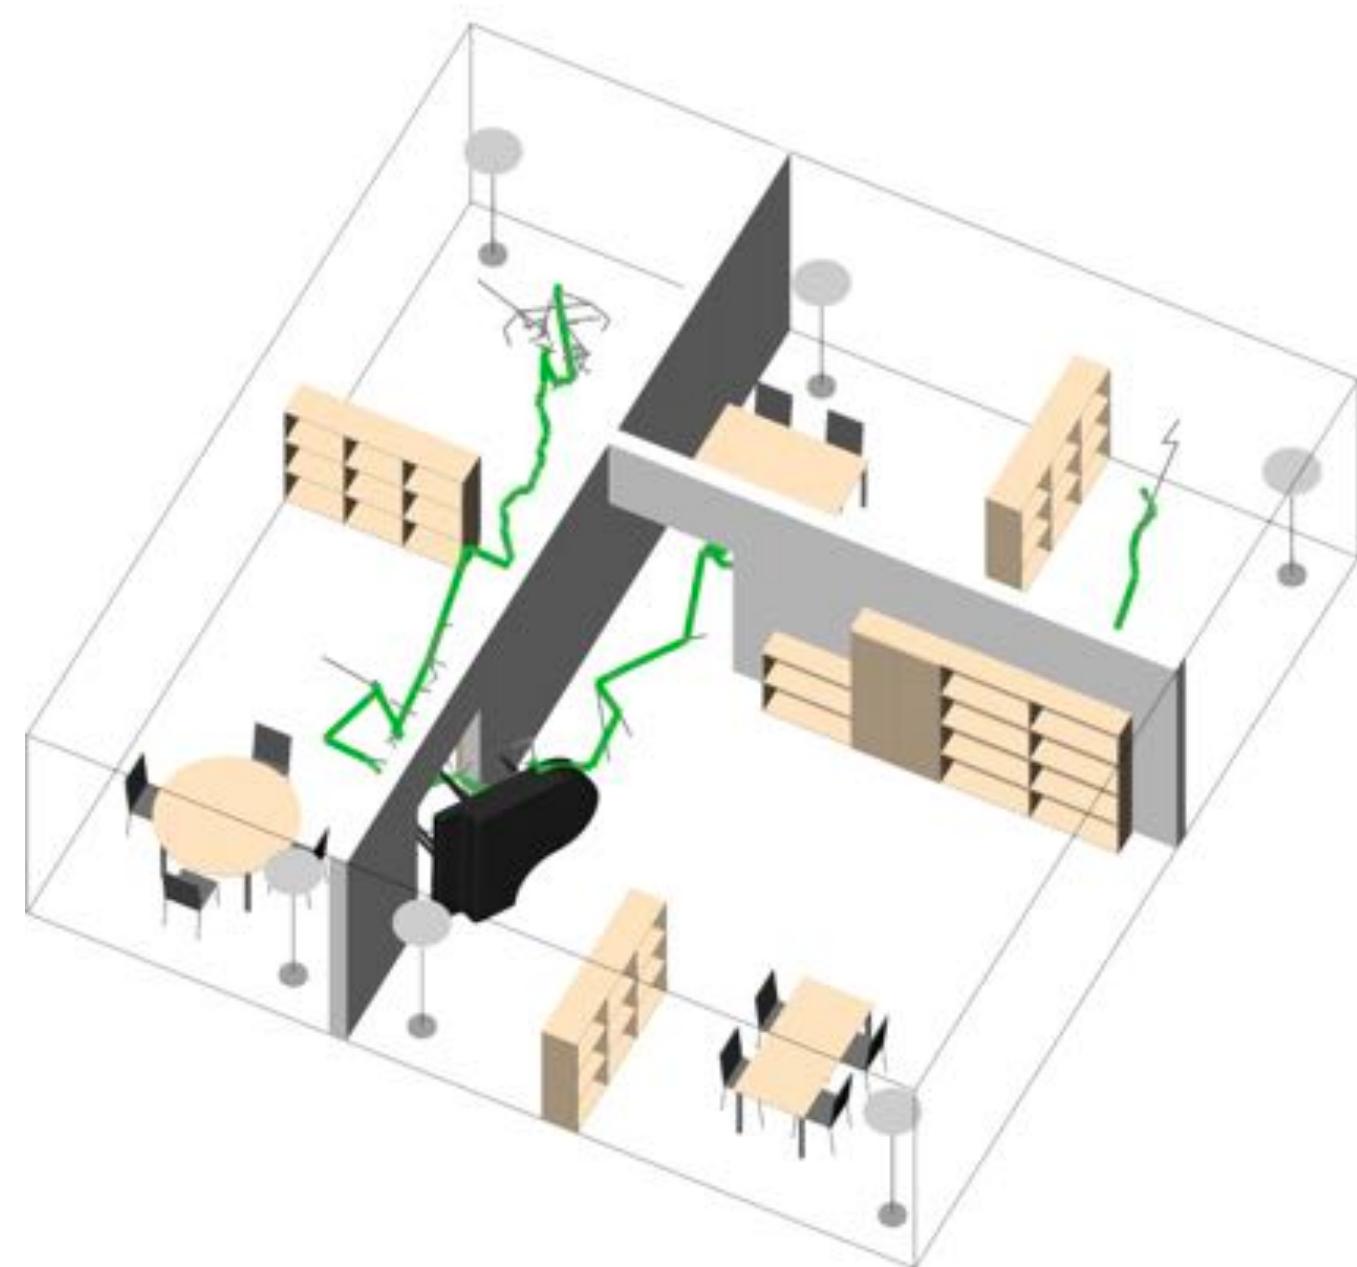
■ LAZYSP

■ RRT + PS



**We formalize anytime search as  
Bayesian Reinforcement Learning**

# The Experienced Piano Movers' Problem



New Piano.  
New House.  
Same Mover.

# Search = Eliminating Paths

## Optimal Substructure

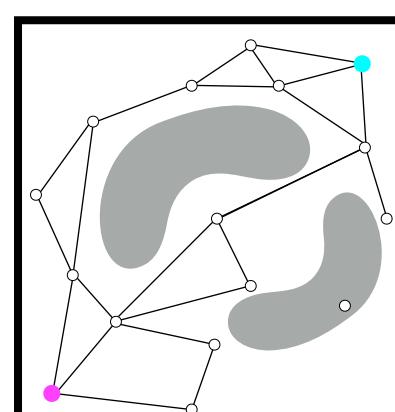
$$f(a) < f(b) \implies f(a \circ x) < f(b \circ x) \forall x$$

*You will never catch up.*

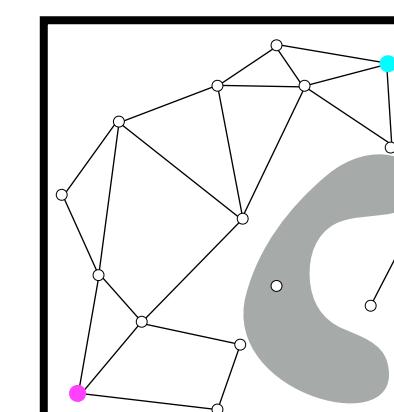
## Posterior Distributions

$$\mathbb{E}_{P(\phi)} \sum_t w(\xi_t)$$

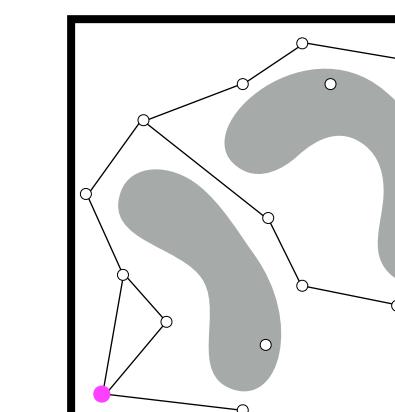
*When you have eliminated the impossible, whatever remains, however improbable, must be the truth.*



$\phi_1$



$\phi_2$



$\phi_3$

# RL = Eliminating Policies

## Optimal Substructure

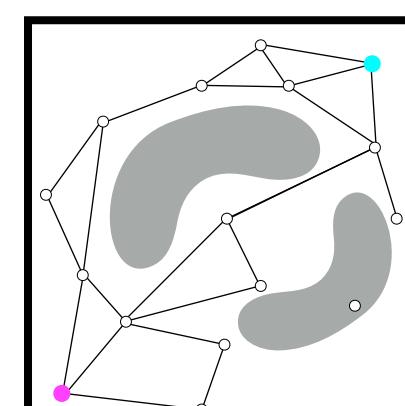
$$f(a) < f(b) \implies f(a \circ x) < f(b \circ x) \forall x$$

*You will never catch up.*

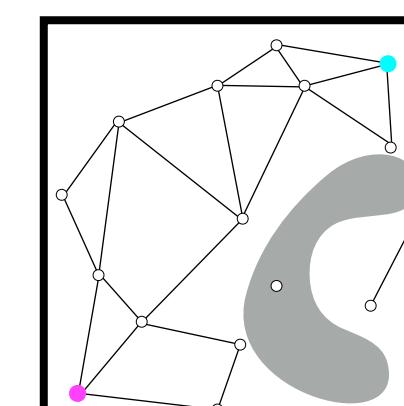
## Posterior Distributions

$$\mathbb{E}_{P(\phi)} \sum_t w(\xi_t)$$

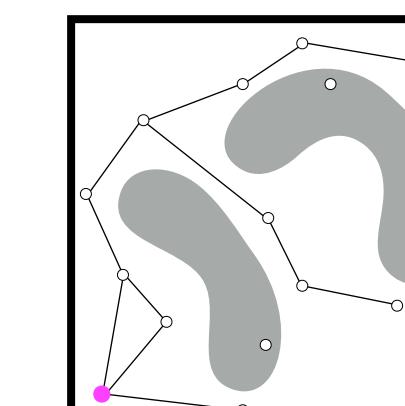
*When you have eliminated the impossible, whatever remains, however improbable, must be the truth.*



$\phi_1$



$\phi_2$



$\phi_3$

Search for Optimal Solutions:  
the Heart of Heuristic Search is Still Beating

Ariel Felner  
ISE Department  
Ben-Gurion University  
ISRAEL  
felner@bgu.ac.il

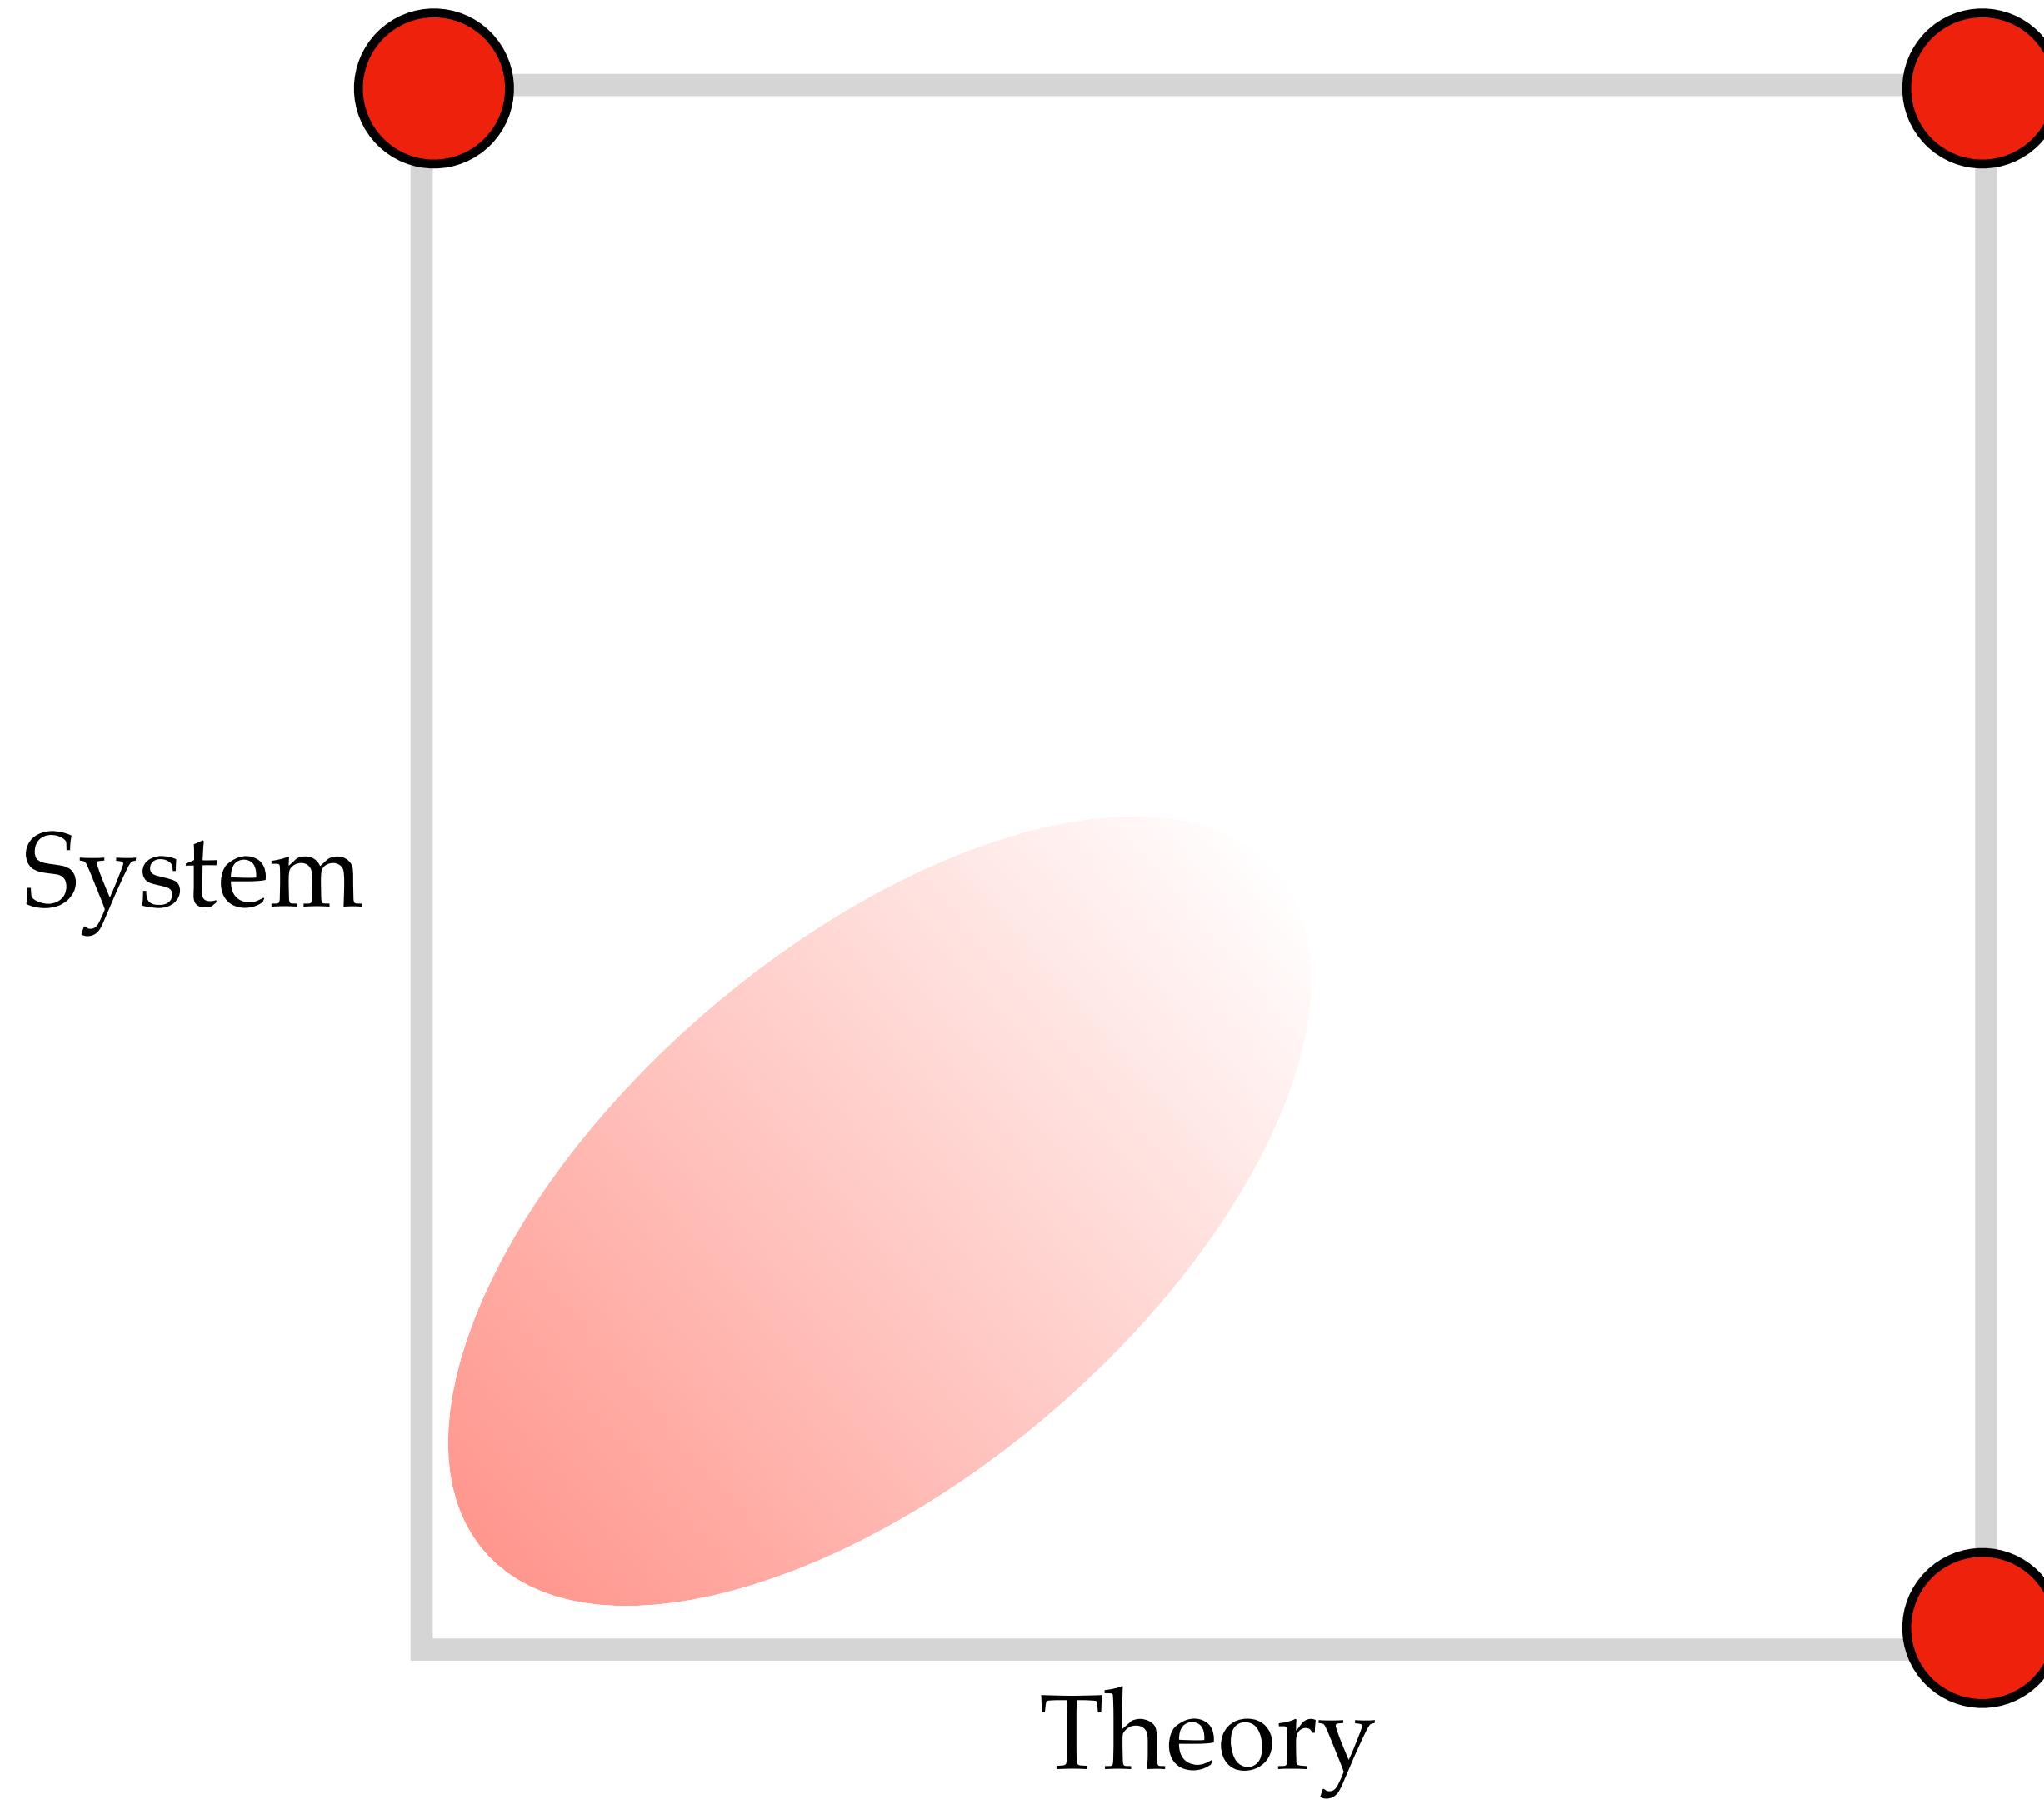


# Exploit Structure Embrace Laziness

Prove some  
Damn Theorems

1

# Aim for the Corners!



*Pareto-Optimal Search over Configuration Space Beliefs for Anytime Motion Planning*, Choudhury et al., IROS 2016.

*A Unifying Formalism for Shortest Path Problems with Expensive Edge Evaluations via Lazy Best-First Search over Paths with Edge Selectors*, Dellin and Srinivasa, ICAPS 2016.

*Near-Optimal Edge Evaluation in Explicit Generalized Binomial Graphs*, Choudhury et al., N\*IPS 2017.

*Bayesian Active Edge Evaluation on Expensive Graphs*, Choudhury et al., IJCAI 2018.

*The Provable Virtue of Laziness in Motion Planning*, Hagtalab et al., ICAPS 2018.

*Lazy Receding Horizon A\* for Efficient Path Planning in Graphs with Expensive-to-Evaluate Edges*, Mandalika et al., ICAPS 2018.

*Leveraging experience in lazy search*, Bhardwaj et al., RSS 2019.

*Generalized Lazy Search for Robot Motion Planning: Interleaving Search and Edge Evaluation via Event-based Toggles*, Mandalika et al., ICAPS 2019.

*Posterior Sampling for Anytime Motion Planning on Graphs with Expensive-to-Evaluate Edges*, Hou et al., ICRA 2020.

*Bayesian Residual Policy Optimization: Scalable Bayesian Reinforcement Learning with Clairvoyant Experts*, Lee et al., arXiv: 2002.03042

## Coauthors

Mohan Bhardwaj, Byron Boots, Sushman Choudhury, Sanjiban Choudhury, Chris Dellin, Nika Hagtalab, Brian Hou, Shervin Javadani, Gilwoo Lee, Simon Mackenzie, Aditya Mandalika, Ariel Procaccia, Oren Salzman, Sebastian Scherer.

## Collaborators

Tim Barfoot, Dmitry Berenson, Jon Gammell, David Hsu, Brad Saund, Rahul Vernwal.

## Smarty Pants

Drew Bagnell, Kostas Bekris, Dan Halperin, Kris Hauser, Sven Koenig, Max Likhachev.

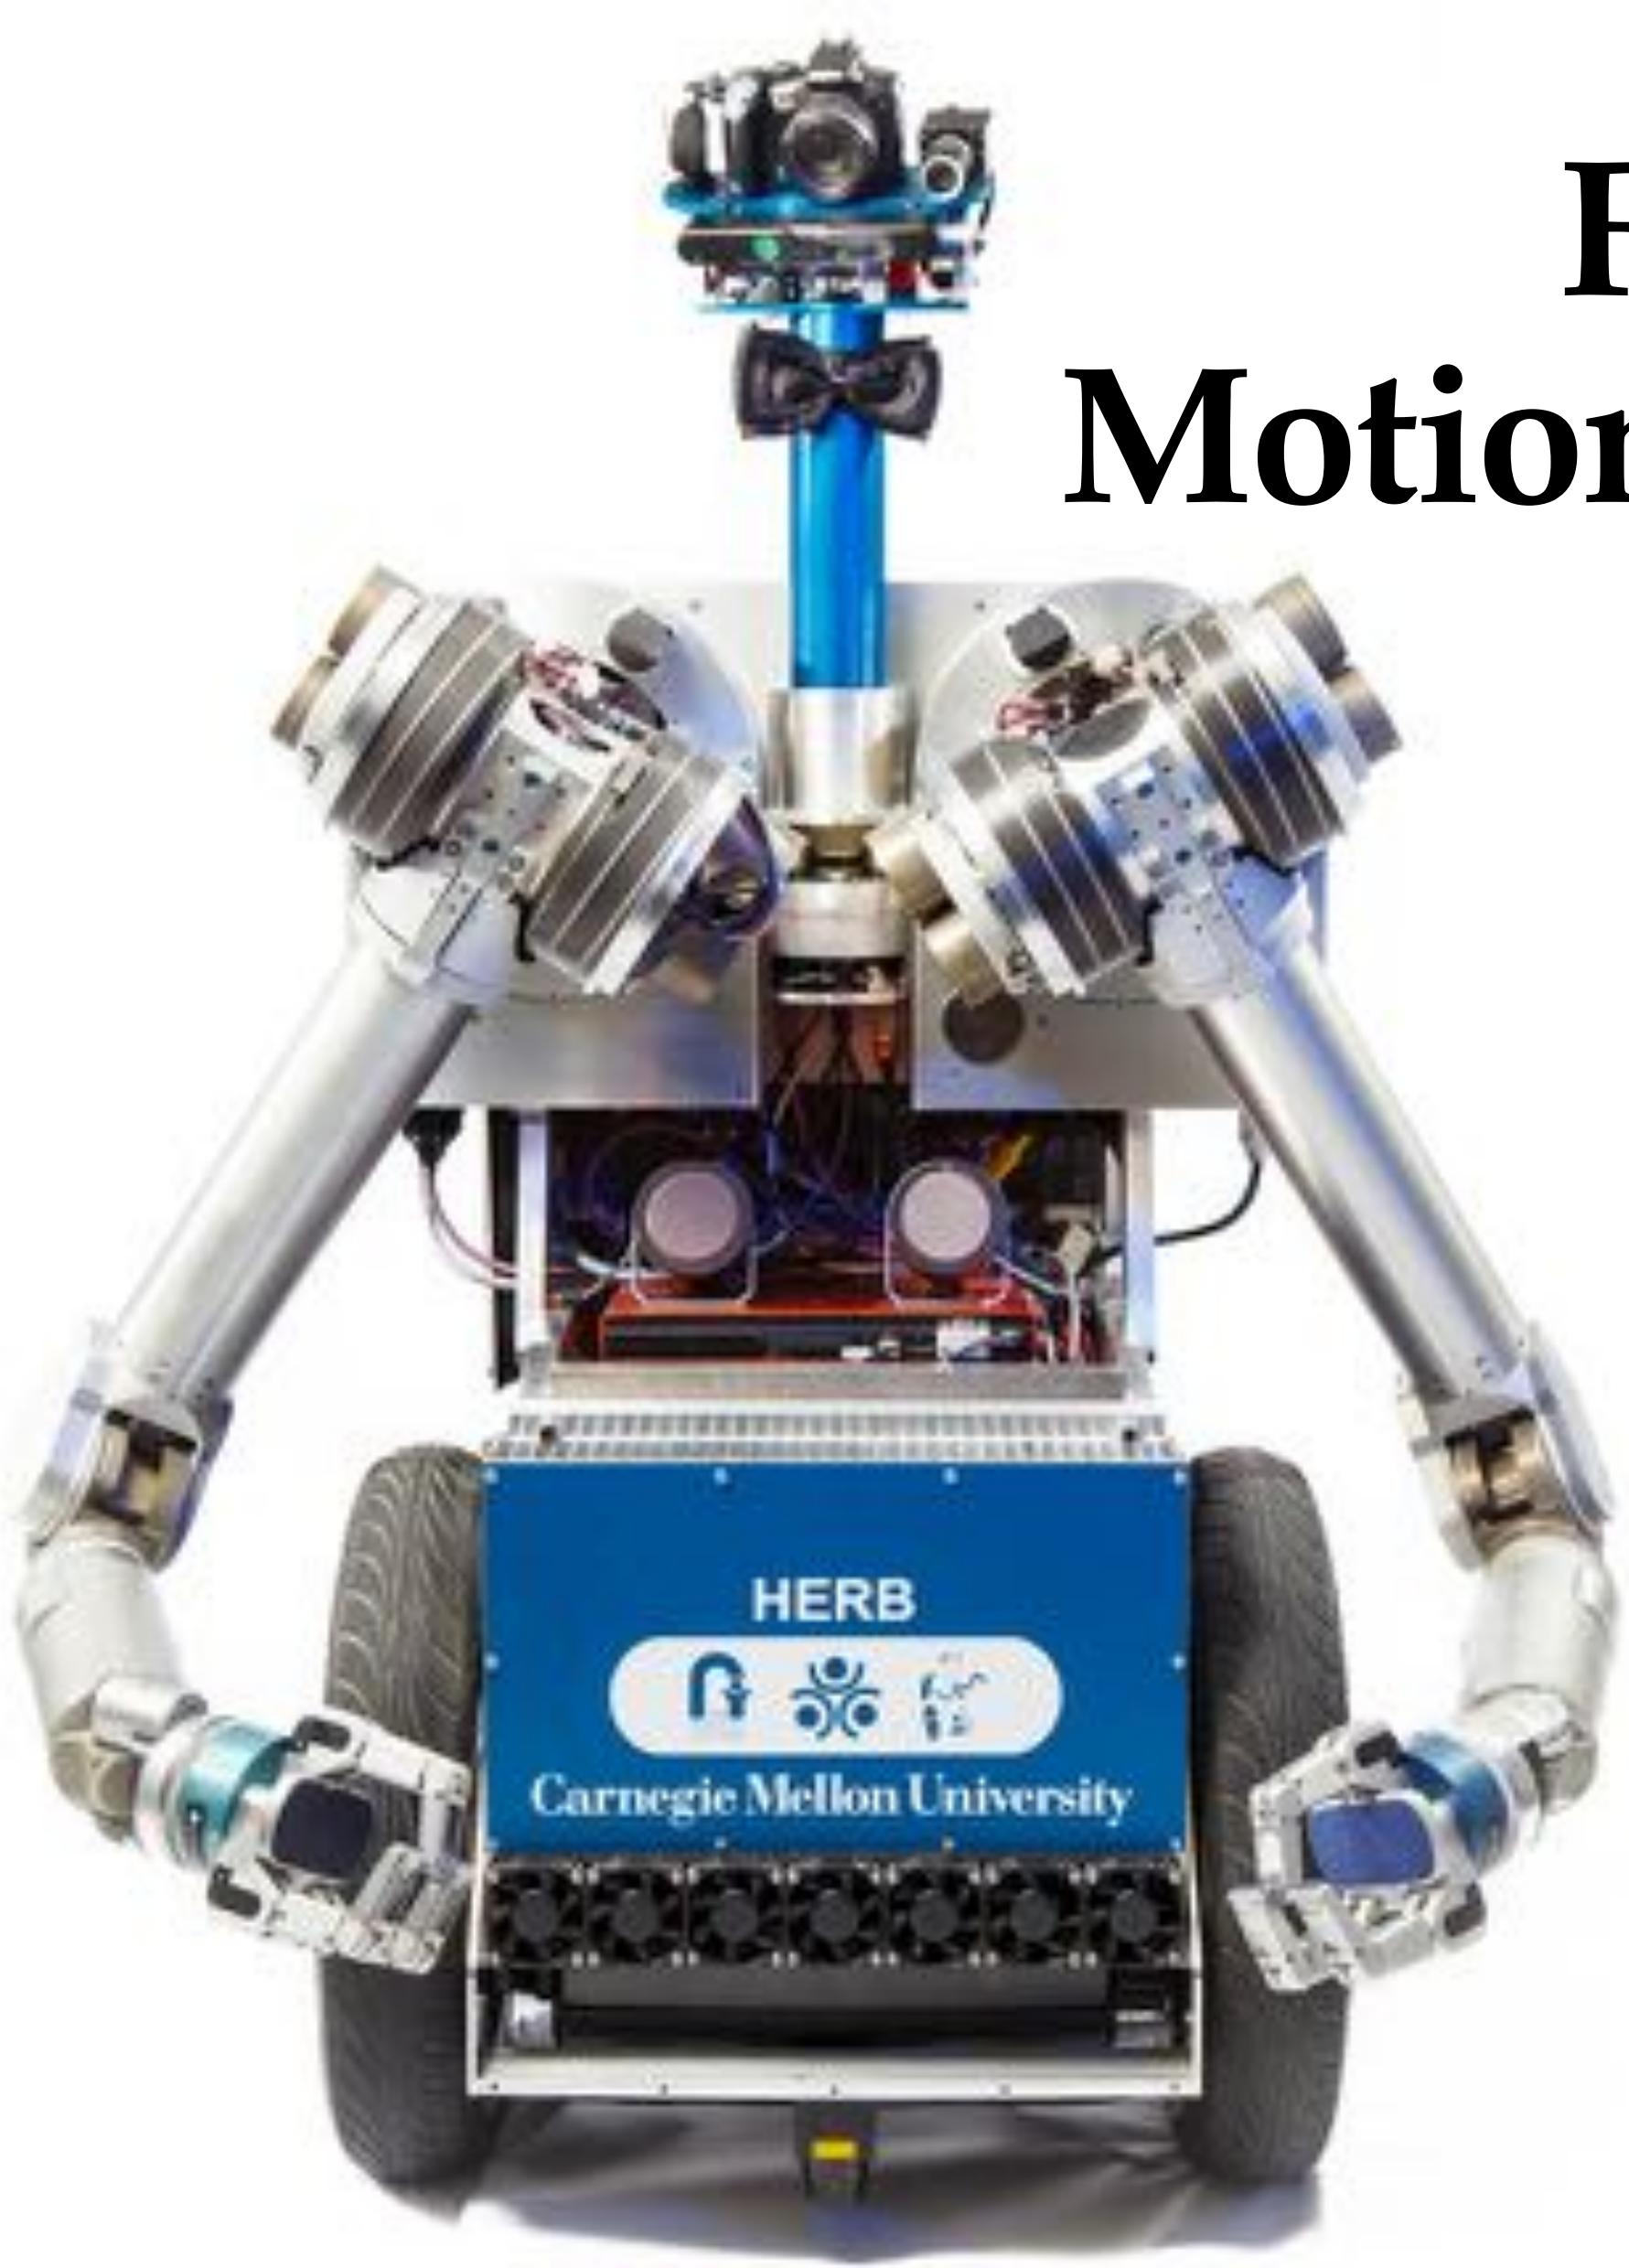
## Funders

Army, DARPA, Honda, NIH, NSF, ONR.

<https://personalrobotics.cs.washington.edu/publications/>

<https://www.amazon.jobs/en/teams/rai>





# Formalizing Connections Between Motion Planning and Machine Learning



Siddhartha Srinivasa  
Boeing Endowed Professor  
University of Washington

**W** PAUL G. ALLEN SCHOOL  
OF COMPUTER SCIENCE & ENGINEERING